

# Modular Monoliths



Simon Brown  
@simonbrown

Big design up front  
vs  
no design up front

Working software  
over  
comprehensive  
documentation

#NoEstimates

# #NoProjects

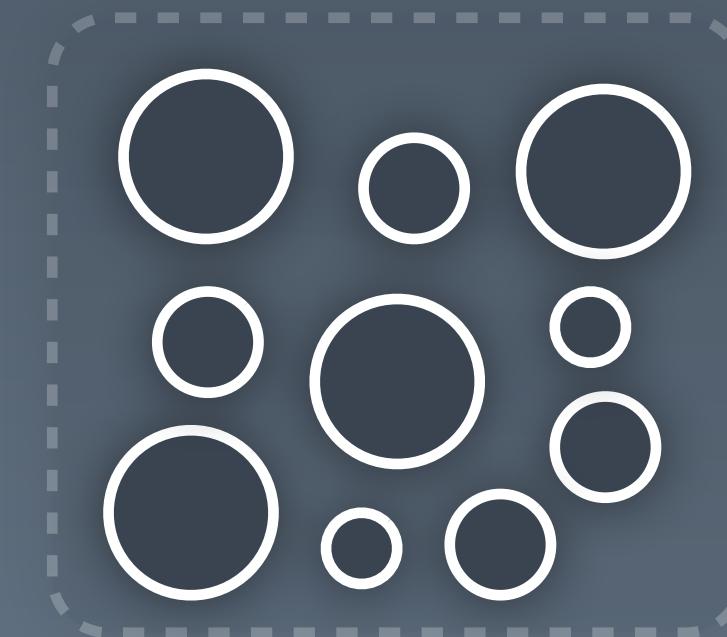
# Monoliths vs microservices

*Monolithic  
architecture*



*Service-based  
architecture*

(SOA, microservices, etc)



The hidden assumption is that all  
monoliths are  
big balls of mud



**Simon Brown**

@simonbrown

I'll keep saying this ... if people can't build monoliths properly, microservices won't help.  
[#qconlondon](#) [#DesignThinking](#) [#Modularity](#)

---

RETWEETS

**244**

LIKES

**100**



---

10:49 AM - 4 Mar 2015



...



**Architect Clippy**

@architectclippy

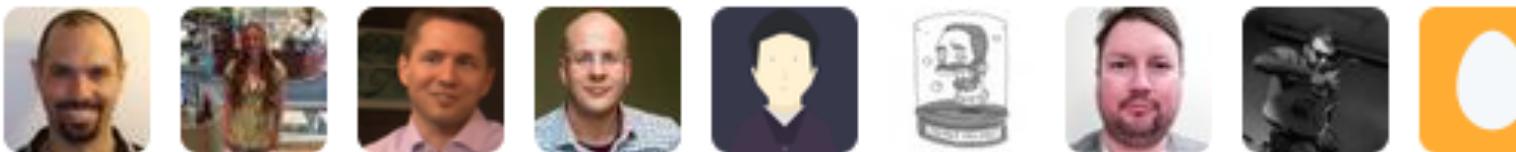
I see you have a poorly structured monolith.  
Would you like me to convert it into a poorly  
structured set of microservices?

RETWEETS

4,031

LIKES

2,345

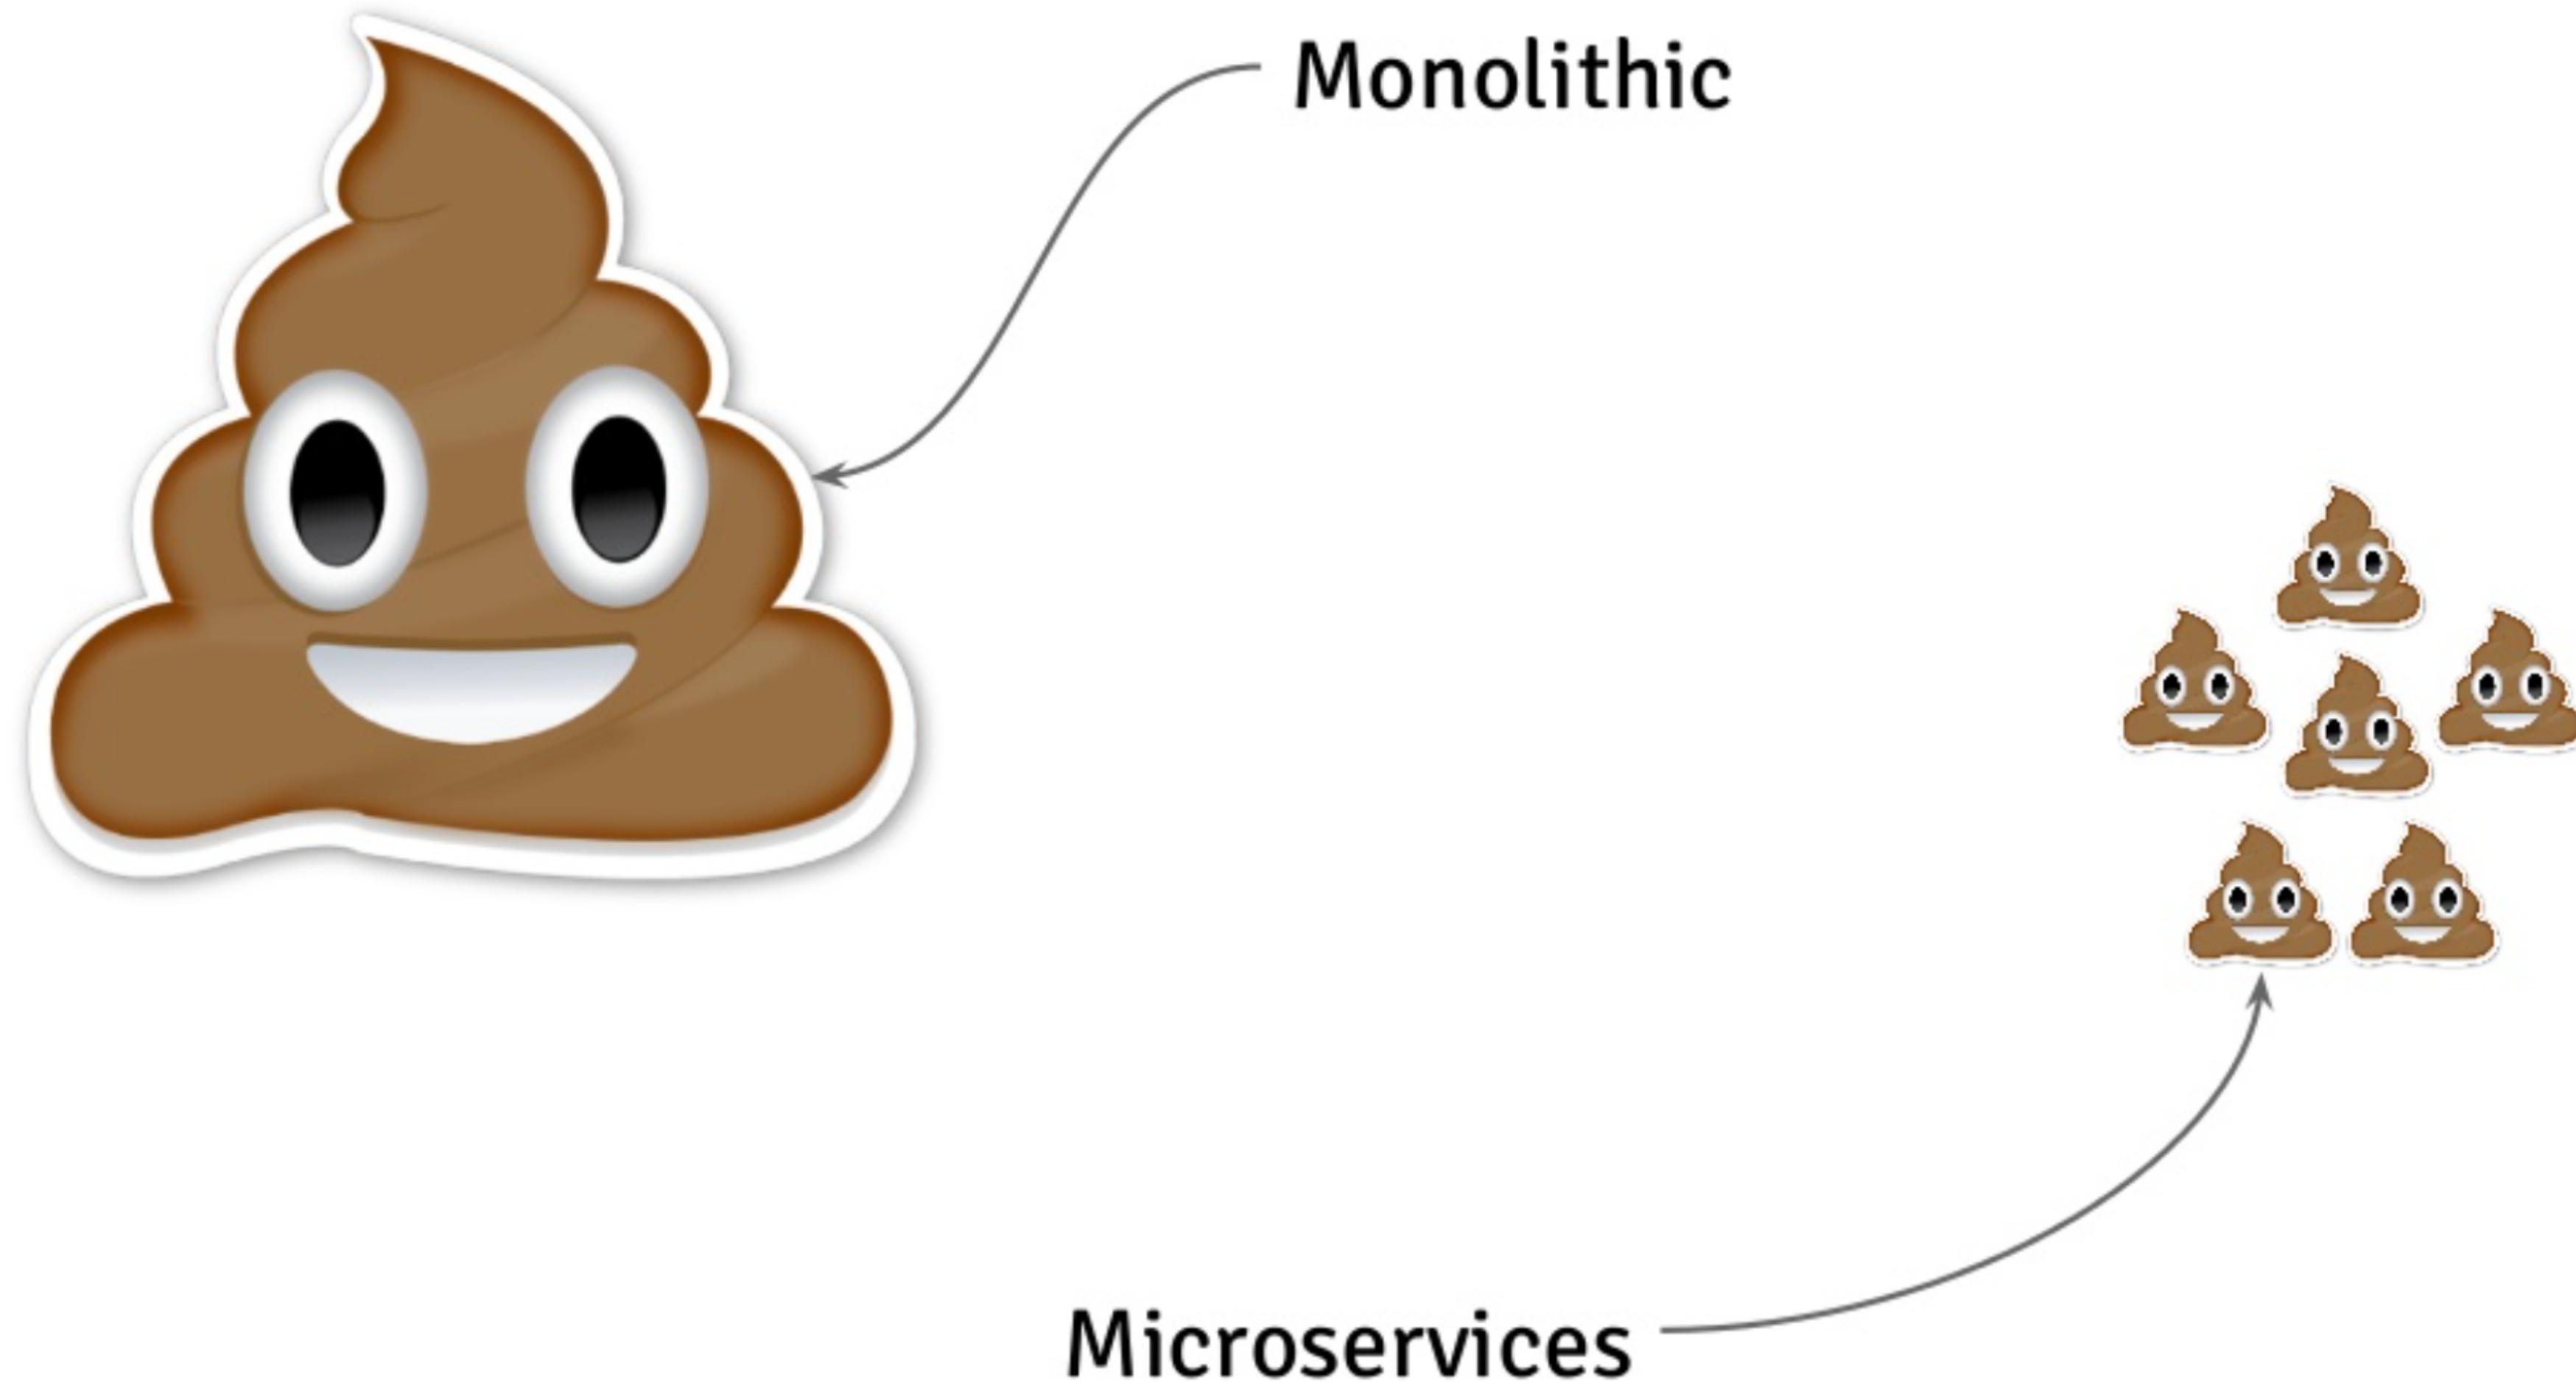


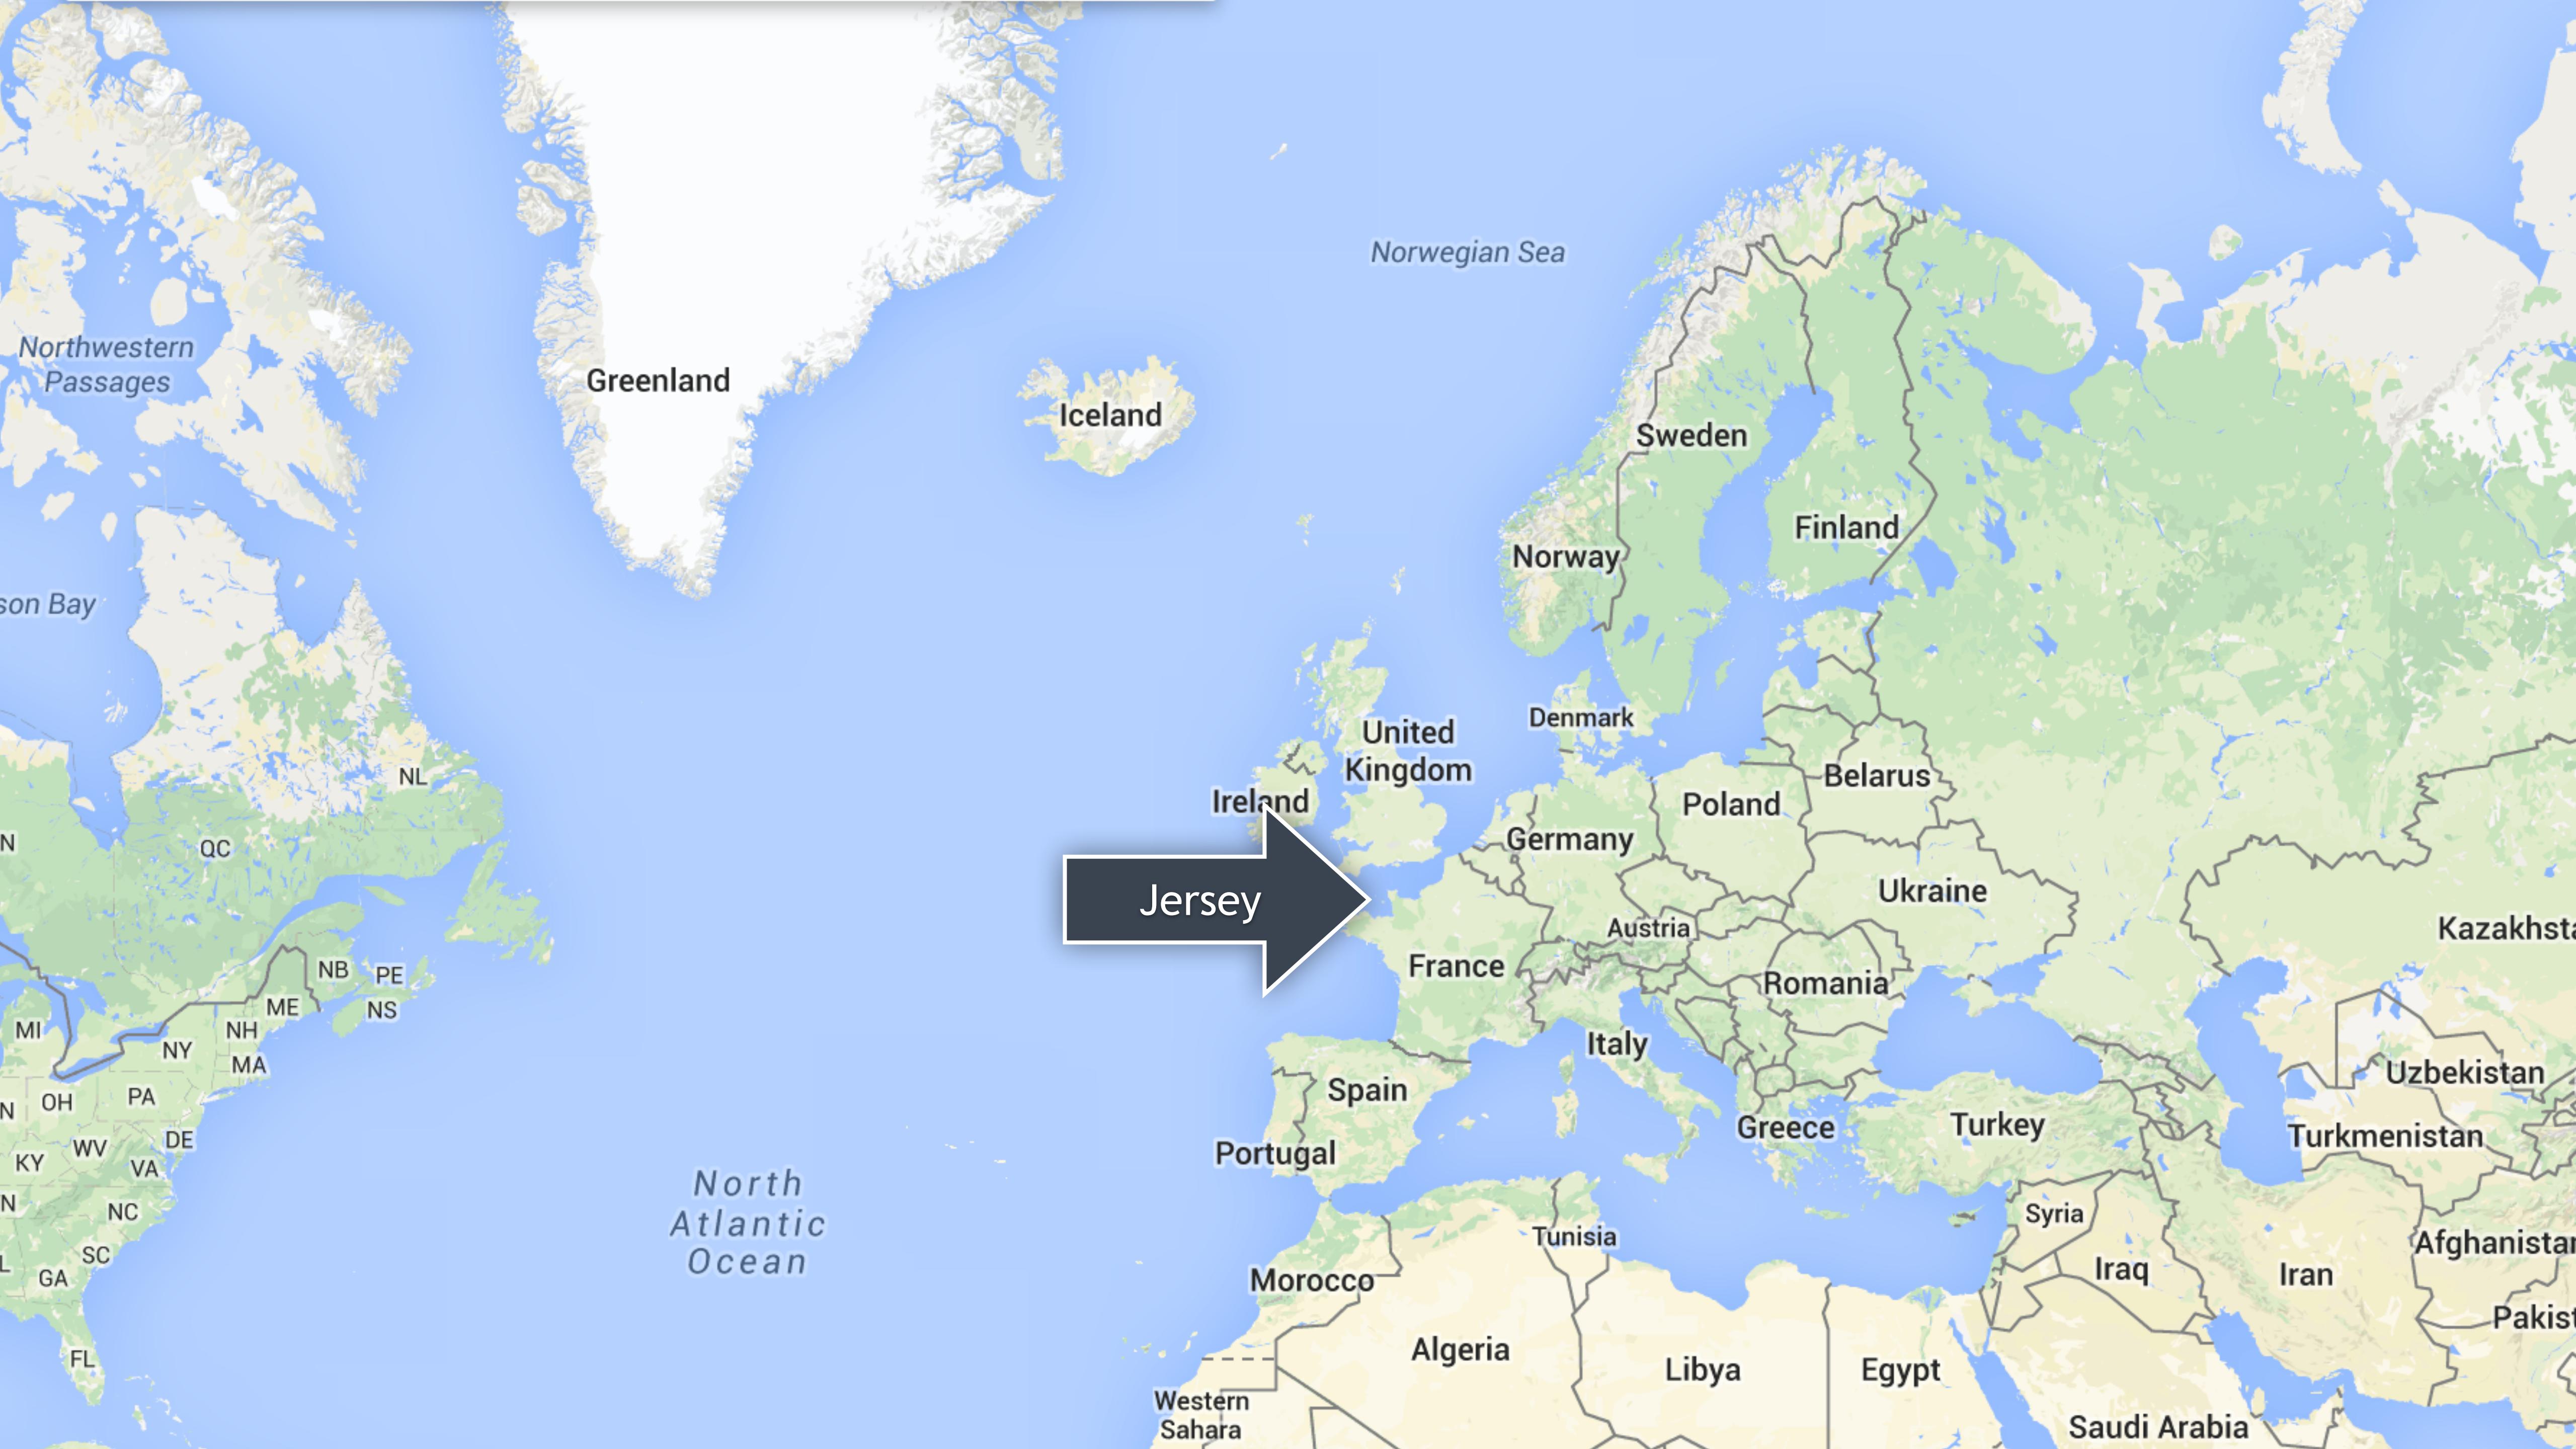
12:59 AM - 24 Feb 2015

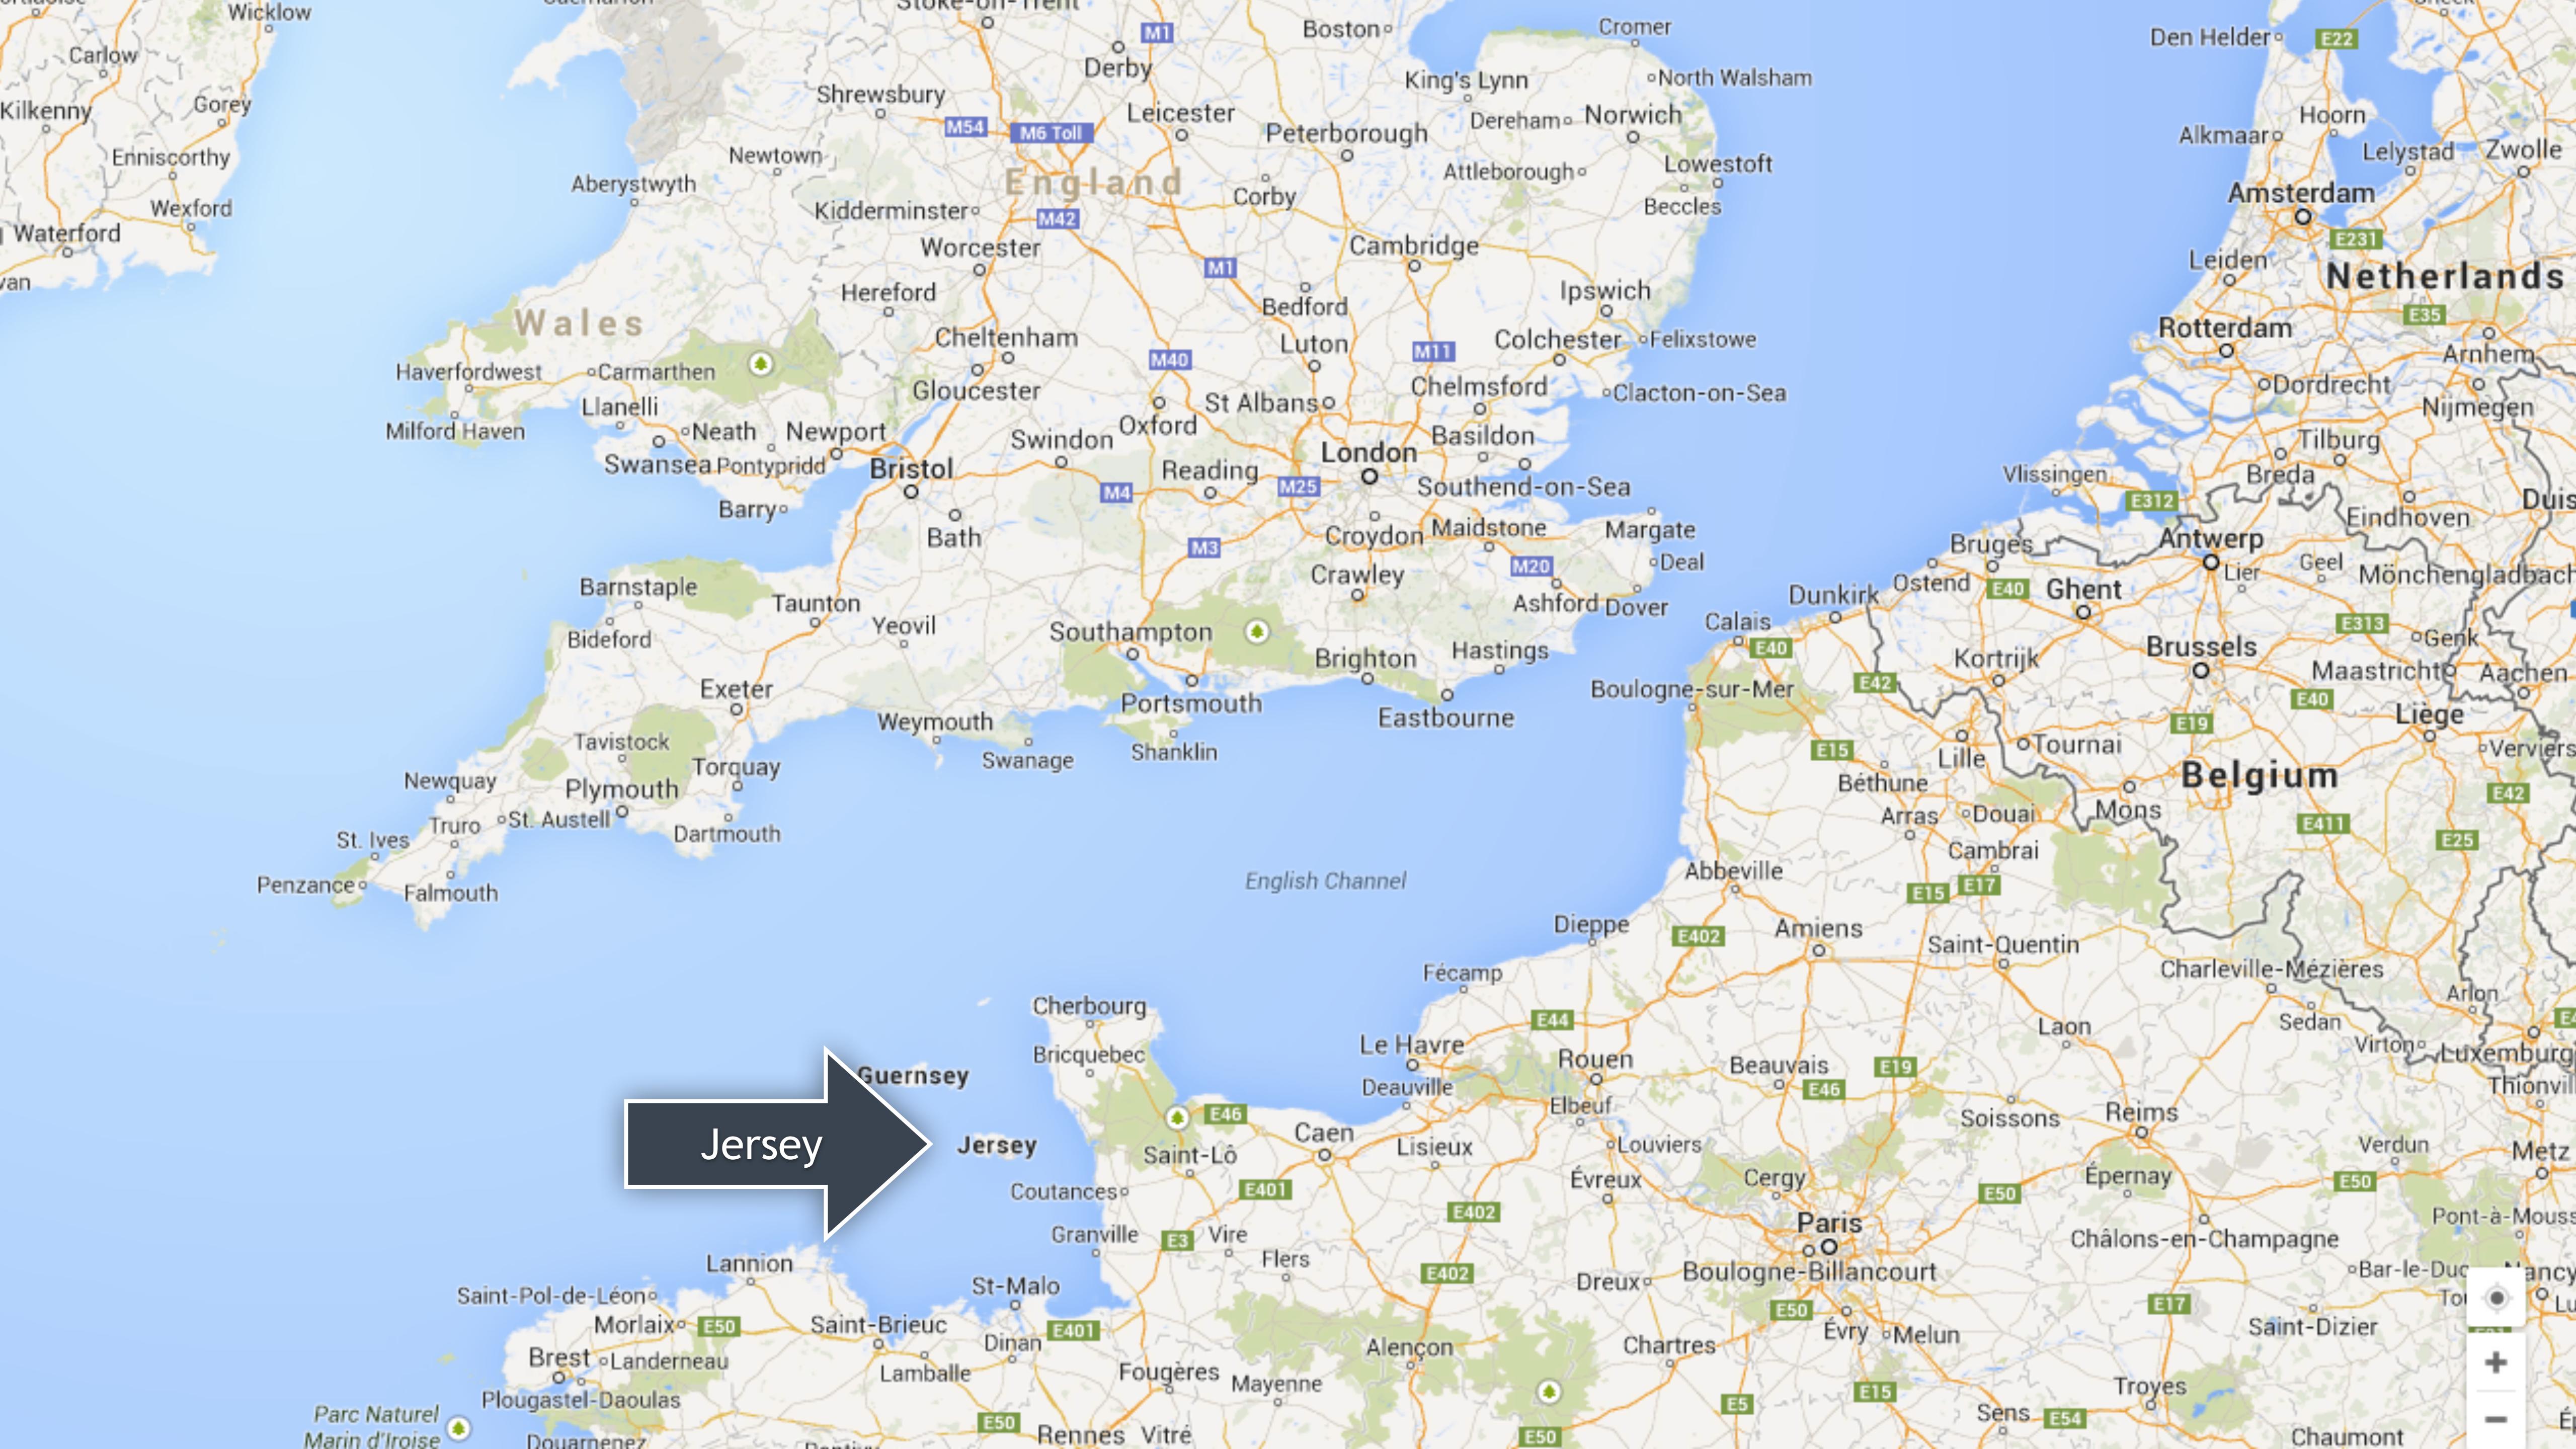


...

# Monolithic vs Microservices

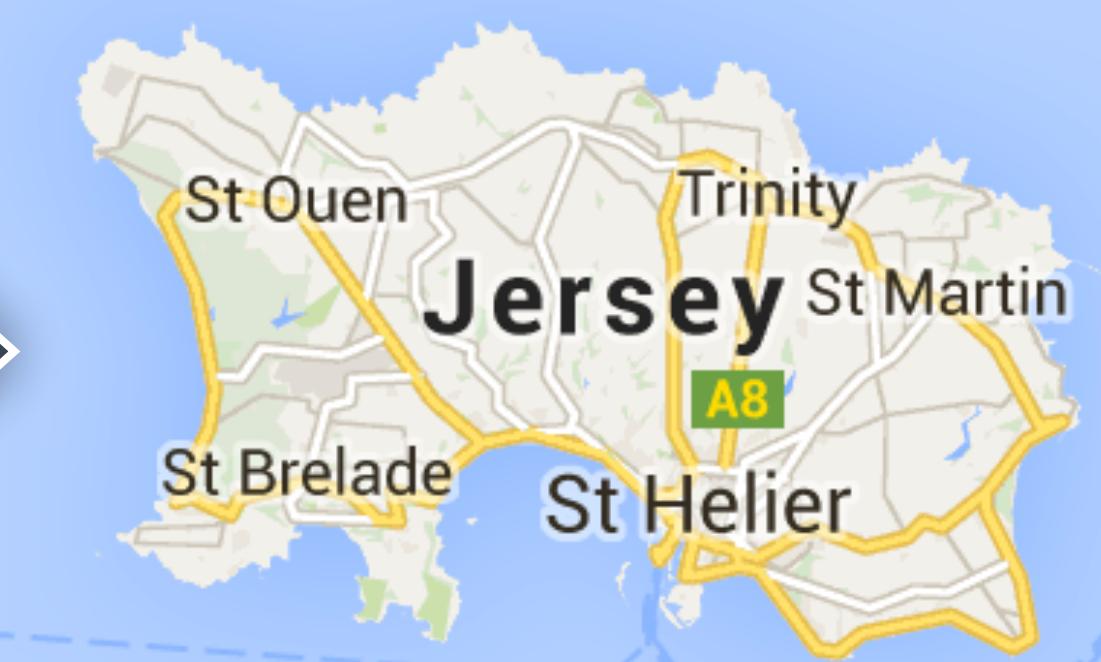


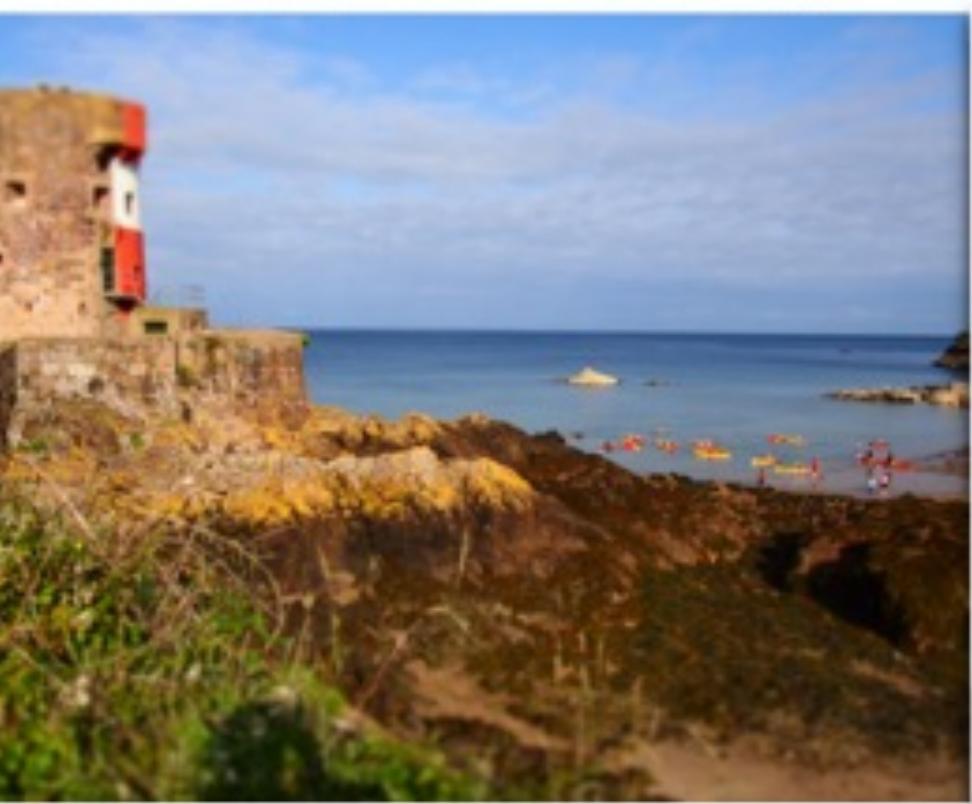






Jersey





I help software teams understand  
software architecture,  
technical leadership and  
the balance with agility



Software architecture  
needs to be more  
**accessible**

coding  
({the})  
architecture

# Software Architecture *for Developers*

Volume 1

---

Technical leadership by **coding**, coaching,  
collaboration and just enough up front design

Simon Brown

coding  
({the})  
architecture

# Software Architecture *for Developers*

Volume 2

---

Visualise, document and explore  
your software architecture

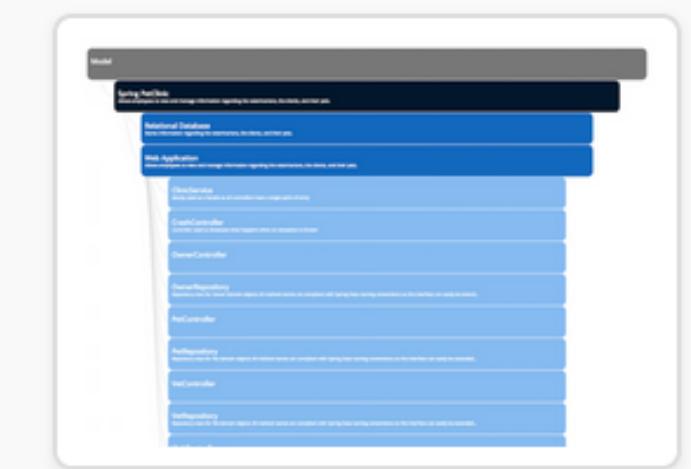
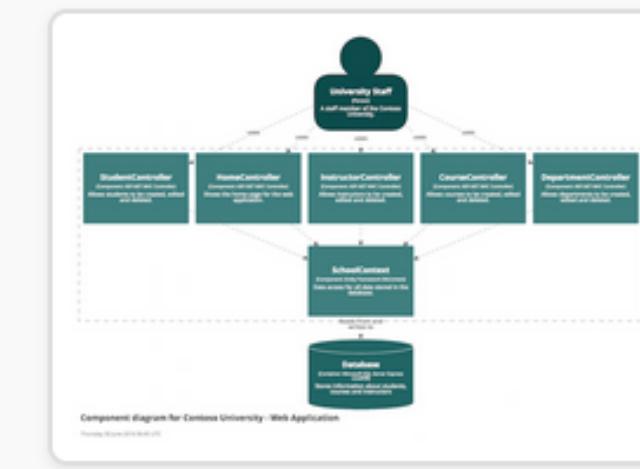
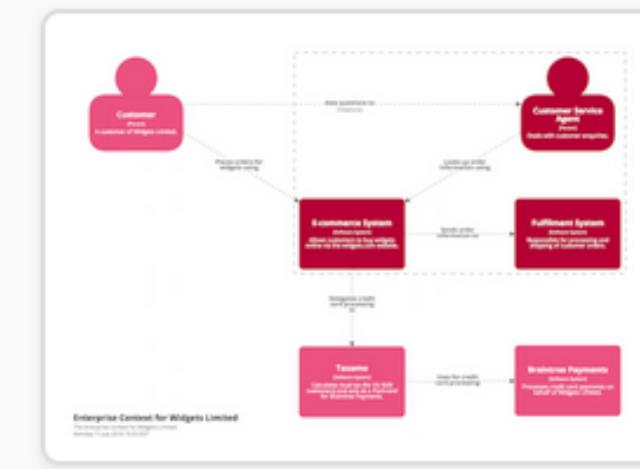
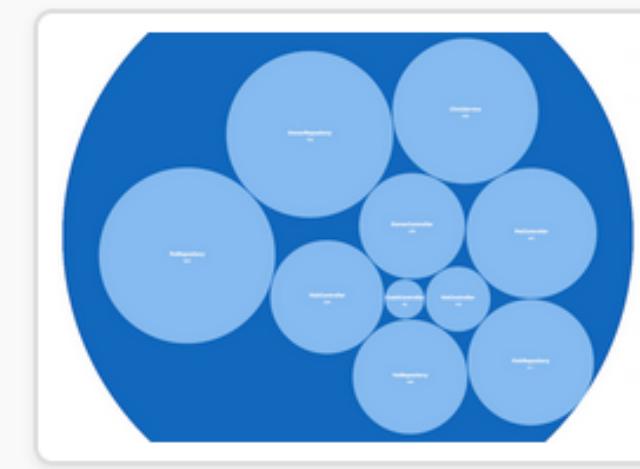
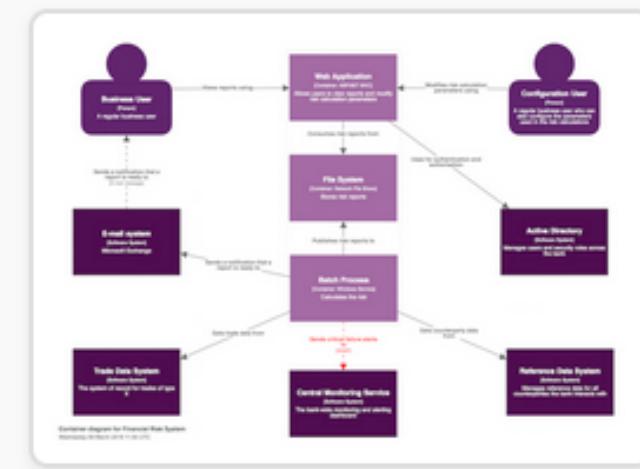
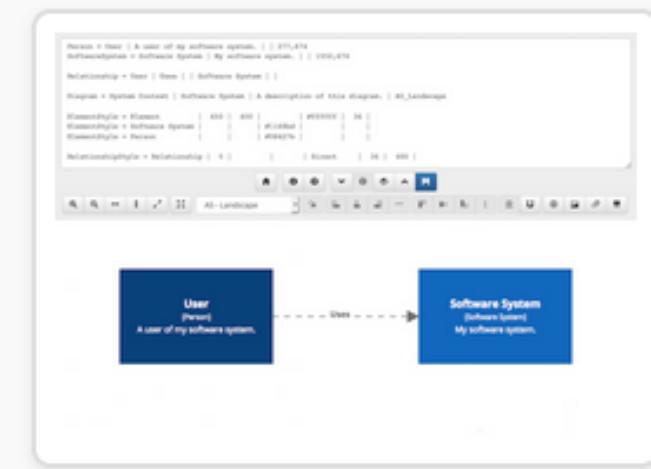
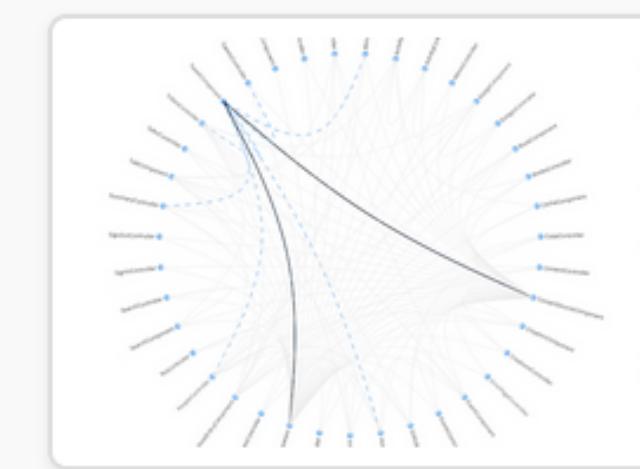
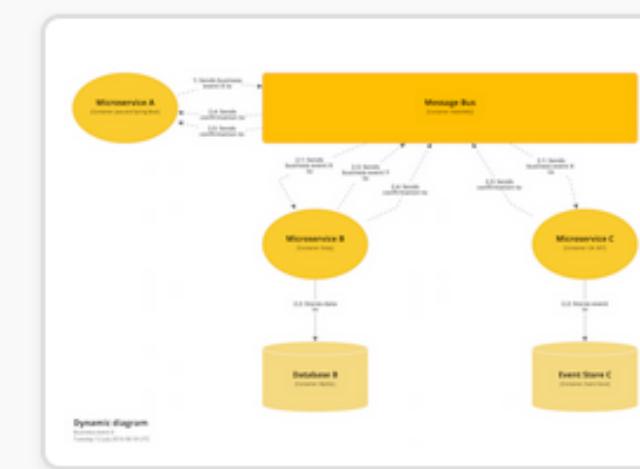
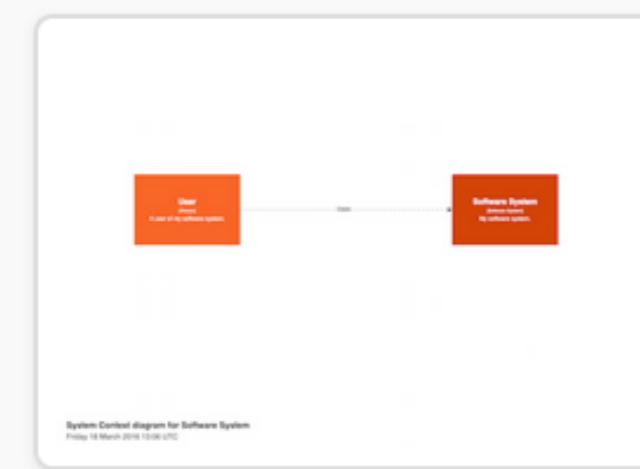
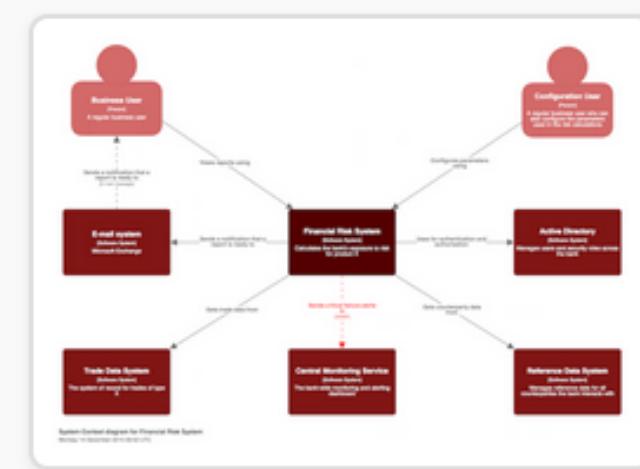
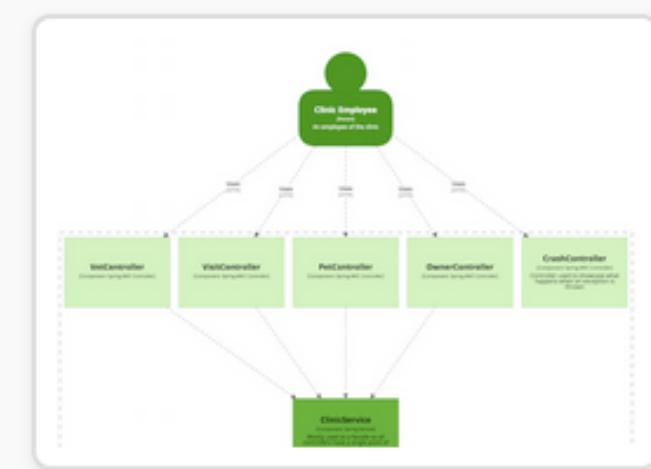
Simon Brown

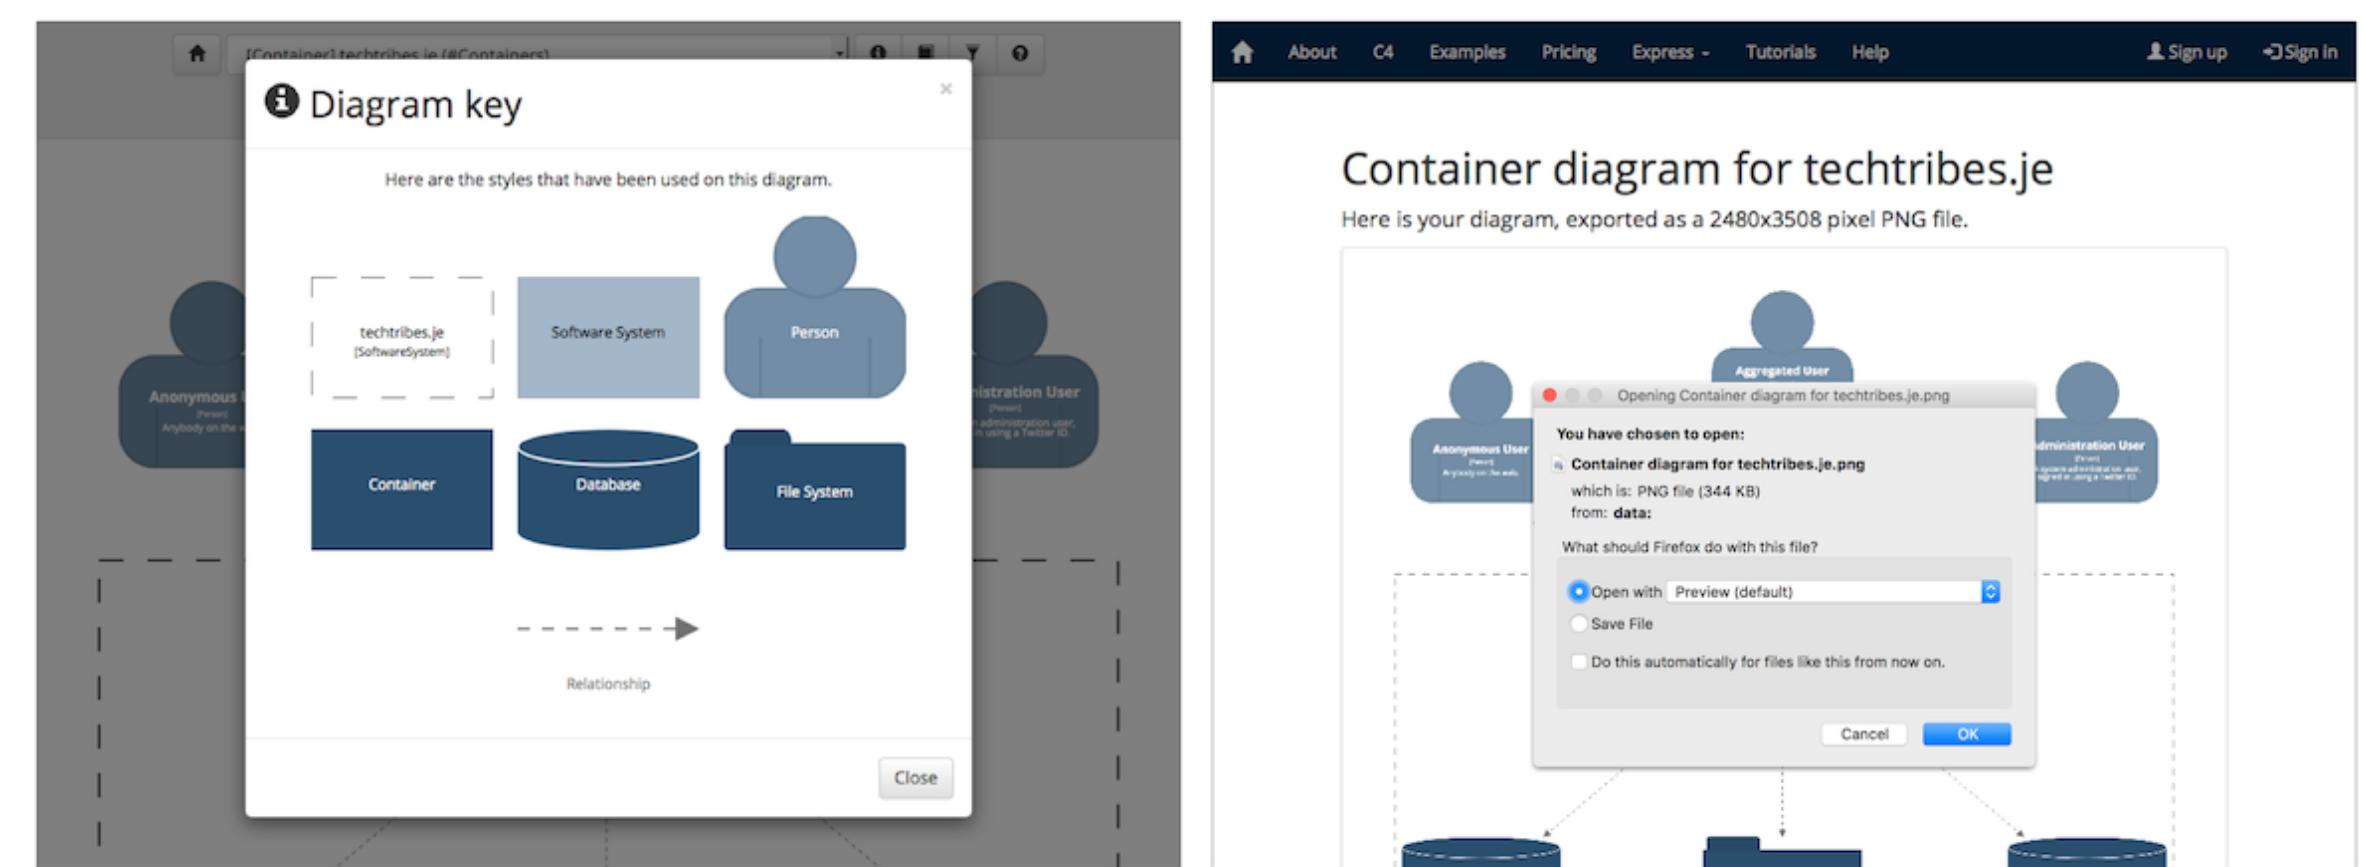
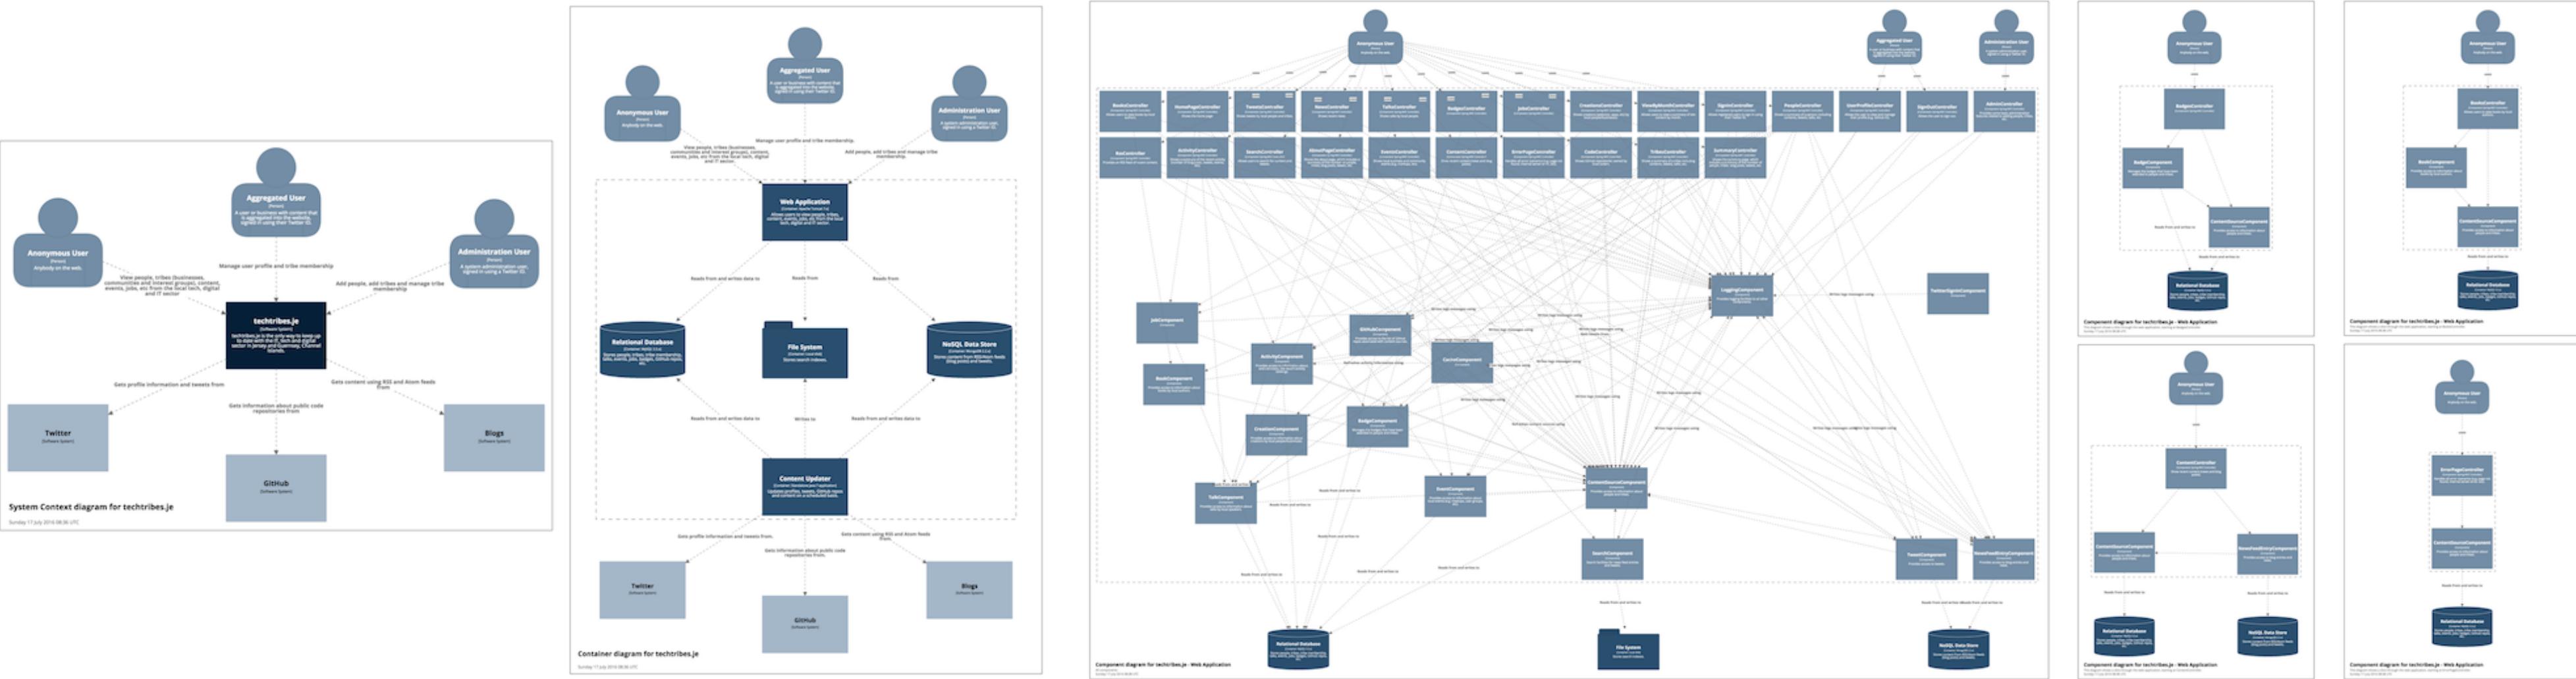




# Structurizr

Visualise, document and explore your software architecture





# Visualise, document and explore your software architecture

**techtribes.je**

**1. Context**

The techtribes.je website provides a way to find people, tribes (businesses, communities, interest groups, etc) and content related to the tech, IT and digital sector in Jersey and Guernsey. At the most basic level, it's a content aggregator for local tweets, news, blog posts, events, talks, jobs and more. Here's a context diagram that provides a visual summary of this:

**2. Quality Attributes**

- Performance
- Scalability
- Security
- Availability
- Internationalization
- Localization
- Browser compatibility

**3. Constraints**

- Budget
- Principles

**4. Principles**

**techtribes.je**

**Context**

This section provides information about the principles adopted for the development of the techtribes.je website.

**Package by component**

To provide a simple mapping of the software architecture into the code, the package structure of the code reflects a "package by component" convention rather than "package by layer".

A component is often a combination of a number of classes in different layers.

**System Context diagram for techtribes.je**

The purpose of the website is to:

1. Consolidate and share local content, helping to promote it inside and outside of the local community.
2. Encourage an open, sharing and learning culture within the local community.

**Users**

The techtribes.je website has three types of user:

1. **Anonymous:** anybody with a web browser can view content on the site.
2. **Authenticated:** people/tribes who have content aggregated into the website can sign-in to the website using their registered Twitter ID (if they have one) to modify some of their basic profile information.
3. **Admin:** people with administrative (super-user) access to the website can manage the people, tribes and

**techtribes.je**

**Principles**

This section provides information about the principles adopted for the development of the techtribes.je website.

**Package by component**

jo.techtribes.service

```

graph TD
    subgraph jo_techtribes_service [jo.techtribes.service]
        direction TB
        subgraph interface_tweetservice [jo.techtribes.service<br><code><interface></interface></code>->TweetService]
            TweetService
        end
        subgraph default_tweetservice [DefaultTweetService]
            DefaultTweetService
        end
        interface_tweetservice --- DefaultTweetService
    end
    interface_tweetservice <-->|jo.techtribes.data| jo_techtribes_data[jo.techtribes.data]
    subgraph jo_techtribes_data [jo.techtribes.data]
        direction TB
        subgraph interface_tweetdao [jo.techtribes.data<br><code><interface></interface>->TweetDao<br/>MongoDBTweetDao</code>]
            interface_tweetdao
            MongoDBTweetDao
        end
        subgraph mongo_db_tweetdao [MongoDBTweetDao]
            MongoDBTweetDao
        end
        interface_tweetdao --- MongoDBTweetDao
    end
    interface_tweetservice <-->|jo.techtribes.component.tweet| jo_techtribes_component_tweet[jo.techtribes.component.tweet]
    jo_techtribes_component_tweet <-->|jo.techtribes.component.tweet| jo_techtribes_data
    interface_tweetdao <-->|jo.techtribes.component.tweet| jo_techtribes_component_tweet

```

**Package by layer**

**Package by layer vs package by component**

This means that the codebase is broken up into a number of components, each of which has:

- A well-defined public interface.
- Strong encapsulation (i.e. all implementation details are package protected where possible).
- A Spring configuration file called component.xml to configure and wire the component together into the rest of the system.

**techtribes.je**

**Operation and Support**

This section provides information about the operational and support aspects of the techtribes.je website.

**Starting MySQL**

MySQL is installed as a service, and should be running after a server restart. You can check this by using the following command:

```
sudo netstat -tulp | grep mysql
```

If you need to start MySQL, you can use the following command:

```
sudo service mysql start
```

**Starting MongoDB**

MongoDB is also installed as a service, and should be running after a server restart. You can check this by using the following commands:

```
sudo netstat -tulp | grep mongo
tail /var/log/mongodb/mongod.log
```

If you need to start MongoDB, you can use the following command:

```
sudo service mongod start
```

**Starting the Web Server**

Apache Tomcat is also installed as a service, and should be running after a server restart. You can check this by using the following commands:

```
ps -Af | grep tomcat
tail /var/lib/tomcat7/logs/catalina.out
```

If you need to start Tomcat, you can use the following command:

```
~/techtribes.je/bin/start-tomcat.sh
```

**Starting the Content Updater**

The Content Updater is a standalone Java process that needs to be started manually after a server restart. You can do this with the following command (where XYZ is the build number):

```
~/techtribes.je/bin/start-updater.sh XYZ
```

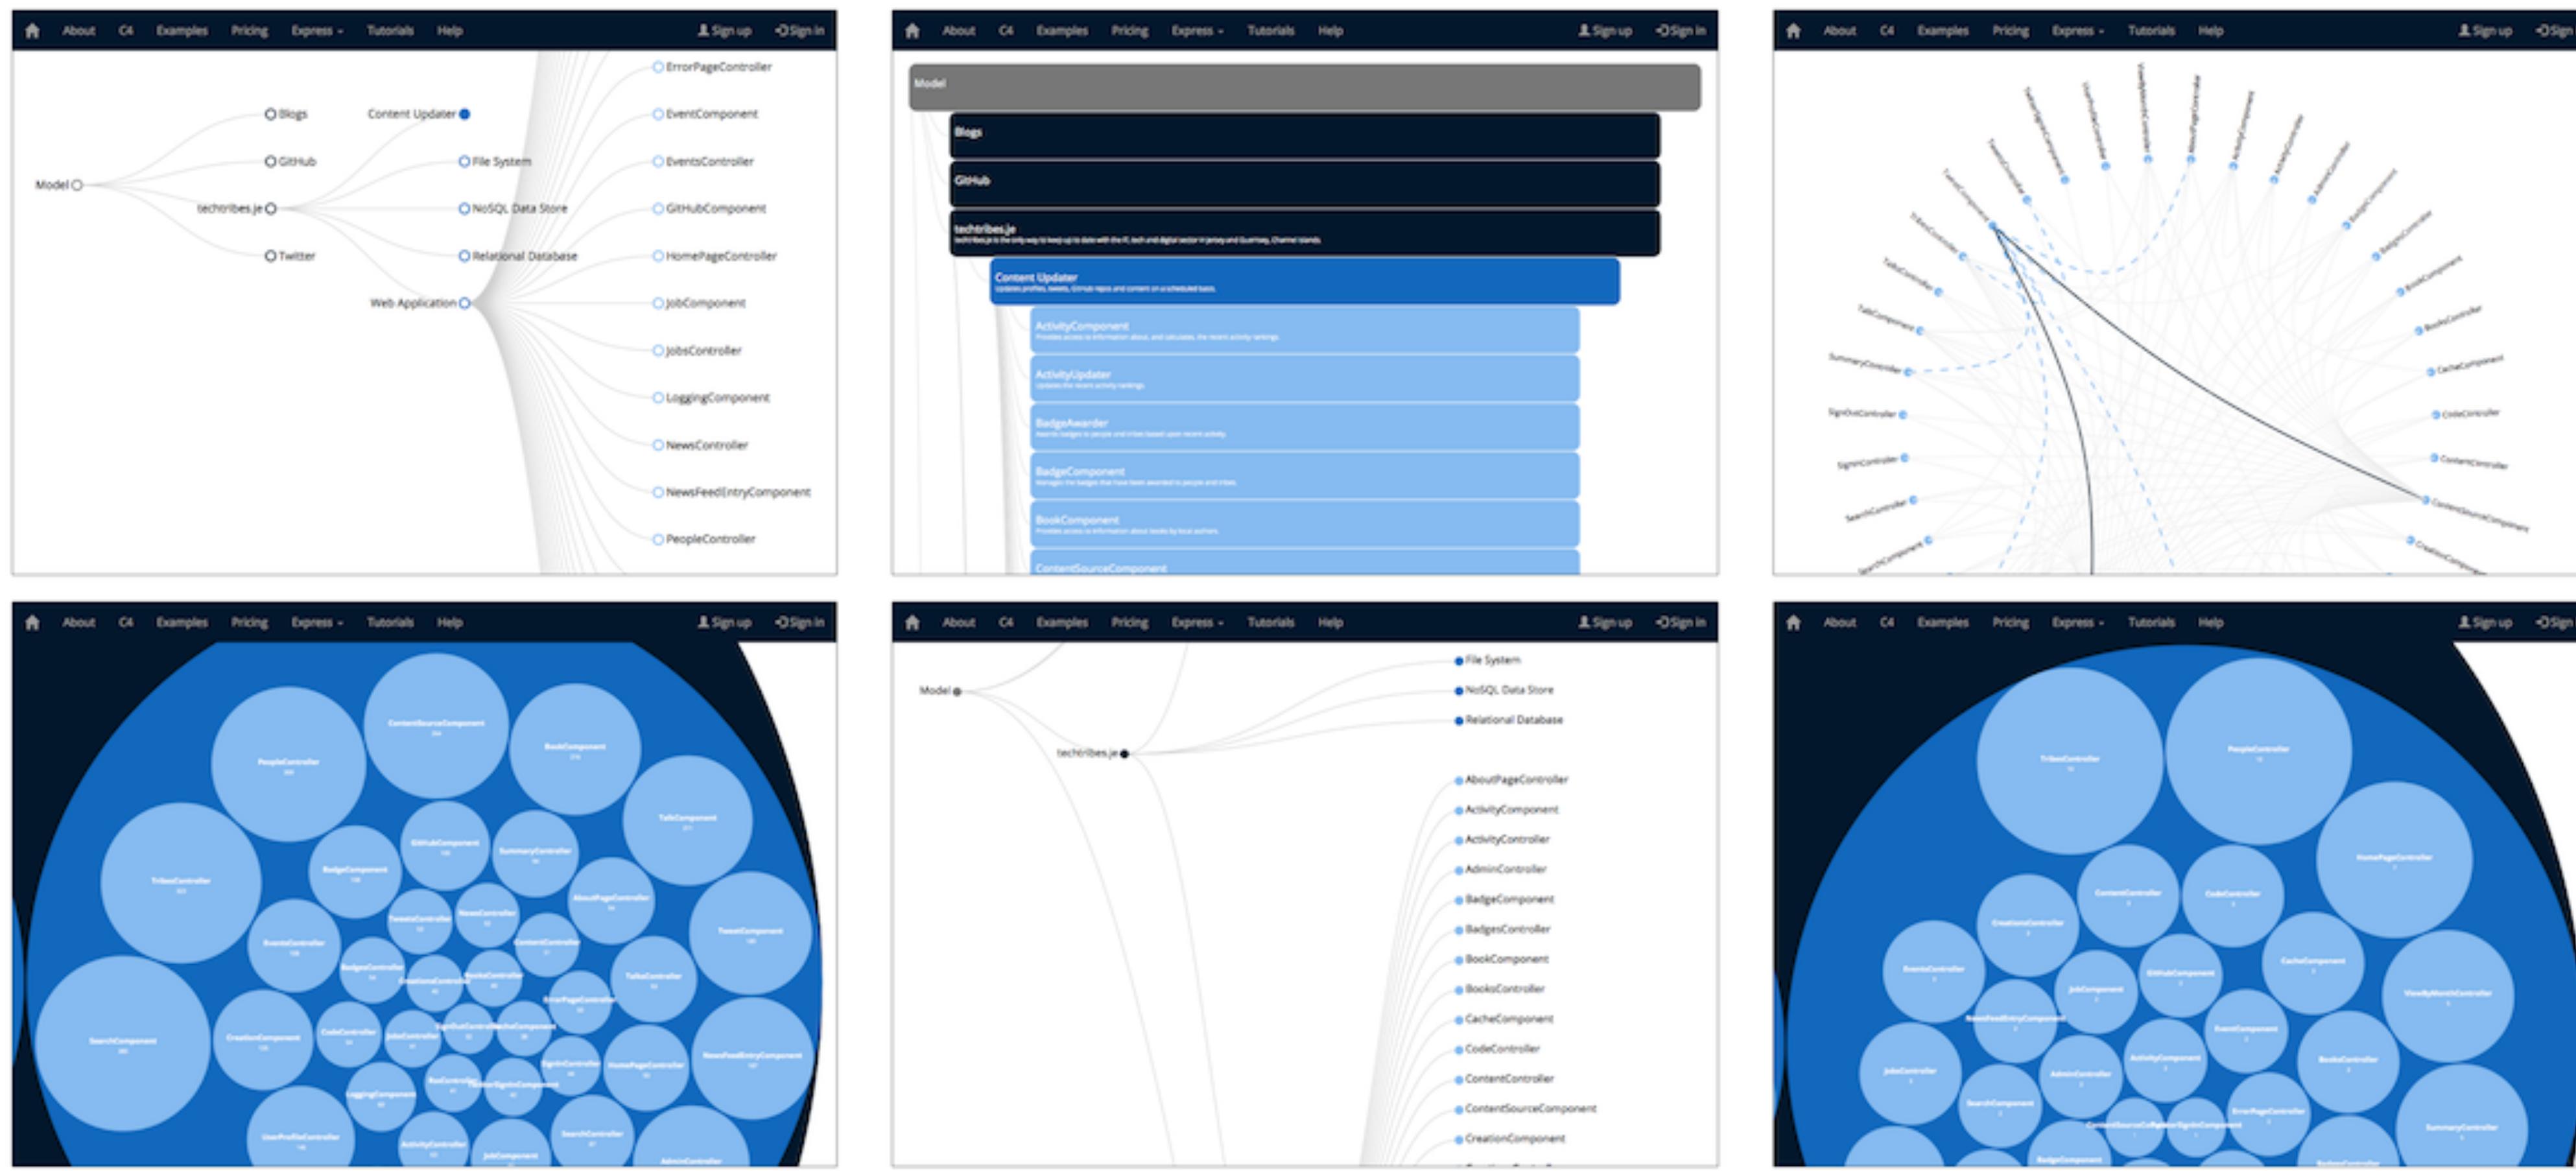
You can check the log file with the following command:

```
~/techtribes.je/bin/updater-log.sh XYZ
```

**Monitoring**

The only monitoring on the techtribes.je website is Pinotom, which is configured to test that the website is still

# Visualise, document and explore your software architecture



# Visualise, document and **explore** your software architecture

```

Workspace workspace = new Workspace("My model", "This is a model of my software system.");
Model model = workspace.getModel();

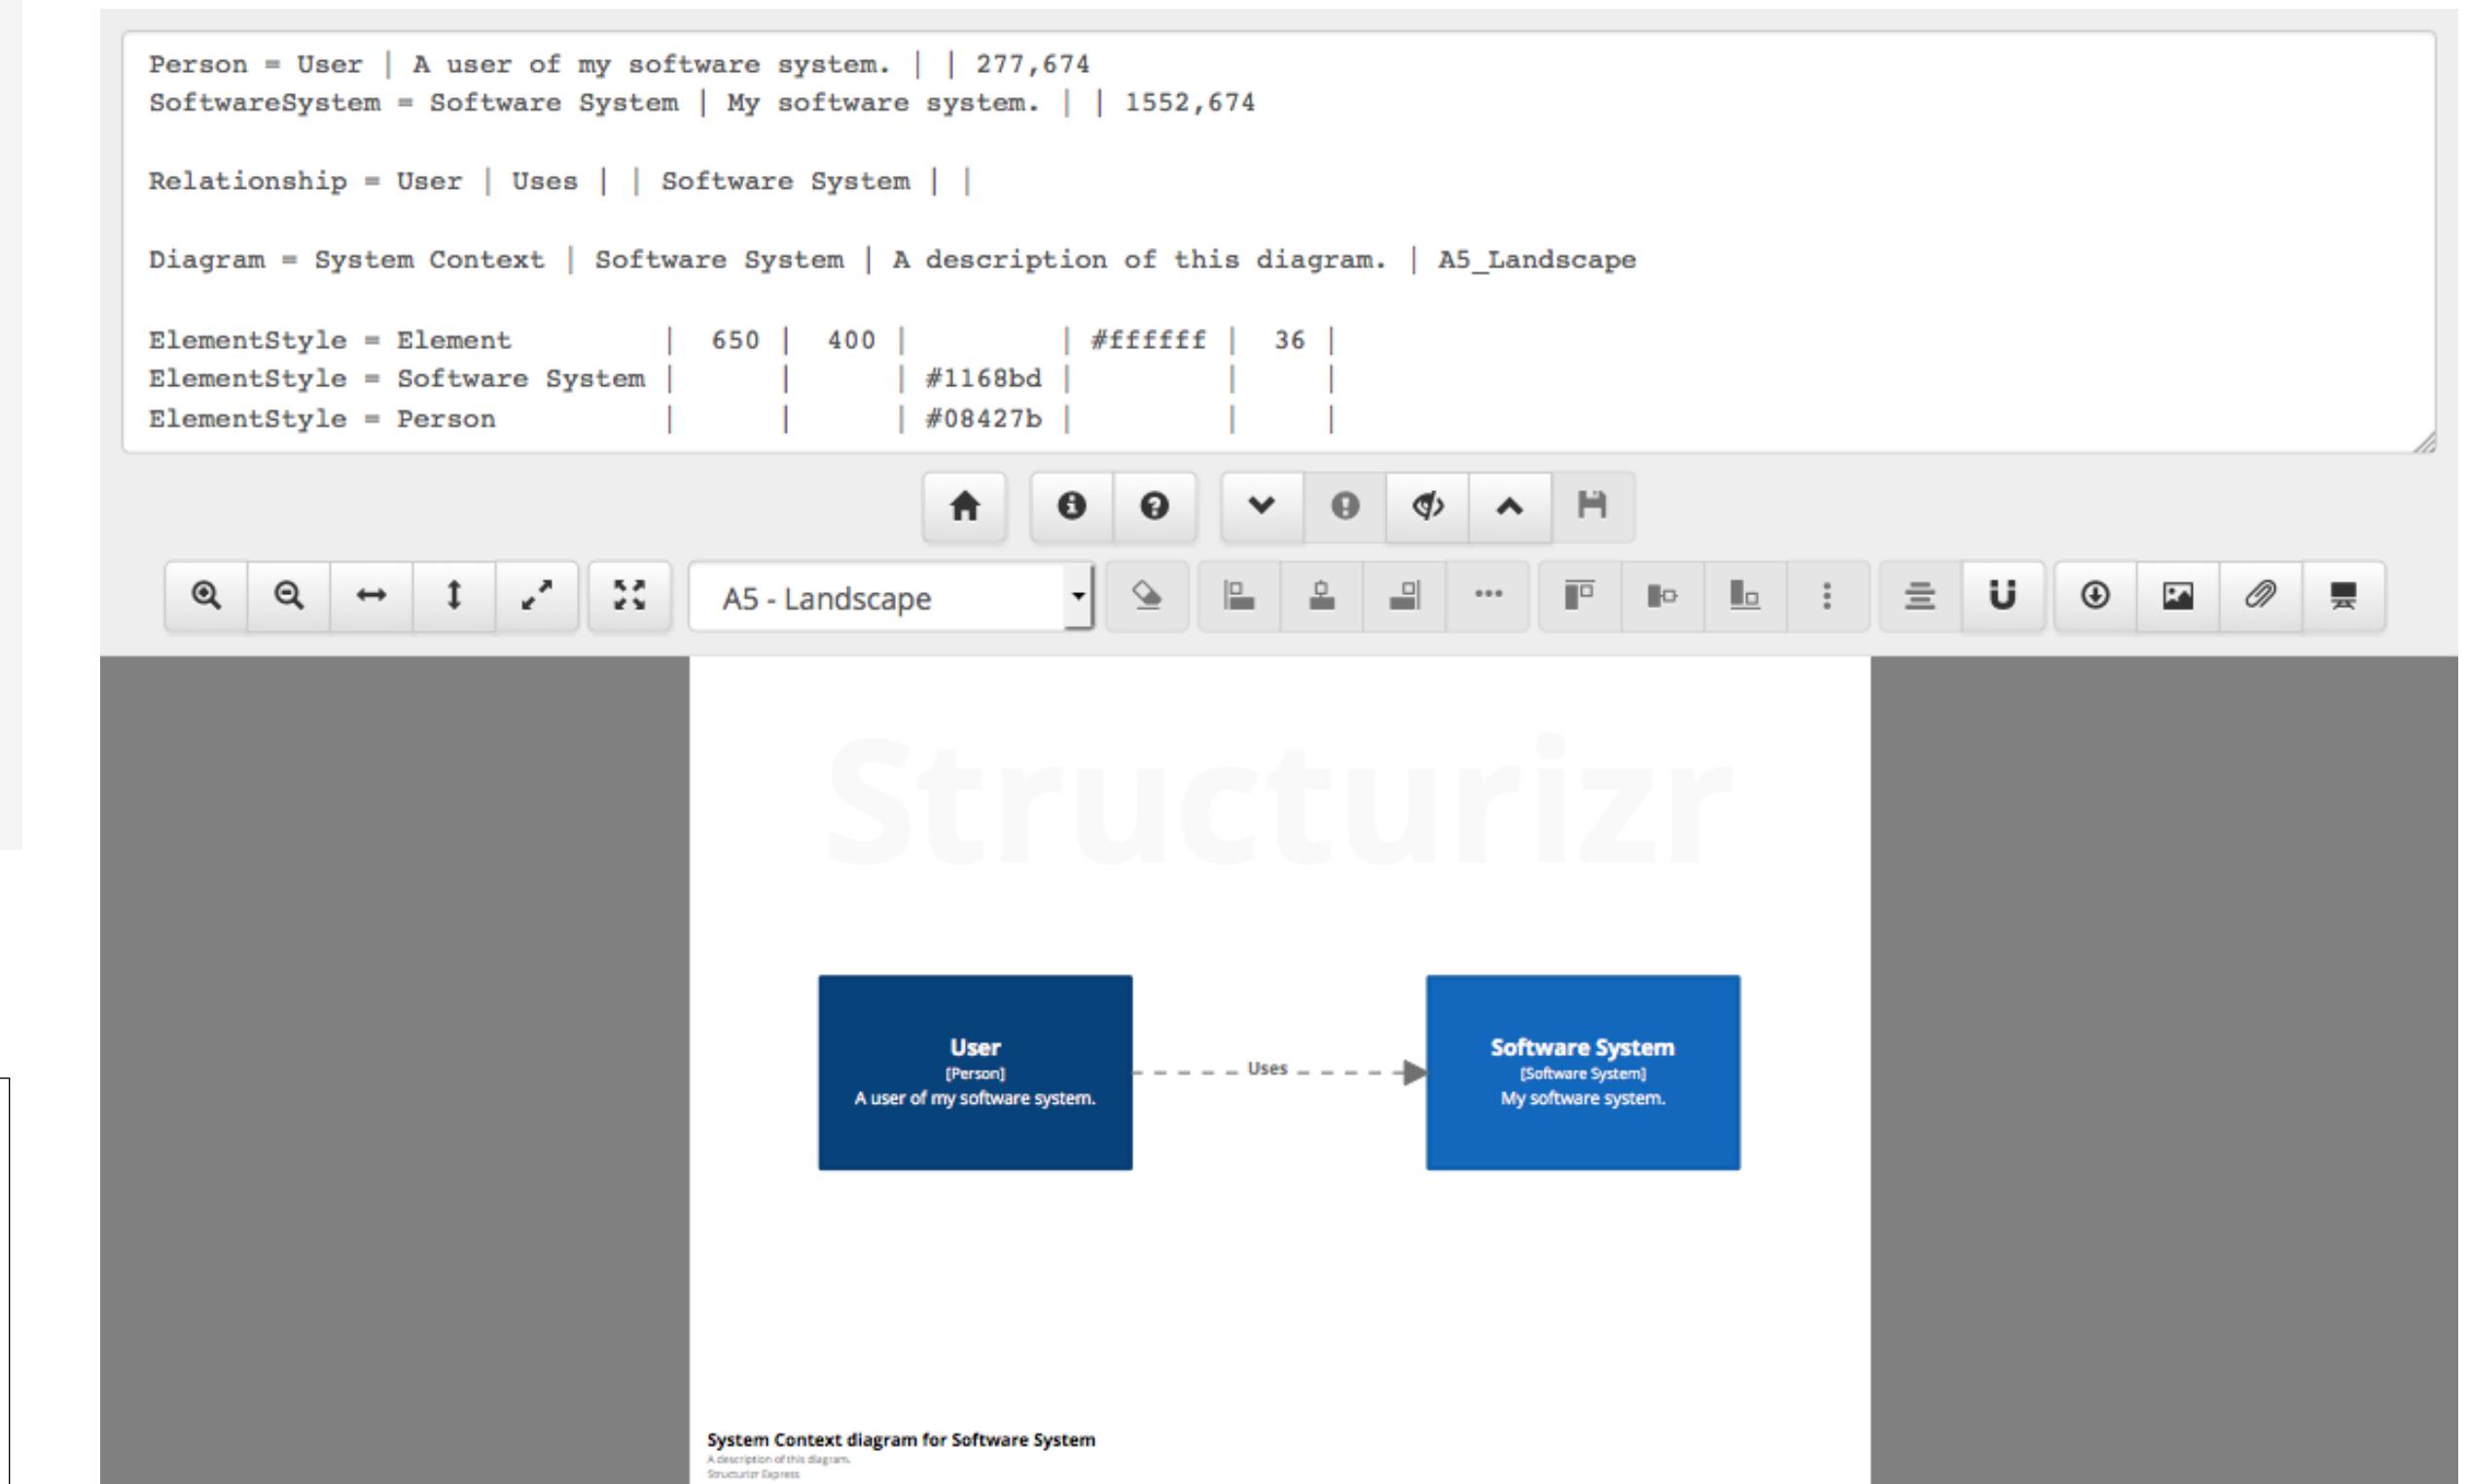
Person user = model.addPerson("User", "A user of my software system.");
SoftwareSystem softwareSystem = model.addSoftwareSystem("Software System", "My software system.");
user.uses(softwareSystem, "Uses");

ViewSet viewSet = workspace.getViews();
SystemContextView contextView = viewSet.createSystemContextView(softwareSystem, "context", "A simple example...");
contextView.addAllSoftwareSystems();
contextView.addAllPeople();

Styles styles = viewSet.getConfiguration().getStyles();
styles.addElementStyle(Tags.SOFTWARE_SYSTEM).background("#1168bd").color("#ffffff");
styles.addElementStyle(Tags.PERSON).background("#08427b").color("#ffffff");

StructurizrClient structurizrClient = new StructurizrClient("key", "secret");
structurizrClient.putWorkspace(1234, workspace);

```

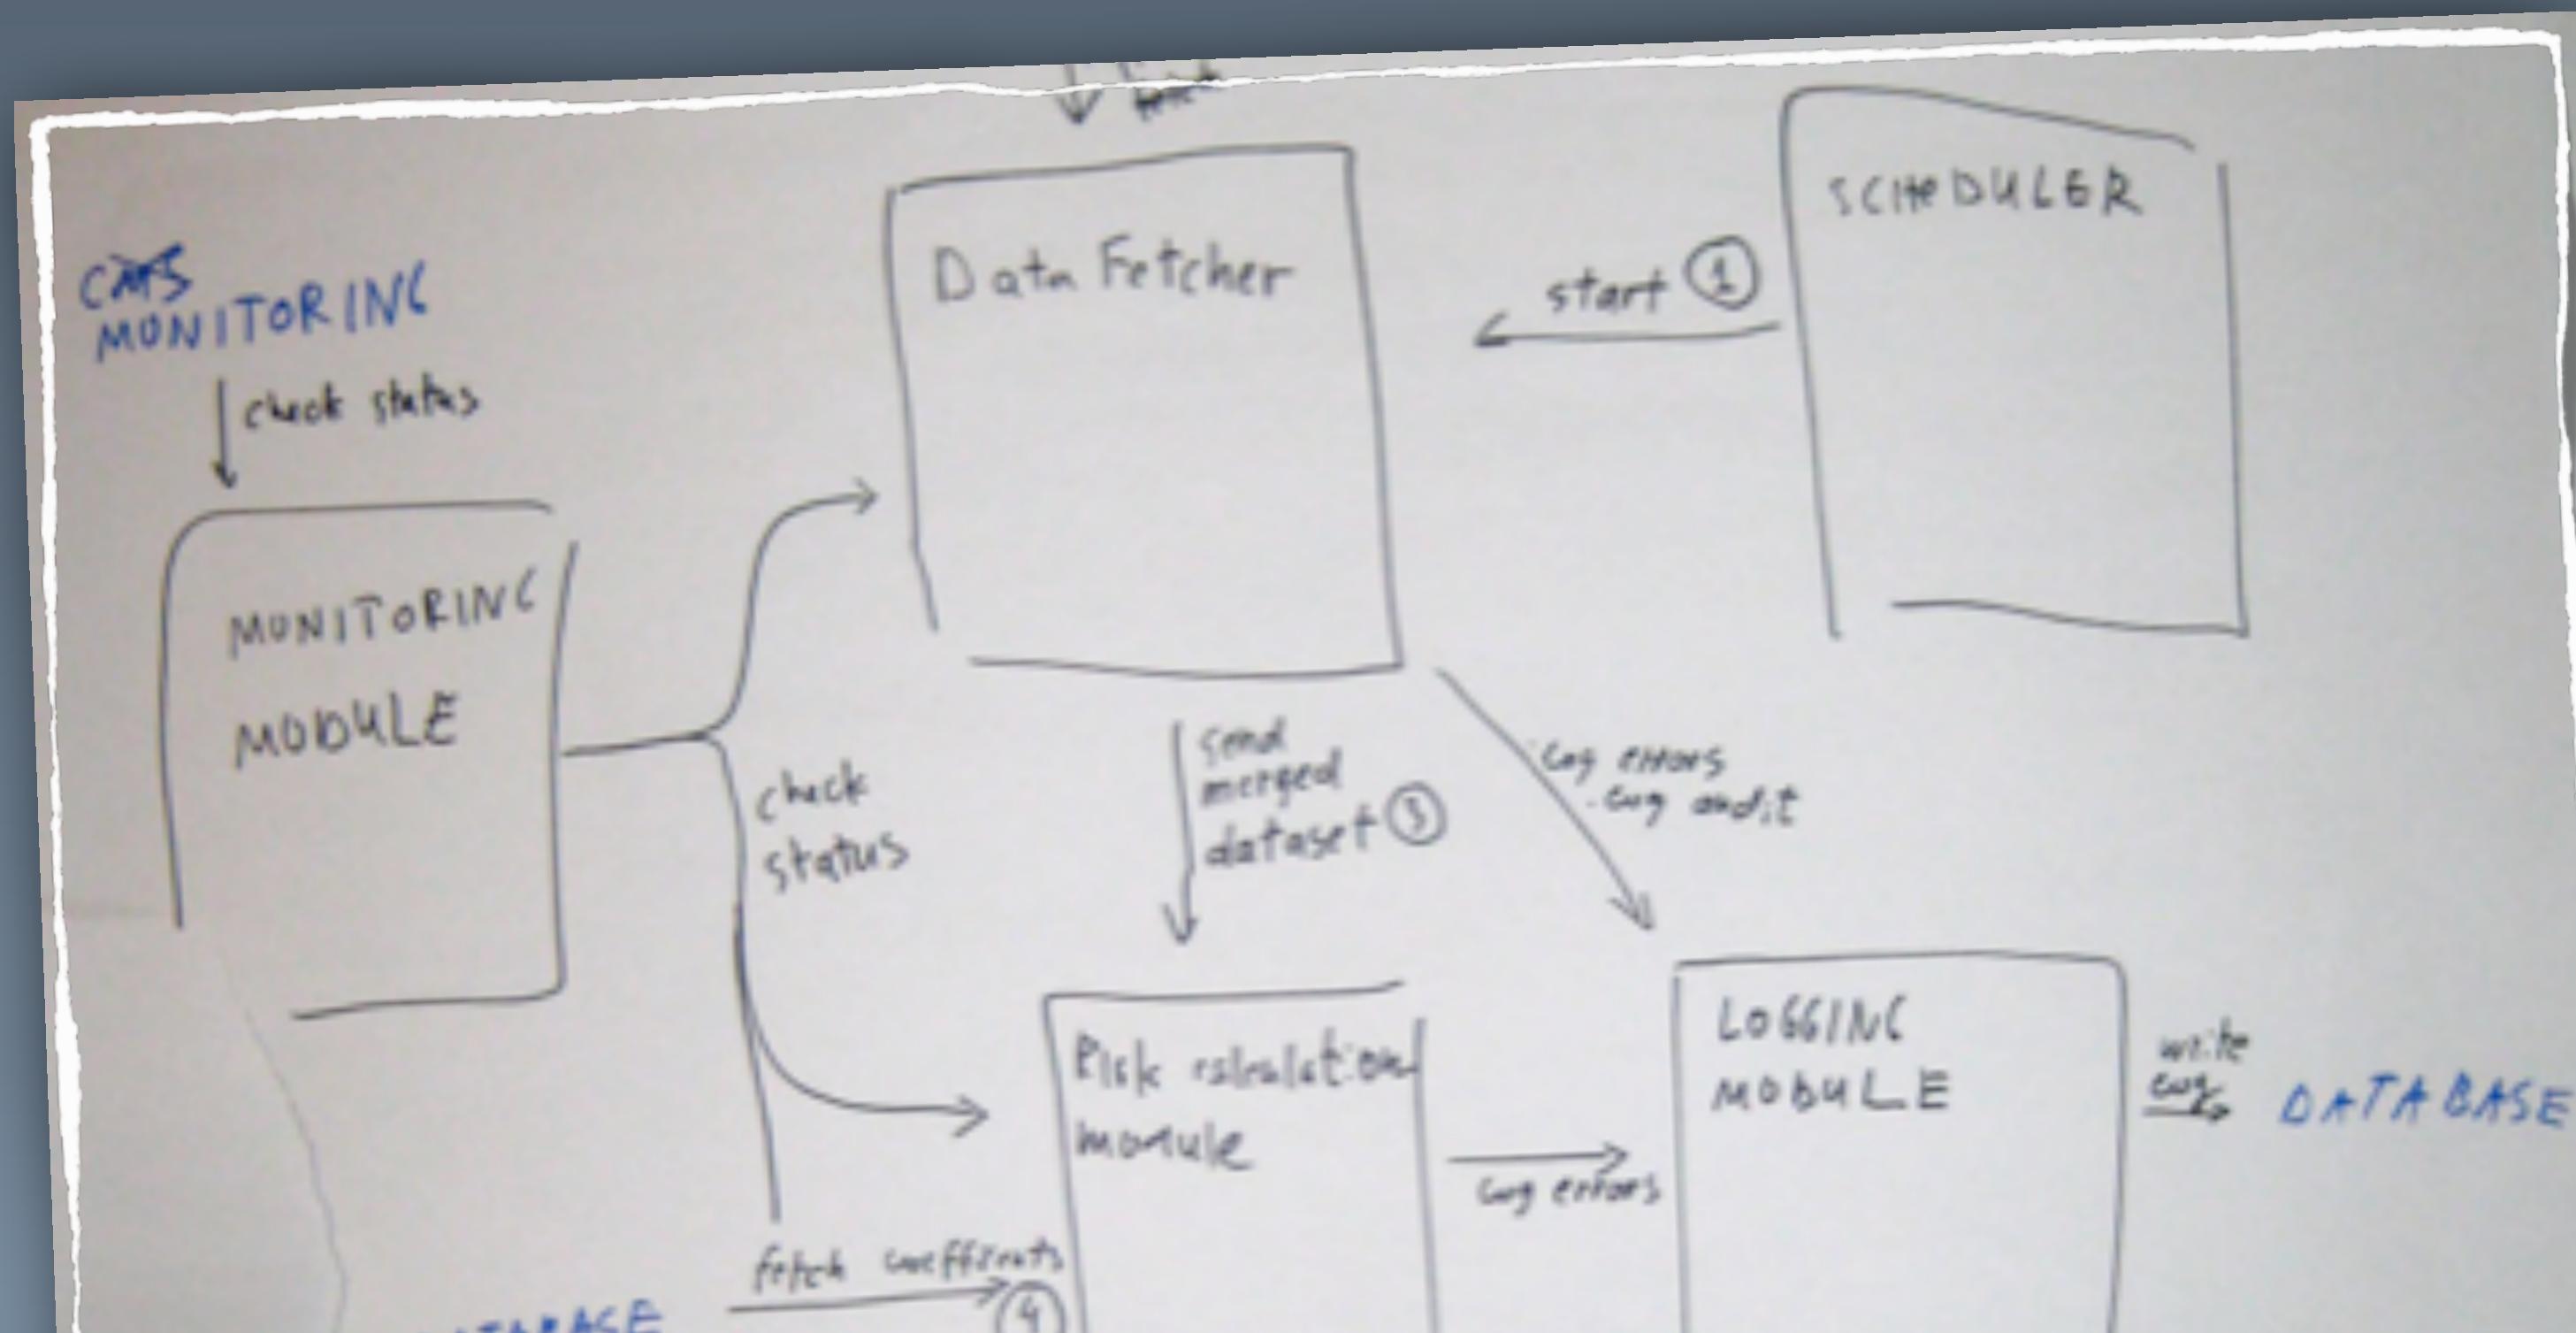


# Software architecture diagrams as **code** or **text**

A well structured  
codebase is easy  
to visualise

# Abstraction

is about reducing detail  
rather than creating a different representation



Abstractions help us  
reason about  
a big and/or complex  
software system

Abstractions  
should reflect the  
code

As an industry, we lack a  
common vocabulary  
with which to think about, describe  
and communicate software architecture



**Container**

web application, application server, standalone application,  
browser, database, file system, etc)

**Container**

(e.g. web application, application server, standalone application,  
browser, database, file system, etc)

**Container**

(e.g. web application, application server, standalone application,  
browser, database, file system, etc)

Component

Component

Component

Class

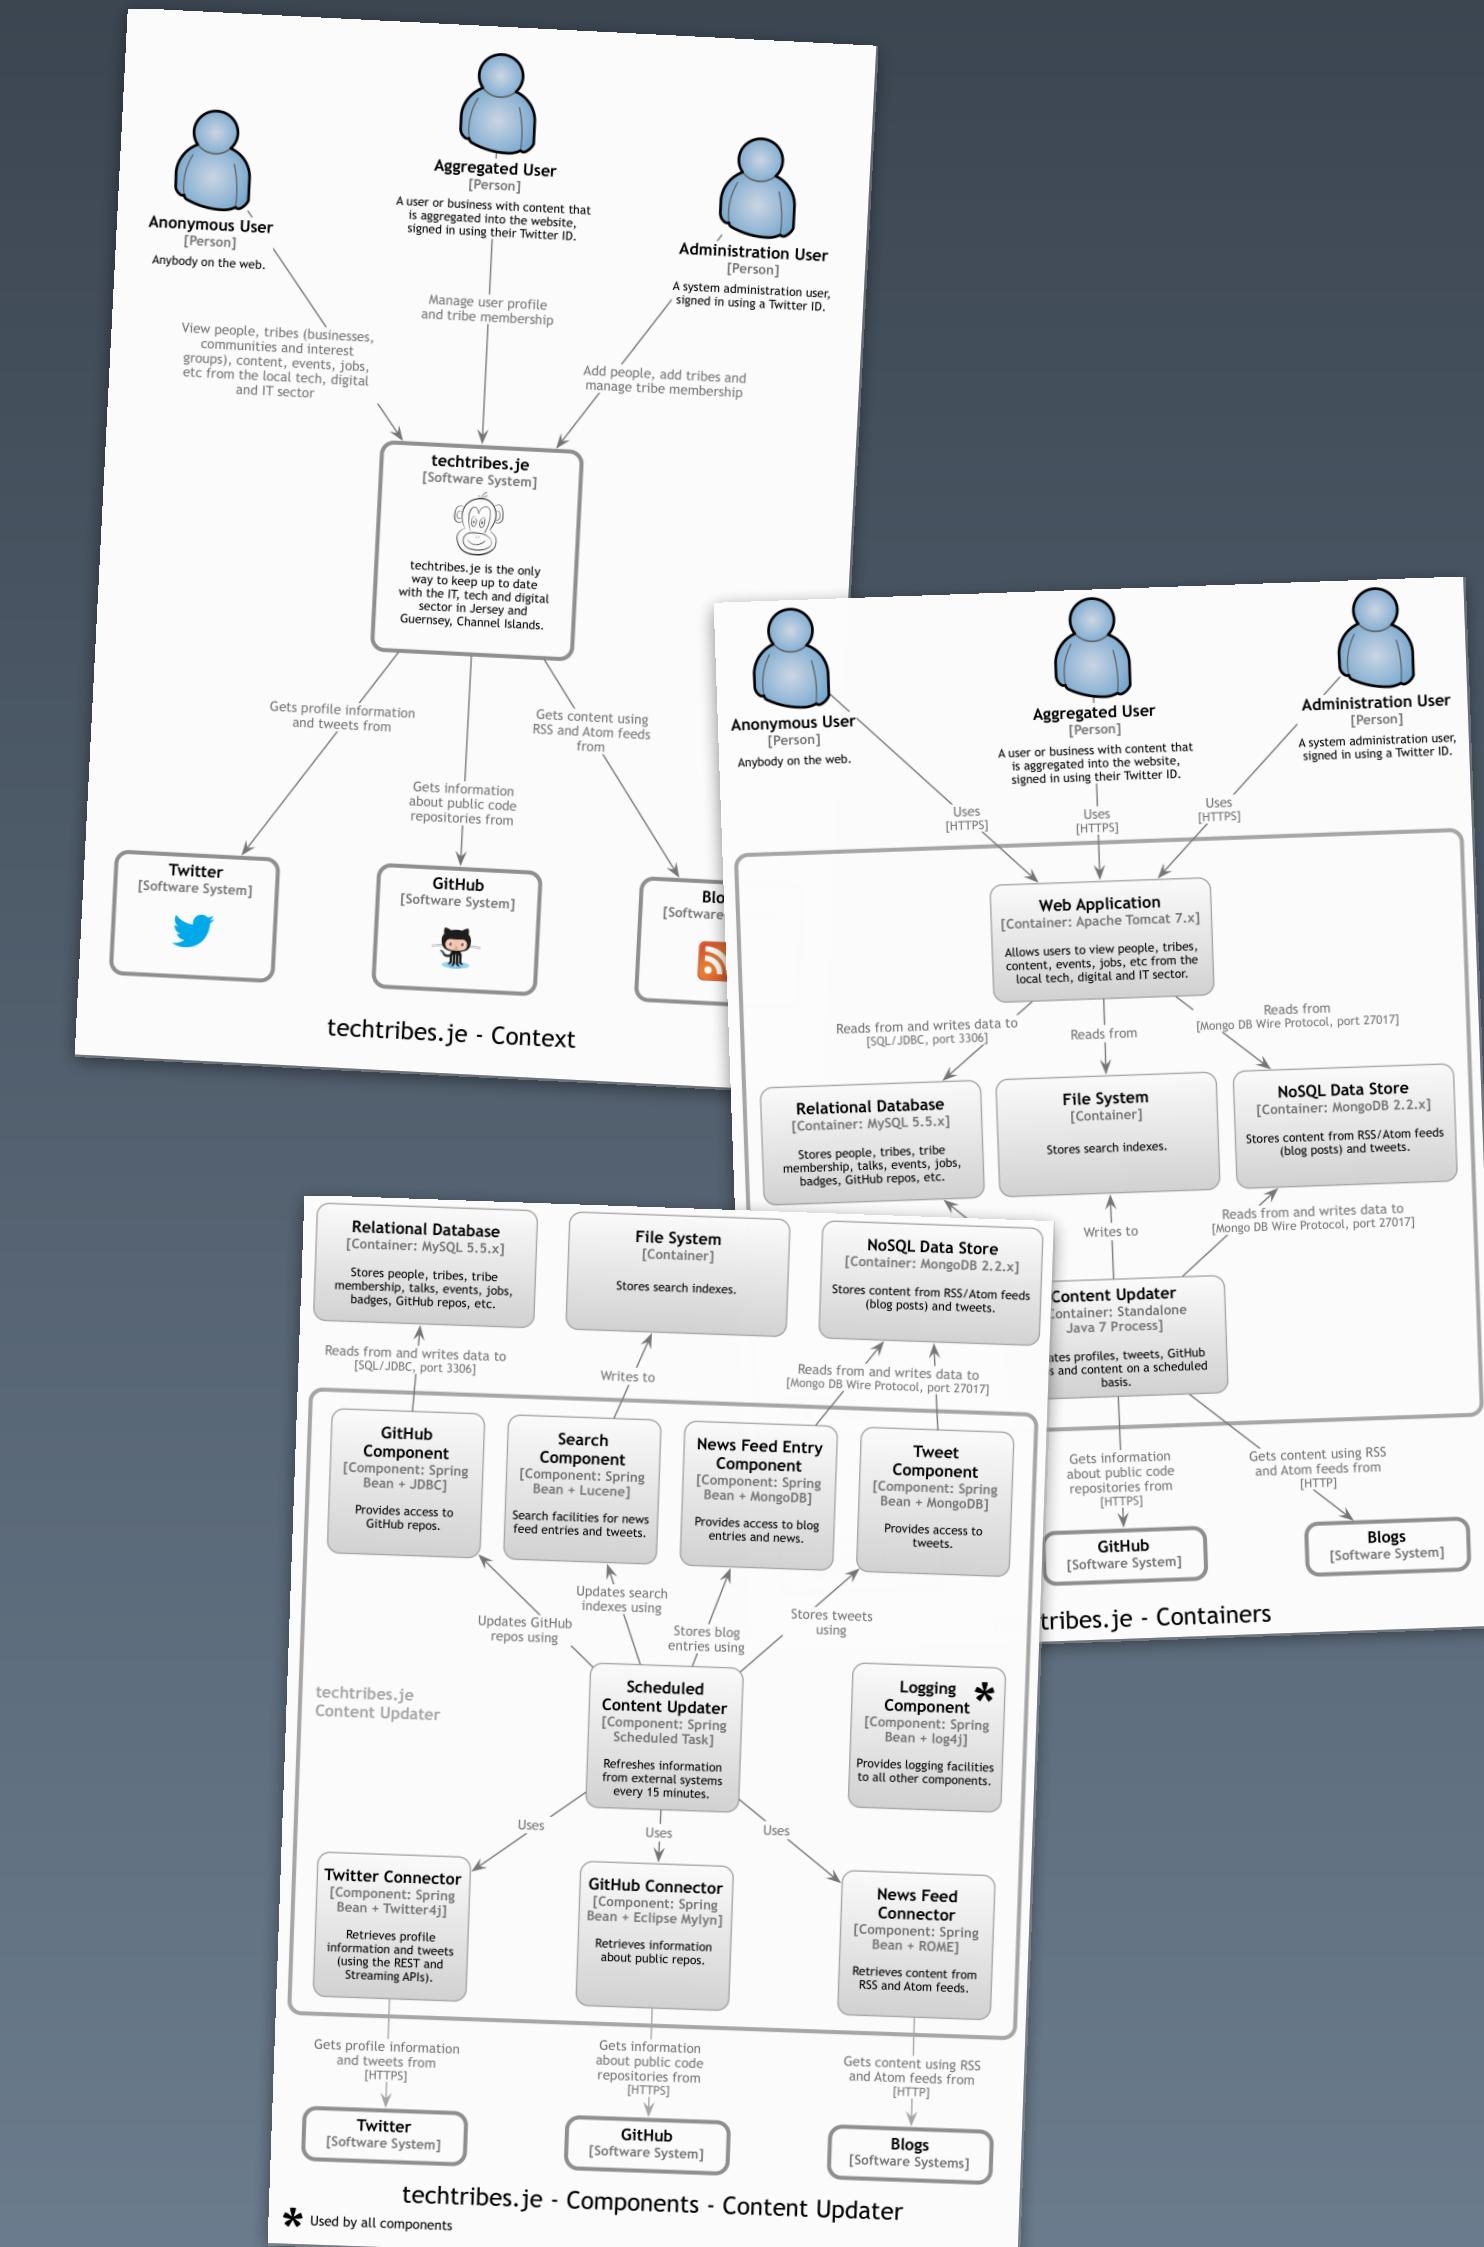
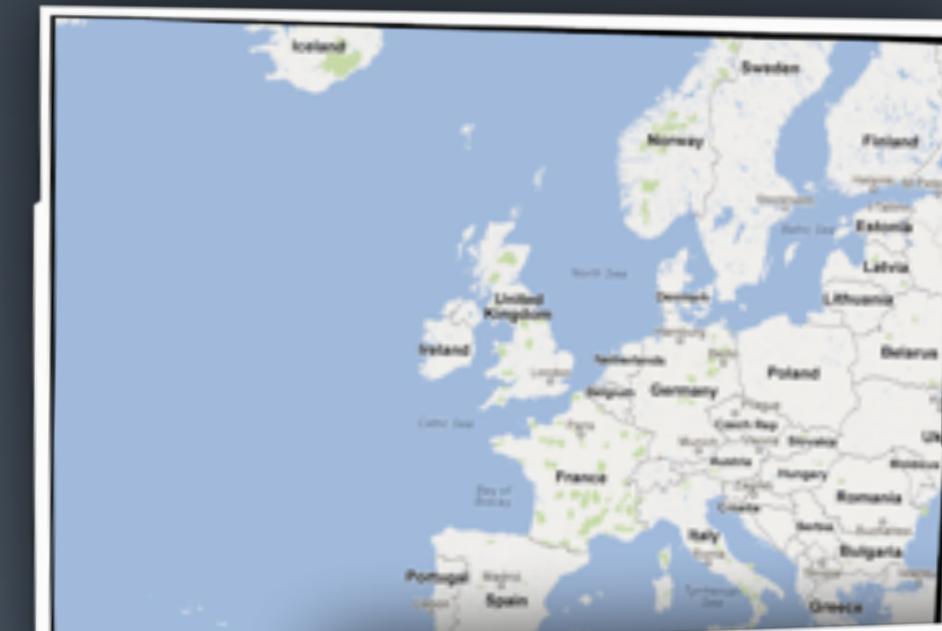
Class

Class

A **software system** is made up of one or more **containers**,  
each of which contains one or more **components**,  
which in turn are implemented by one or more **classes**.

# C4

Context, Containers, Components  
and Classes (or Code)



# Diagrams are maps that help a team navigate a complex codebase

The screenshot shows the Techtribes website interface. At the top, there's a navigation bar with links to News, Events, Talks, Content, Tweets, Code, People, Tribes, and Jobs. Below the navigation is a search bar with placeholder text "Find me people who know about..." and "Search...". A small "Talk me on GitHub" button is in the top right corner.

**Most active people:** Shows three small profile pictures and a grid of ten smaller ones below.

**Most active business tribes:** Shows icons for Property, C5, and RBS.

**Most active community tribes:** Shows icons for various local organizations.

## News

- C5 Alliance plans Microsoft events in Channel Islands**

Channel Island cloud provider, C5 Alliance are organising two breakfast events in both Jersey and Guernsey, named 'Leveraging Microsoft Technologies for Regulatory Compliance'. The breakfast briefings are due to include demonstrations of the latest Microsoft technologies and how they are combined. The briefing will cover Microsoft CRM process driven forms, SharePoint Workflow & Collaboration and SQL Server Data Warehousing technology. C5 Alliance, who work with a number of clients, both financial and...

Posted Today
- Jersey residents set to have choice in fibre broadband**

Sure customers will soon be able to access Jersey's fibre network following the reaching of an agreement between Sure and JT that finalises the commercial arrangements for access to the network. The agreement means that JT has gone some way to fulfilling the second condition of the eight that were set out in the States of Jersey's funding arrangements for the network, as agreed by the Treasury Minister, Senator Philip Oxo. "This is excellent news for our broadband customers who have been extremely pati...

Posted Yesterday
- Logicalis Group taking over Jersey cloud provider**

Logicalis Group, the international IT solutions and managed services provider, has announced the acquisition of Jersey's iConsult Limited, a privately owned Jersey company and provider of desktop and mail hosted solutions to the small medium businesses (SMB) market within the Channel Islands. Through their data facility in Jersey the company services over 800 users on the islands, mainly in the financial and professional services sectors. Their main offering is a hosted desktop solution, using primarily ...

Posted 18 Oct 2013

[More...](#)

## Local events

2014 2013 2012

- Ivan Nikkhoo - Growth Funding**

Topics of discussion will be: Growth capital, Funding cycles, Investment decisions plus Valuation and exits. With over 29 years of industry experience in various senior capacities internationally, Ivan is a Managing Director at Siemer & Associates and a...

Pomme d'Or Hotel, St Helier, Jersey  
29 Oct 2013 at 17:30
- Tech Tribes Talks**

The third set of Tech Tribes talks are ready to rock your world! After a very successful July event at the Royal Yacht we've decided to go back for our October talks. We have a great line up of speakers and we take great pleasure in inviting you to atte...

The Royal Yacht, St Heller  
24 Oct 2013 at 17:30
- The Internet of Everything and Gigabit Jersey**

'The Internet of everything' Currently, there are an estimated 10 to 15 billion 'things' connected to the Internet and this is predicted to grow to 50 billion by 2020. How will this change our lives? What infrastructure will we need? What opportun...

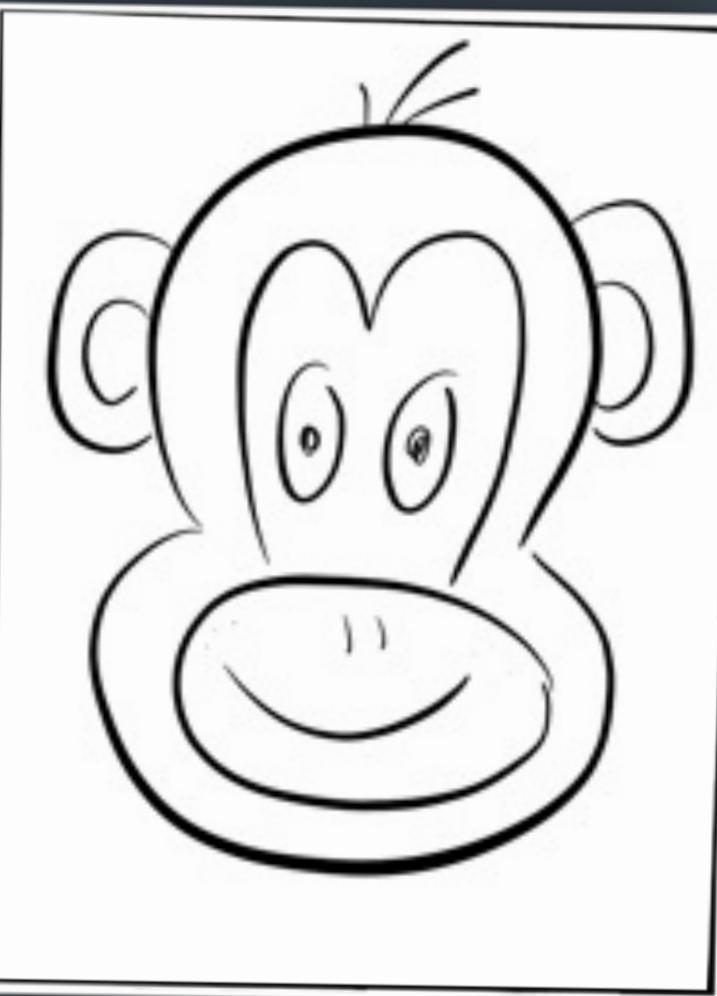
The Grand Hotel, St Heller, Jersey  
Tomorrow at 17:15

## Talks by local speakers

2014 2013 2012

- Ted talk**
- Agile software**
- Software architecture and the balance with agility**

# techtribes.je

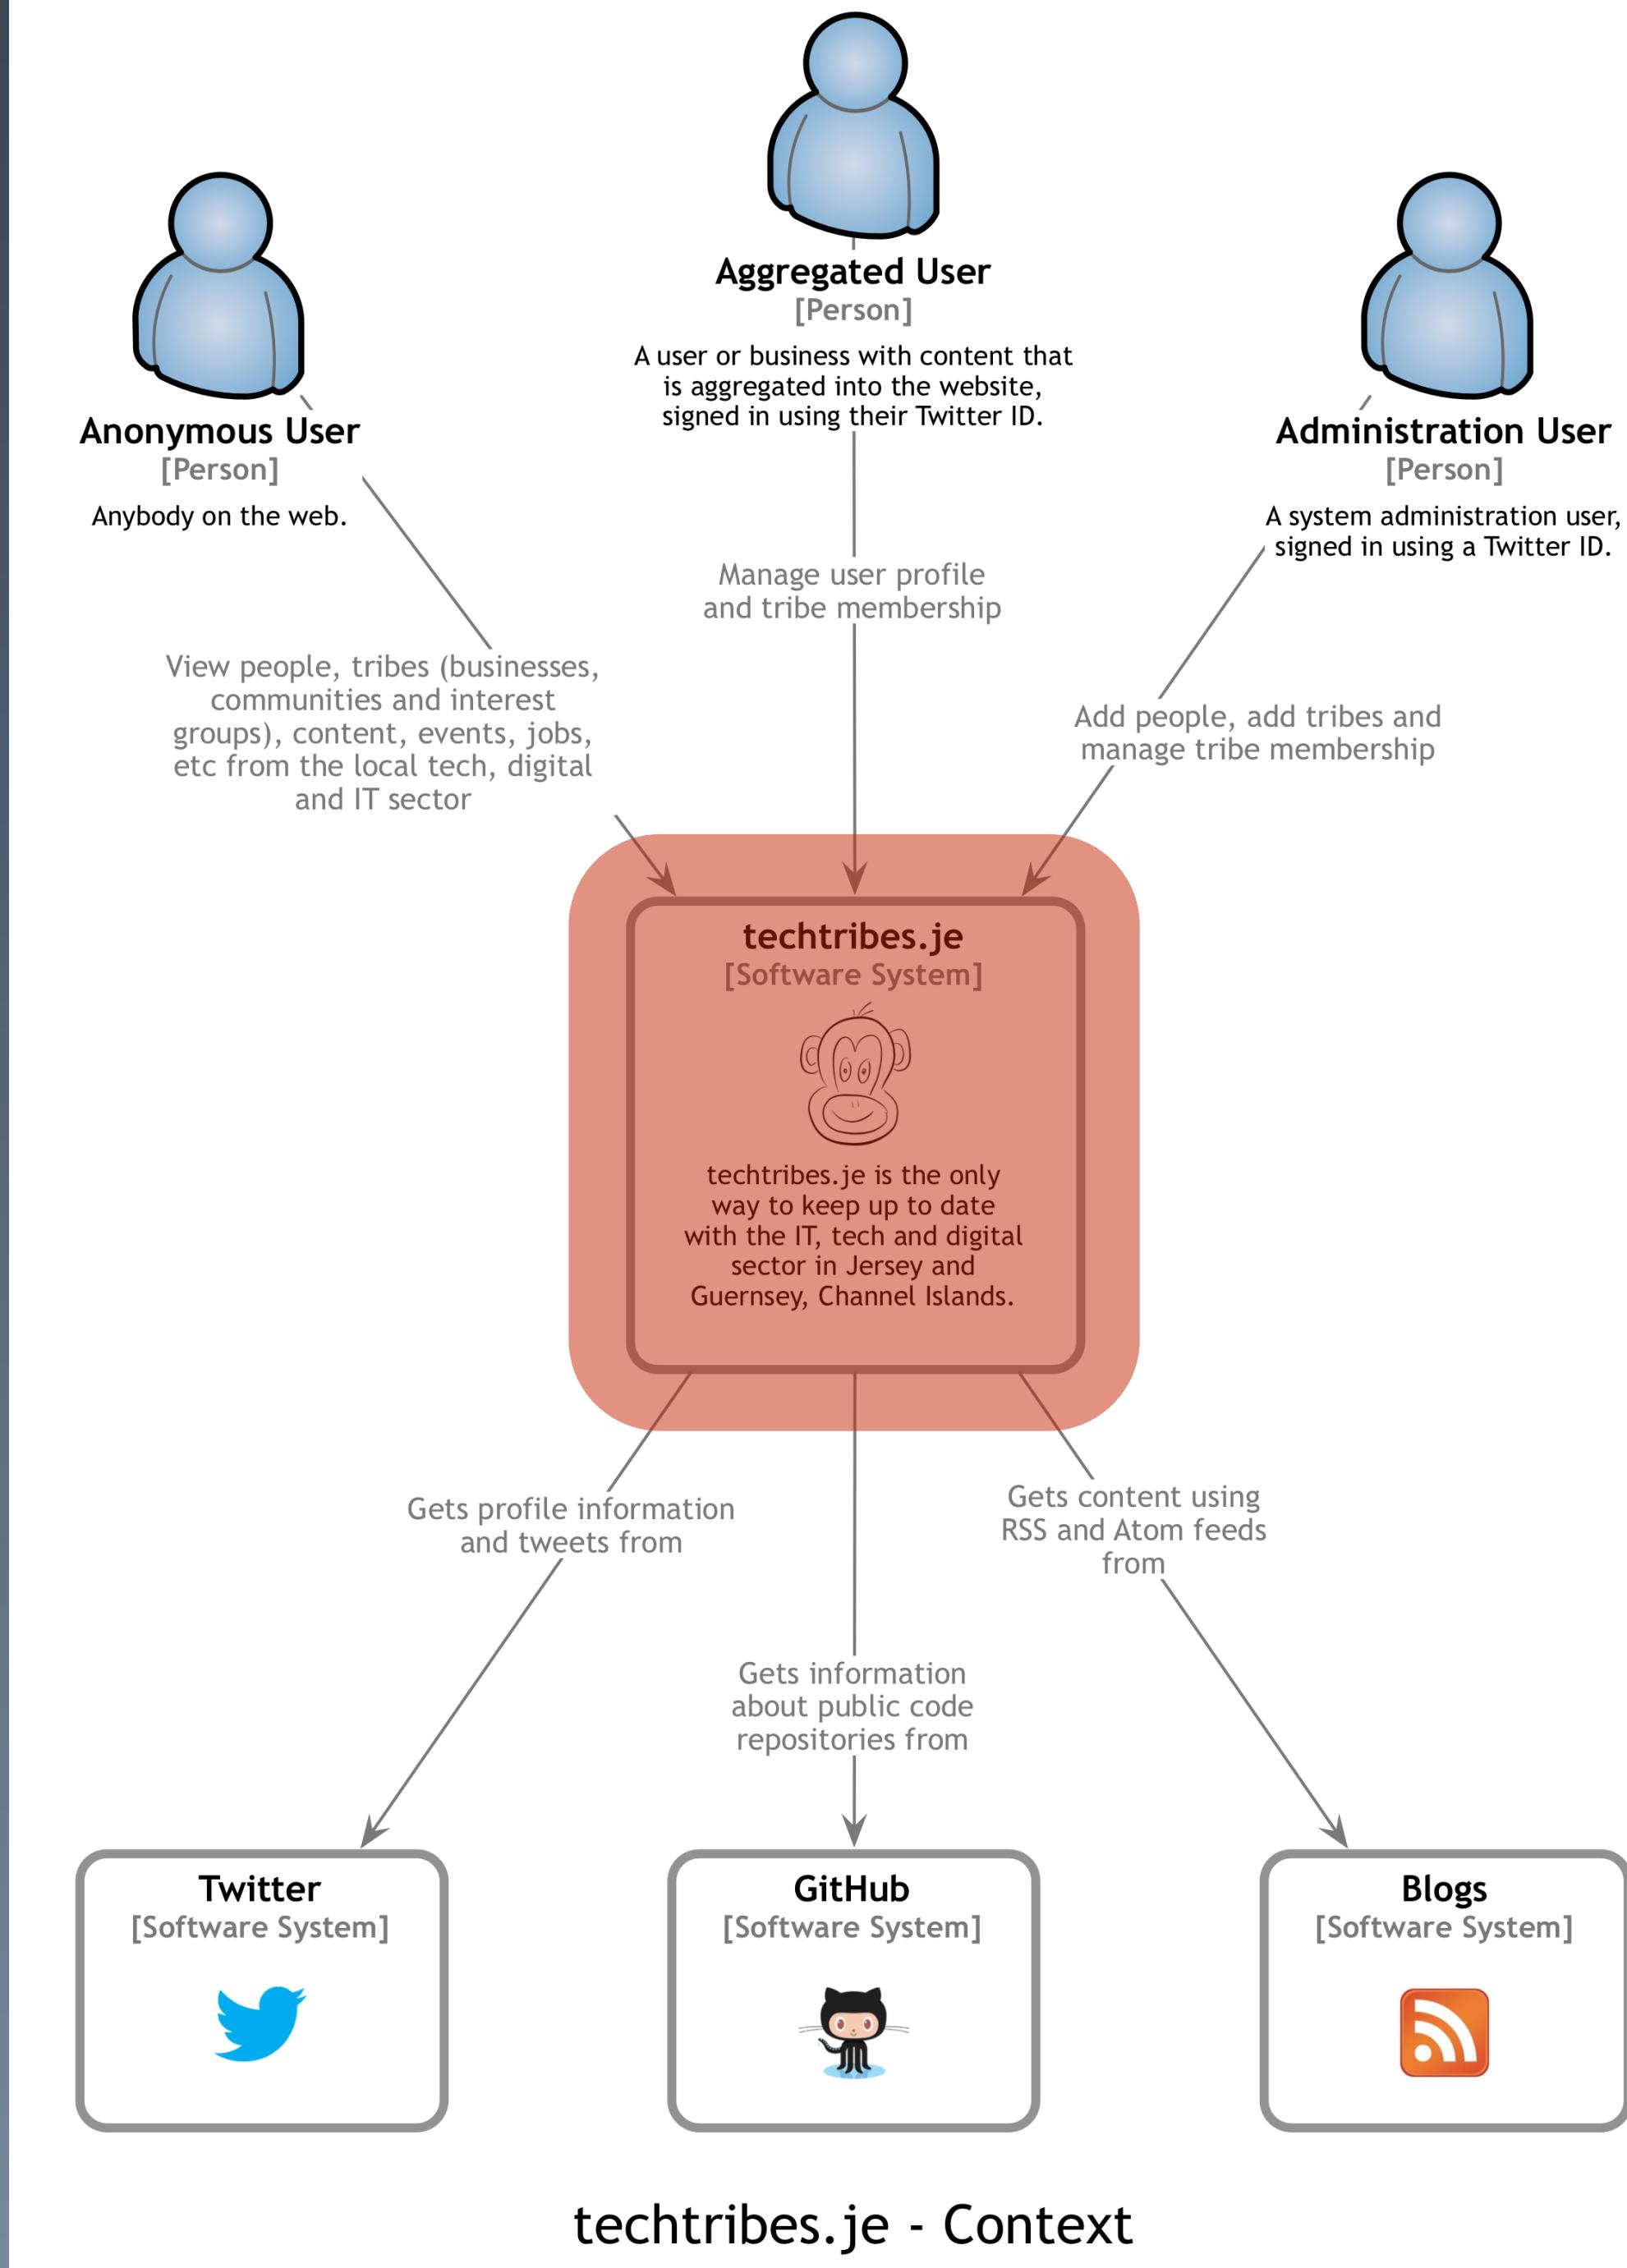


# Context diagram (level 1)

## Container diagram (level 2)

### Component diagram (level 3)

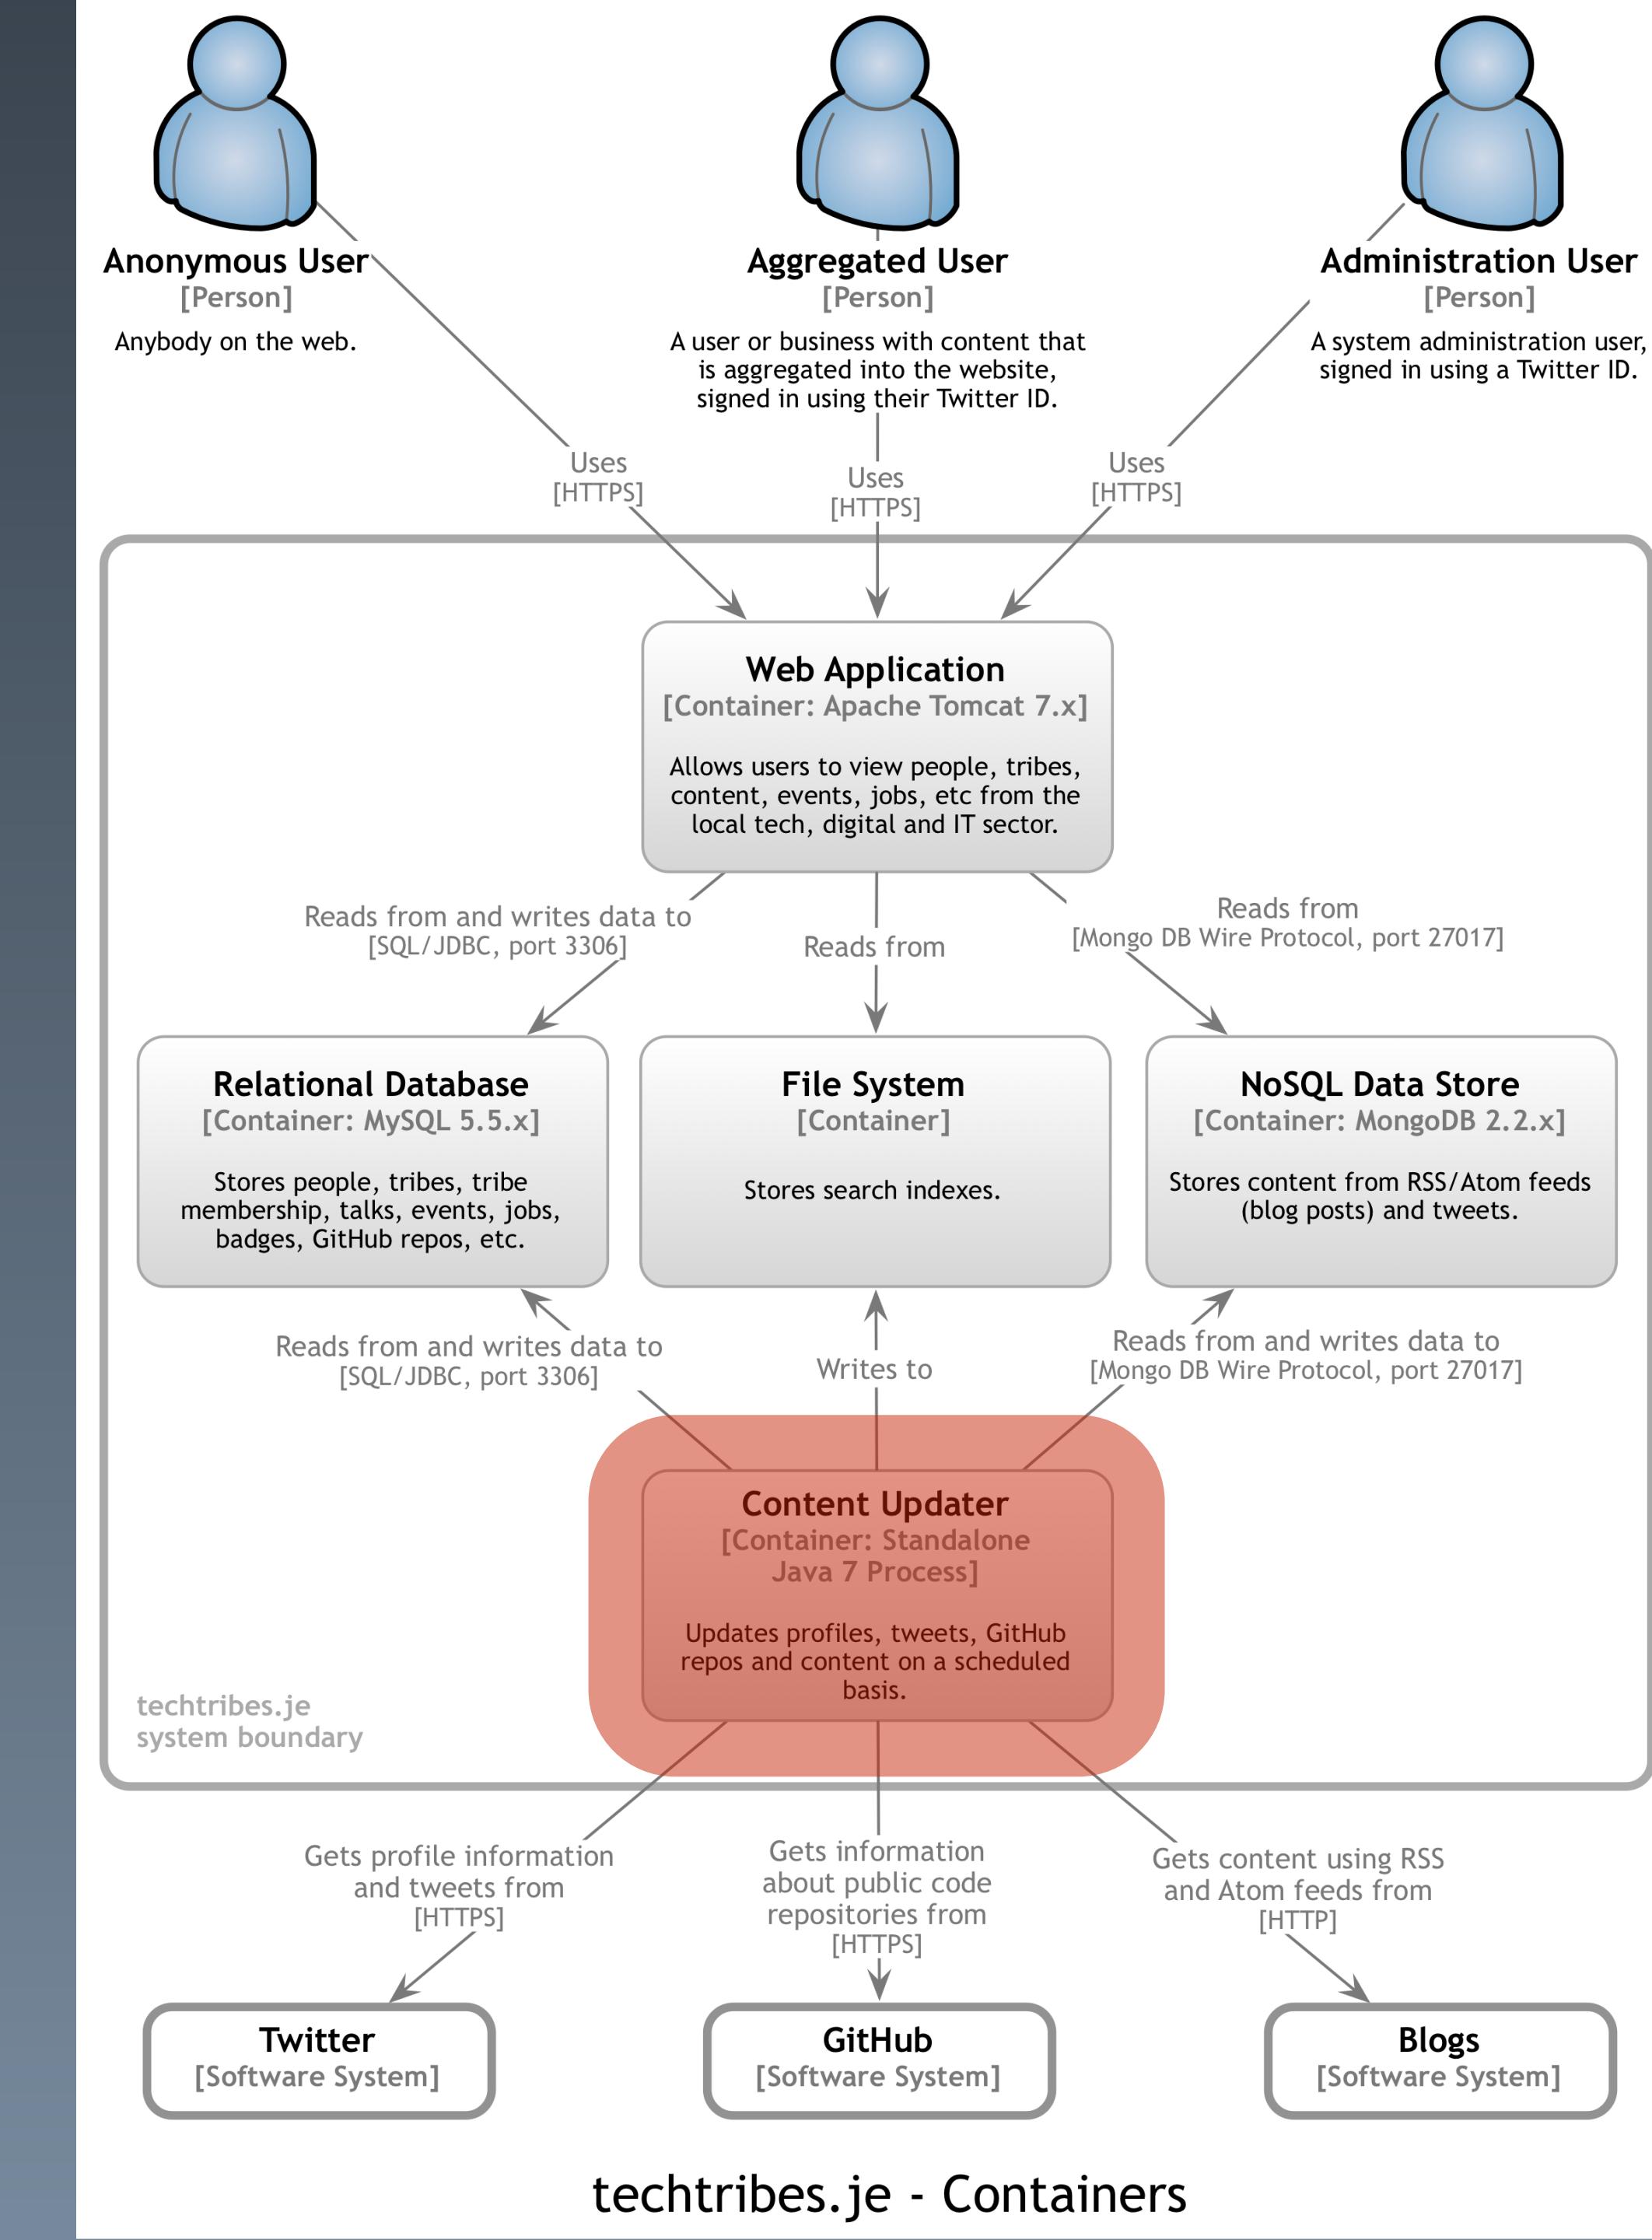
#### Class diagram (level 4)



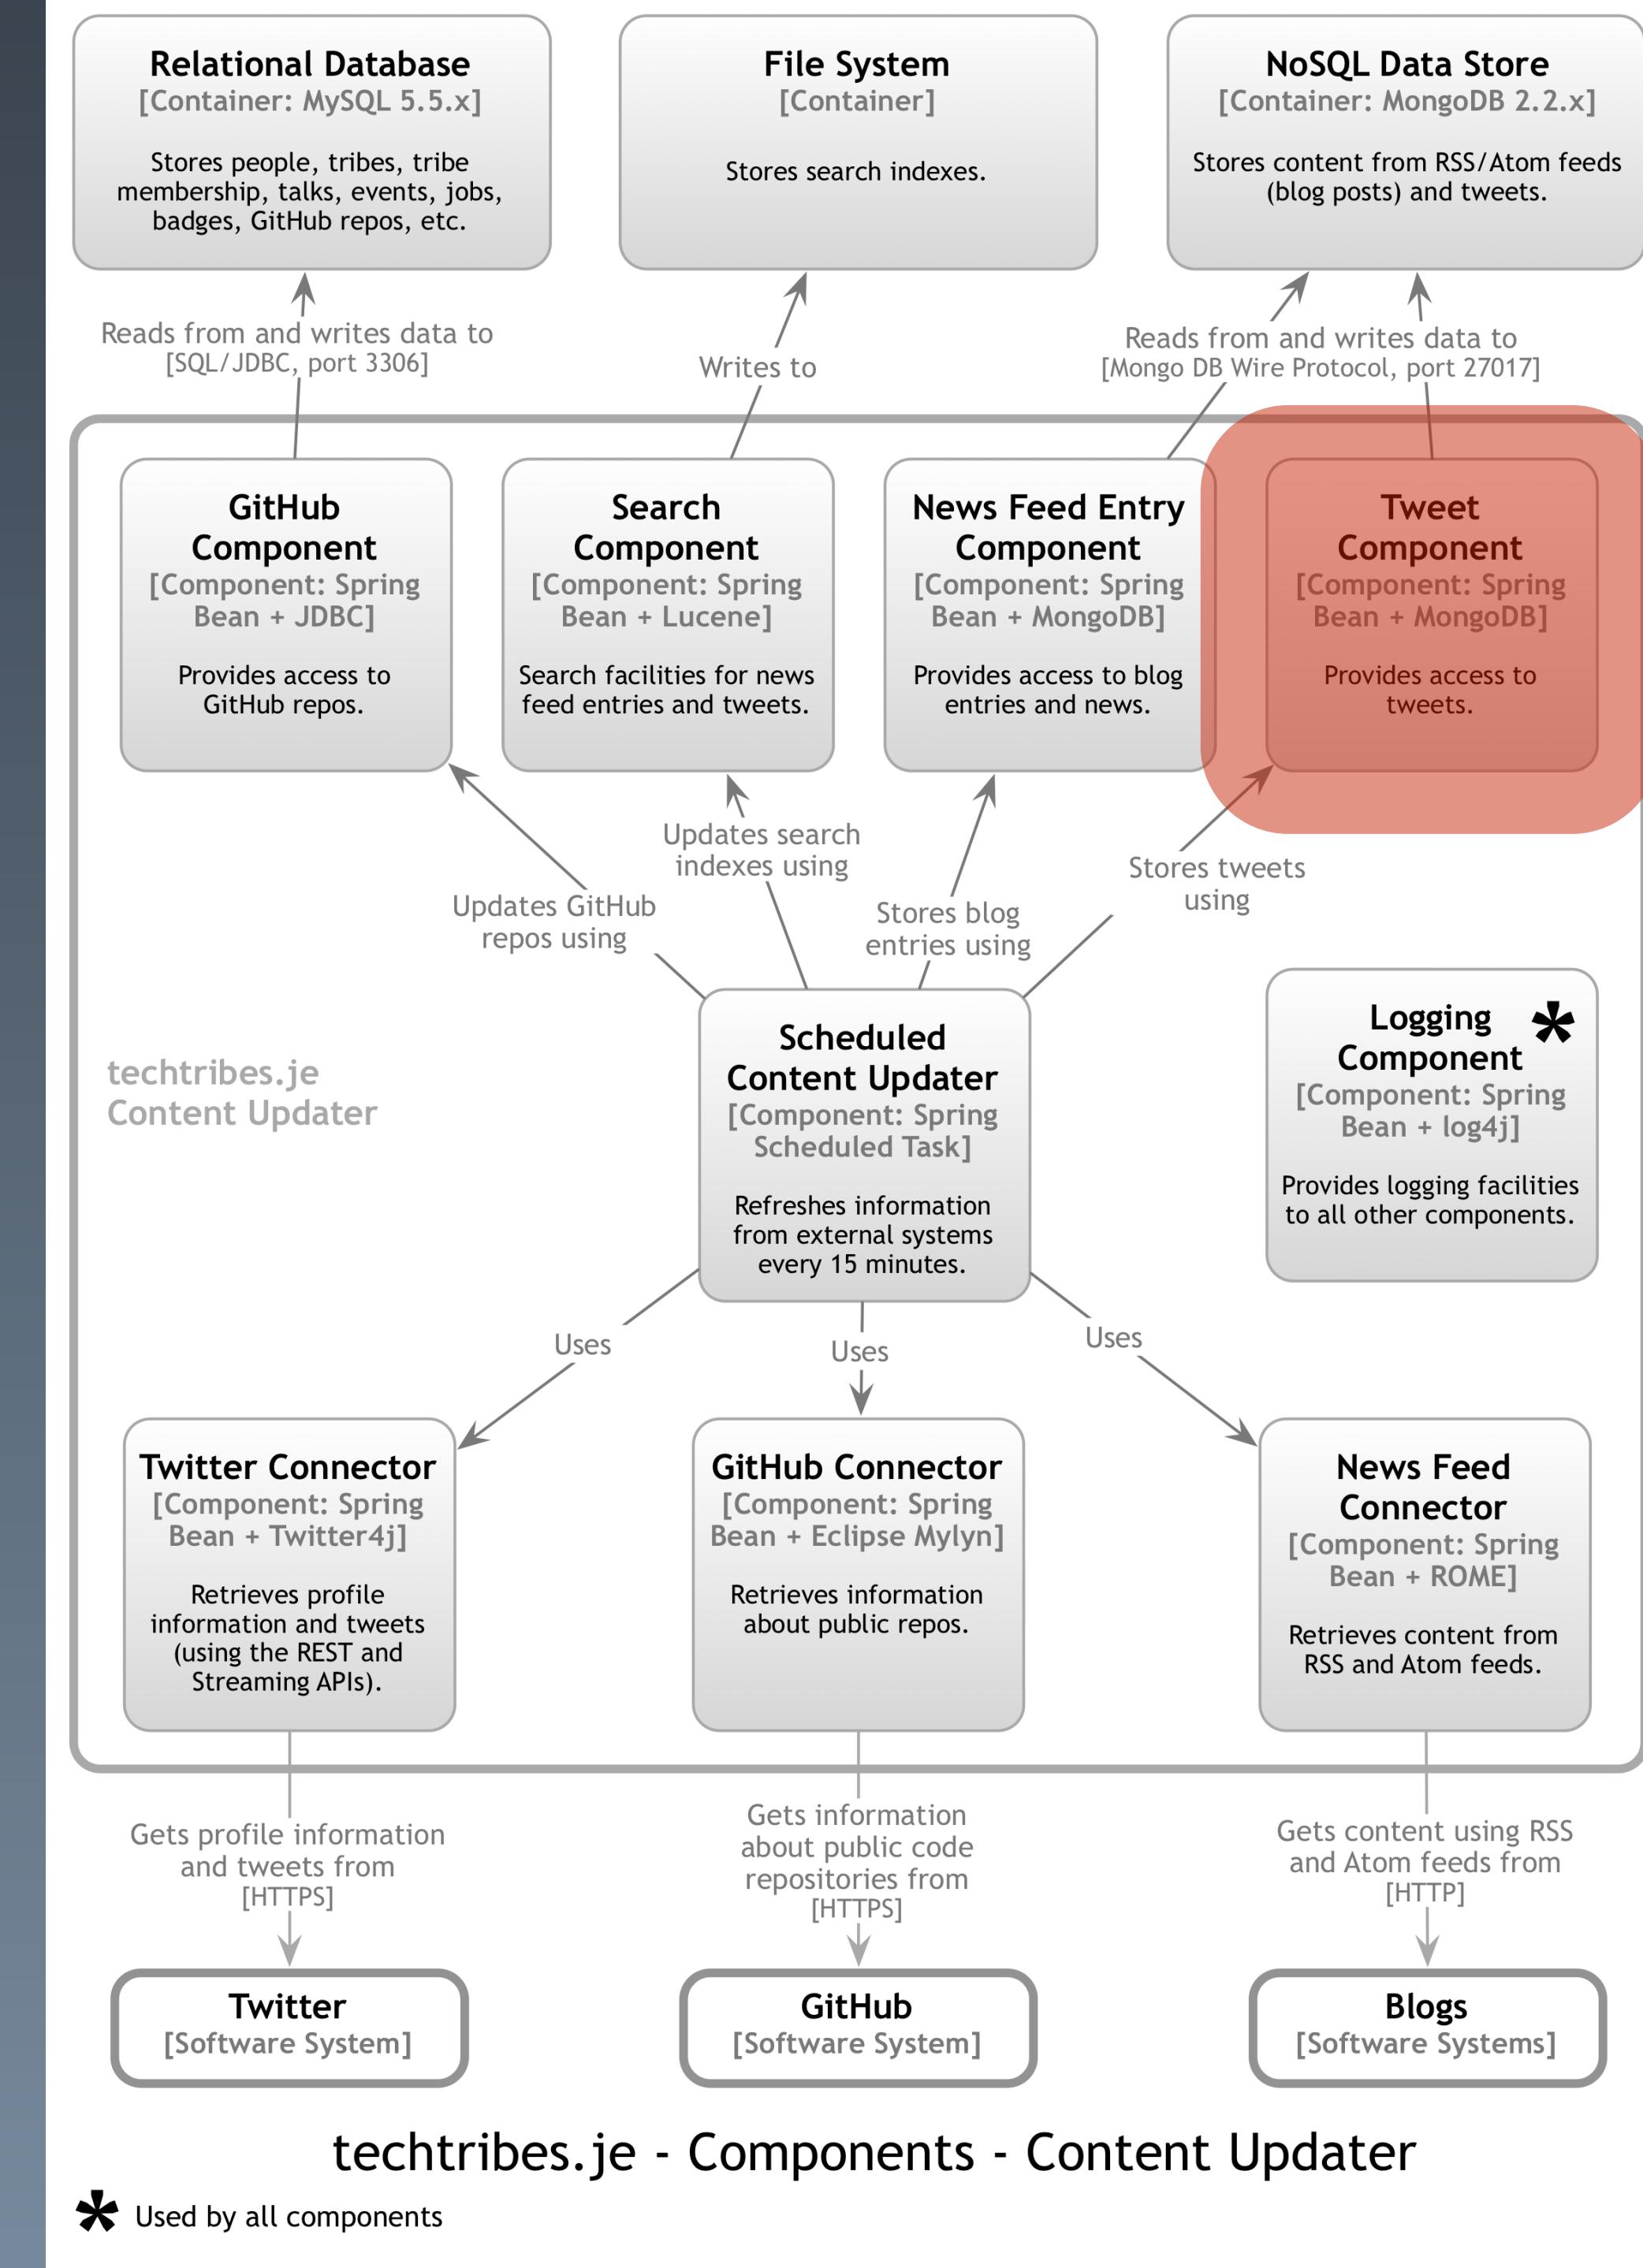
# Container diagram (level 2)

## Component diagram (level 3)

### Class diagram (level 4)



# Component diagram (level 3)



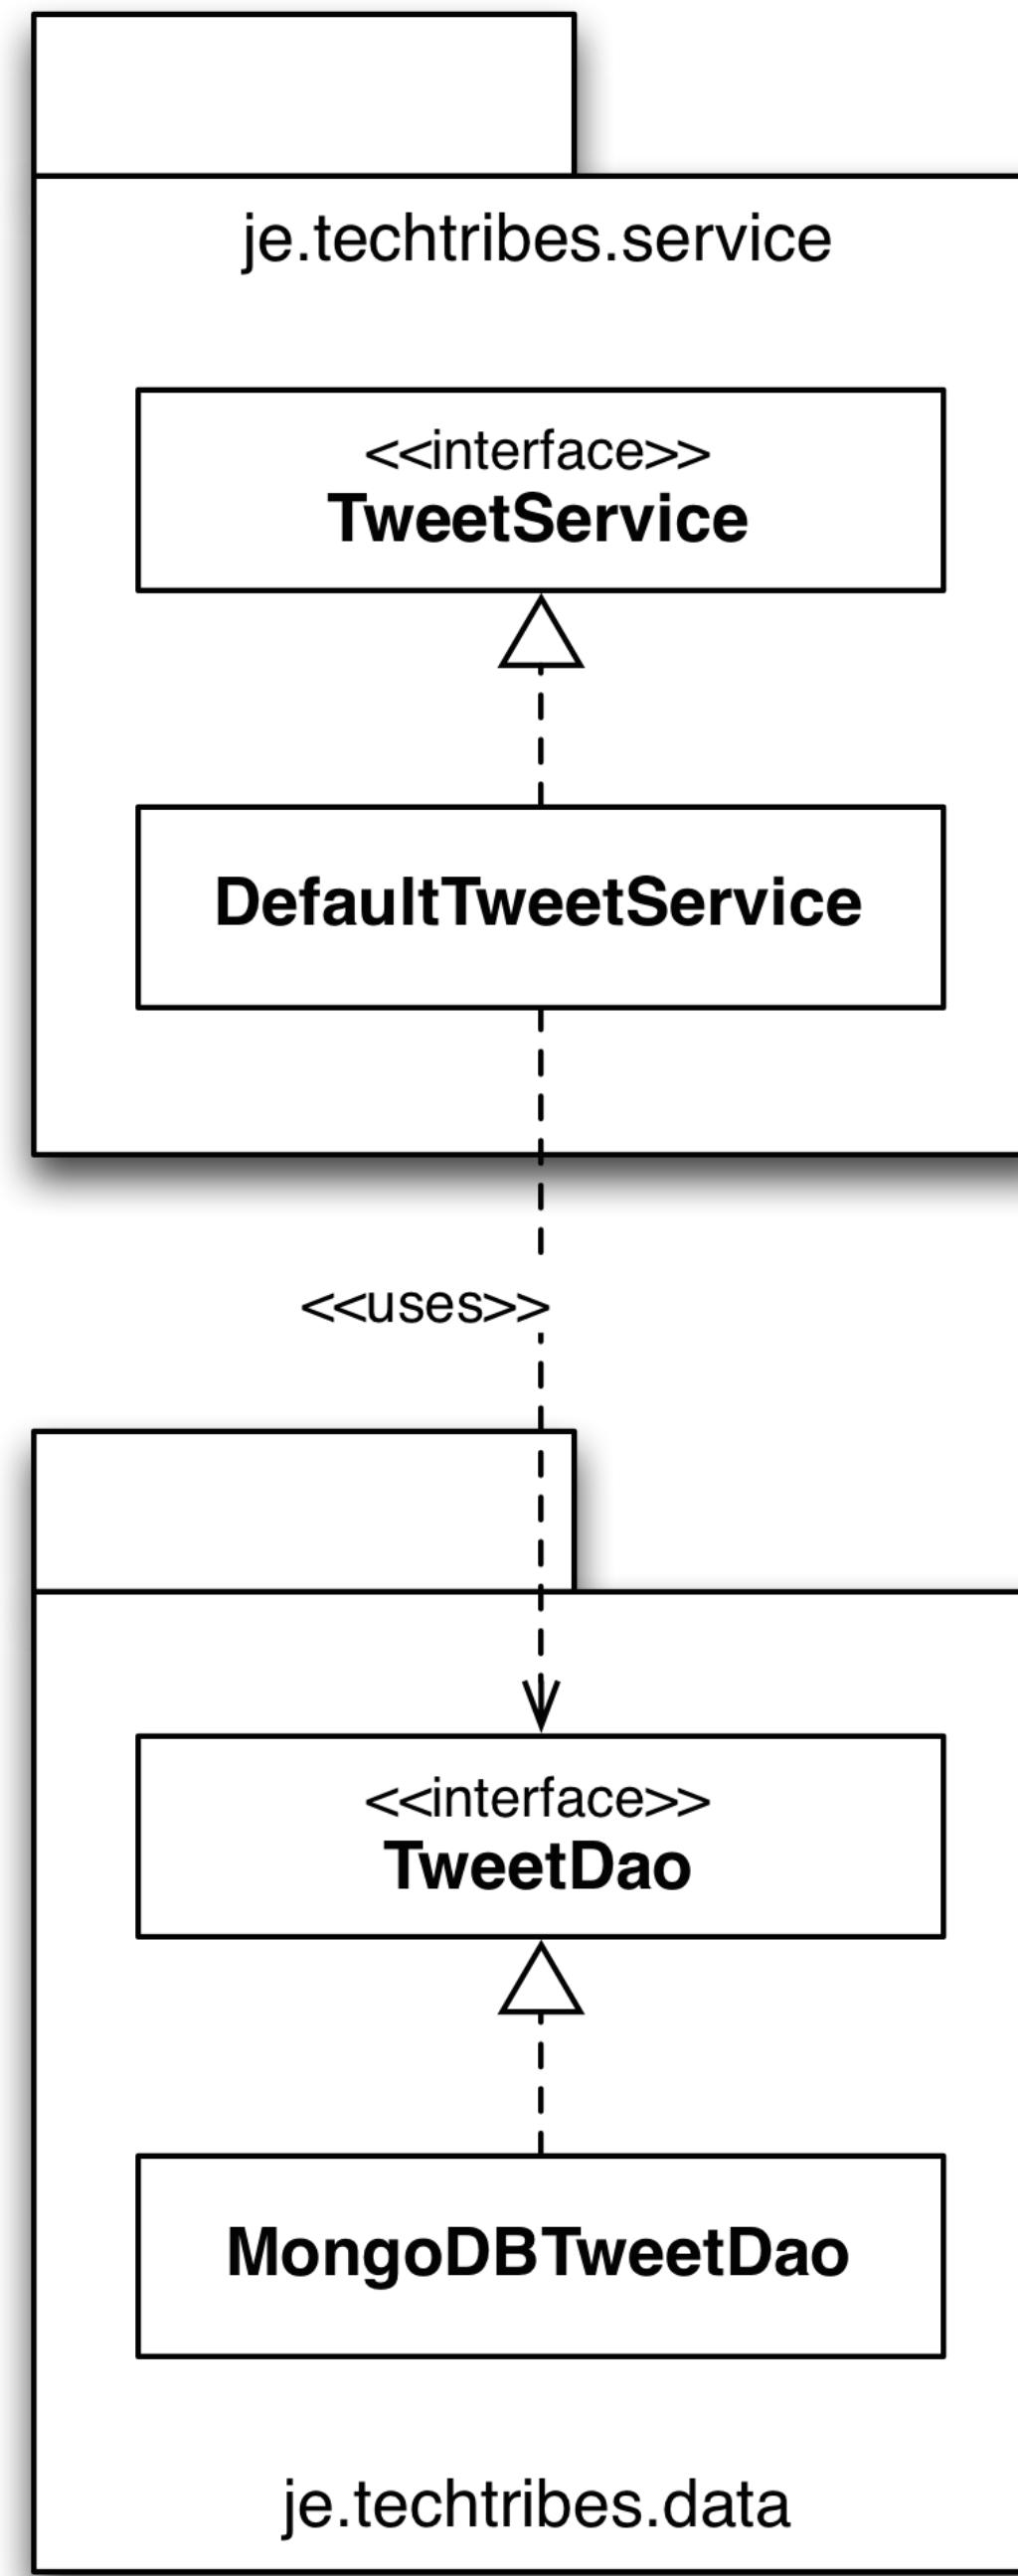
**\*** Used by all components

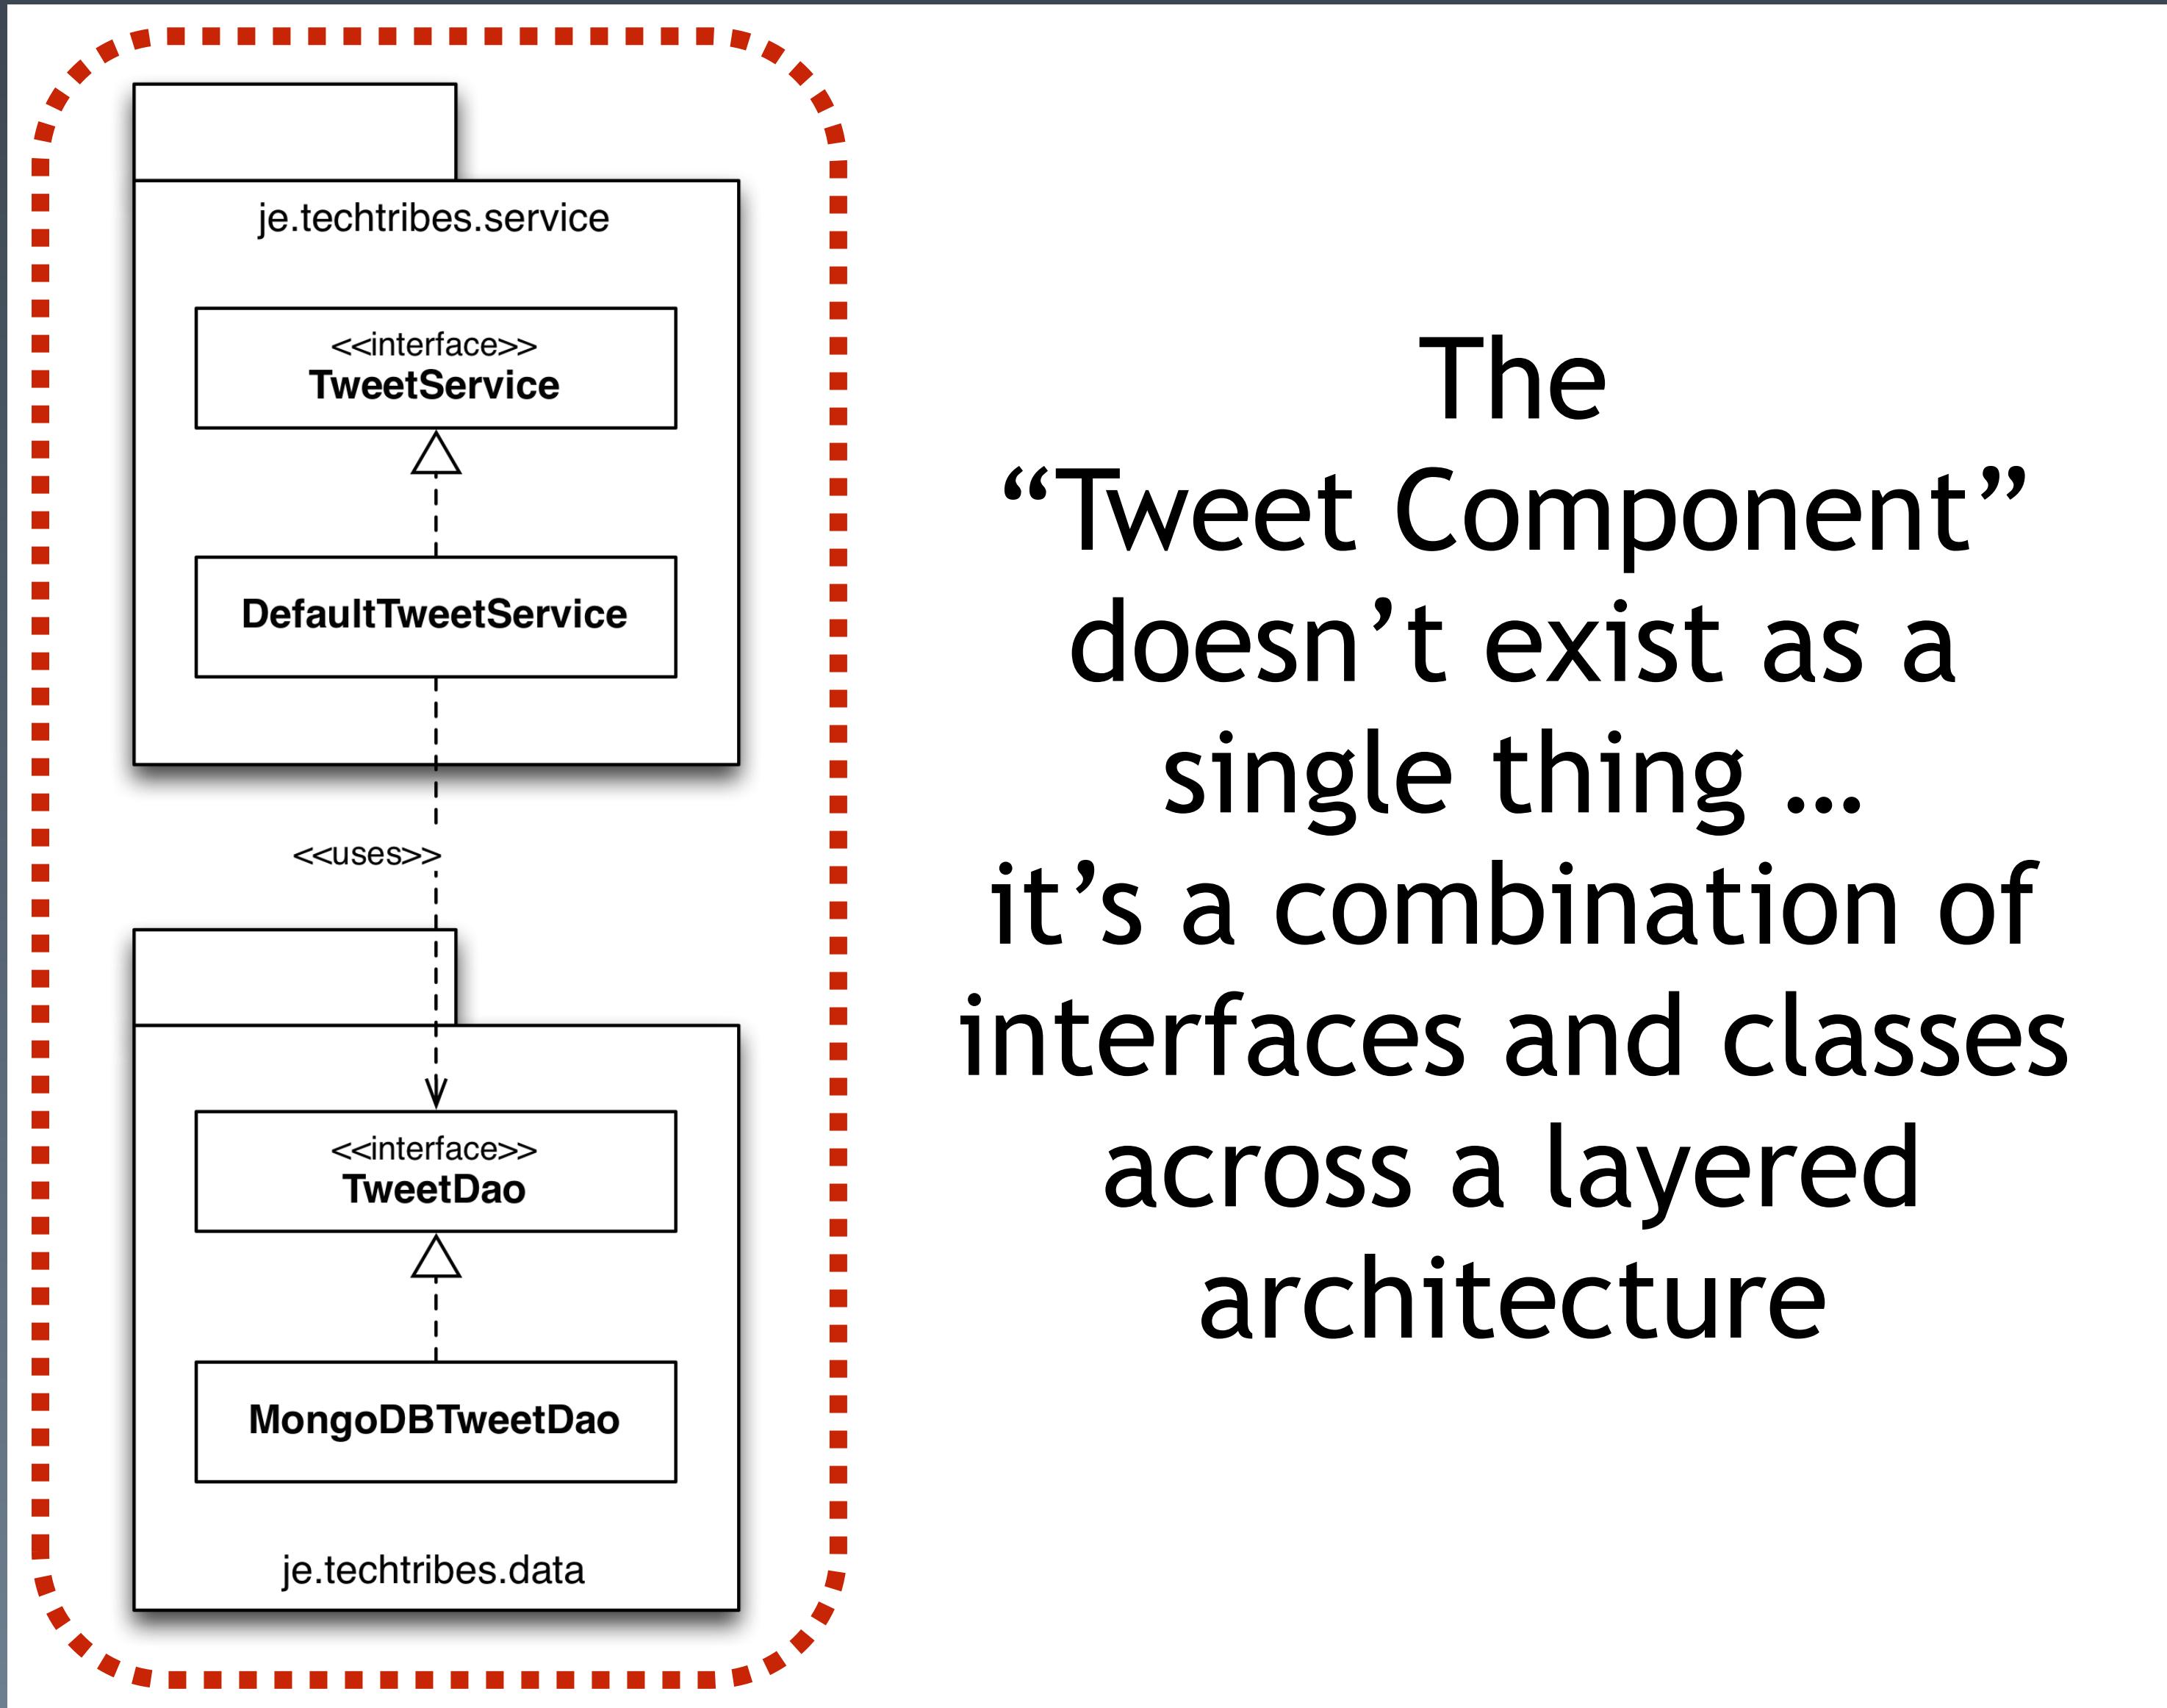
# Context diagram (level 1)

# Container diagram (level 2)

# Component diagram (level 3)

# Class diagram (level 4)



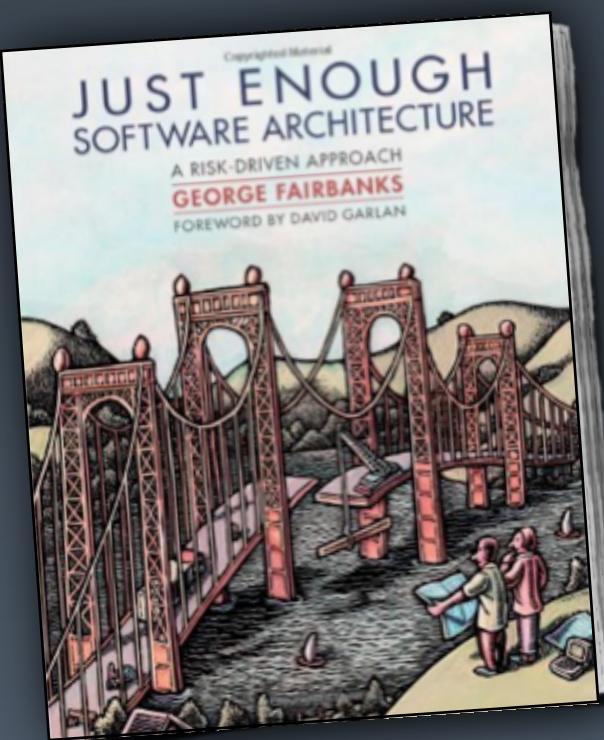


The  
“Tweet Component”  
doesn’t exist as a  
single thing ...  
it’s a combination of  
interfaces and classes  
across a layered  
architecture

Where's my component?

*“But the component  
exists conceptually”*

Does your code reflect the  
abstractions  
that you think about?



# “the model-code gap”

**Model-code gap.** Your architecture models and your source code will not show the same things. The difference between them is the *model-code gap*. Your architecture models include some abstract concepts, like components, that your programming language does not, but could. Beyond that, architecture models include intensional elements, like design decisions and constraints, that cannot be expressed in procedural source code at all.

Consequently, the relationship between the architecture model and source code is complicated. It is mostly a refinement relationship, where the extensional elements in the architecture model are refined into extensional elements in source code. This is shown in Figure 10.3. However, intensional elements are not refined into corresponding elements in source code.

Upon learning about the model-code gap, your first instinct may be to avoid it. But reflecting on the origins of the gap gives little hope of a general solution in the short term: architecture models help you reason about complexity and scale because they are abstract and intensional; source code executes on machines because it is concrete and extensional.

# Software Reflexion Models: Bridging the Gap between Source and High-Level Models\*

Gail C. Murphy and David Notkin

Dept. of Computer Science & Engineering  
University of Washington  
Box 352350  
Seattle WA, USA 98195-2350  
[{gmurphy,notkin}@cs.washington.edu](mailto:{gmurphy,notkin}@cs.washington.edu)

## Abstract

Software engineers often use high-level models (for instance, box and arrow sketches) to reason and communicate about an existing software system. One problem with high-level models is that they are almost always inaccurate with respect to the system's source code. We have developed an approach that helps an engineer use a high-level model of the structure of an existing software system as a lens through which to see a model of that system's source code. In particular, an engineer defines a high-level model and specifies how the model maps to the source. A tool then computes a software reflexion model that shows where the engineer's high-level model agrees with and where it differs from a model of the source.

The paper provides a formal characterization of reflexion models, discusses practical aspects of the approach, and relates experiences of applying the approach and tools to a number of different systems. The illustrative example used in the paper describes the application of reflexion models to NetBSD, an implementation of Unix comprised of 250,000 lines of C code. In only a few hours, an engineer computed several reflexion models that provided him with a useful, global overview of the structure of the NetBSD virtual memory subsystem. The approach has also been applied to aid in the understanding and experimental reengineering of the Microsoft Excel spreadsheet product.

\*This research was funded in part by the NSF grant CCR-8858804 and a Canadian NSERC post-graduate scholarship.

<sup>0</sup>Permission to make digital/hard copies of all or part of this material without fee is granted provided that the copies are not made or distributed for profit or commercial advantage, the ACM copyright/server notice, the title of the publication and its date appear, and notice is given that copyright is by permission of the Association for Computing Machinery, Inc. (ACM). To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

SIGSOFT '95 Washington, D.C., USA  
©1995 ACM 0-89791-716-2/95/0010...\$3.50

Kevin Sullivan

Dept. of Computer Science  
University of Virginia  
Charlottesville VA, USA 22903  
[sullivan@cs.virginia.edu](mailto:sullivan@cs.virginia.edu)

## 1 Introduction

Software engineers often think about an existing software system in terms of high-level models. Box and arrow sketches of a system, for instance, are often found on engineers' whiteboards. Although these models are commonly used, reasoning about the system in terms of such models can be dangerous because the models are almost always inaccurate with respect to the system's source.

Current reverse engineering systems derive high-level models from the source code. These derived models are useful because they are, by their very nature, accurate representations of the source. Although accurate, the models created by these reverse engineering systems may differ from the models sketched by engineers; an example of this is reported by Wong et al. [WTMS95].

We have developed an approach, illustrated in Figure 1, that enables an engineer to produce sufficiently accurate high-level models in a different way. The engineer defines a high-level model of interest, extracts a source model (such as a call graph or an inheritance hierarchy) from the source code, and defines a declarative mapping between the two models. A *software reflexion model* is then computed to determine where the engineer's high-level model does and does not agree with the source model.<sup>1</sup> An engineer interprets the reflexion model and, as necessary, modifies the input to iteratively compute additional reflexion models.

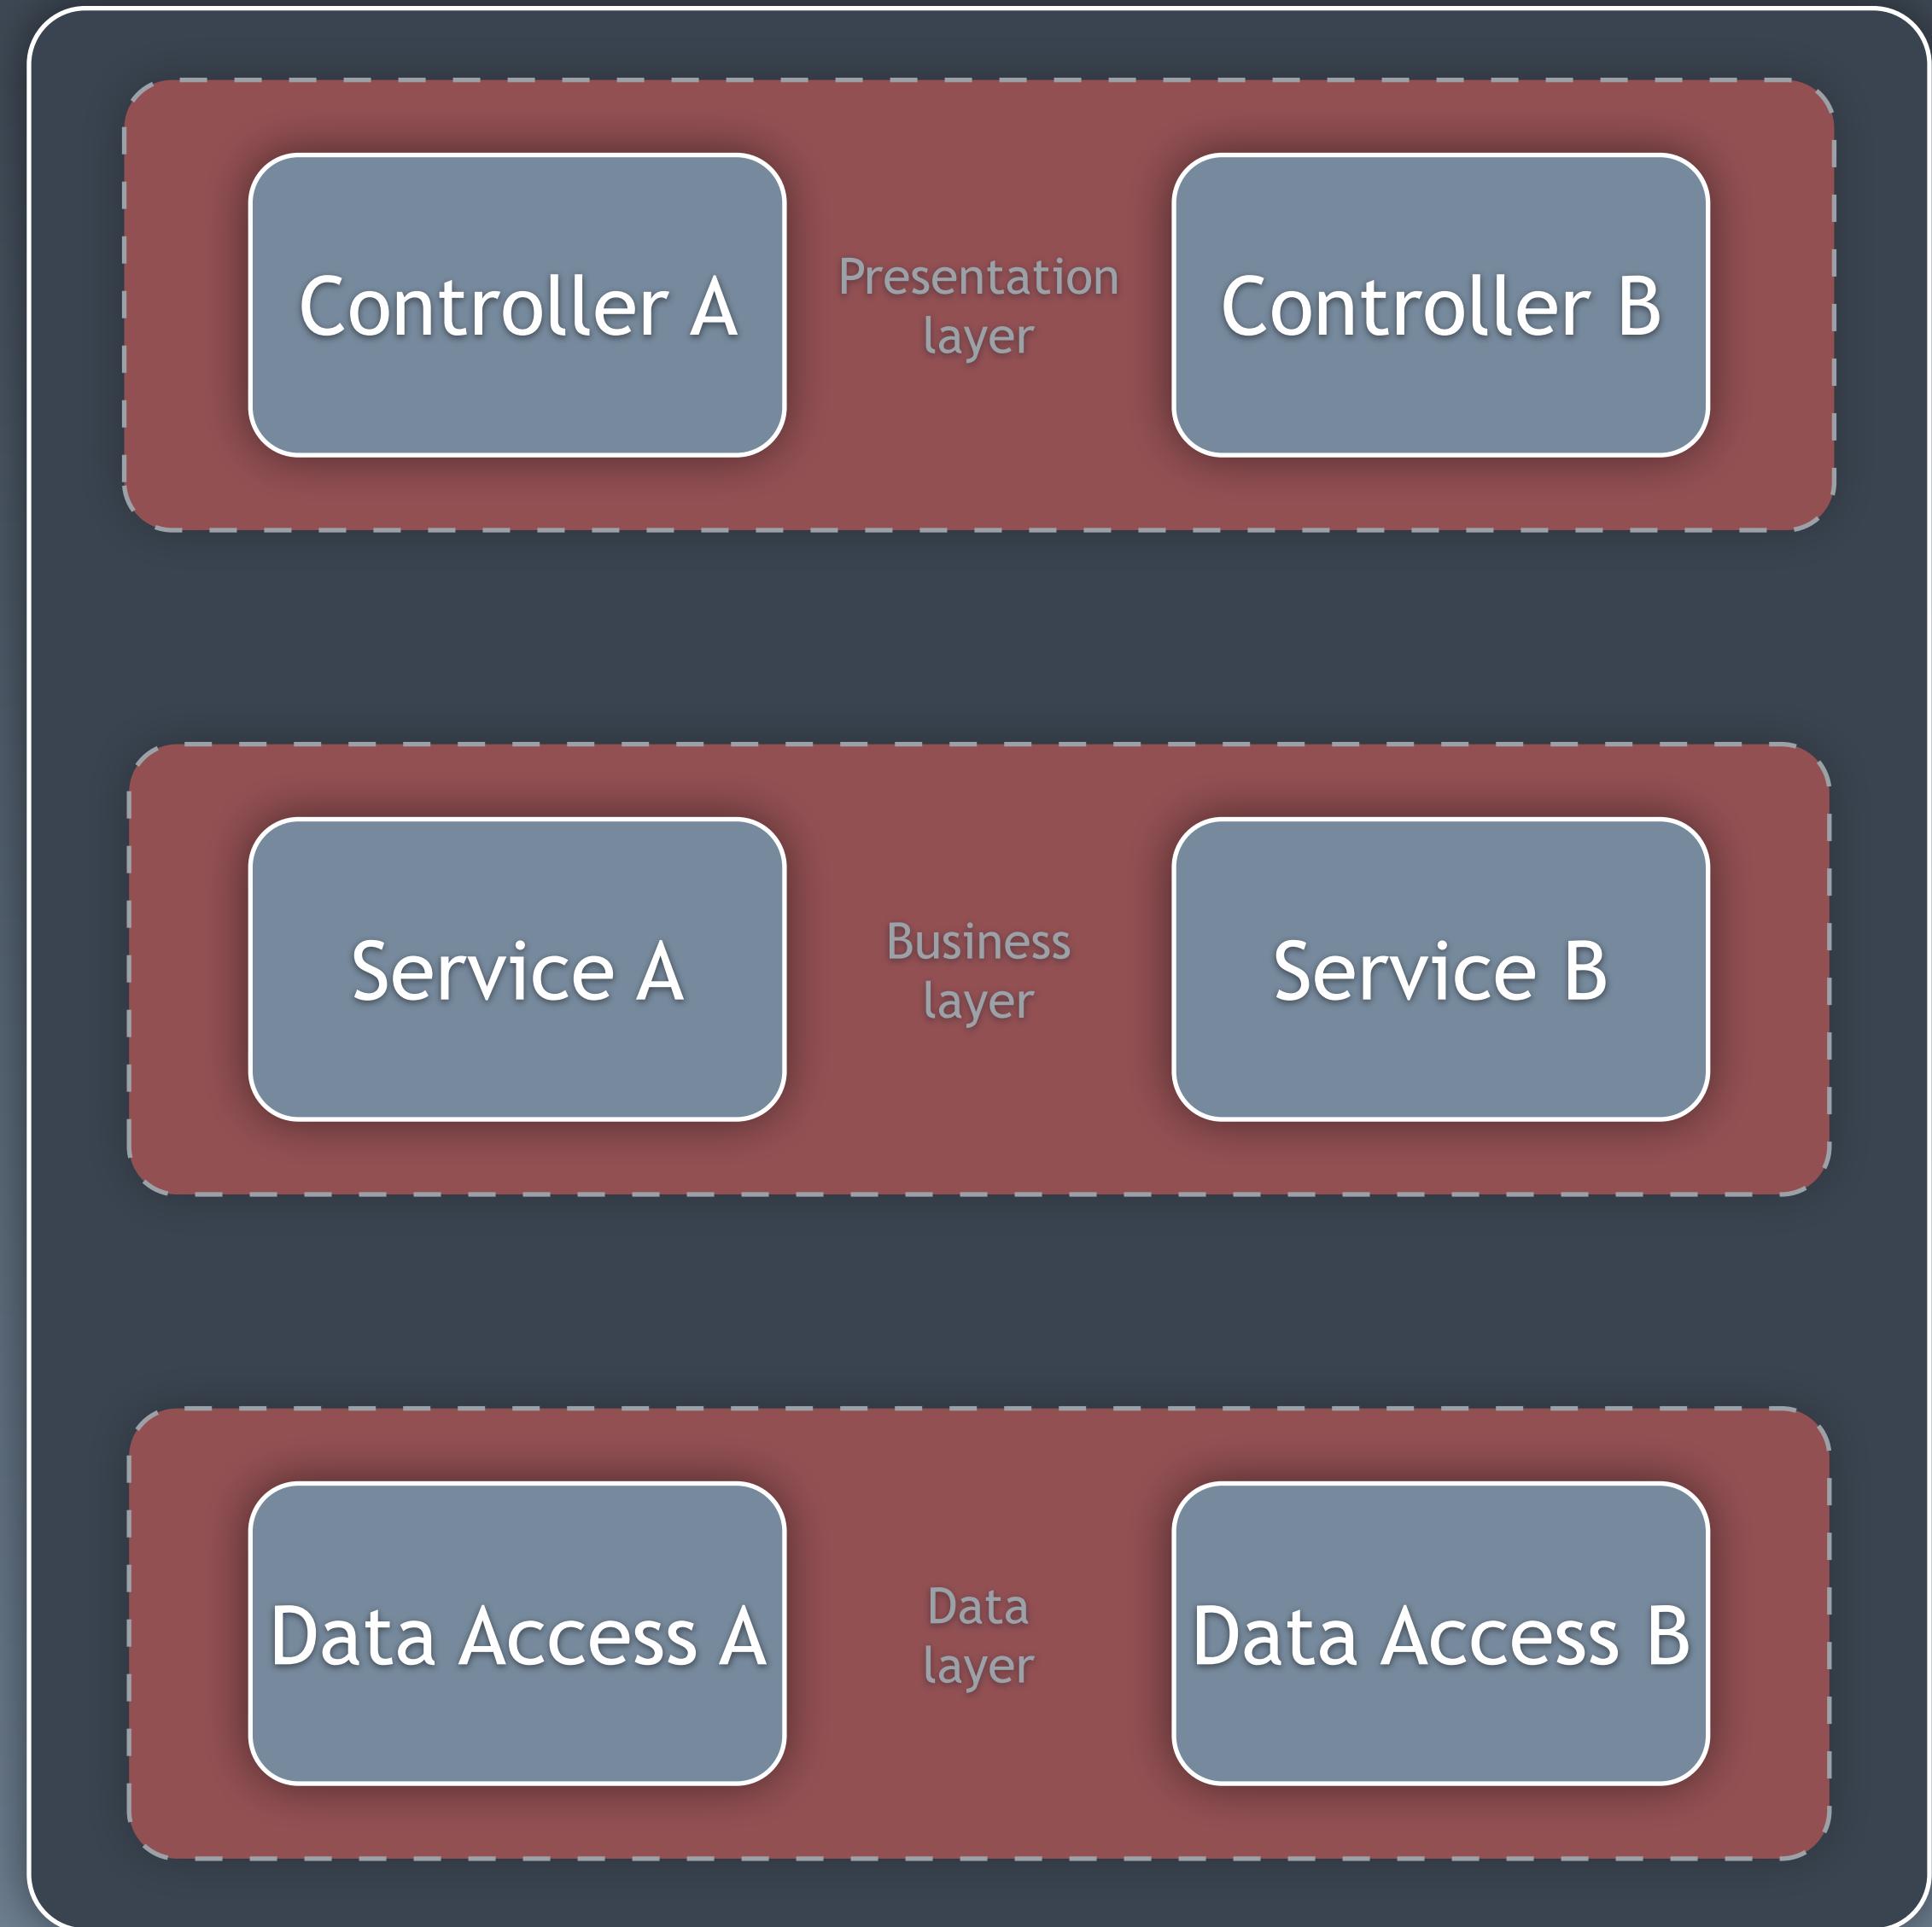
<sup>1</sup>The old English spelling differentiates our use of "reflexion" from the field of reflective computing [Smi84].

# 1 Introduction

Software engineers often think about an existing software system in terms of high-level models. Box and arrow sketches of a system, for instance, are often found on engineers' whiteboards. Although these models are commonly used, reasoning about the system in terms of such models can be dangerous because the models are almost always inaccurate with respect to the system's source.

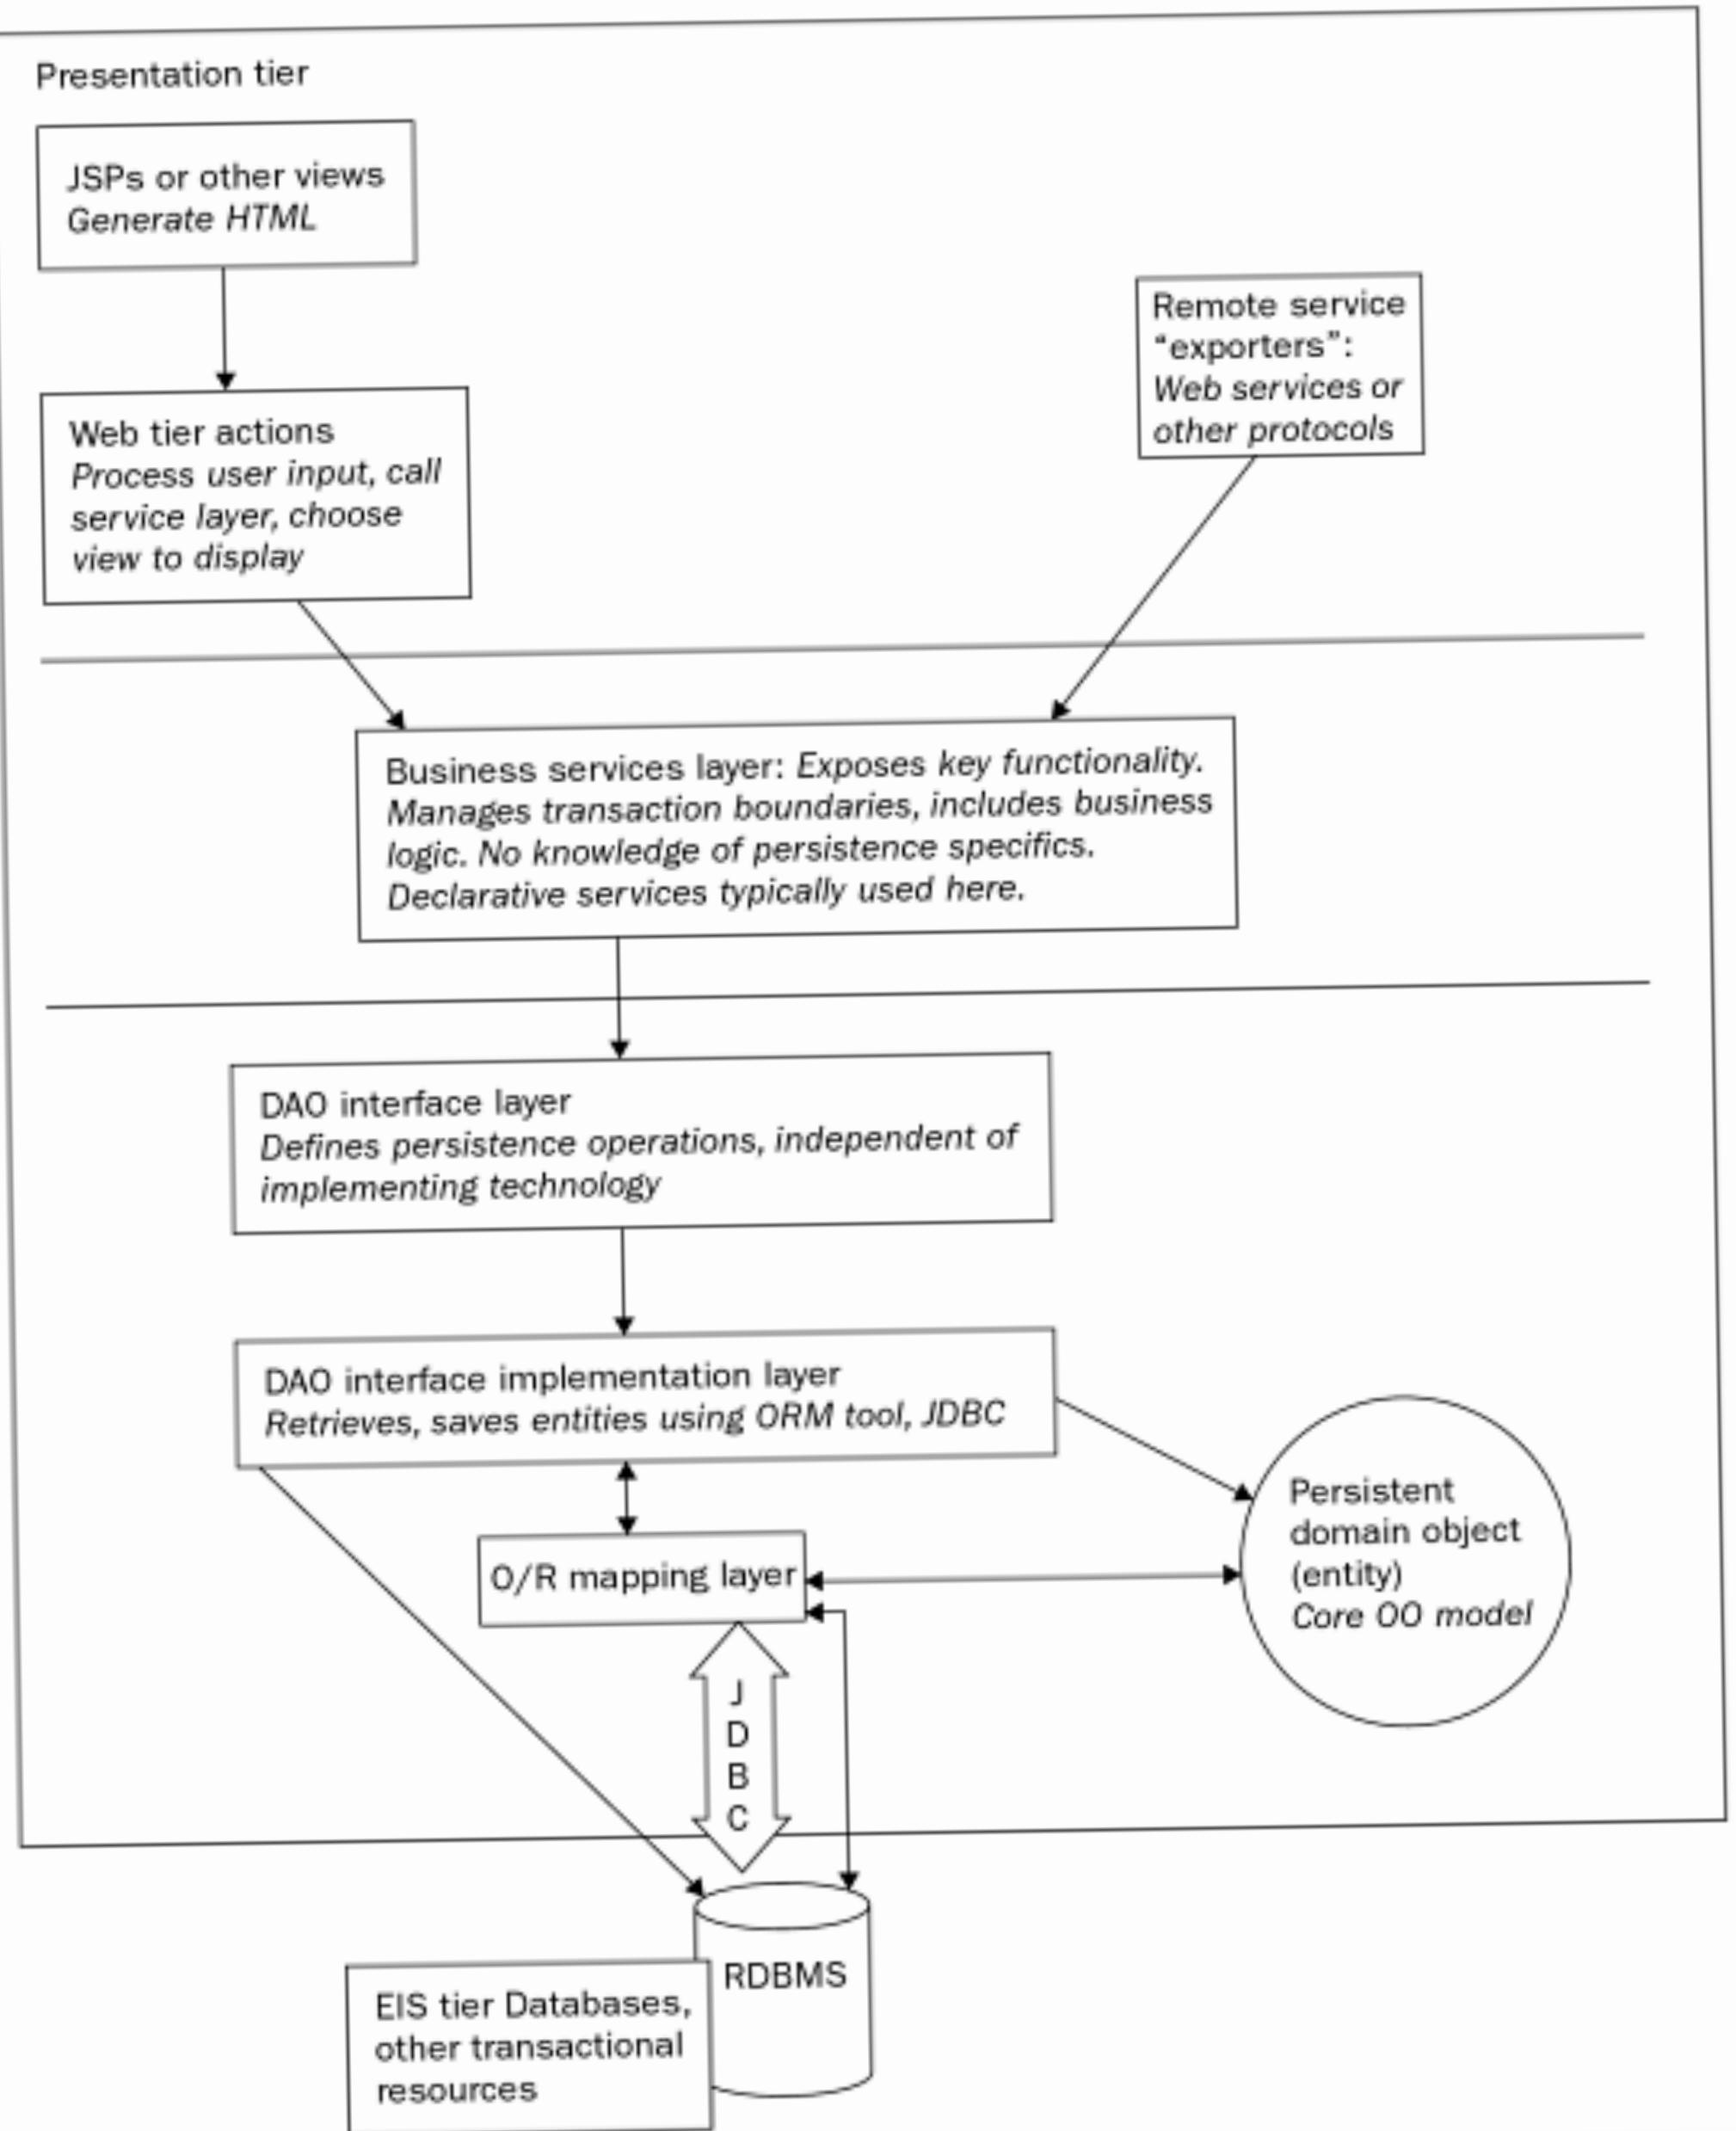
Current reverse engineering systems derive high-level models from the source code. These derived models are useful because they are, by their very nature, accurate representations of the source. Although accurate, the models created by these reverse engineering systems may differ from the models sketched by engineers; an exam-

We often think  
in components  
but write classes



Package by layer (horizontal slicing)

## Chapter 1



## Introducing the Spring Framework

Let's summarize each layer and its responsibilities, beginning closest to the database or other enterprise resources:

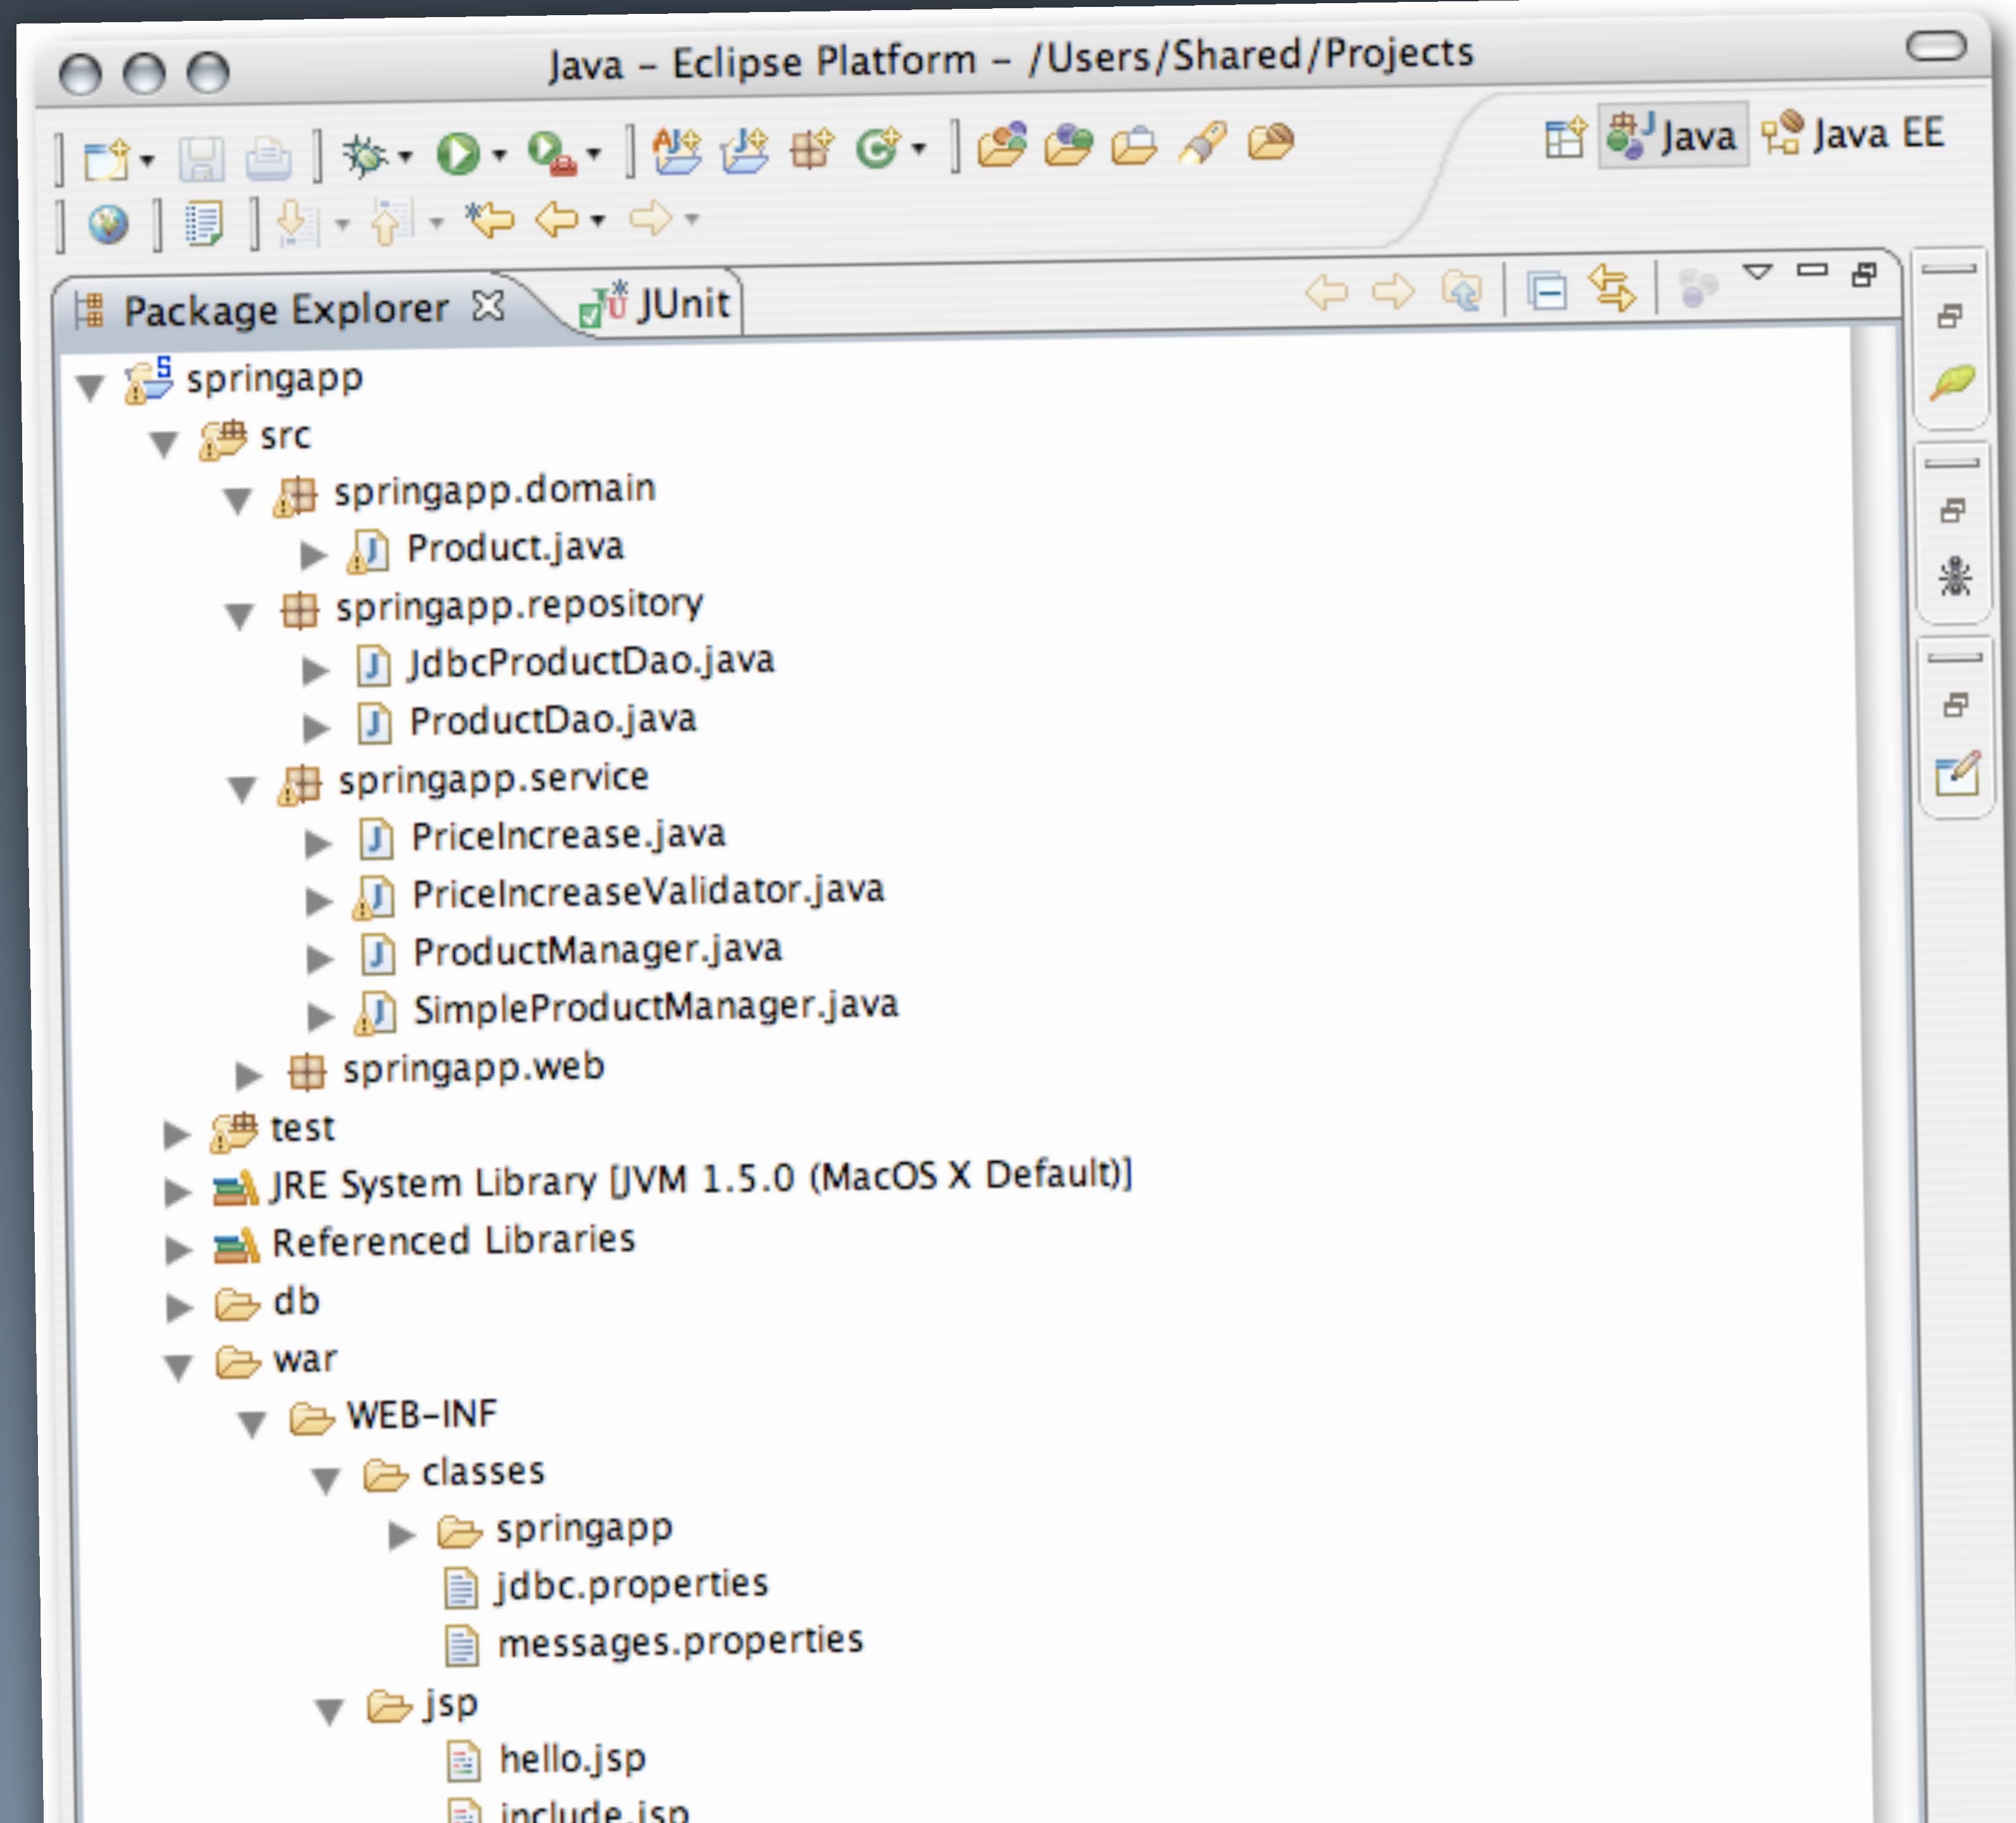
- ❑ **Presentation layer:** This is most likely to be a web tier. This layer should be as thin as possible. It should be possible to have alternative presentation layers—such as a web tier or remote web services facade—on a single, well-designed middle tier.
- ❑ **Business services layer:** This is responsible for transactional boundaries and providing an entry point for operations on the system as a whole. This layer should have no knowledge of presentation concerns, and should be reusable.
- ❑ **DAO interface layer:** This is a layer of interfaces *independent of any data access technology* that is used to find and persist persistent objects. This layer effectively consists of *Strategy* interfaces for the Business services layer. This layer should not contain business logic. Implementations of these interfaces will normally use an O/R mapping technology or Spring's JDBC abstraction.
- ❑ **Persistent domain objects:** These model real objects or concepts such as a bank account.
- ❑ **Databases and legacy systems:** By far the most common case is a single RDBMS. However, there may be multiple databases, or a mix of databases and other transactional or non-transactional legacy systems or other enterprise resources. The same fundamental architecture is applicable in either case. This is often referred to as the *EIS (Enterprise Information System)* tier.

In a J2EE application, all layers except the EIS tier will run in the application server or web container. Domain objects will typically be passed up to the presentation layer, which will display data they contain, *but not modify them*, which will occur only within the transactional boundaries defined by the business services layer. Thus there is no need for distinct Transfer Objects, as used in traditional J2EE architecture.

In the following sections we'll discuss each of these layers in turn, beginning closest to the database.

Spring aims to decouple architectural layers, so that each layer can be modified as

**Spring aims to decouple architectural layers, so that each layer can be modified as far as possible without impacting other layers. No layer is aware of the concerns of the layer above; as far as possible, dependency is purely on the layer immediately below. Dependency between layers is normally in the form of interfaces, ensuring that coupling is as loose as possible.**



java - Standard project/package structure of a j2ee web application - Stack Overflow  
stackoverflow.com/questions/5878774/standard-project-package-structure-of-a-j2ee-web-application

## 1 Answer

active   oldest   votes

If you are using maven, it's best to follow the standard maven project layout. You can get maven to generate this structure for you by doing,

**8**

```
mvn archetype:generate
```

and select spring-mvc-jpa-archetype from the list of choices

This will give you a package structure like,

```
pom.xml
src
  main
    java
      mygroup
        controller
          HomeController.java
          PersonController.java
        dao
          PersonDao.java
        model
          Person.java
    resources
      db.properties
      log4j.xml
      META-INF
        persistence.xml
  webapp
    index.html
    META-INF
      context.xml
      MANIFEST.MF
    resources
      css
        screen.css
    WEB-INF
      spring
        app
          controllers.xml
          servlet-context.xml
          db.xml
          root-context.xml
      views
```

Related

- 2 How do I use a custom authentication mechanism for a Java web application with Spring Security?
- 0 Is there a sample web project of integrating Spring-MVC and Spring-Data-JPA?
- 1 Reusing DAOs in another web application
- 0 Java Web Application Warming
- 3 Is a parallel Spring-MVC application possible with a non-spring web app?
- 2 Adding a web interface (Spring MVC) to existing Java application
- 0 package strucuture for Spring MVC project with multiple sub projects using maven
- 0 How to add a setup

Technical Operations Engineer, Devops, Cloud Systems and AWS  
Voyanta  
London, United Kingdom

Linux Service Engineers  
Shazam  
London, United Kingdom

3 x Senior Engineer (Ruby) 3 month+ contract  
We are Friday Limited  
London, United Kingdom

More jobs near London...

MvcMovie - Microsoft Visual Studio Express... Quick Launch (Ctrl+Q)

FILE EDIT VIEW PROJECT BUILD DEBUG TEAM TOOLS TEST WINDOW HELP

Internet Explorer Debug

HomeController.cs

MvcMovie.Controllers.HomeCont Index()

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Web;
using System.Web.Mvc;

namespace MvcMovie.Controllers
{
    public class HomeController : Controller
    {
        public ActionResult Index()
        {
            ViewBag.Message = "Modify this template";

            return View();
        }

        public ActionResult About()
        {
            ViewBag.Message = "Your app description here.";

            return View();
        }
    }
}
```

Solution Explorer

Search Solution Explorer (Ctrl+Shift+F)

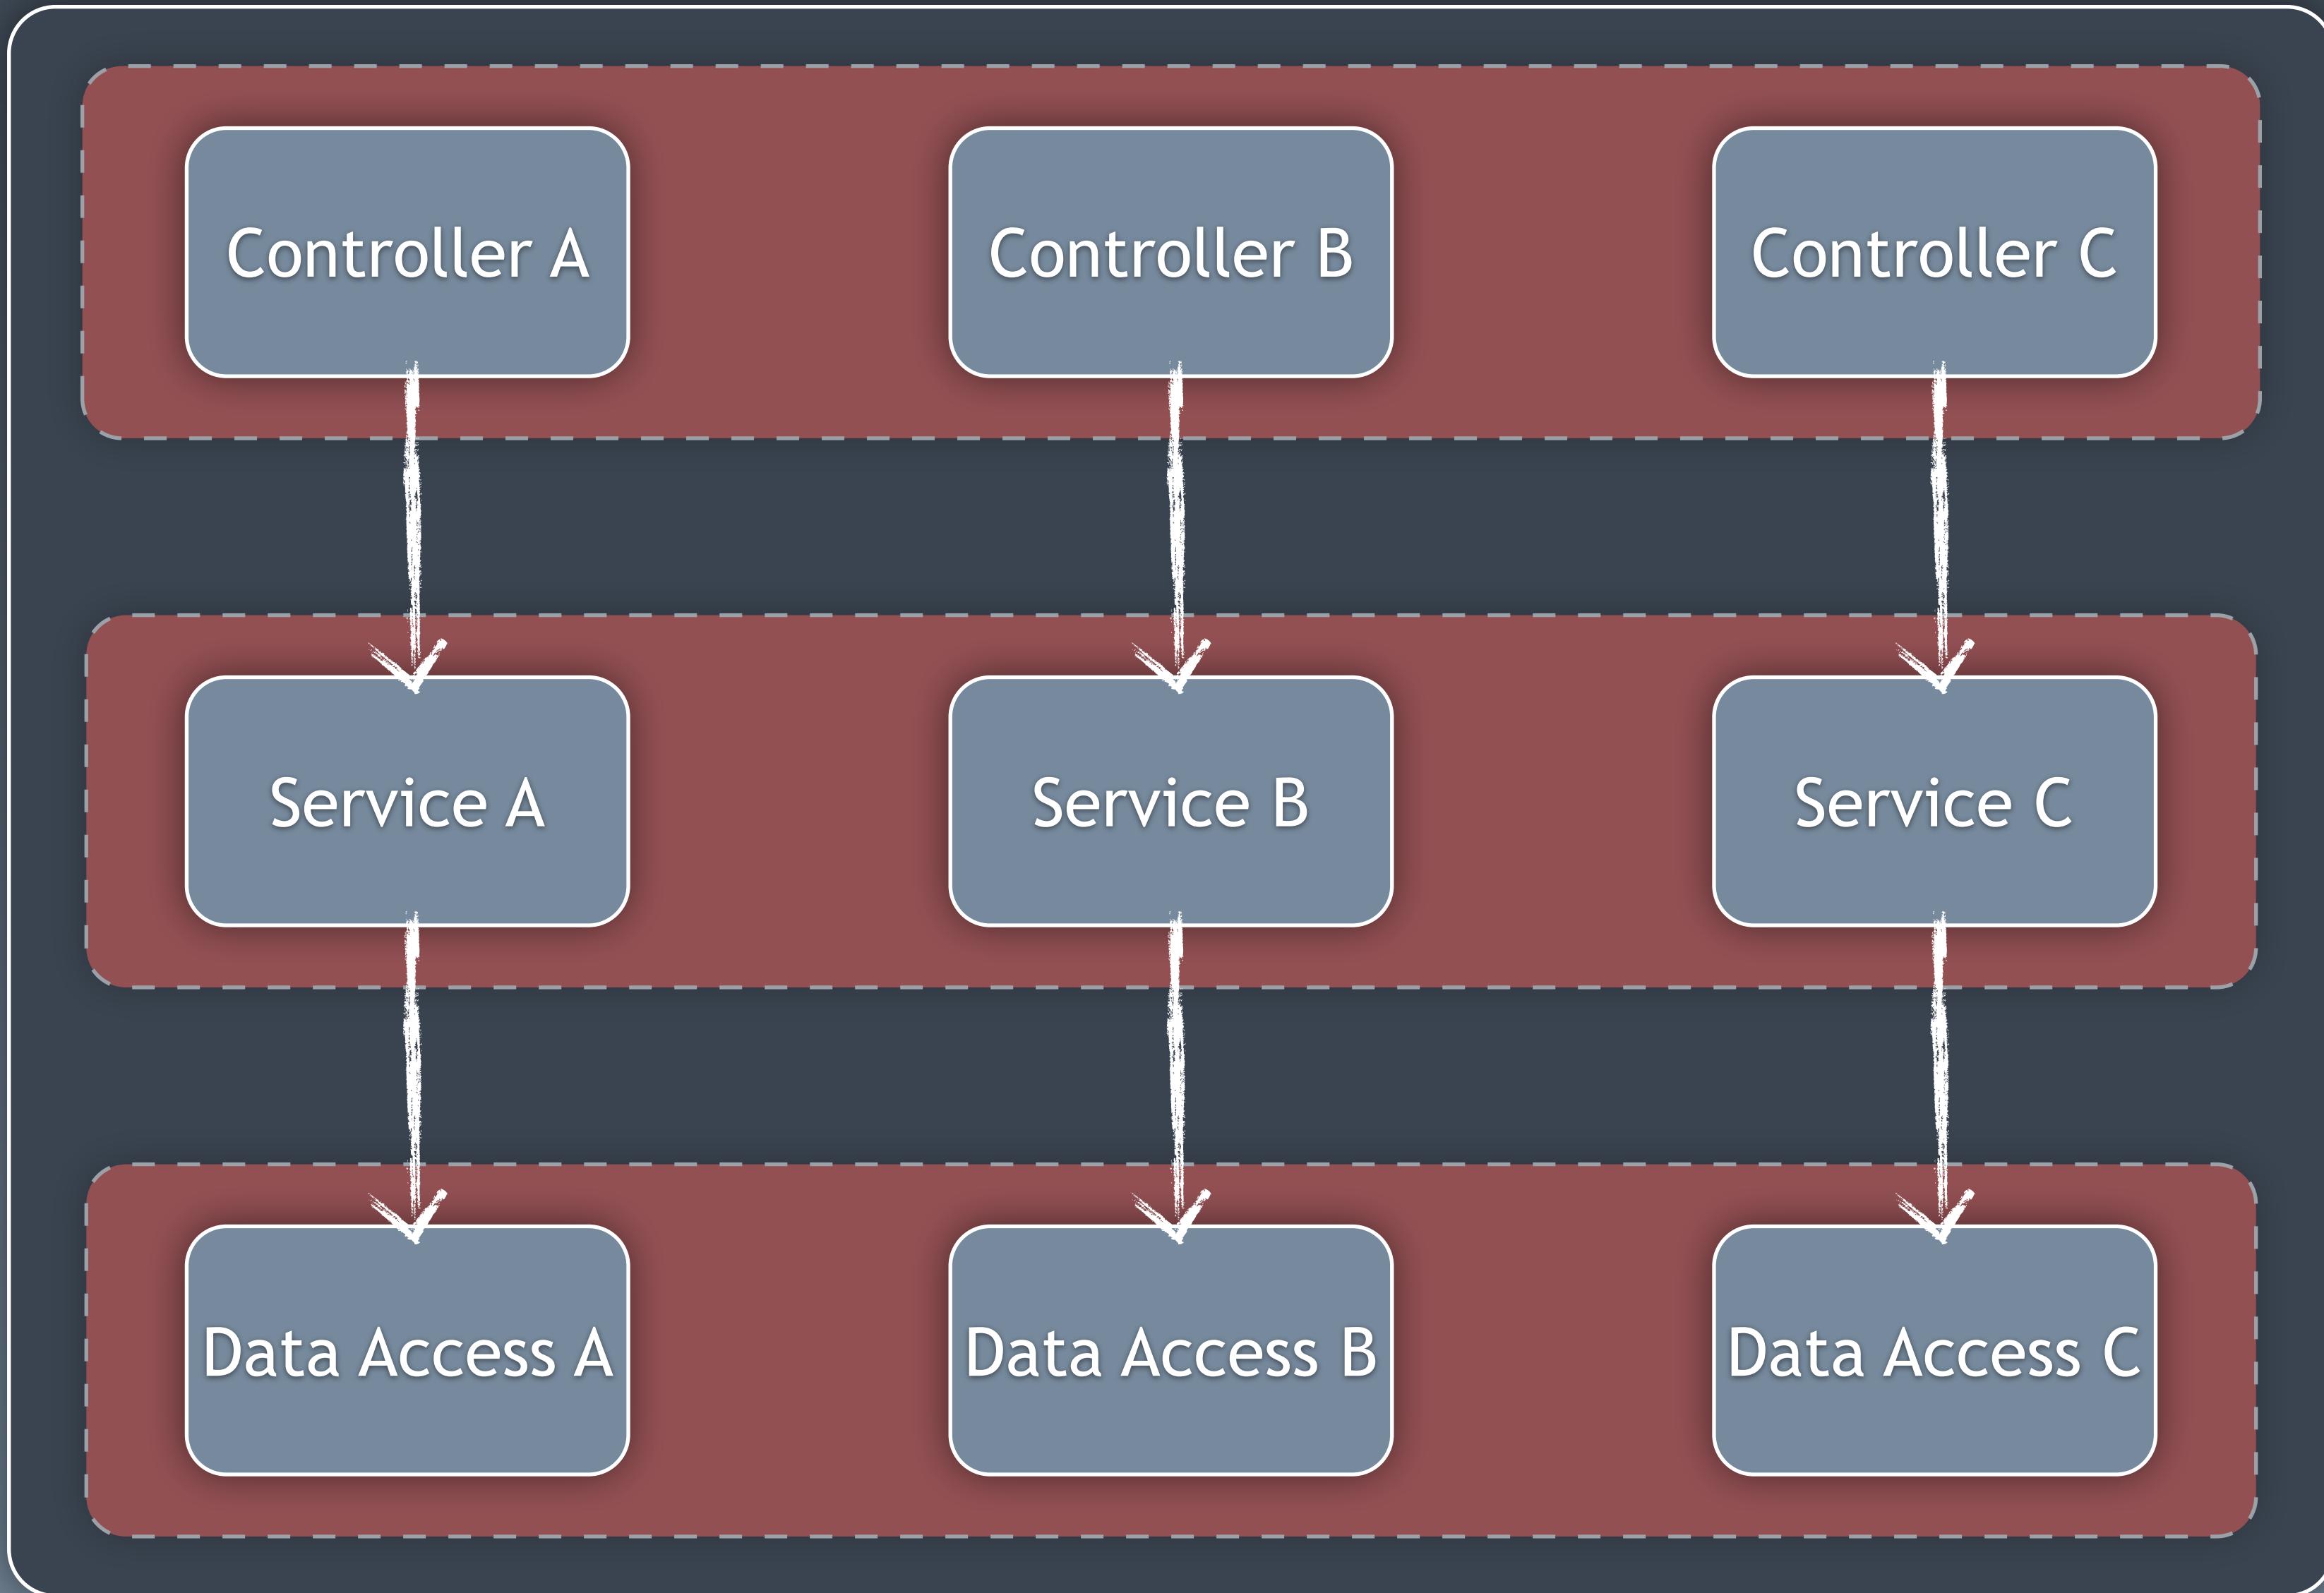
Solution 'MvcMovie' (1 project)

MvcMovie

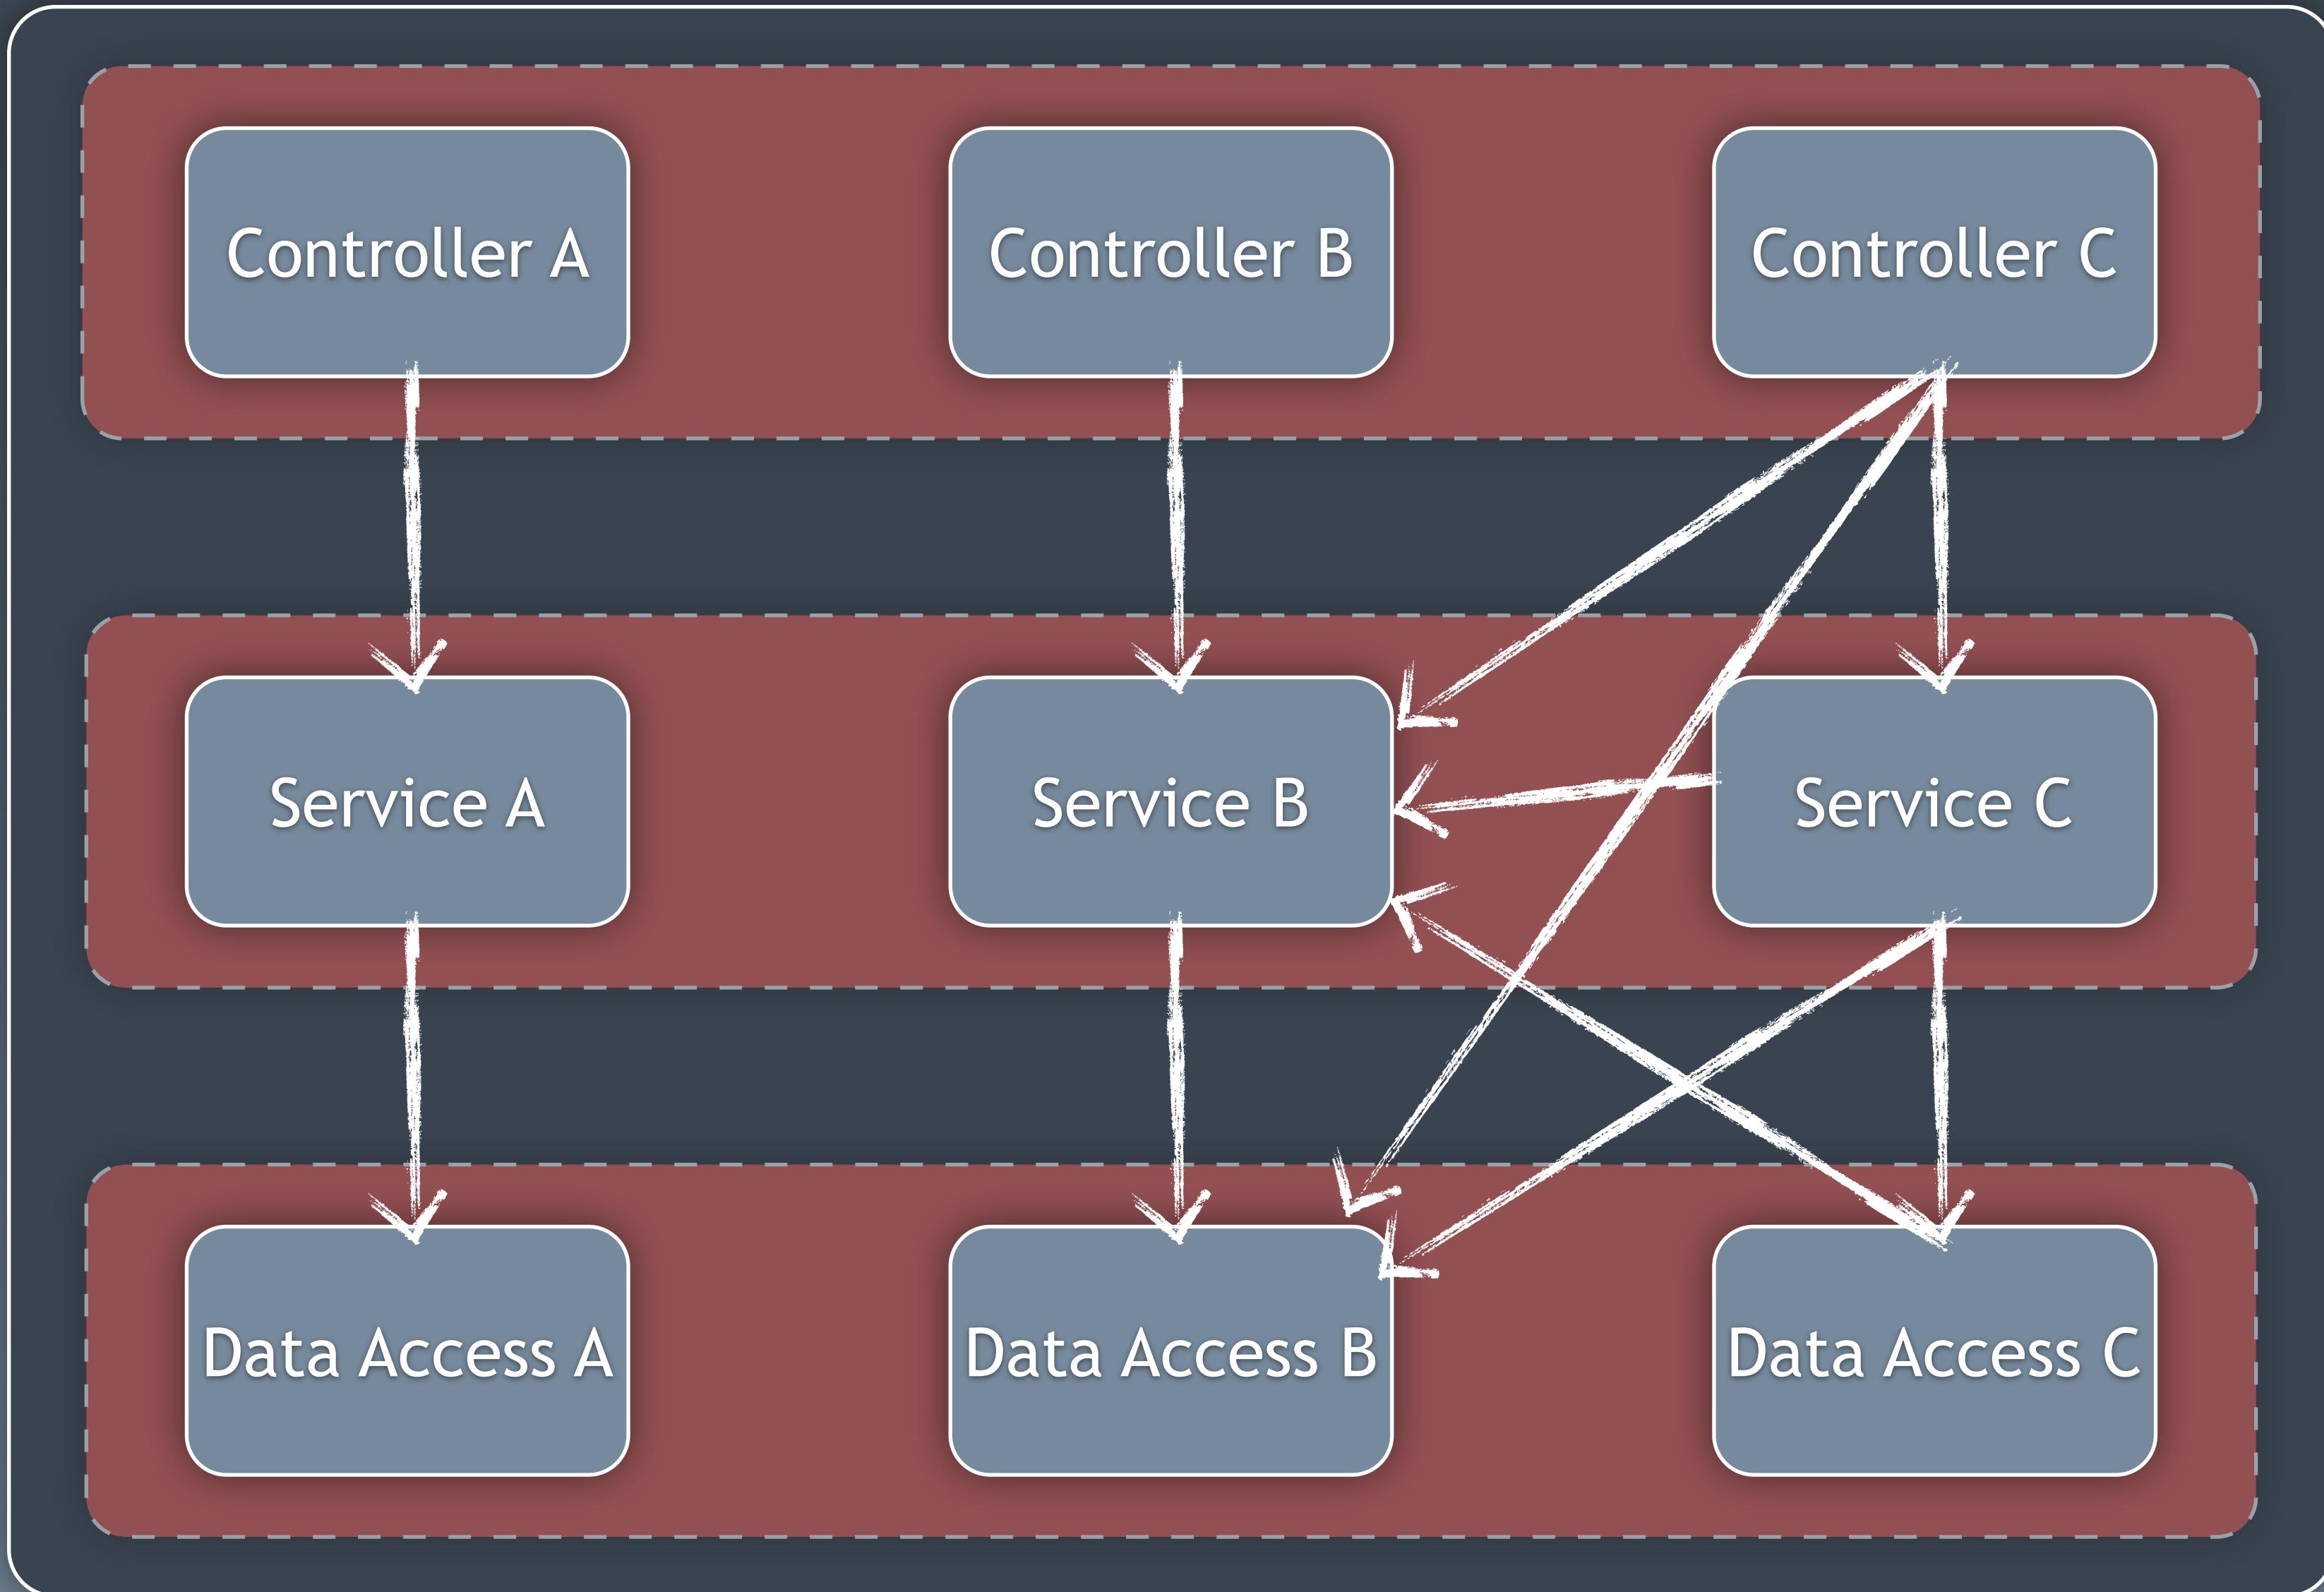
- Properties
- References
- App\_Data
- App\_Start
- Content
- Controllers
- Filters
- Images
- Models
- Scripts
- Views
- favicon.ico
- Global.asax
- packages.config
- Web.config

100 %

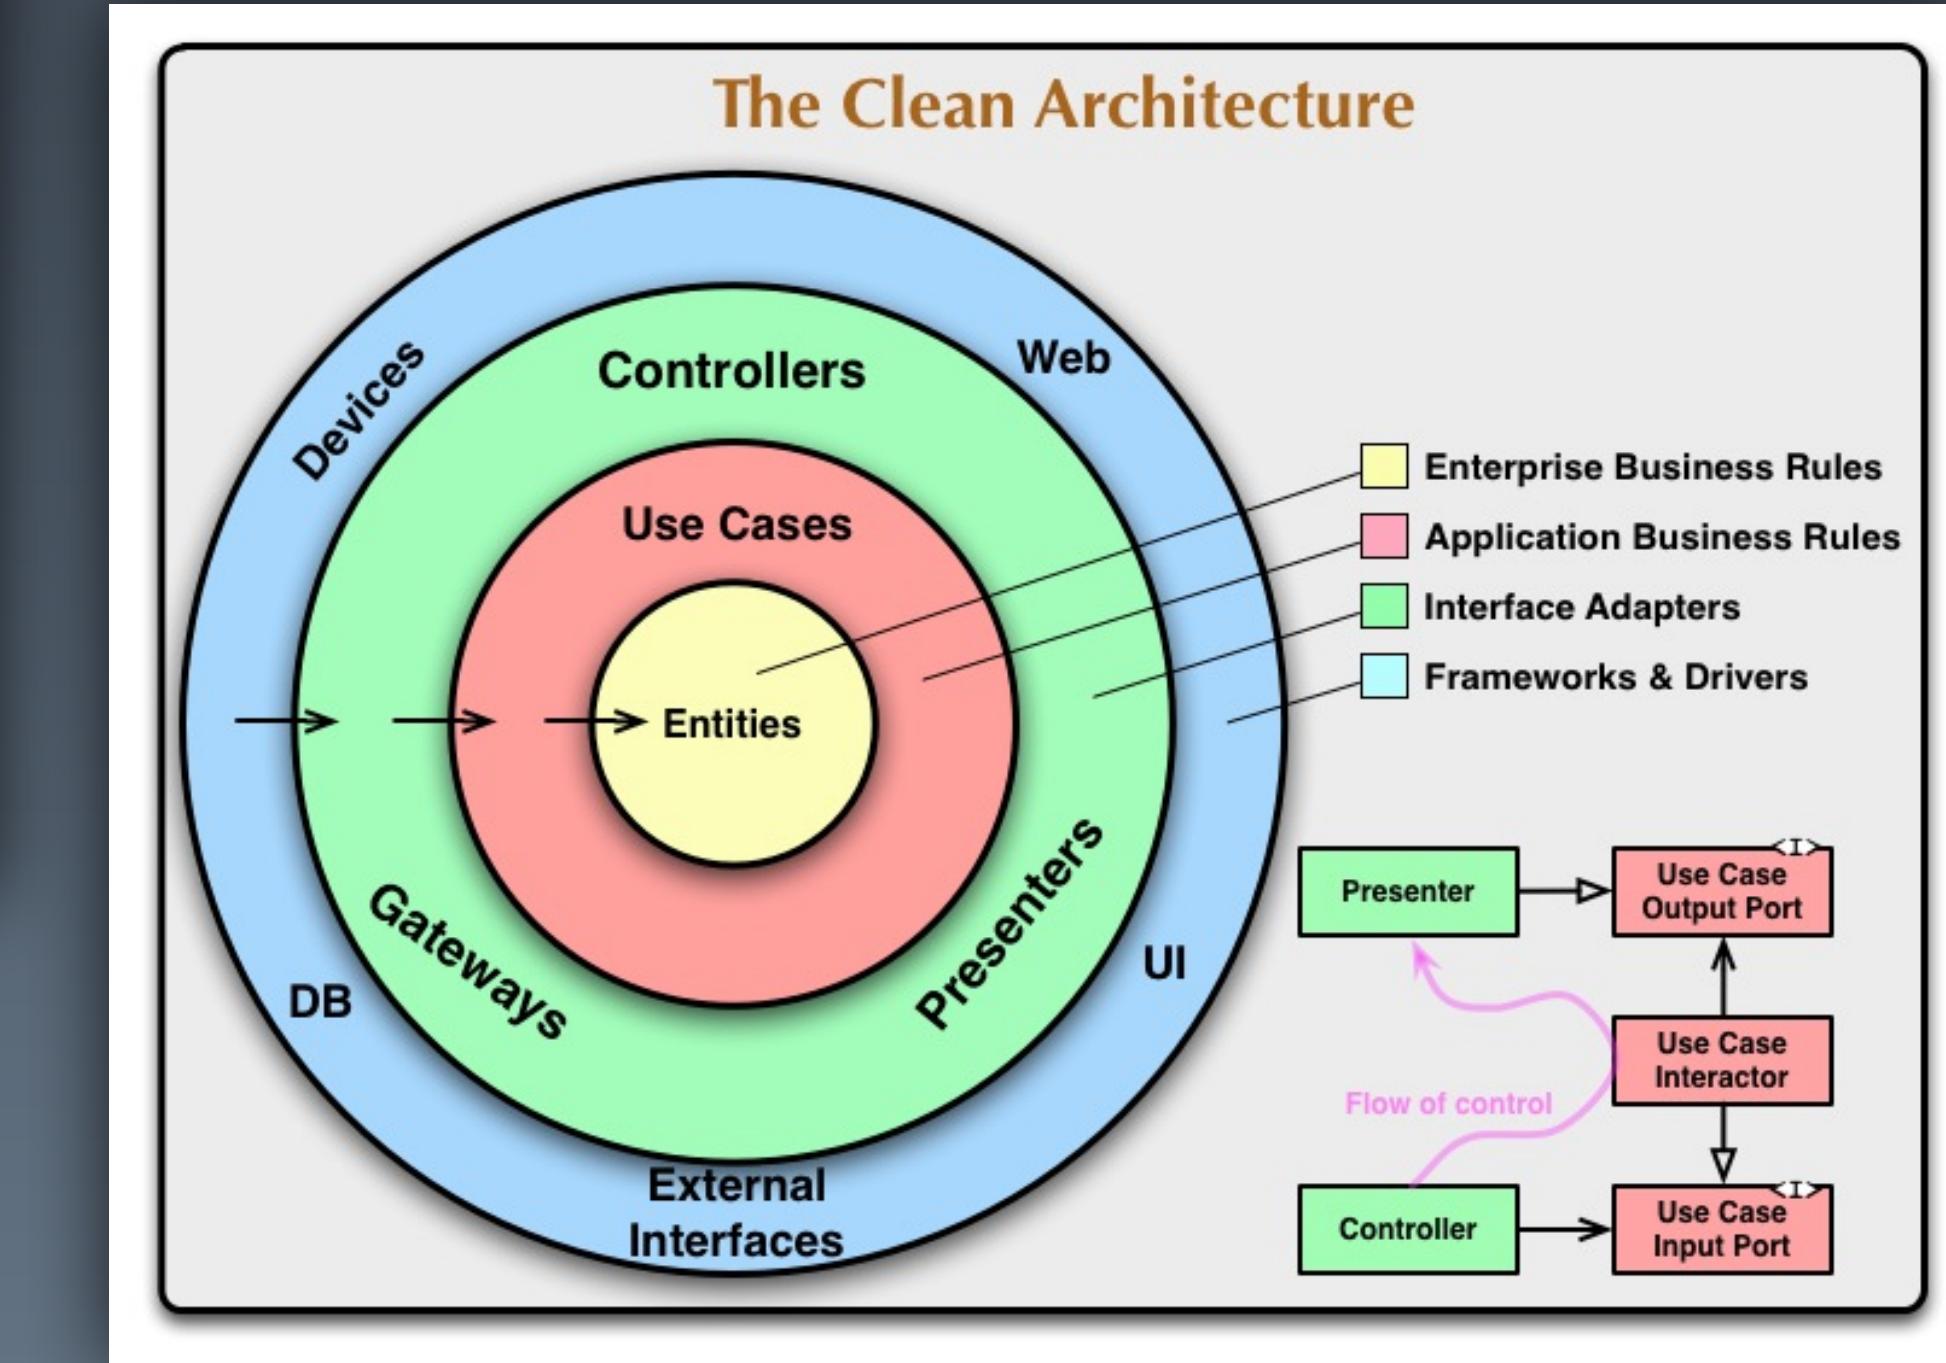
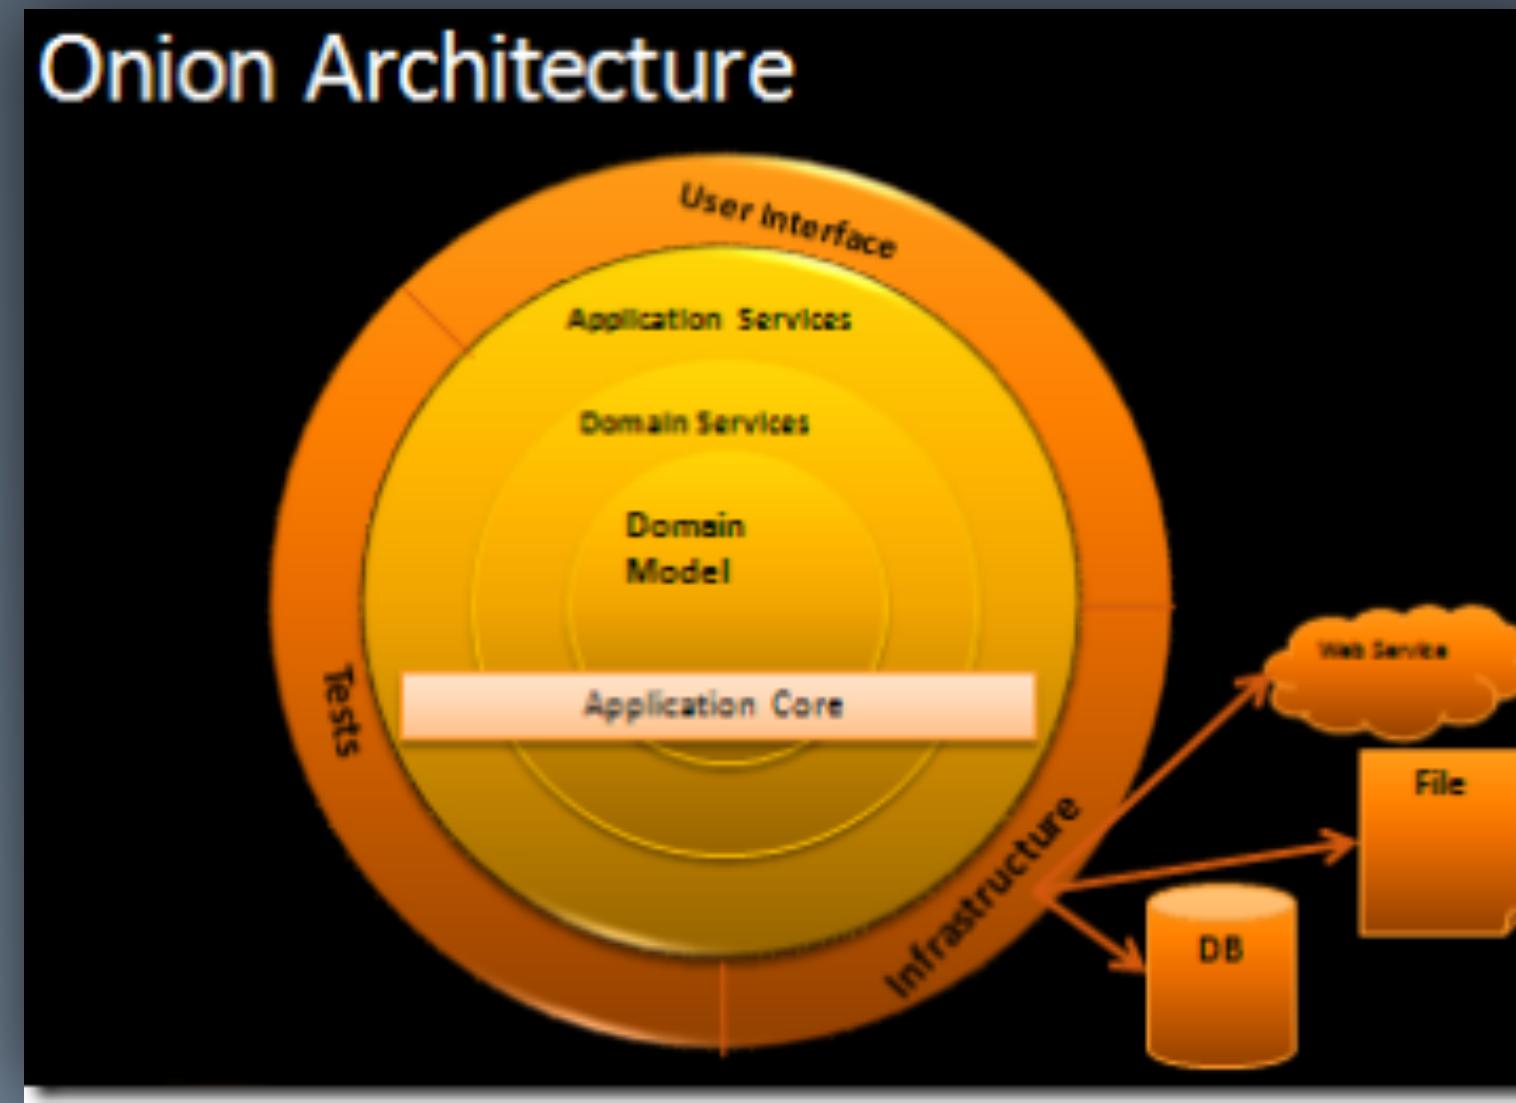
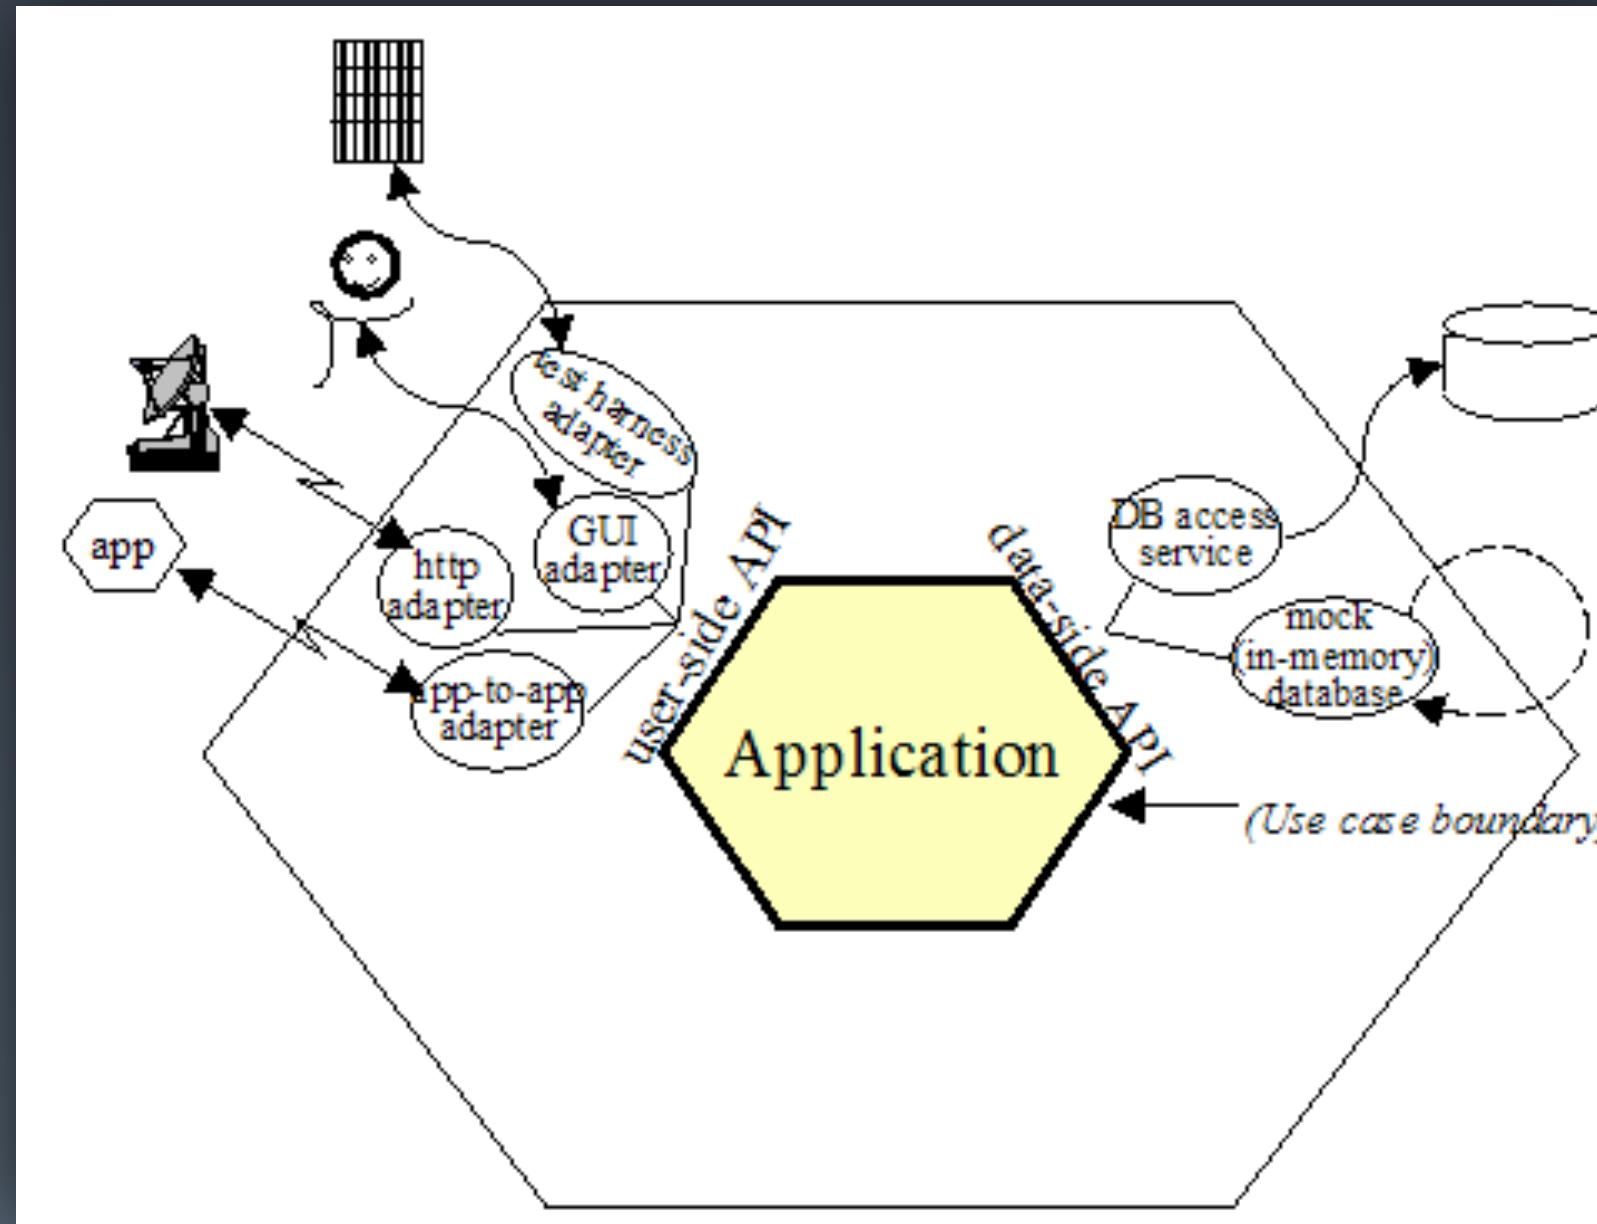
Ready Ln 1 Col 1 Ch 1 INS



Package by layer (horizontal slicing)

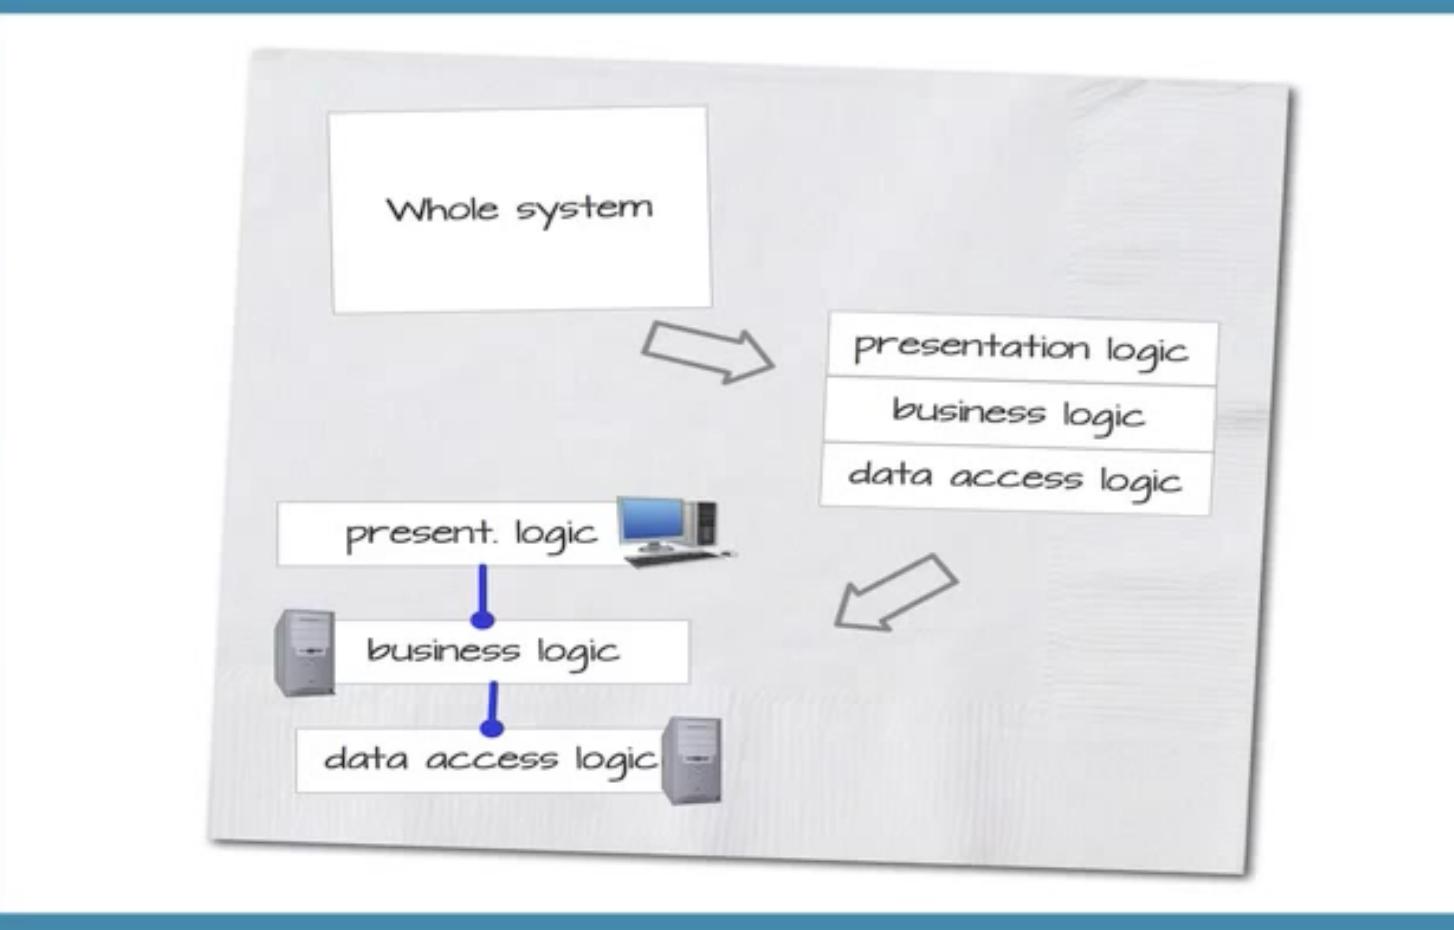


Package by layer (horizontal slicing)



Hexagons, onions,  
ports & adapters, etc

Should layers be  
considered  
harmful?



SOFTWARE  
ARCHITECT  
14 – 17 OCTOBER 2014 | HOTEL RUSSELL, LONDON

Let software design come to life using software cells -  
Ralf Westphal



DevWeek Events



403

230 views

+ Add to

Share

More



4



0

OSI layers  
vs  
software  
layers  
(~13 minutes)

Cargo cult programming can also refer to the results of applying a design pattern or coding style blindly without understanding the reasons behind that design principle.

[https://en.wikipedia.org/wiki/Cargo\\_cult\\_programming](https://en.wikipedia.org/wiki/Cargo_cult_programming)

Is architectural  
layering in  
software another  
cargo cult?

Are layers significant  
structural  
elements  
or just an  
implementation  
detail?



## LINKS

[Blogroll](#)[Subscribe via RSS](#)

## AUTHORS

[Uncle Bob](#)[Paul Pagel](#)[Micah Martin](#)[Eric Smith](#)[Doug Bradbury](#)[Colin Jones](#)[Mike Jansen](#)[Jim Suchy](#)[Craig Demyanovich](#)[Dariusz Pasciak](#)[Kevin Buchanan](#)[Eric Meyer](#)[Myles Megyesi](#)[Dave Moore](#)[Chris Peak](#)[Patrick Gombert](#)[Margaret Pagel](#)[Steve Kim](#)[Sandro Padin](#)[Kevin Liddle](#)[Malcolm Newsome](#)

# Screaming Architecture

[Uncle Bob](#) → 30 Sep 2011 + [Architecture](#)[Share](#)[Tweet](#)[G+ Share](#)

Imagine that you are looking at the blueprints of a building. This document, prepared by an architect, tells you the plans for the building. What do these plans tell you?

If the plans you are looking at are for a single family residence, then you'll likely see a front entrance, a foyer leading to a living room and perhaps a dining room. There'll likely be a kitchen a short distance away, close to the dining room. Perhaps a dinette area next to the kitchen, and probably a family room close to that. As you looked at those plans, there'd be no question that you were looking at a *house*. The architecture would *scream: house*.

Or if you were looking at the architecture of a library, you'd likely see a grand entrance, an area for check-in-out clerks, reading areas, small conference rooms, and gallery after gallery capable of holding bookshelves for all the books in the library. That architecture would *scream: Library*.

So what does the architecture of your application scream? When you look at the top level directory structure, and the source files in the highest level package; do they scream: **Health Care System**, or **Accounting System**, or **Inventory Management System**? Or do they scream: **Rails**, or **Spring/Hibernate**, or **ASP**?

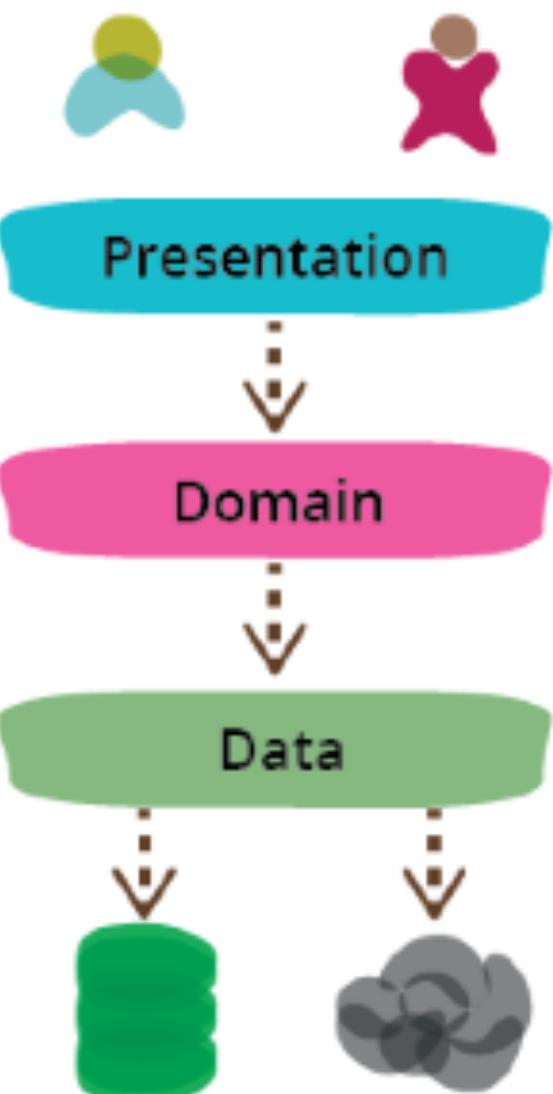
**The Theme of an Architecture**

## PresentationDomainDataLayering



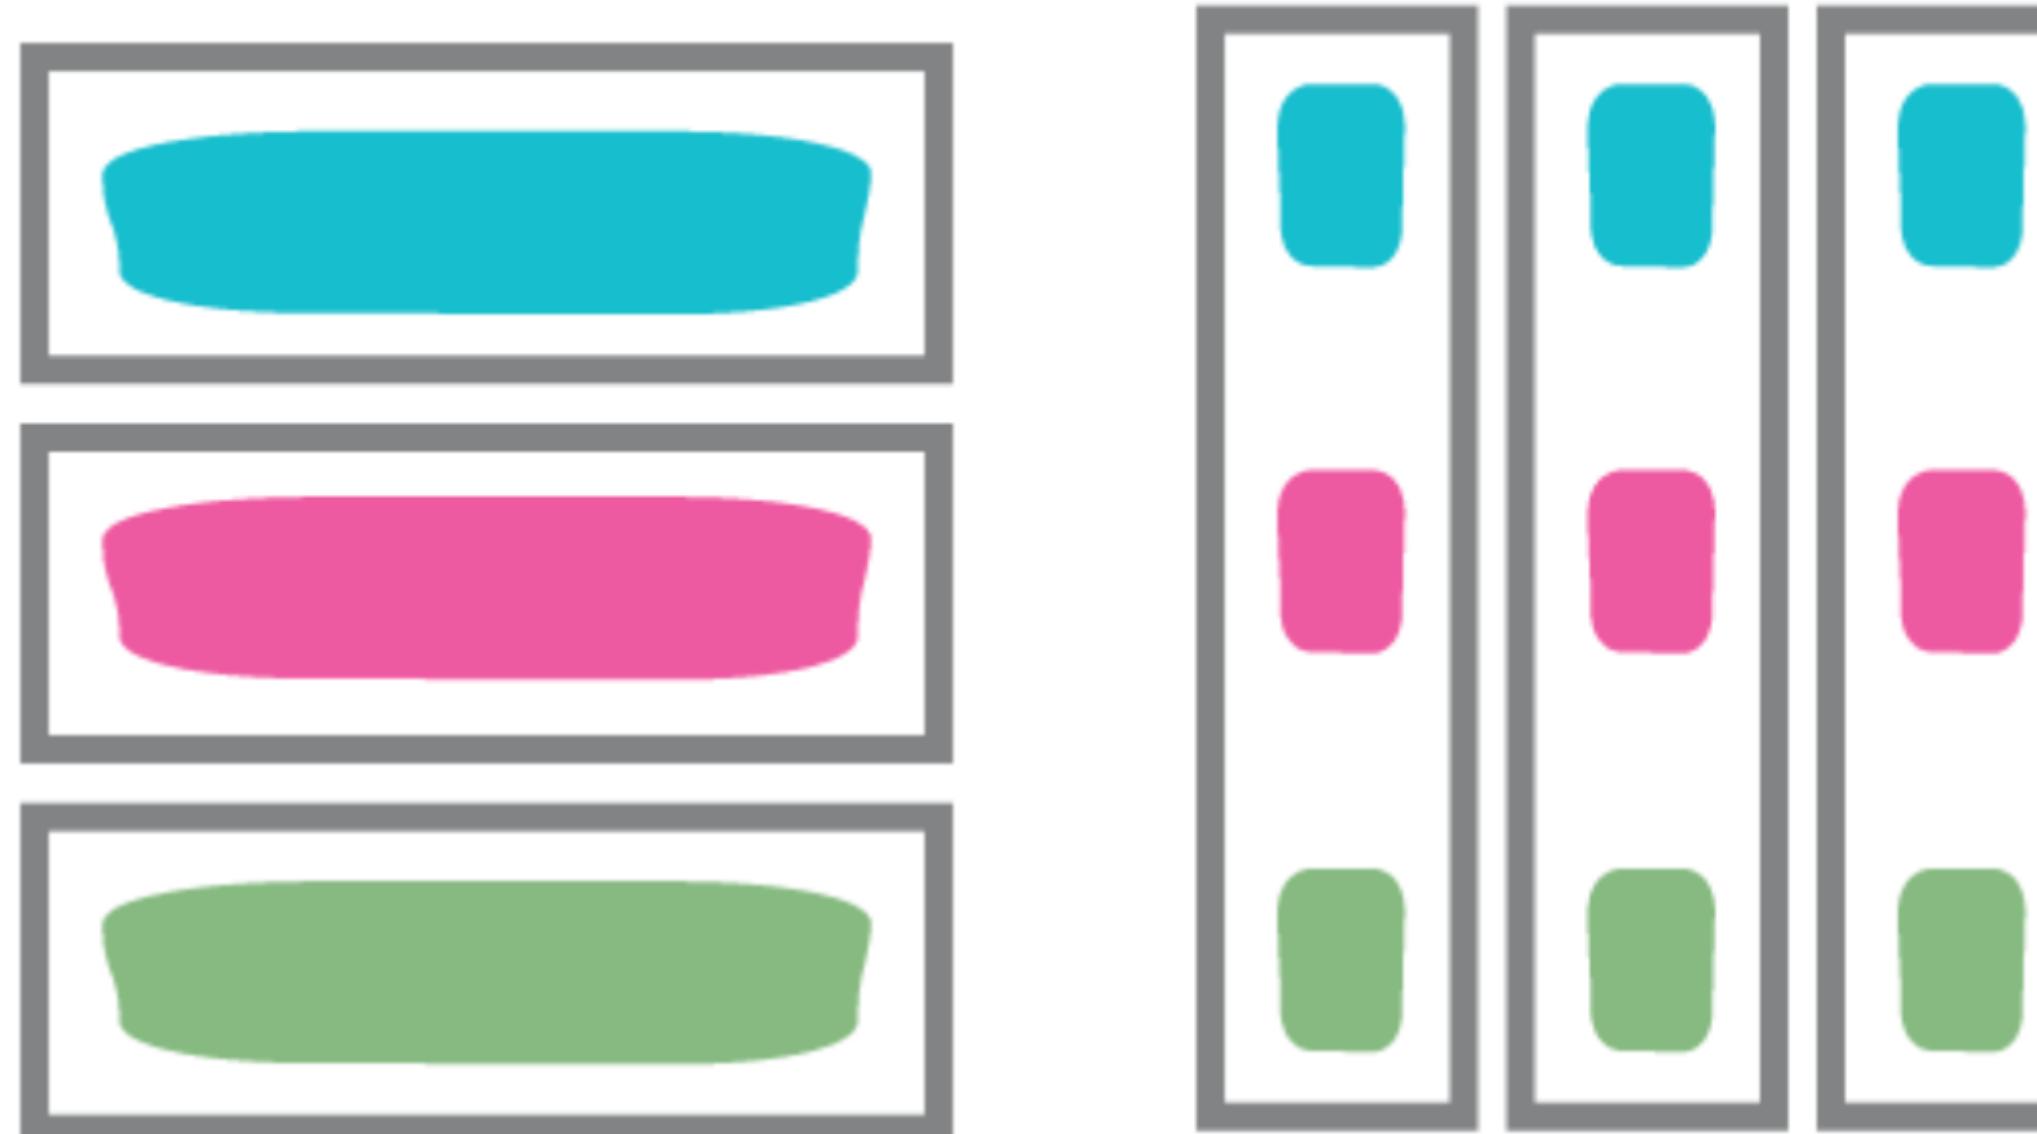
Martin Fowler  
26 August 2015

One of the most common ways to modularize an information-rich application is to separate it into three broad layers: presentation (UI), domain logic, and data access. So you often see web applications split into a web layer that knows about handling http requests and rendering HTML, a business logic layer that contains validations and calculations, and a data access layer that sorts out how to manage persistent data from a database or remote services.



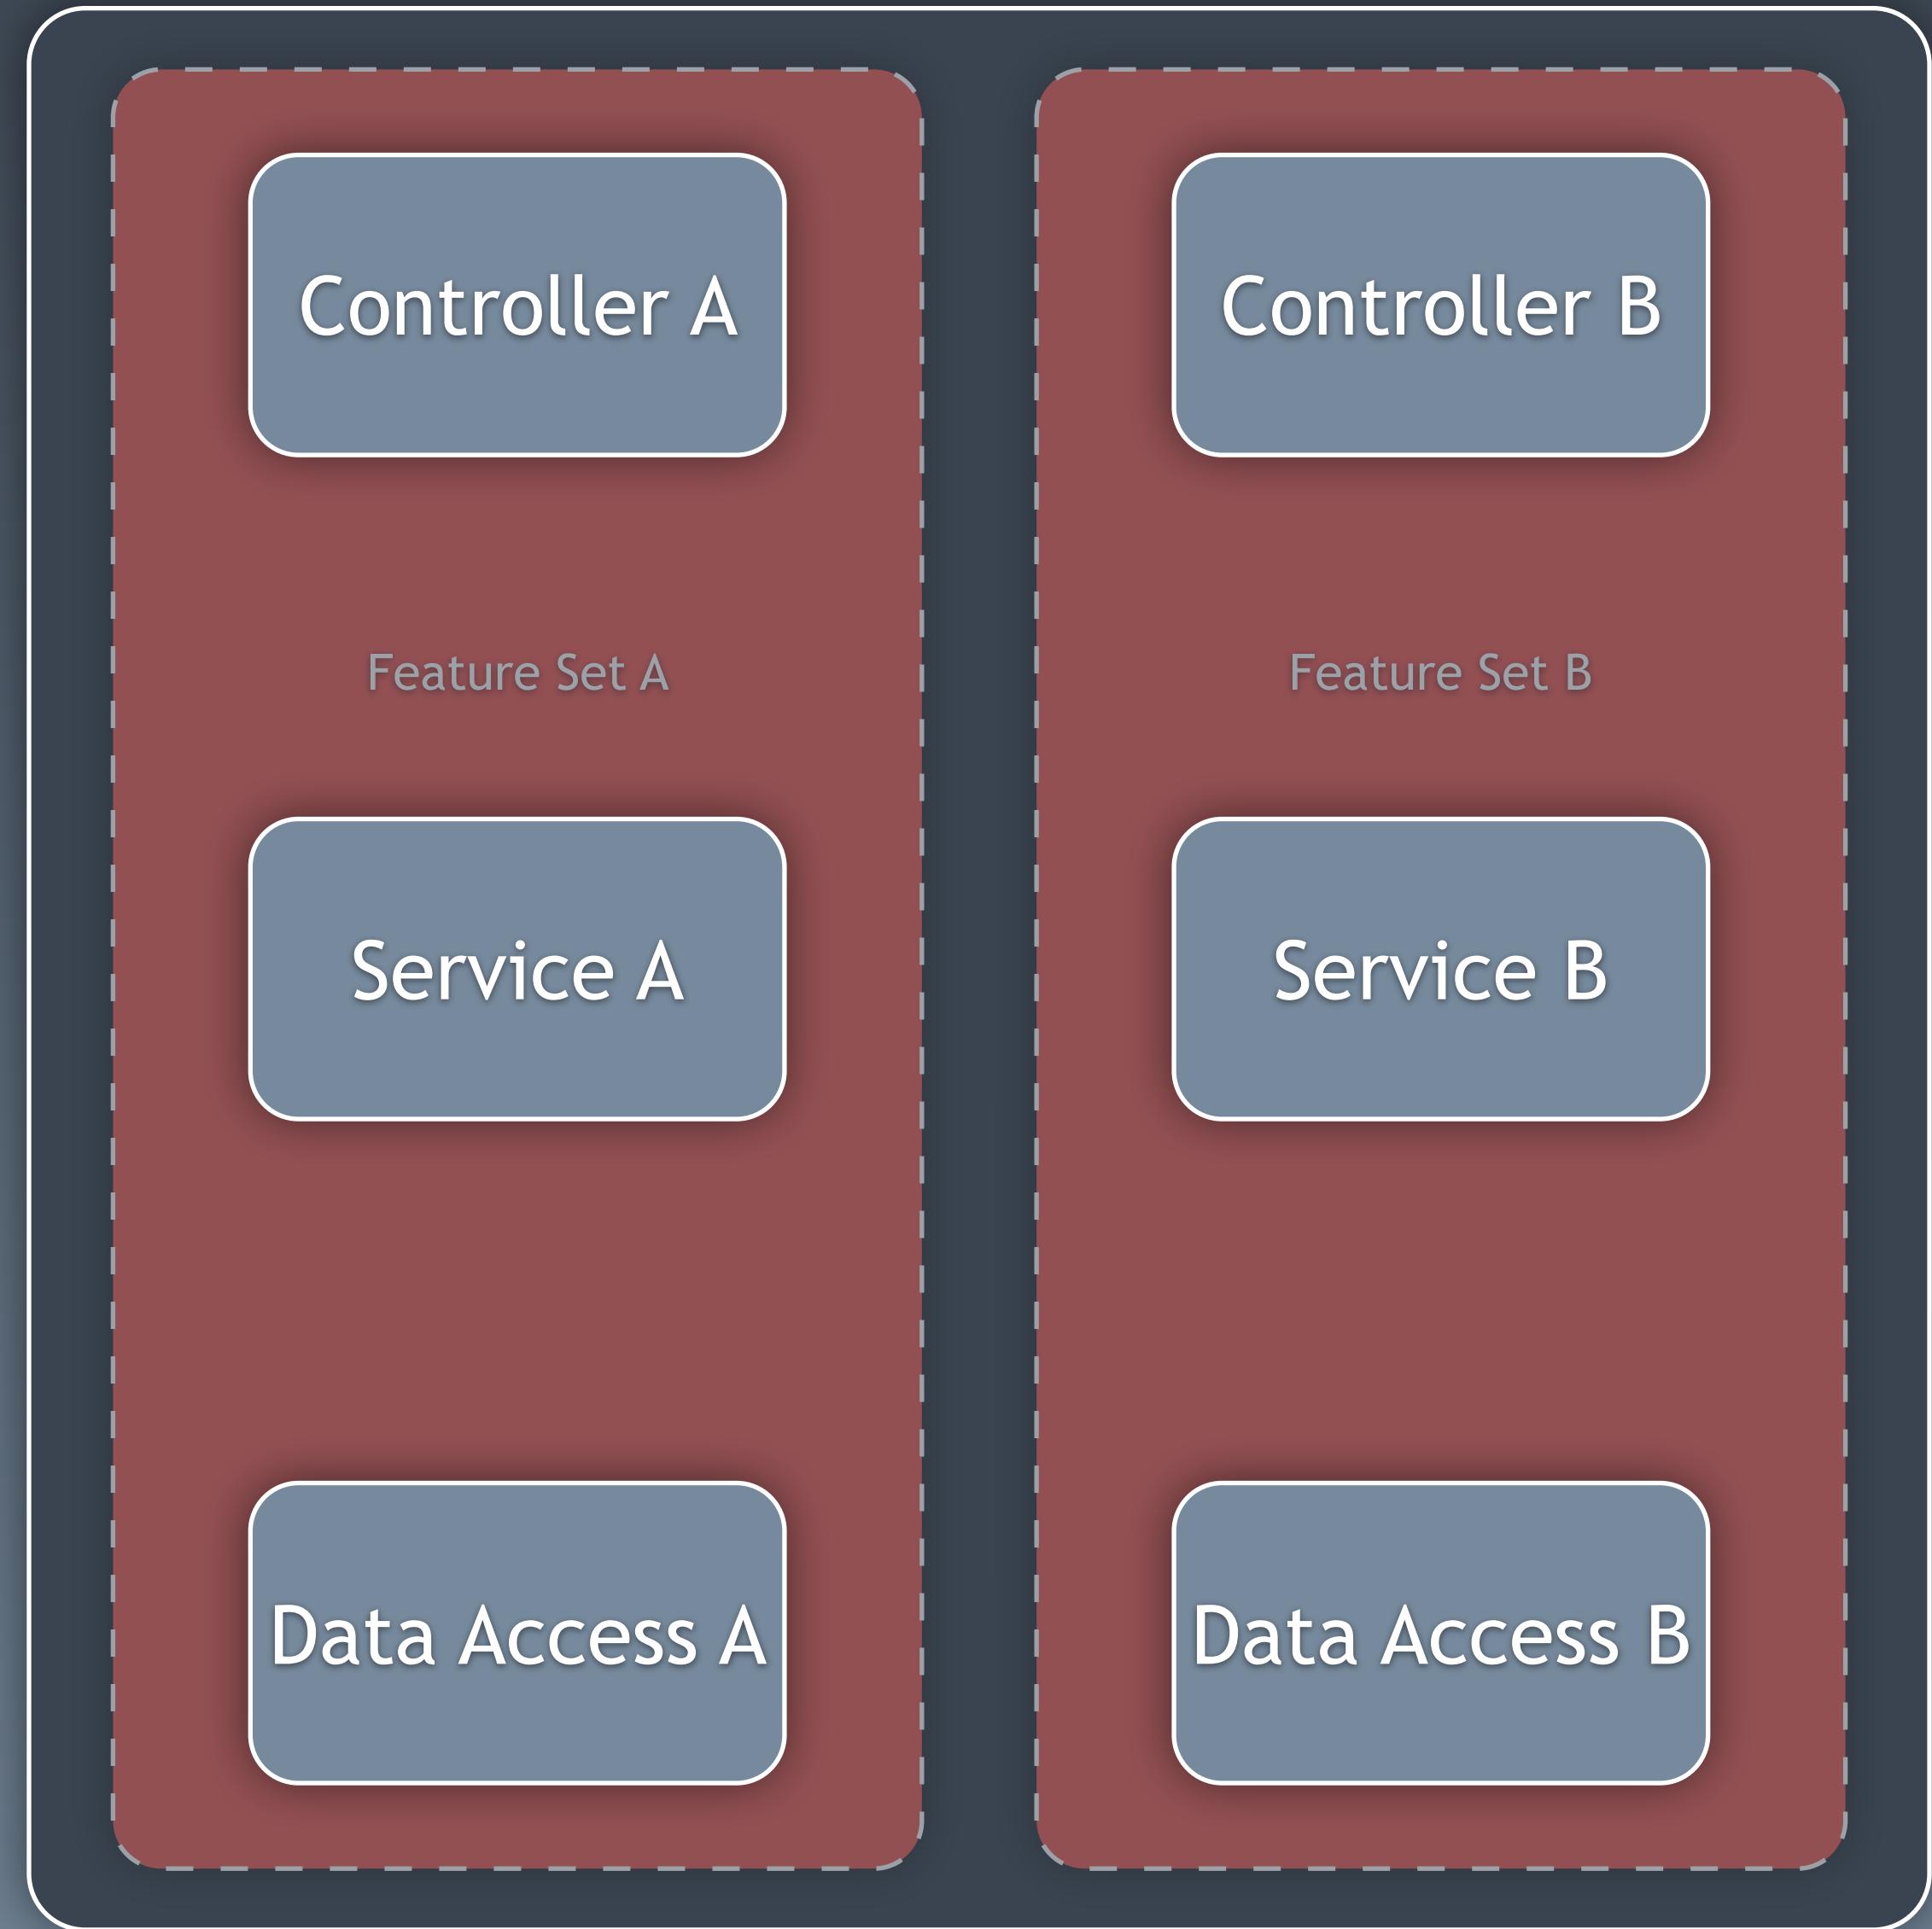
On the whole I've found this to be an effective form of modularization, and one that I regularly use and encourage. Its main advantage (for me) is that it allows me to reduce the scope of a

Although presentation-domain-data separation is a common approach, it should only be applied at a relatively small granularity. As an application grows, each layer can get sufficiently complex on its own that you need to modularize further. When this happens it's usually not best to use presentation-domain-data as the higher level of modules. Often frameworks encourage you to have something like view-model-data as the top level namespaces; that's ok for smaller systems, but once any of these layers gets too big you should split your top level into domain oriented modules which are internally layered.

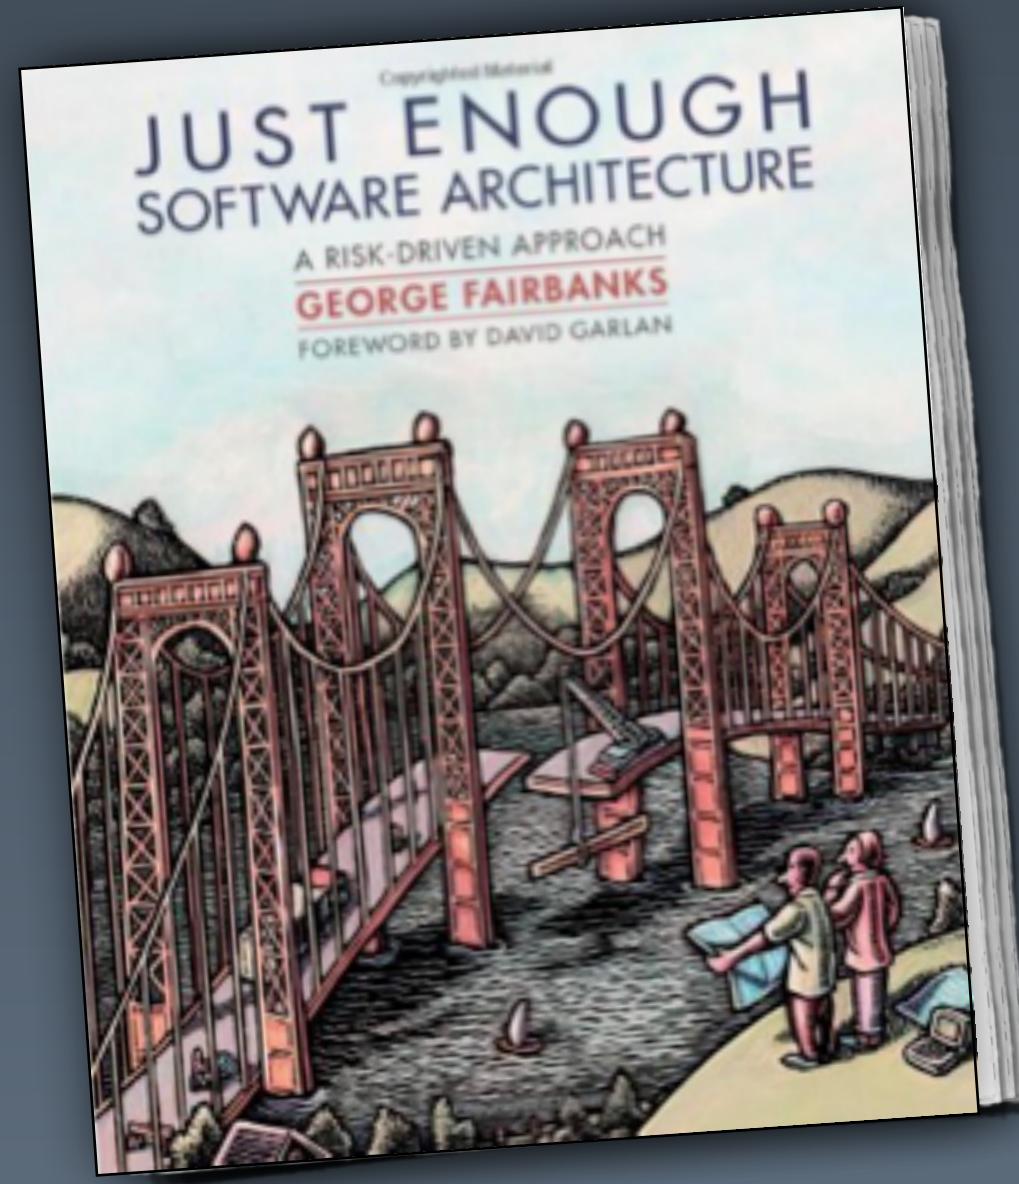


*Don't use layers as the top level modules in a complex application...*

*... instead make your top level modules be full-stack*



Package by feature (vertical slicing)



“architecturally-evident  
coding style”

# Architecturally-evident coding styles include:

Annotations/attributes (@Component, [Component], etc)

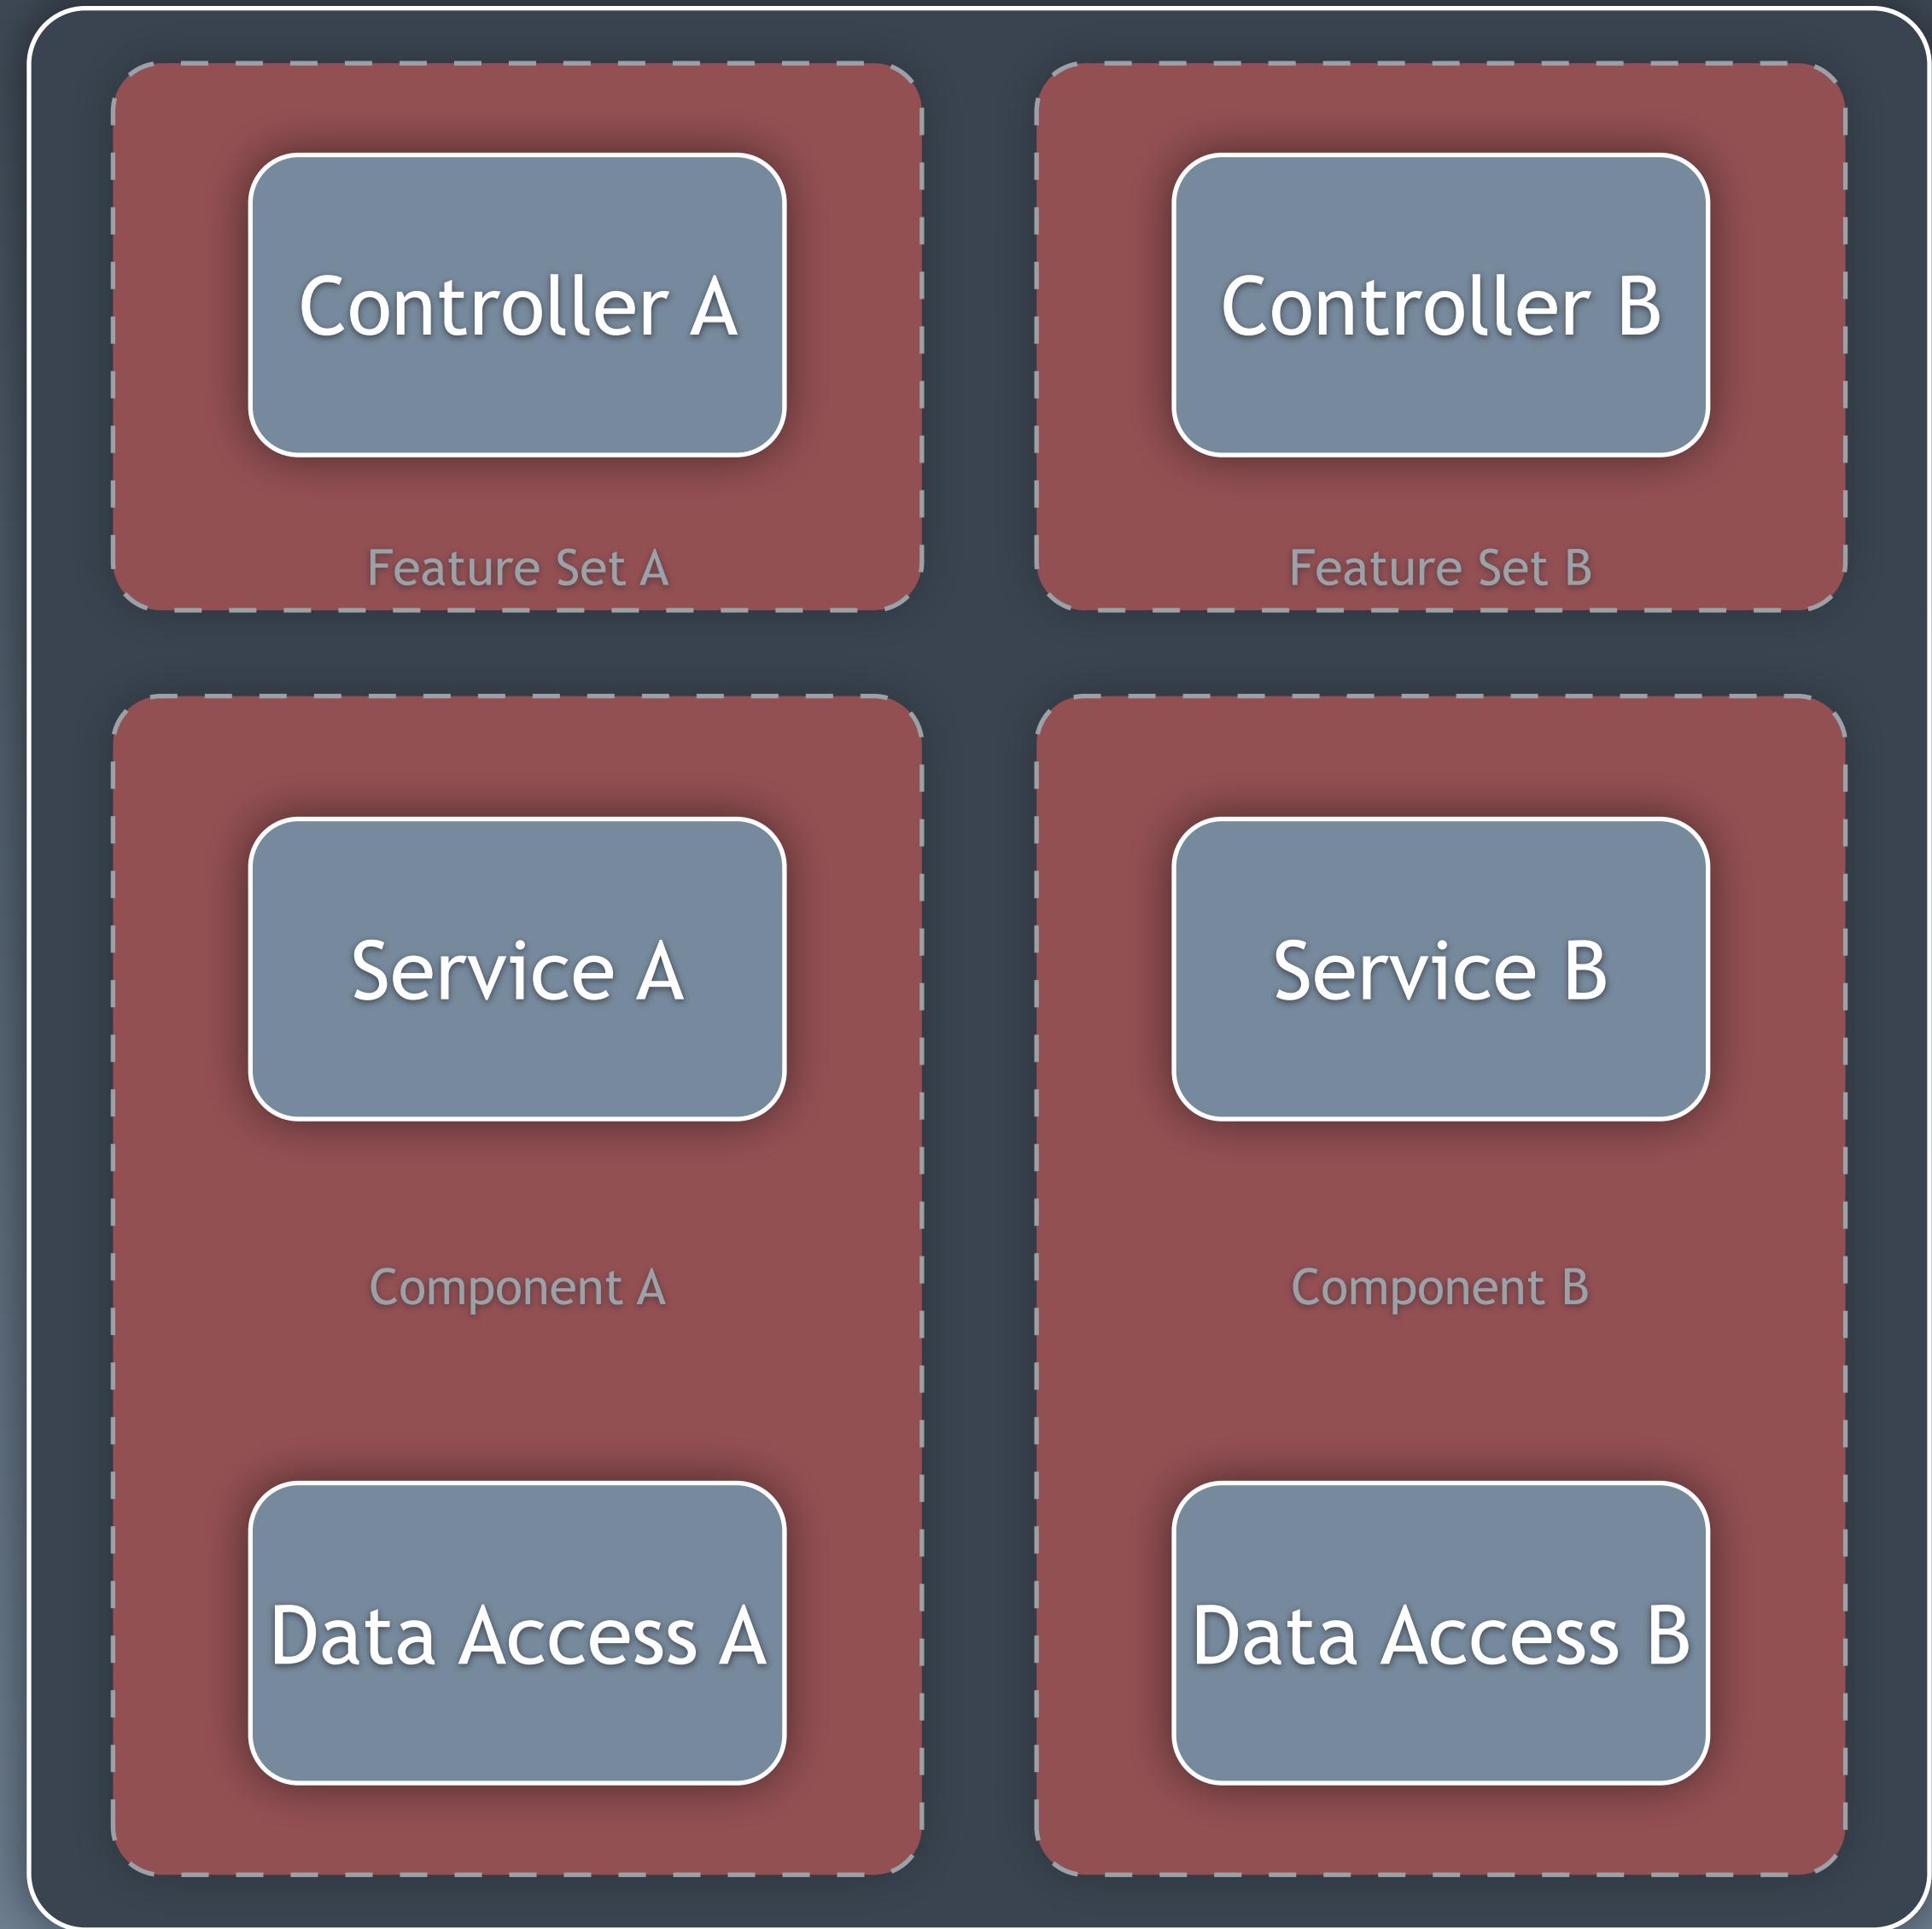
Naming conventions (\*Service)

Namespacing/packaging (com.mycompany.system.components.\*)

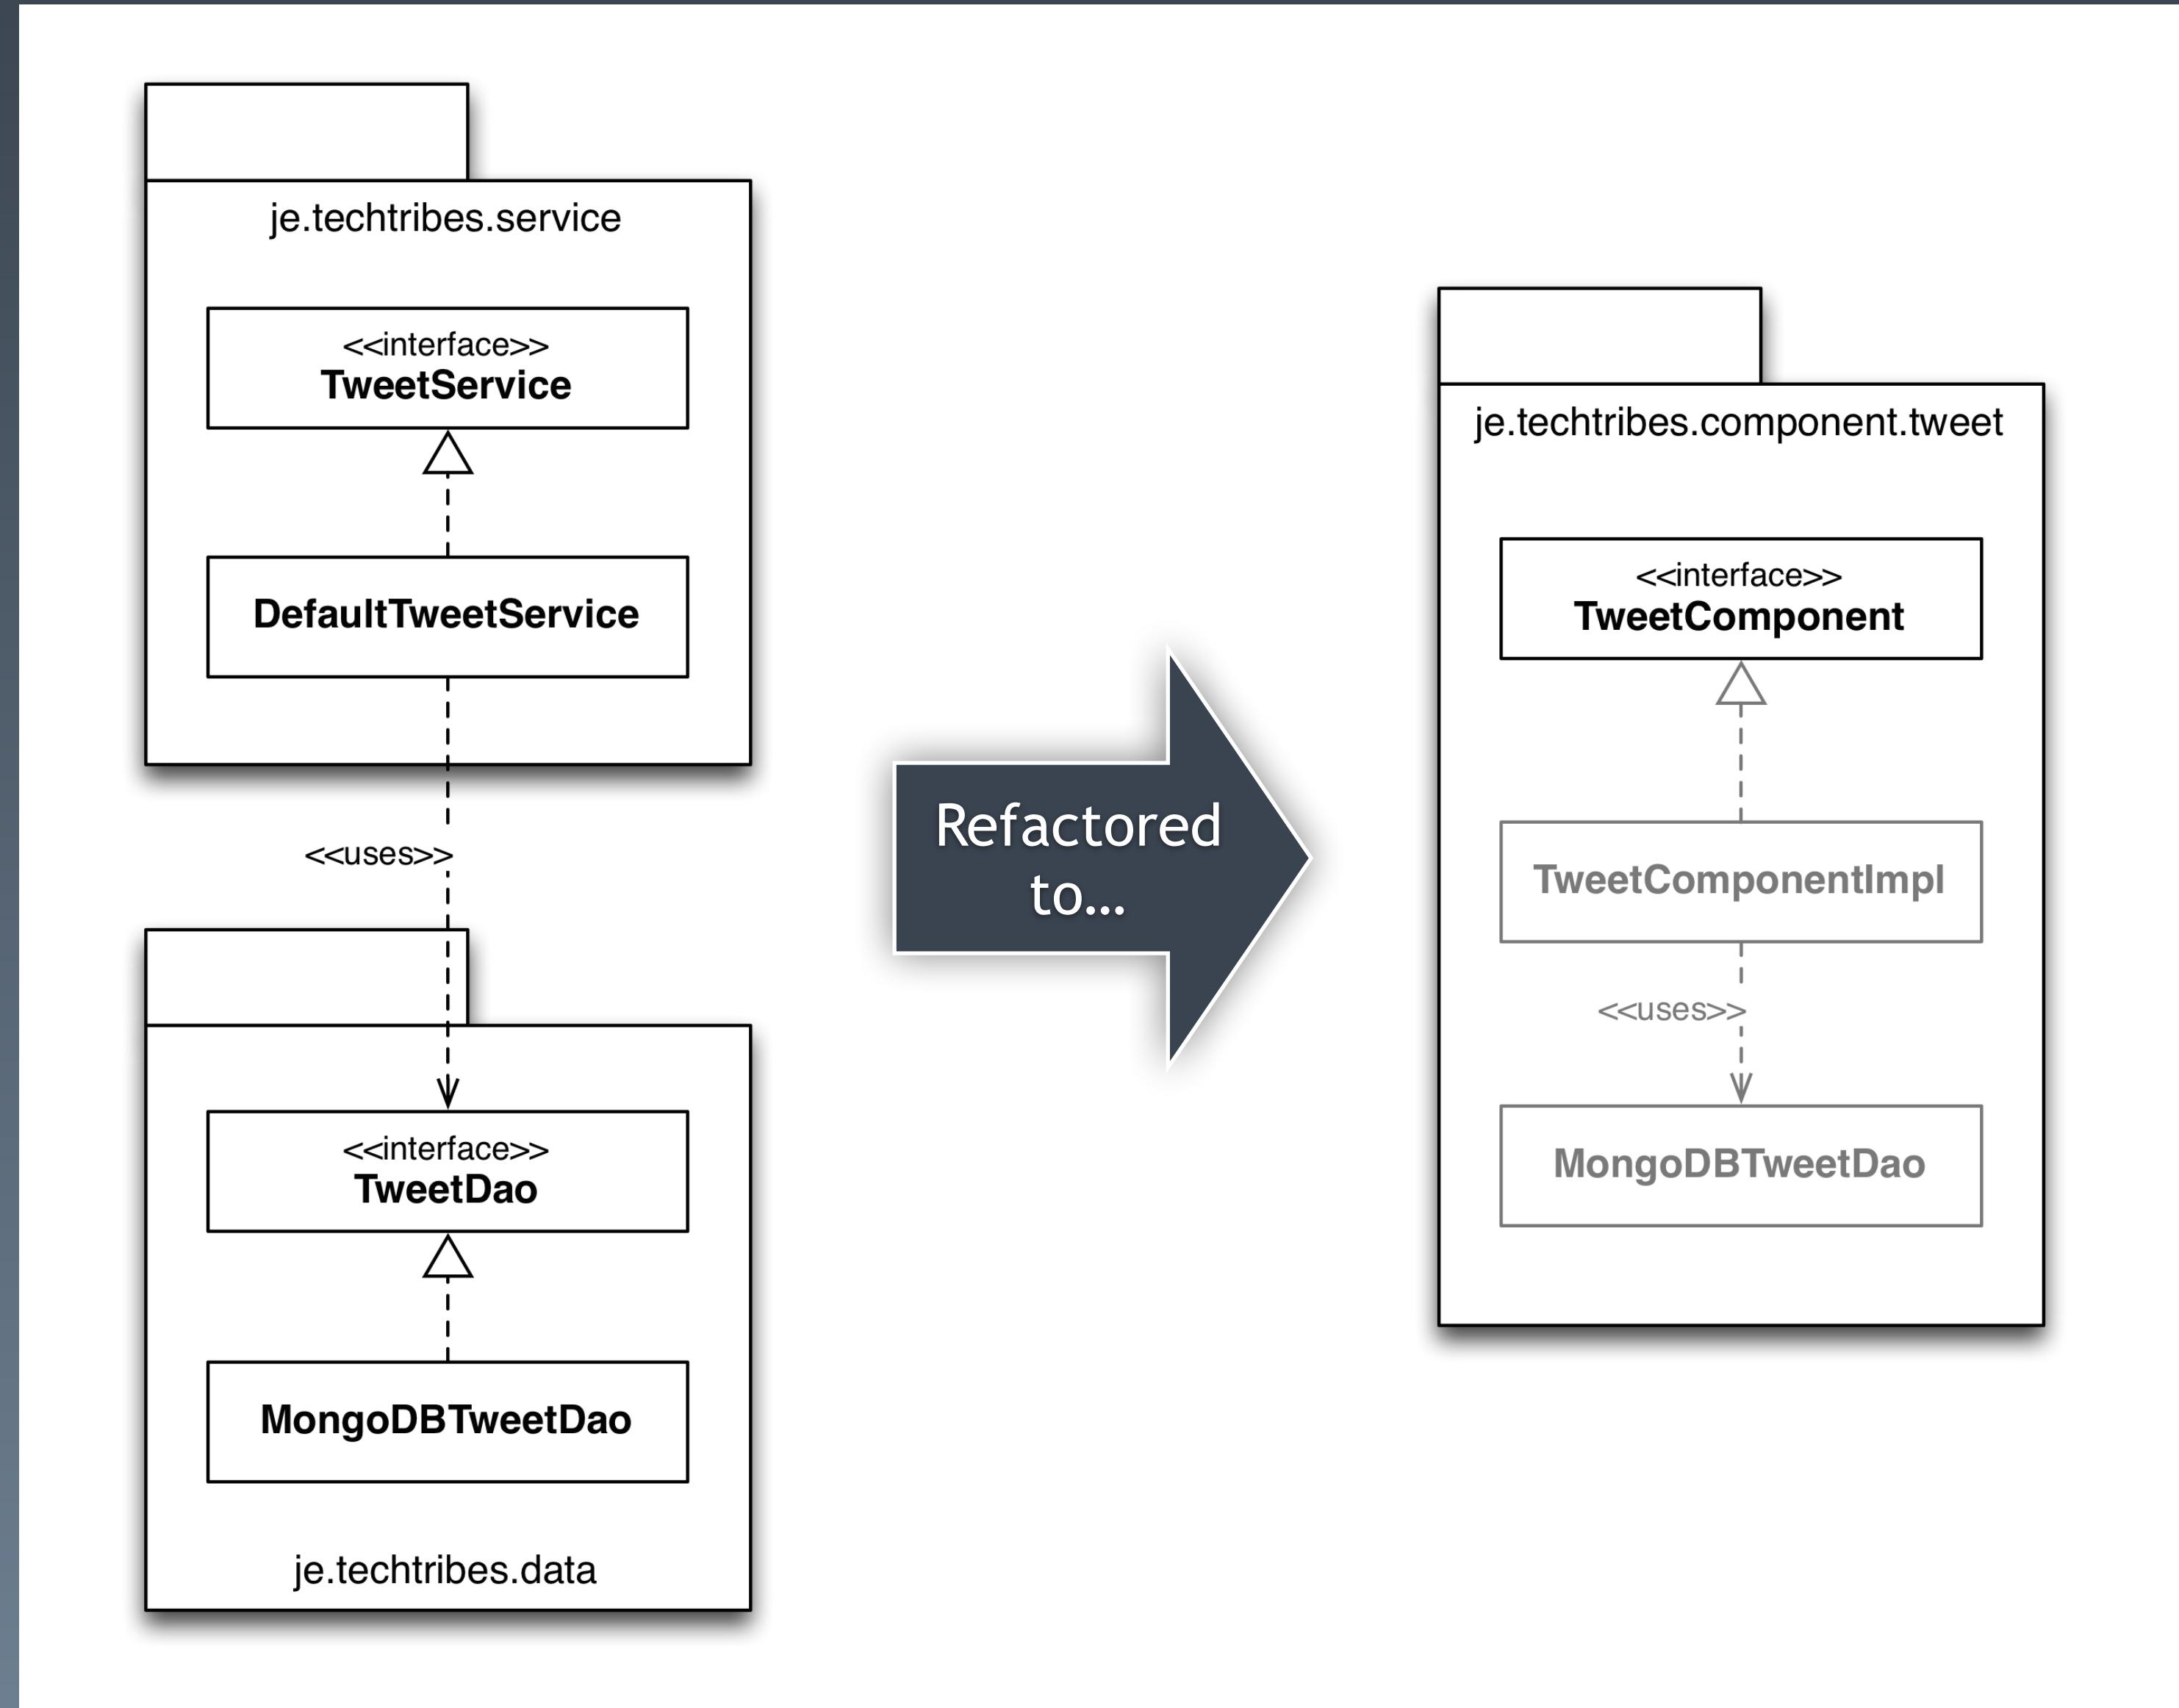
Maven modules, OSGi modules, Java 9 and

Jigsaw, JavaScript module patterns,

ECMAScript 6 modules, microservices, etc



Package by component



An example of an  
architecturally-evident coding style

The screenshot shows a Java code editor and a file browser. The code editor on the right displays the `TweetComponent` interface. The file browser on the left shows the project structure of `techtribes`.

```
package je.techtribes.component.tweet;

import ...

/**
 * Provides access to tweets.
 */
@Component
public interface TweetComponent {

    /**
     * Gets the most recent tweets by page number.
     */
    List<Tweet> getRecentTweets(int page, int pageSize);

    /**
     * Gets the most recent tweets for a given person.
     */
    List<Tweet> getRecentTweets(ContentSource contentSource);

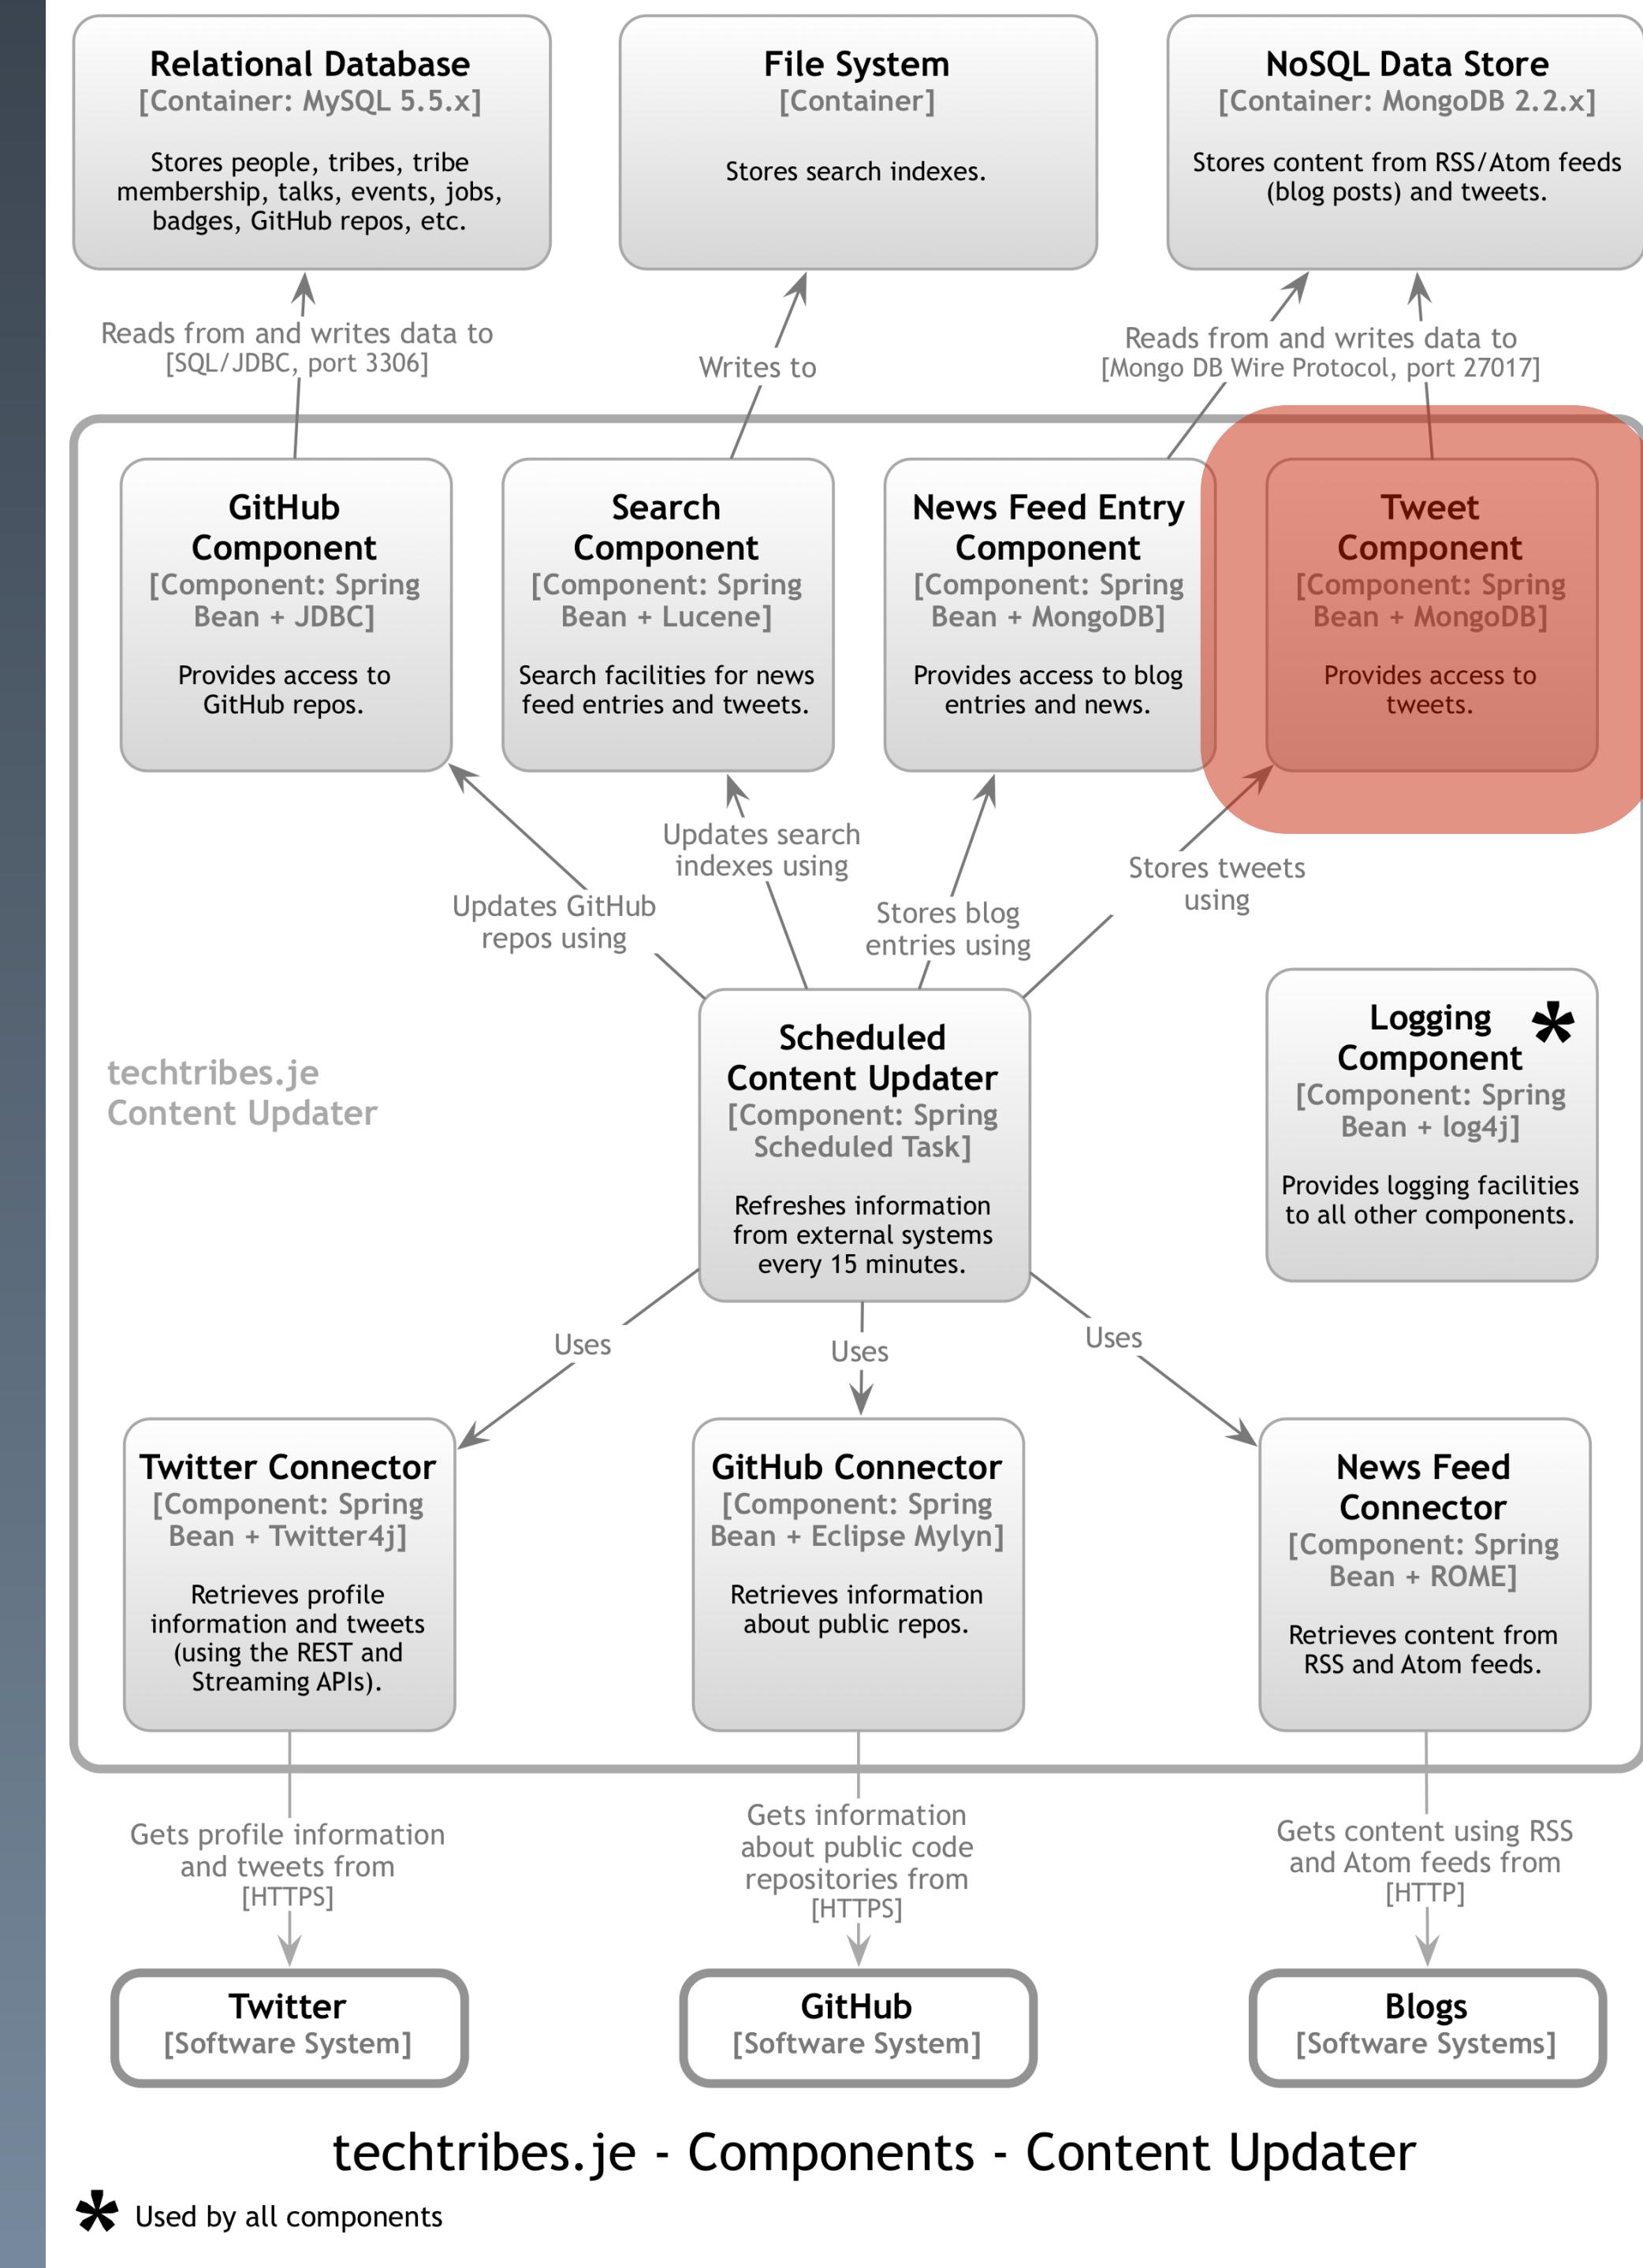
    /**
     * Gets the most recent tweets for a given collection.
     */
    List<Tweet> getRecentTweets(Collection<ContentSource> contentSources);

    /**
     * Calculates how many tweets a given person or tribe has.
     */
    long calculateTotalTweets(ContentSource contentSource);
}
```

**Project Tree:**

- Project
- techtribes [techtribes] (~sandbox/te...)
  - docs
  - lib
    - production
    - runtime
    - test
  - structurizr
  - techtribes-core
    - sql
    - src
      - je
        - techtribes
          - component
          - activity
          - badge
          - book
          - contentsource
          - creation
          - event
          - github
          - log
          - newsfeedentry
          - search
        - talk
        - tweet
          - component.xml
          - MongoDbTweetDao
          - TweetComponent**
          - TweetComponentImpl
          - TweetException
        - domain
        - util
      - config.xml
    - test
    - techtribes-core.iml
    - techtribes-updater
    - techtribes-util
    - techtribes-web
      - .gitignore
      - build\_number

# Component diagram (level 3)



**\*** Used by all components

# Context diagram

(level 1)

# Container diagram

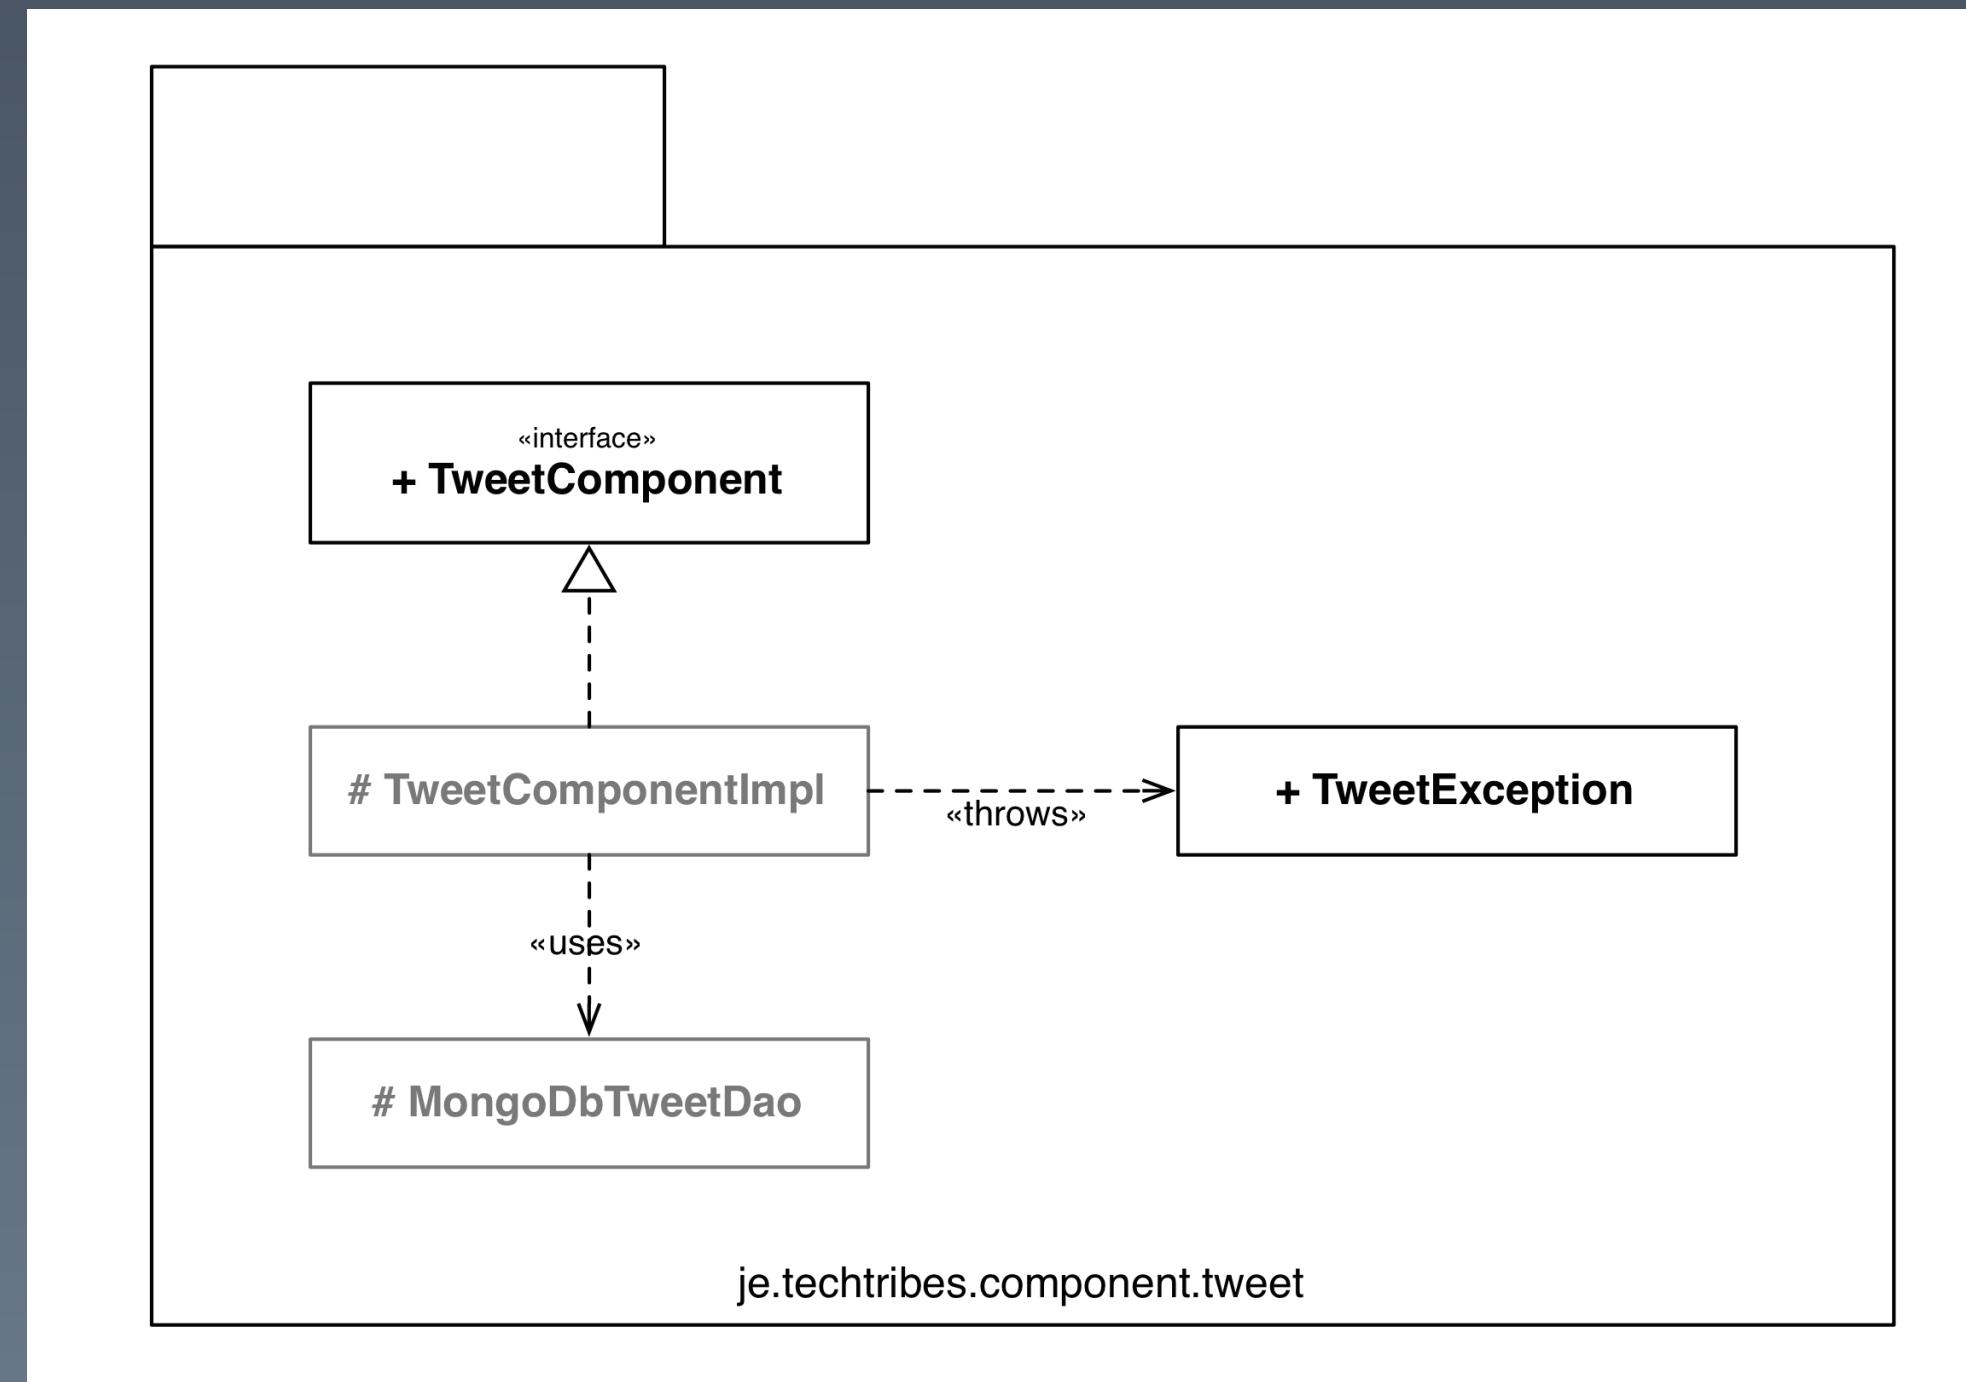
(level 2)

# Component diagram

(level 3)

# Class diagram

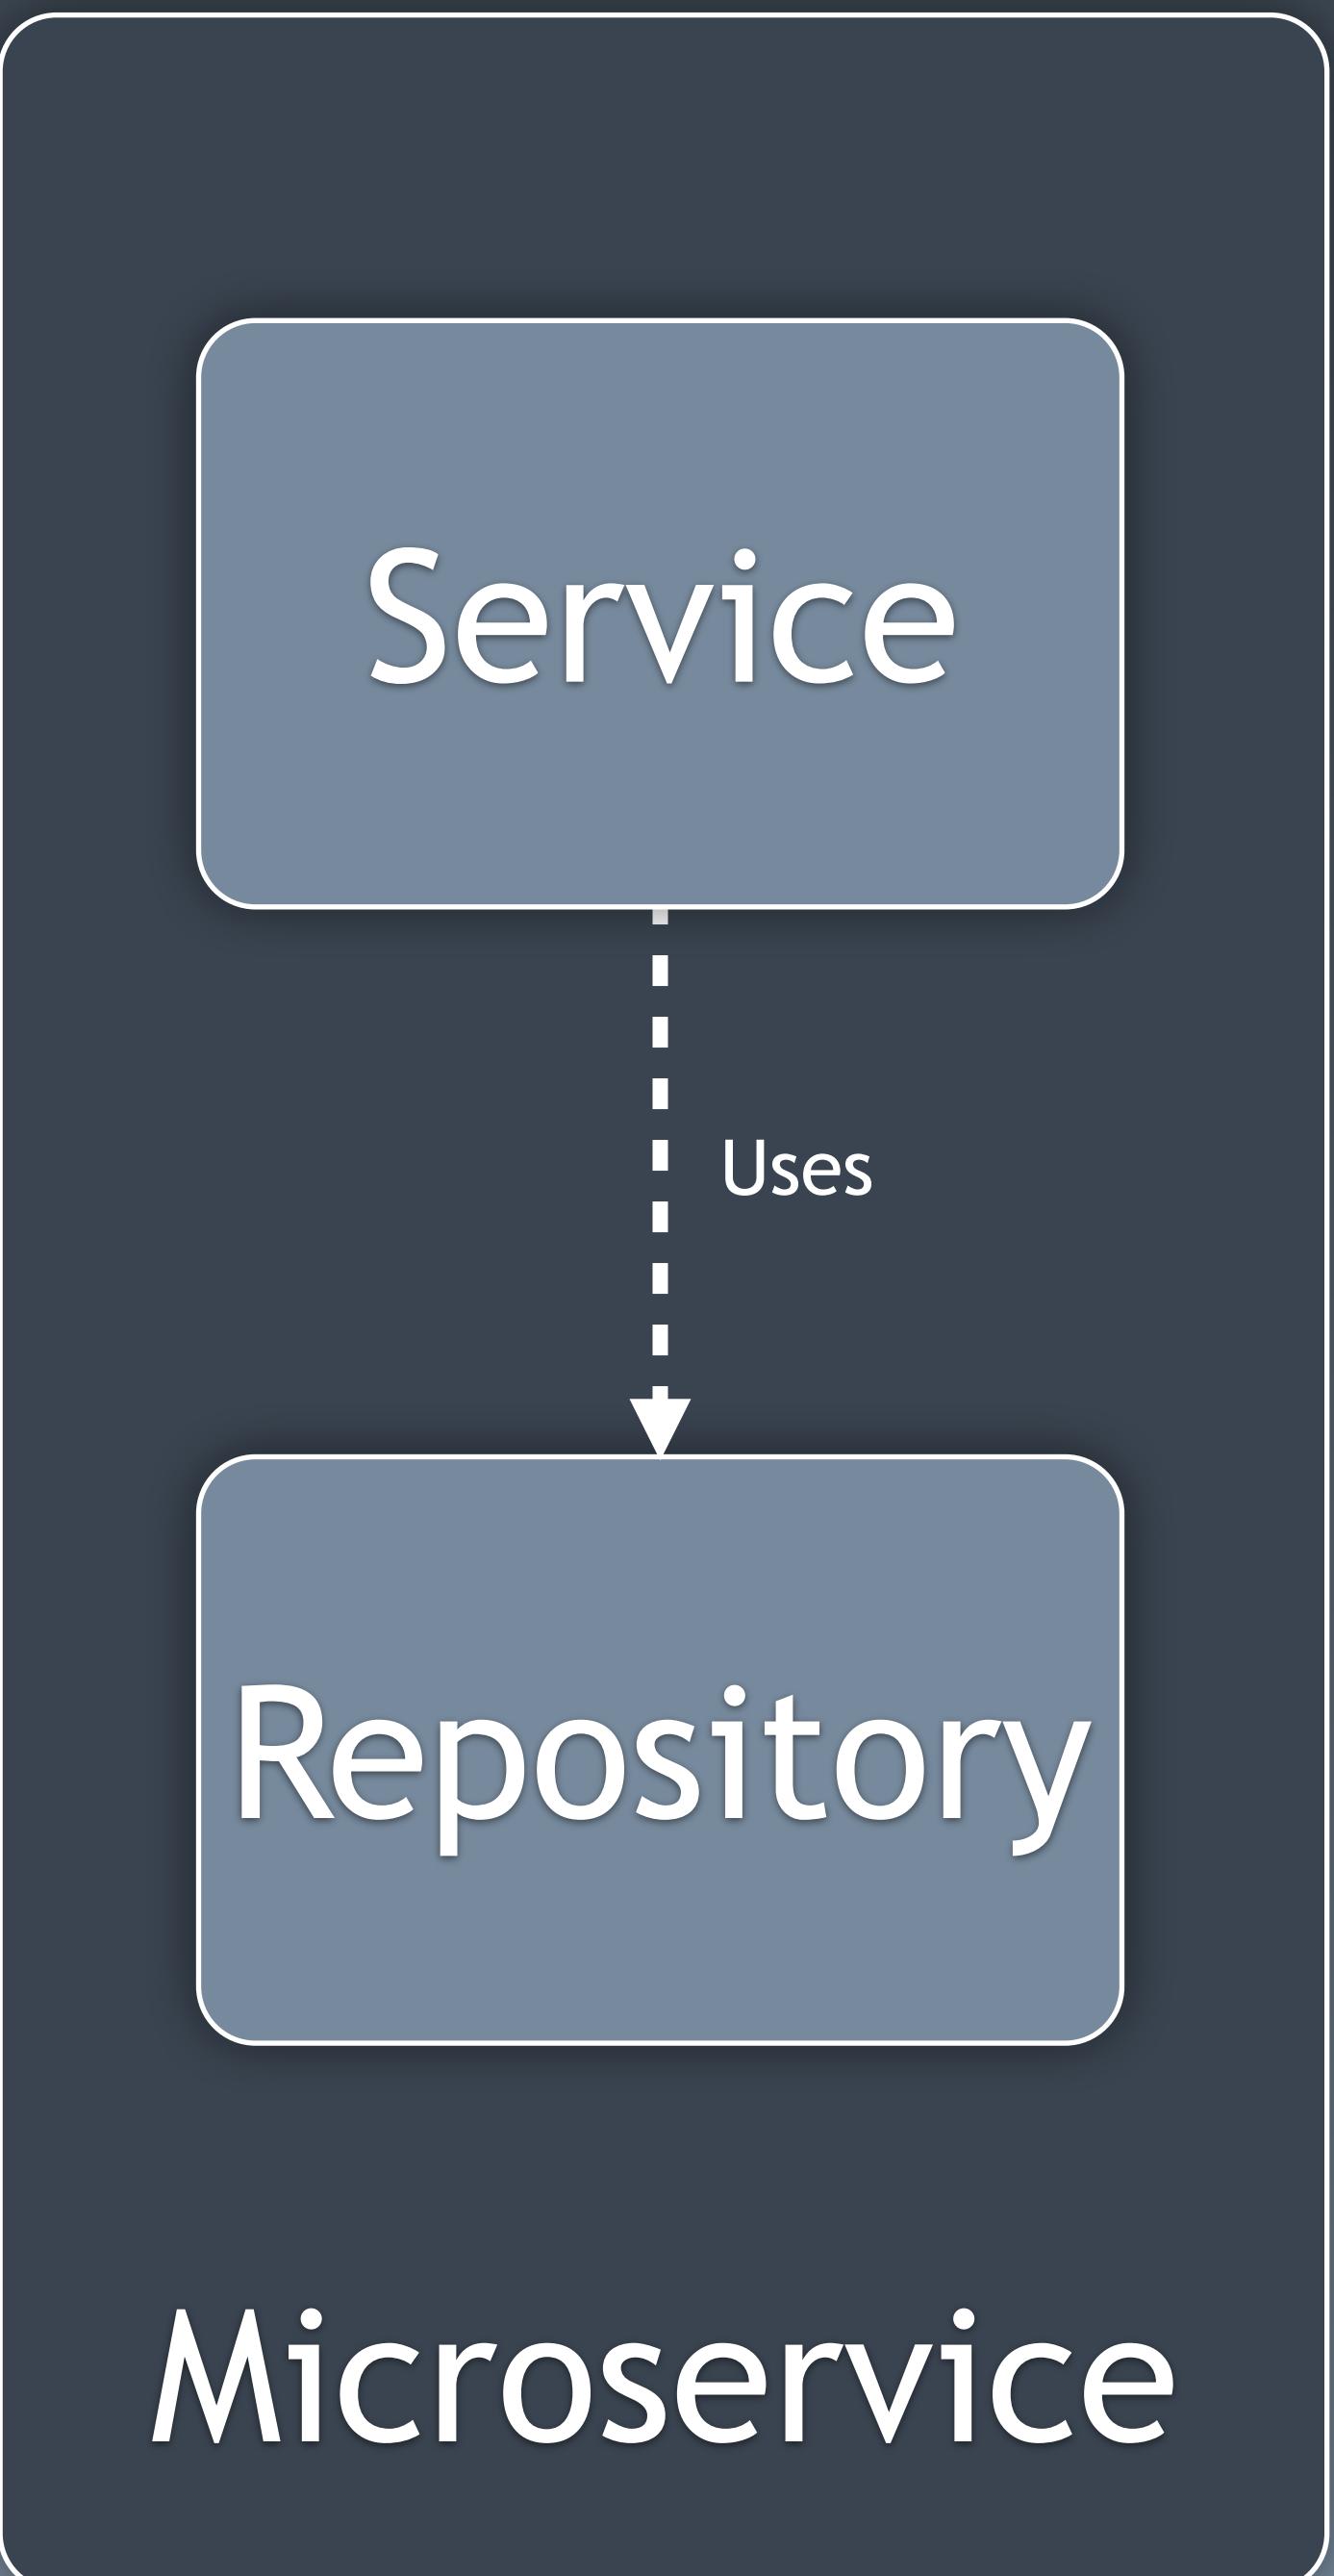
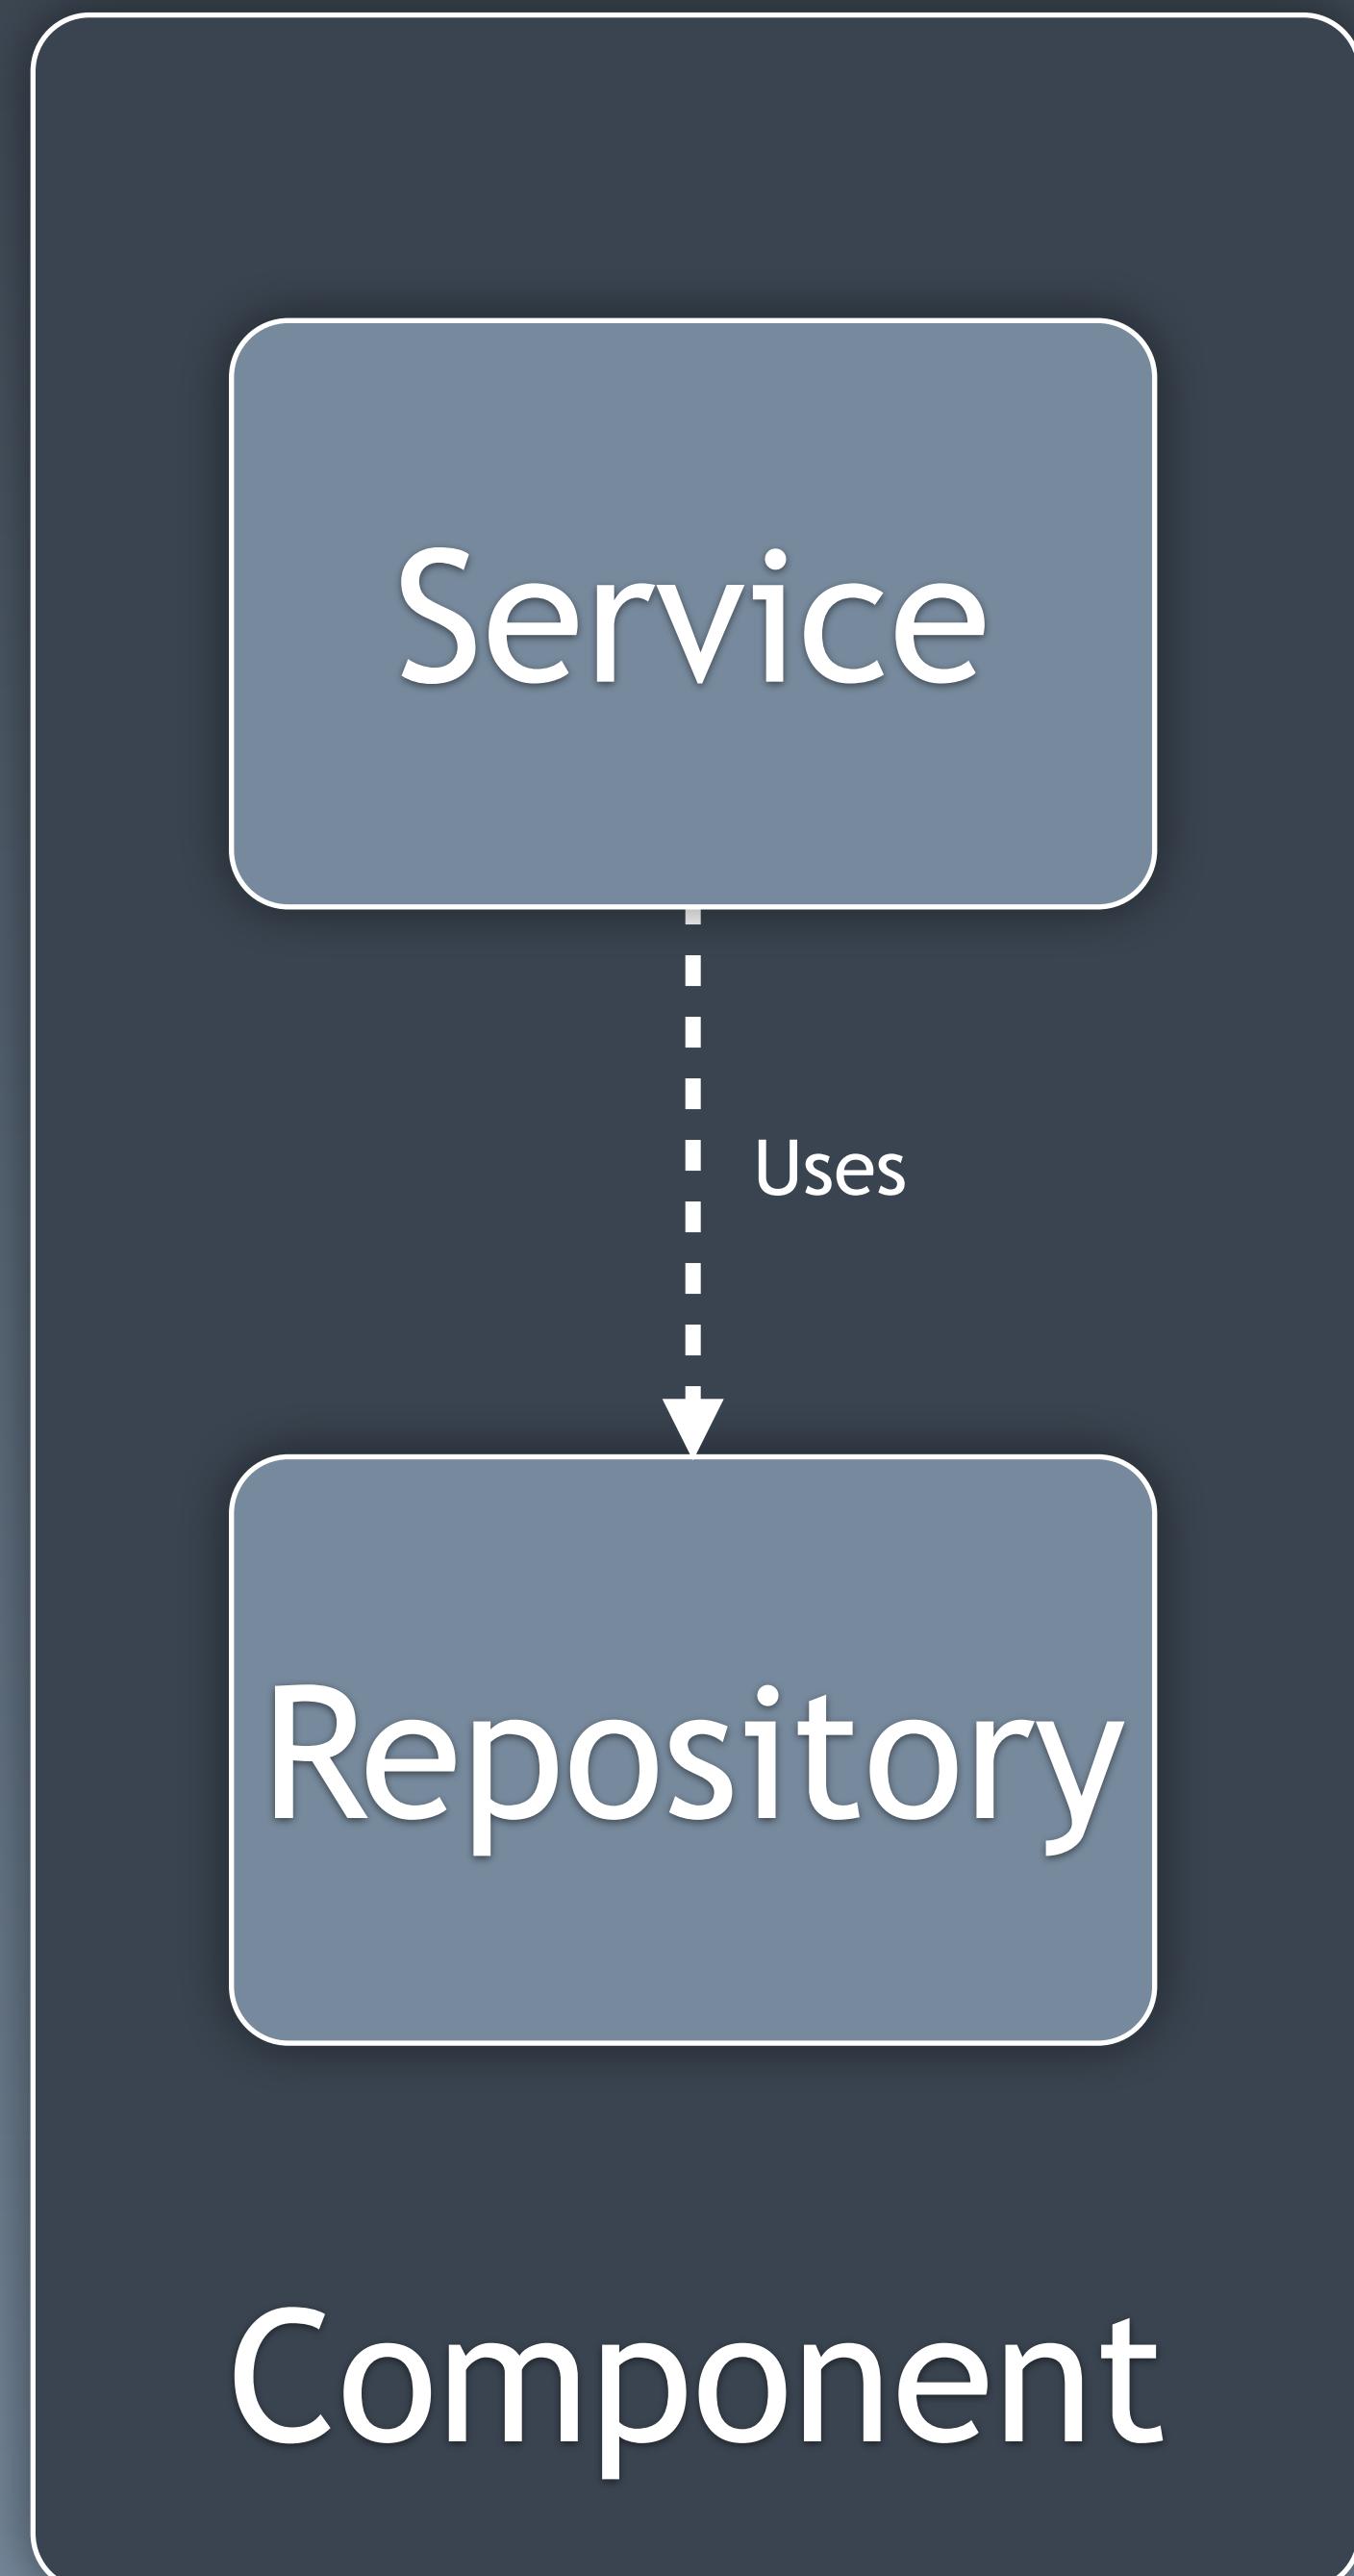
(level 4)



# Modularity as a principle

# Impermeable boundaries

Apply  
component-based  
design and  
service-orientation  
to a monolithic  
application



# Design decisions

How do you  
design  
software?

What is your  
decomposition  
strategy?



**WIKIPEDIA**  
The Free Encyclopedia

Main page  
Contents  
Featured content  
Current events  
Random article  
Donate to Wikipedia  
Wikipedia store

Interaction  
Help  
About Wikipedia  
Community portal  
Recent changes

Article **Talk**

Read **Edit** View history

Search



Not logged in [Talk](#) [Contributions](#) [Create account](#) [Log in](#)

# Decomposition (computer science)

From Wikipedia, the free encyclopedia

**Decomposition** in computer science, also known as **factoring**, is breaking a complex problem or system into parts that are easier to conceive, understand, program, and maintain.

## Contents [hide]

- 1 Overview
- 2 Decomposition topics
  - 2.1 Decomposition paradigm
  - 2.2 Decomposition diagram
- 3 See also
- 4 References
- 5 External links

## Decomposition paradigm [ edit ]

A decomposition paradigm in computer programming is a strategy for organizing a program as a number of parts, and it usually implies a specific way to organize a program text. Usually the aim of using a decomposition paradigm is to optimize some metric related to program complexity, for example the modularity of the program or its maintainability.

Most decomposition paradigms suggest breaking down a program into parts so as to minimize the static dependencies among those parts, and to maximize the **cohesiveness** of each part. Some popular decomposition paradigms are the procedural, modules, abstract data type and **object oriented** ones.

[Printable version](#)

More generally, **functional decomposition** in computer science is a technique for mastering the complexity of the function of a model. A functional model of a system is thereby replaced by a series of functional models of subsystems.<sup>[3]</sup>

Languages

Azerbaijani



---

# On the Criteria To Be Used in Decomposing Systems into Modules

D.L. Parnas  
Carnegie-Mellon University

## Expected Benefits of Modular Programming

The benefits expected of modular programming are: (1) managerial—development time should be shortened because separate groups would work on each module with little need for communication; (2) product flexibility—it should be possible to make drastic changes to one module without a need to change others; (3) comprehensibility—it should be possible to study the system one module at a time. The whole system can therefore be better designed because it is better understood.

# Interaction style

(synchronous, asynchronous,  
commands, messages, events, etc)

Where do you put  
domain  
classes?

Shared code,  
and utilities?

Shared database tables,  
database schemas  
and foreign keys

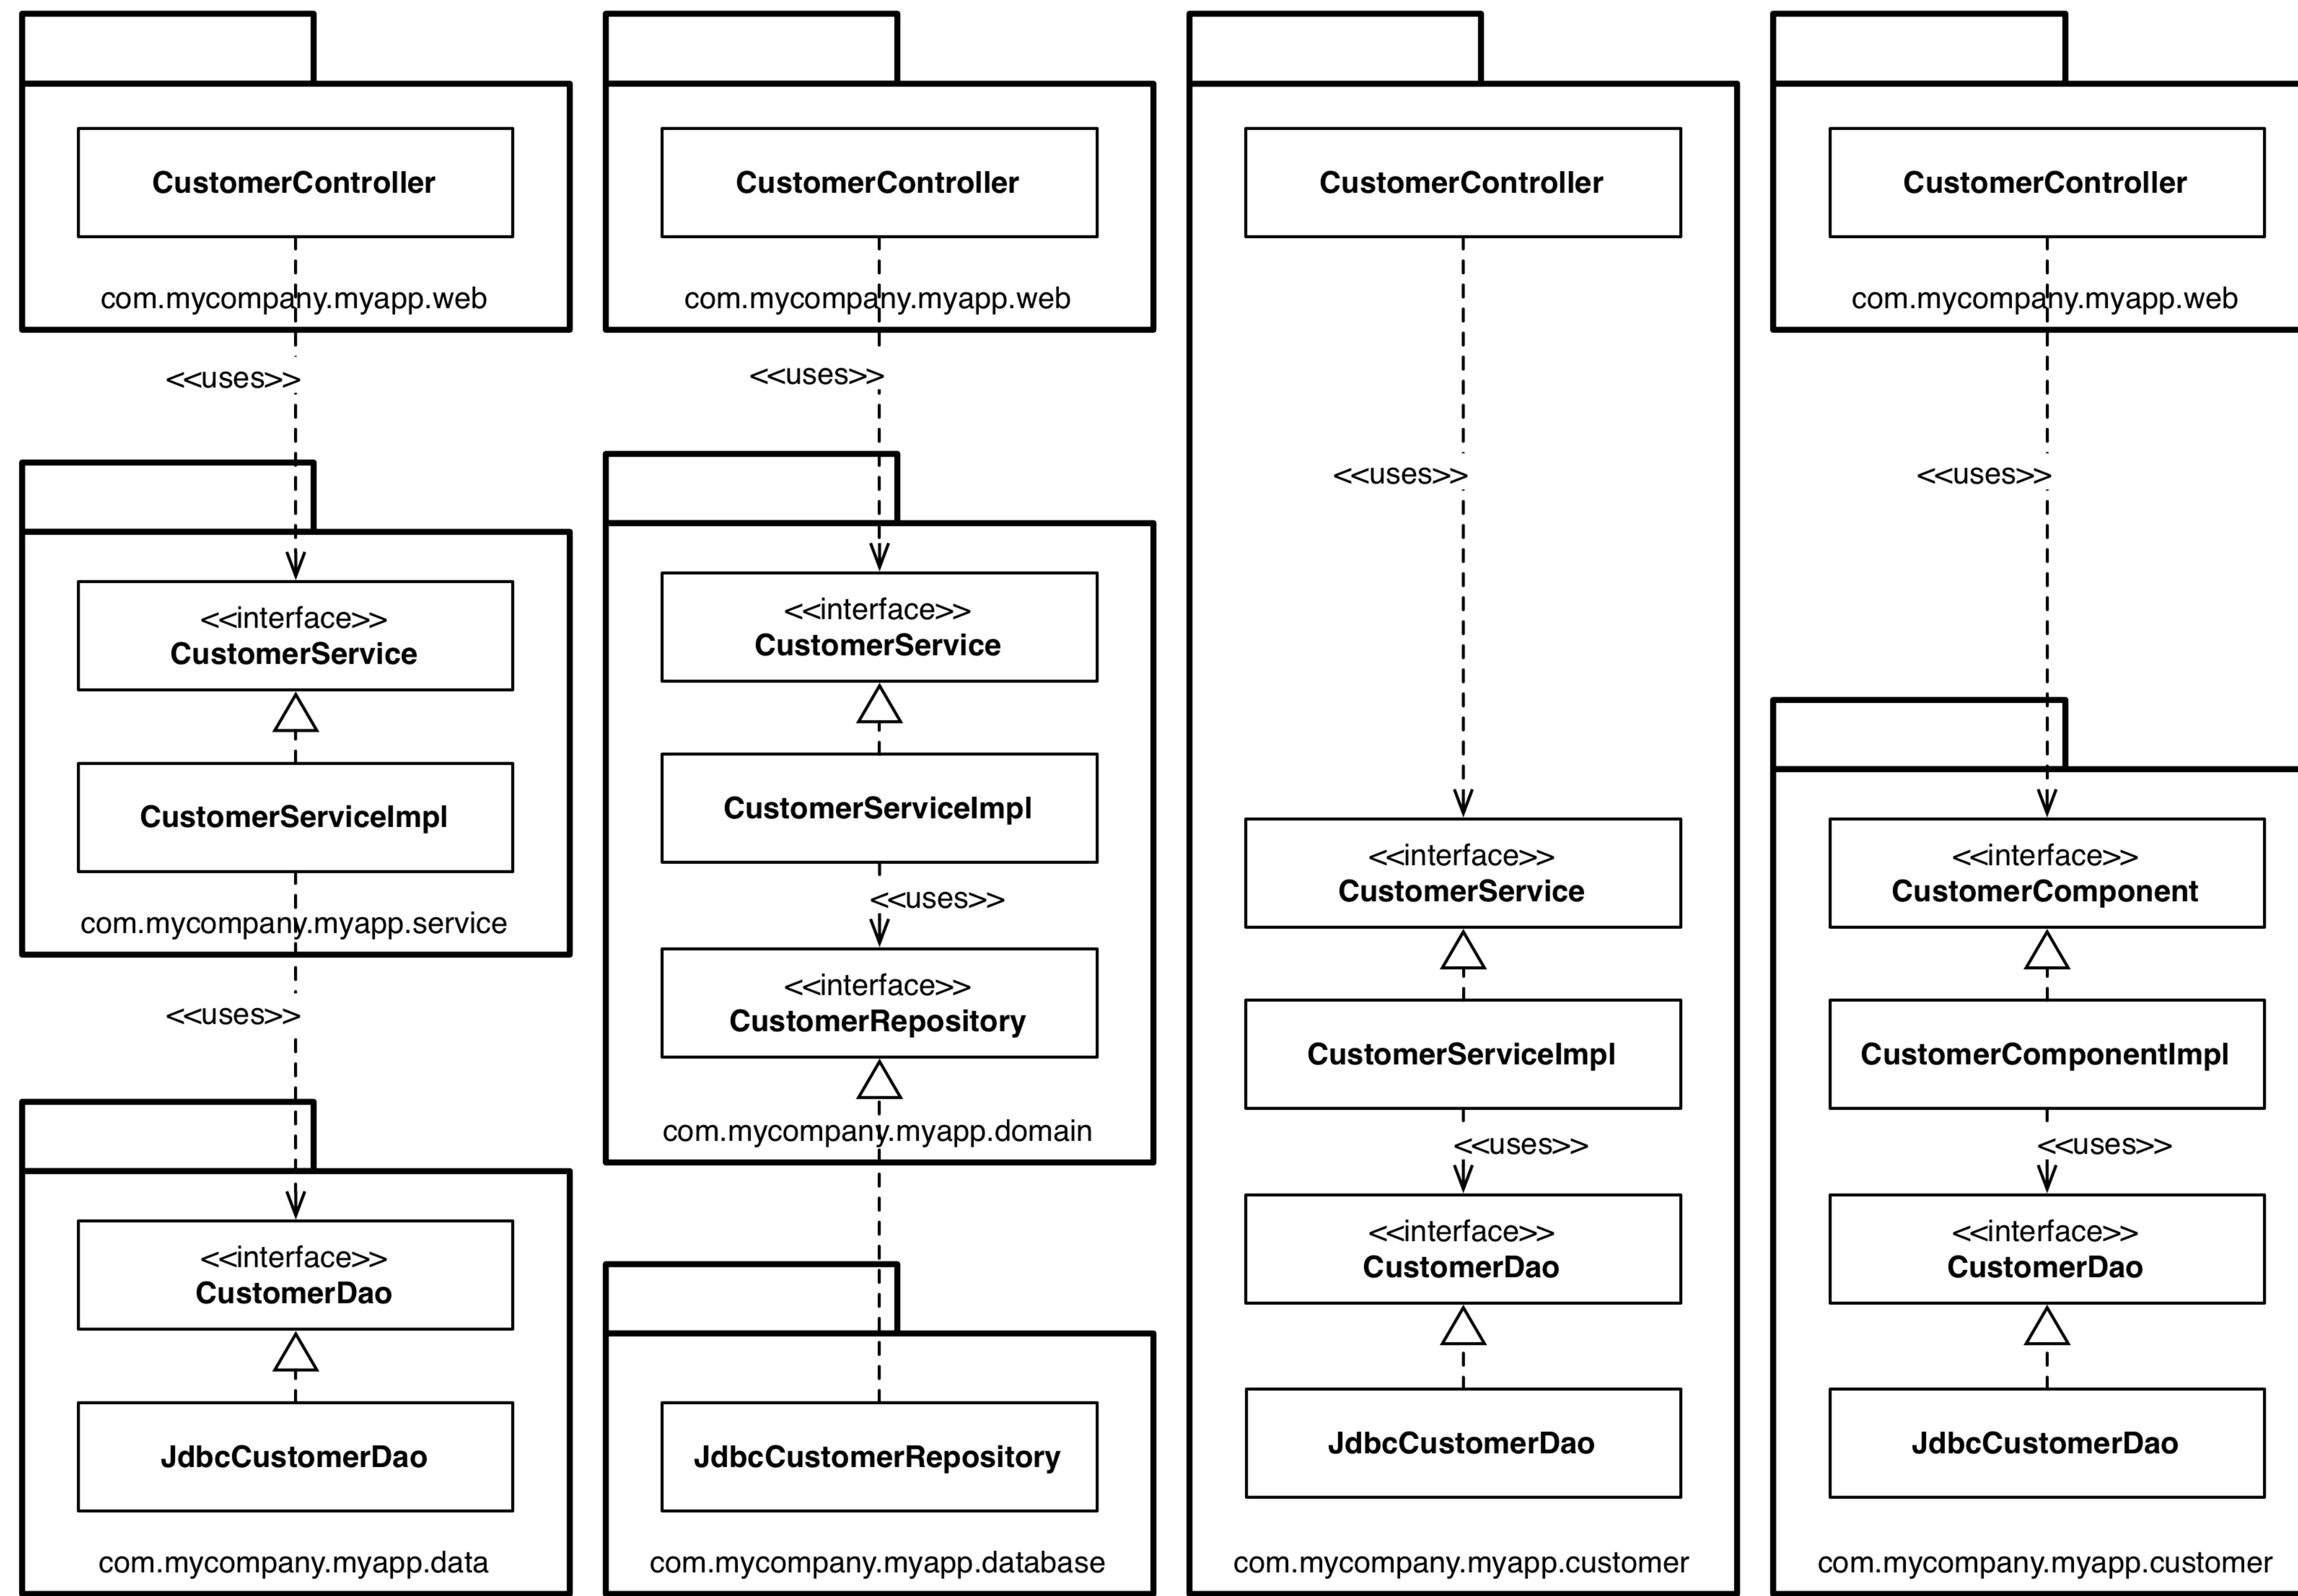
or

polyglot persistence?

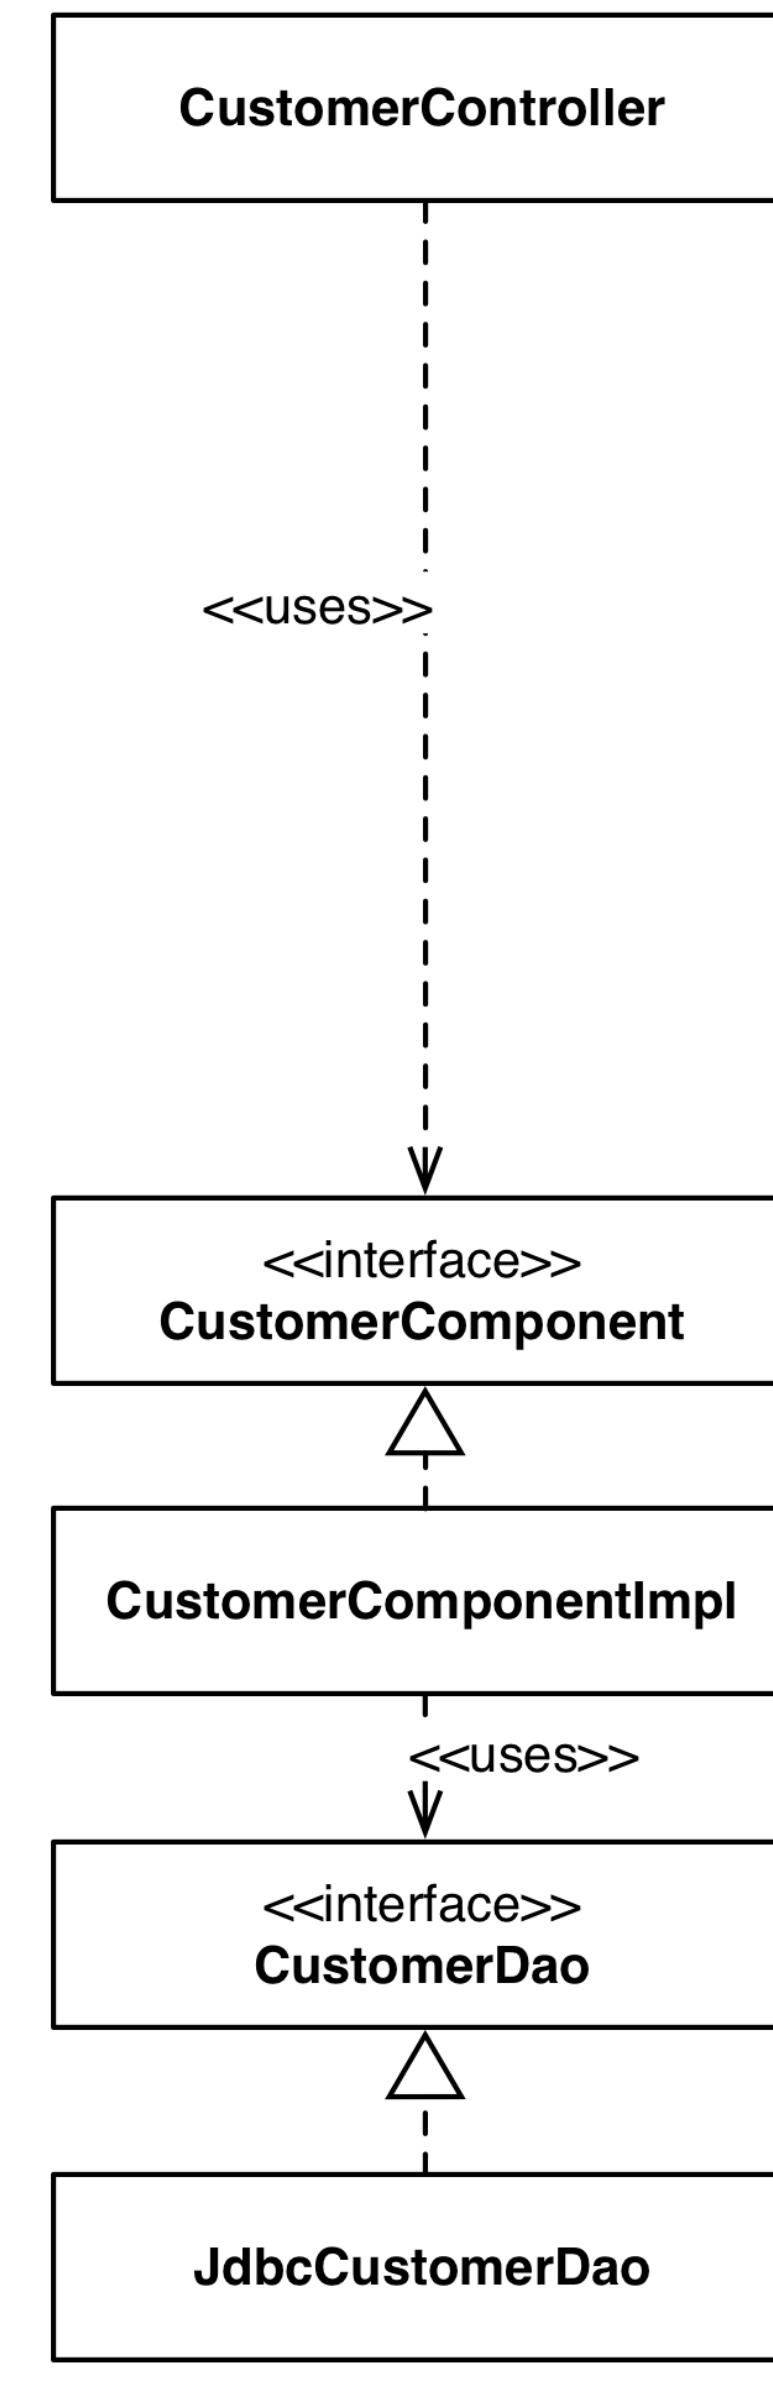
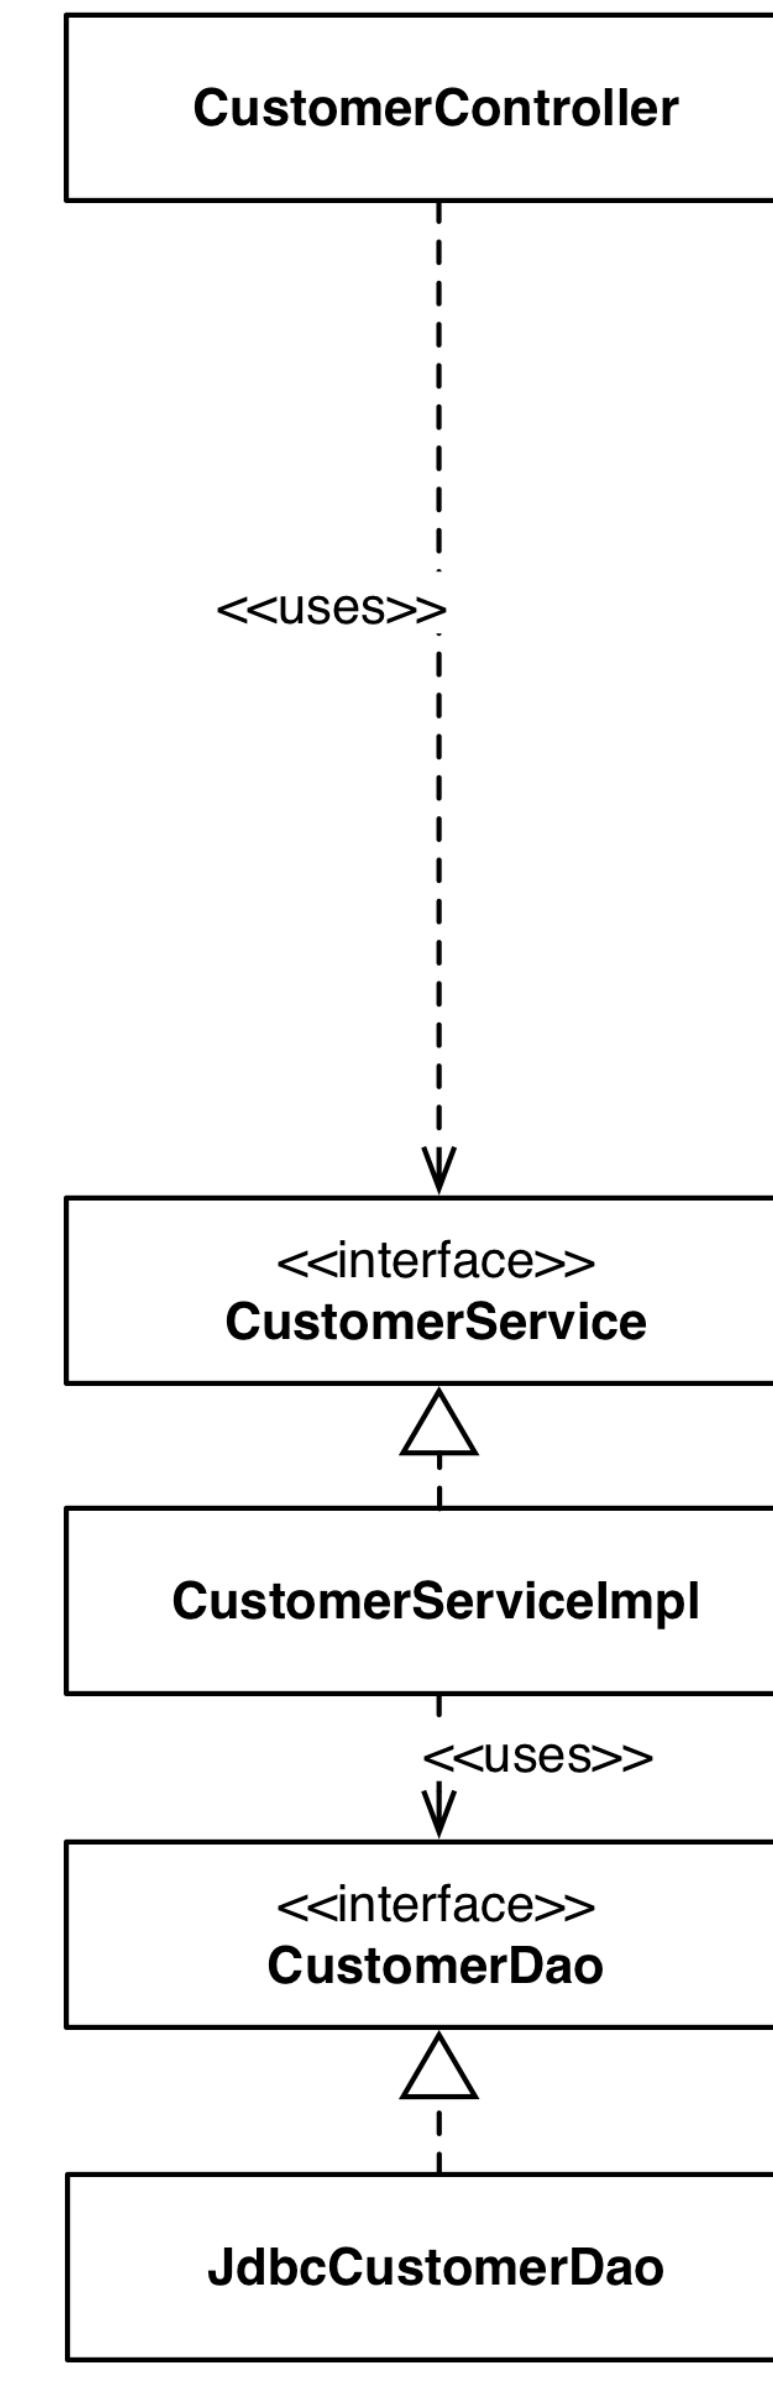
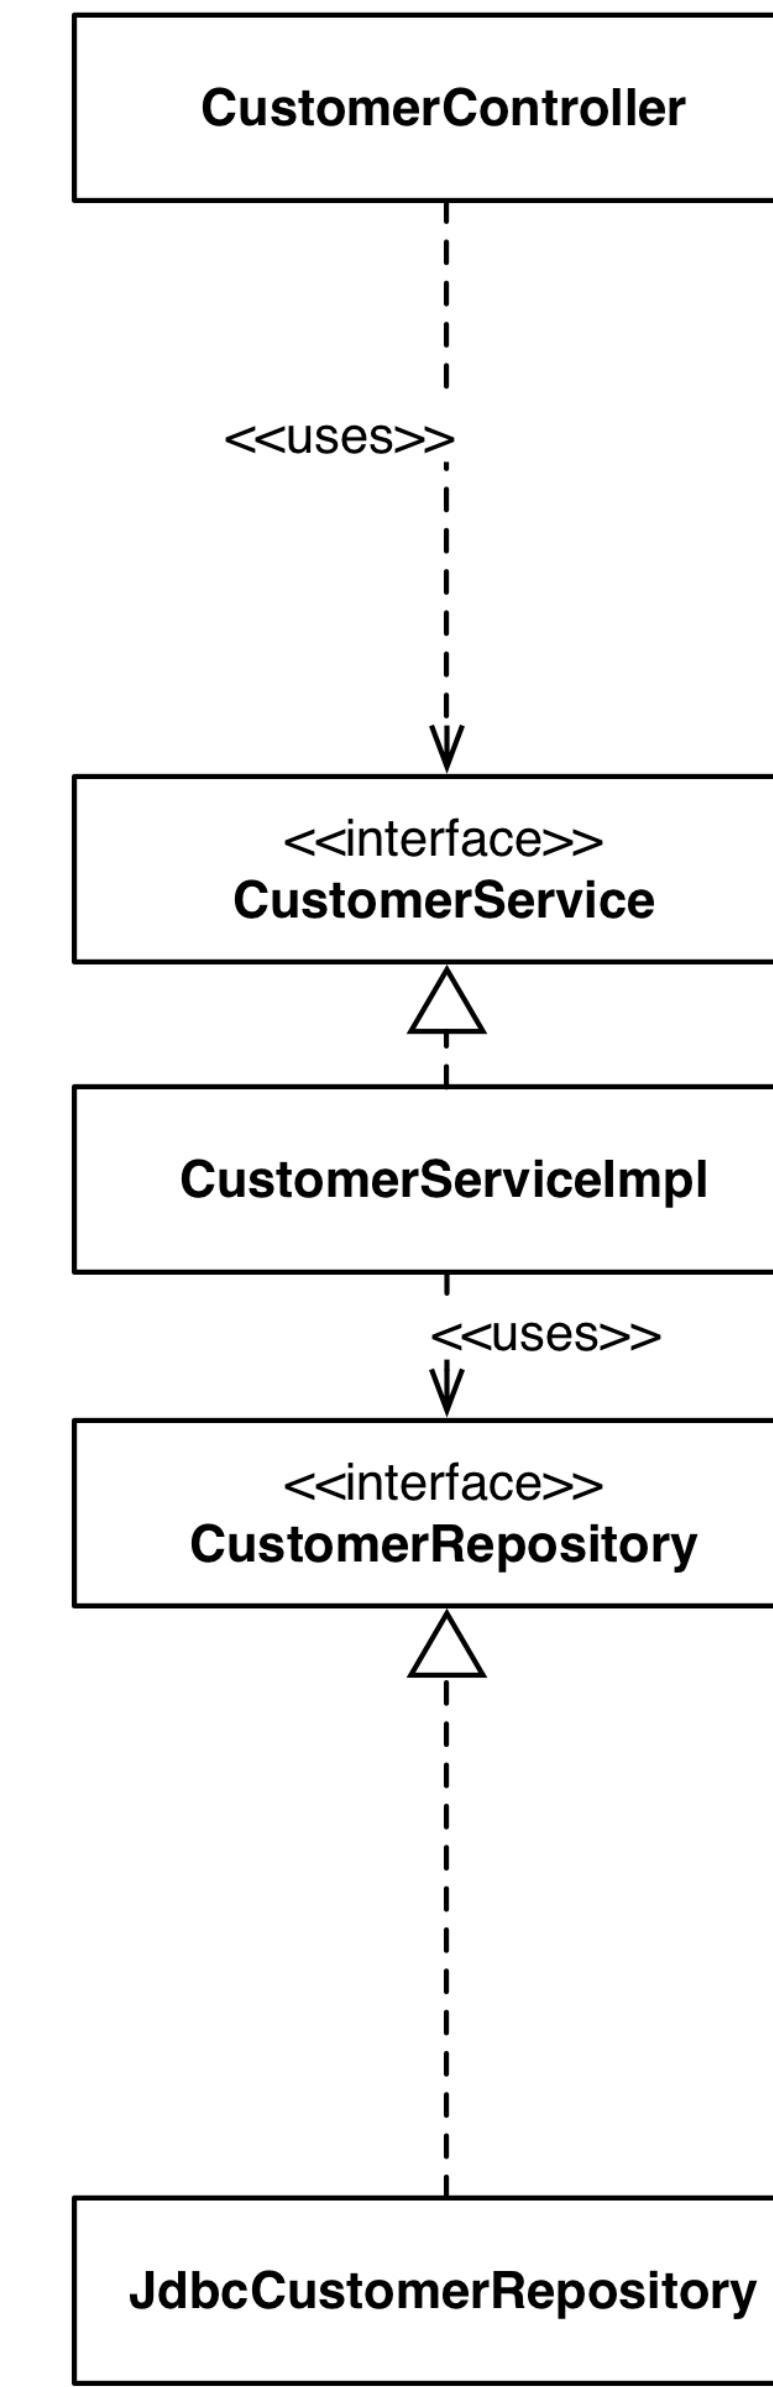
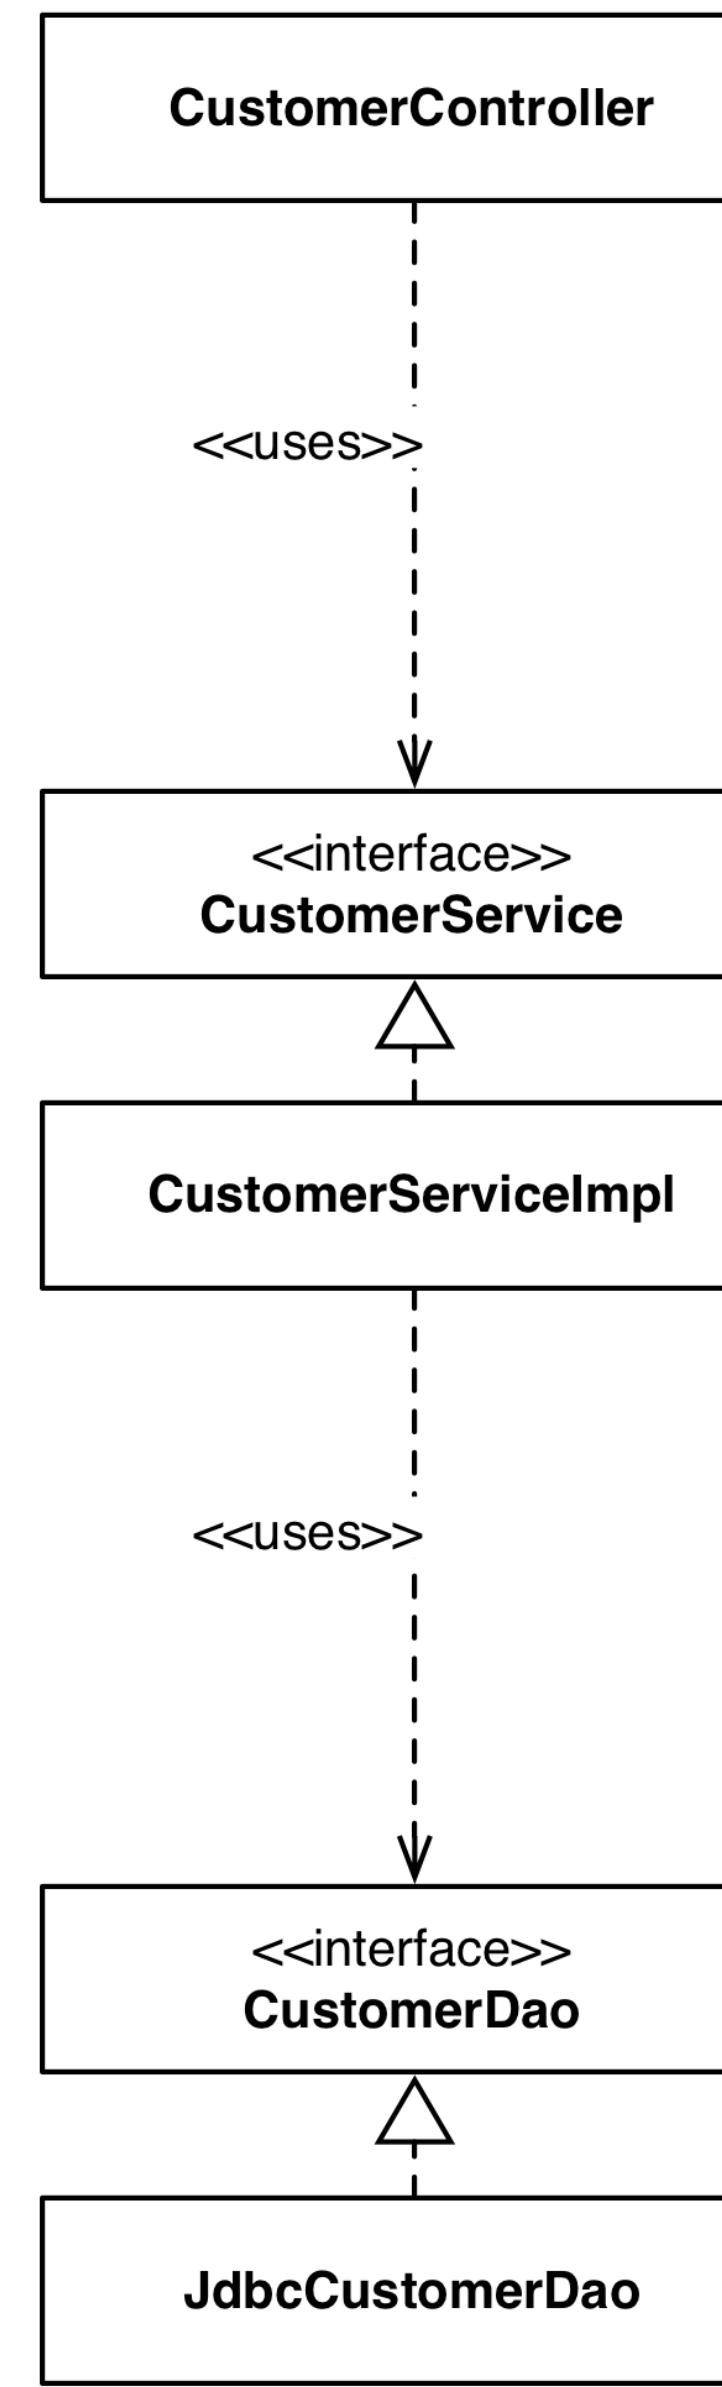
Conscious  
decisions  
rather than cargo cult

# Implementation details

But the devil is in the  
implementation details...

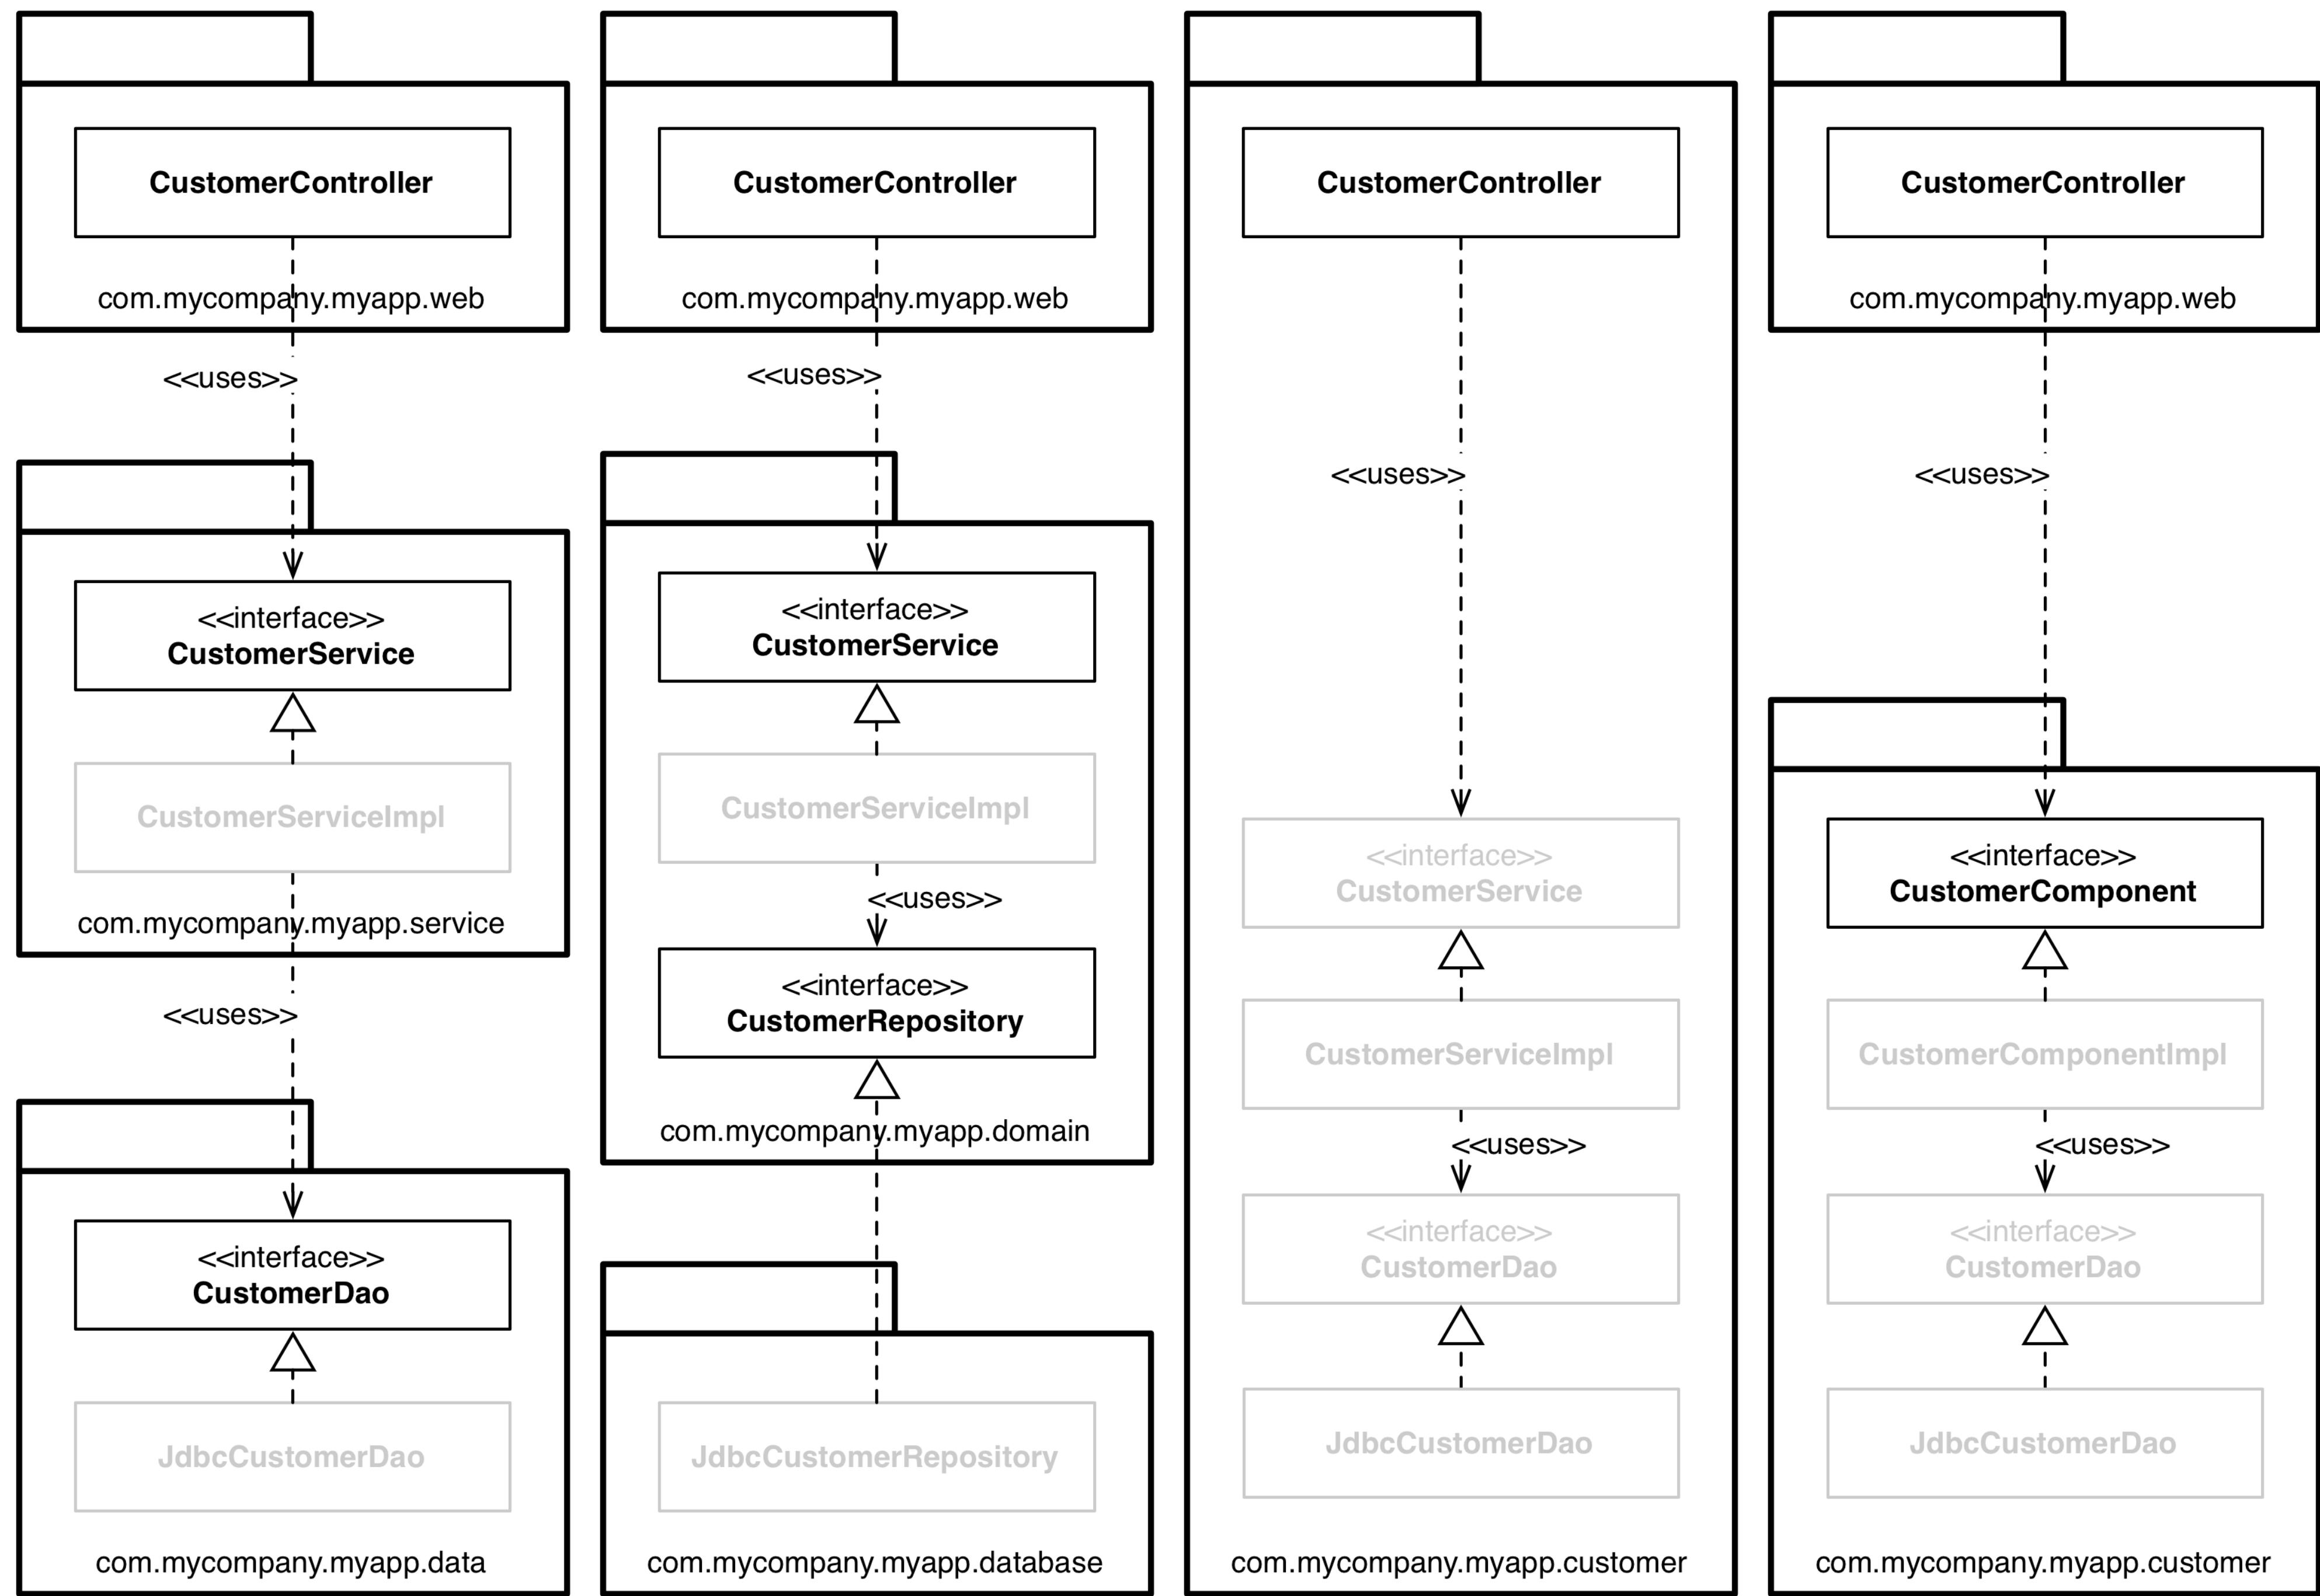


If all types are public,  
Java packages are about  
organisation of code  
rather than encapsulation



The four implementation strategies  
are conceptually different,  
but syntactically the same

Let's bring back the packages and  
mark (fade) those types that  
can be made more restrictive



# Testing

“In the early days of computing when computers were slow, unit tests gave the developer more immediate feedback about whether a change broke the code instead of waiting for system tests to run. Today, with cheaper and more powerful computers, that argument is less persuasive.”



TIME TESTED.  
TESTING IMPROVED.  
[www.RBCS-US.com](http://www.RBCS-US.com)

## Why Most Unit Testing is Waste

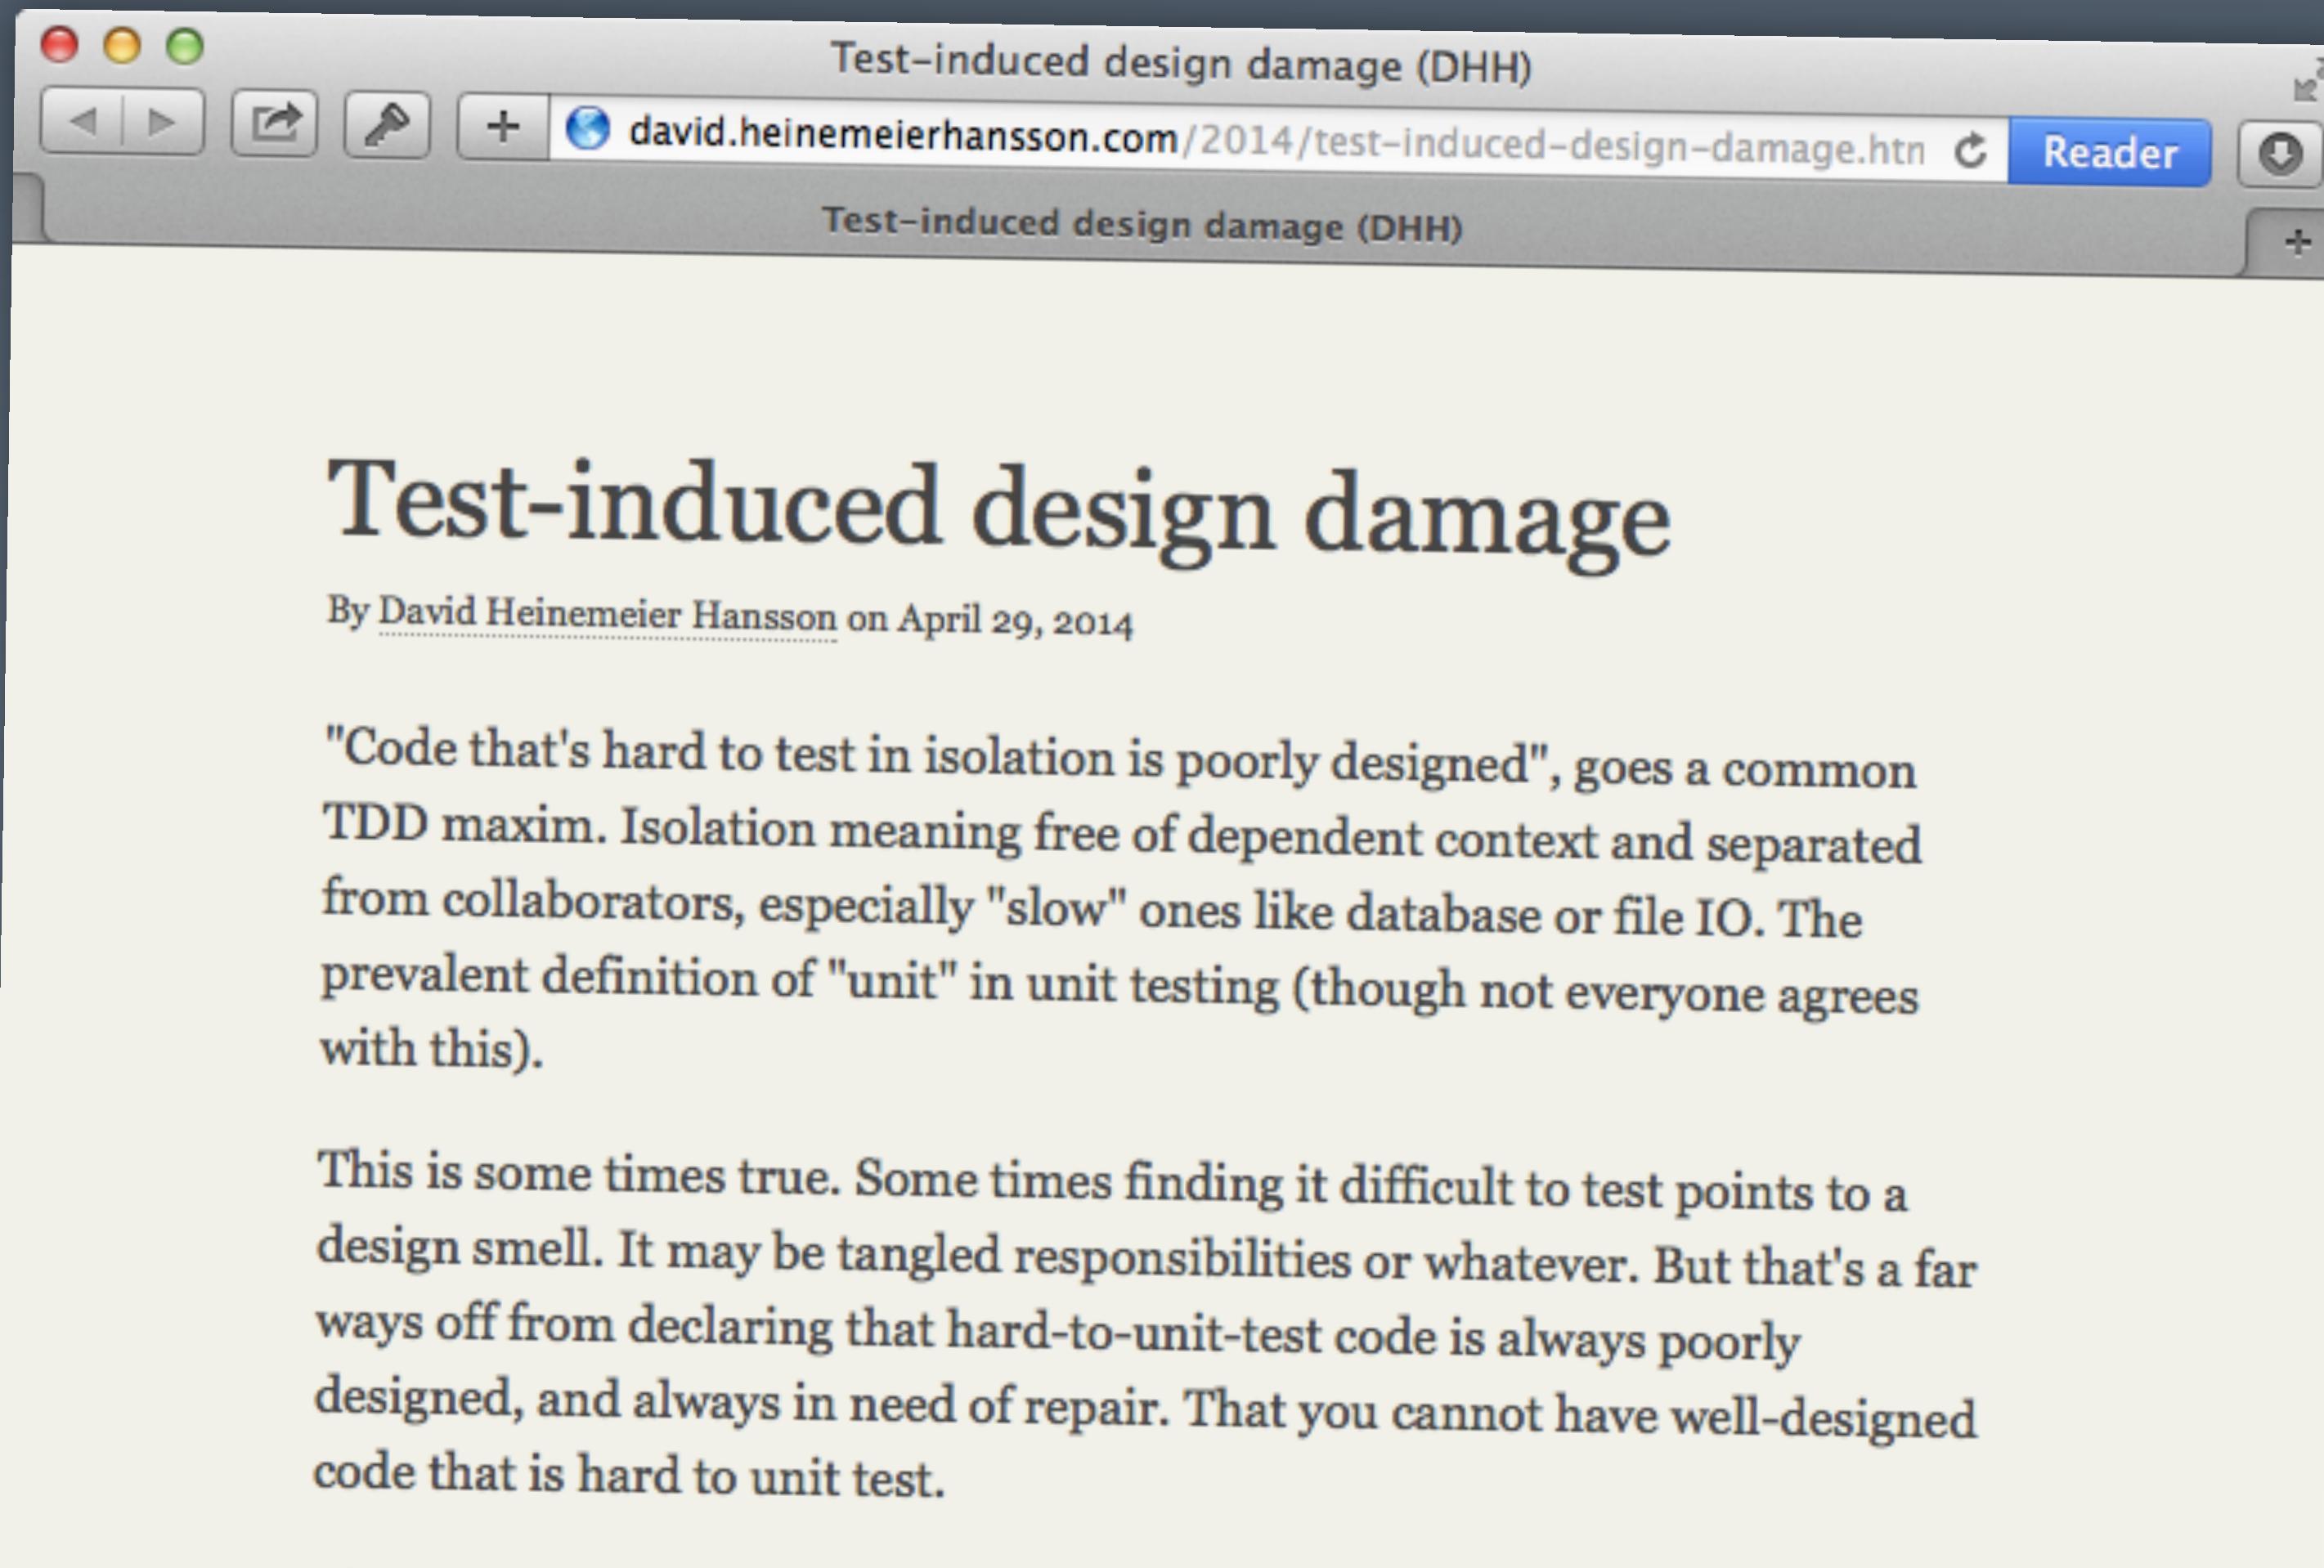
By James O Coplien

### 1.1 Into Modern Times

---

Unit testing was a staple of the FORTRAN days, when a function was a function and was sometimes worthy of functional testing. Computers computed, and functions and procedures represented units of computation. In those days the dominant

# “do not let your tests drive your design”



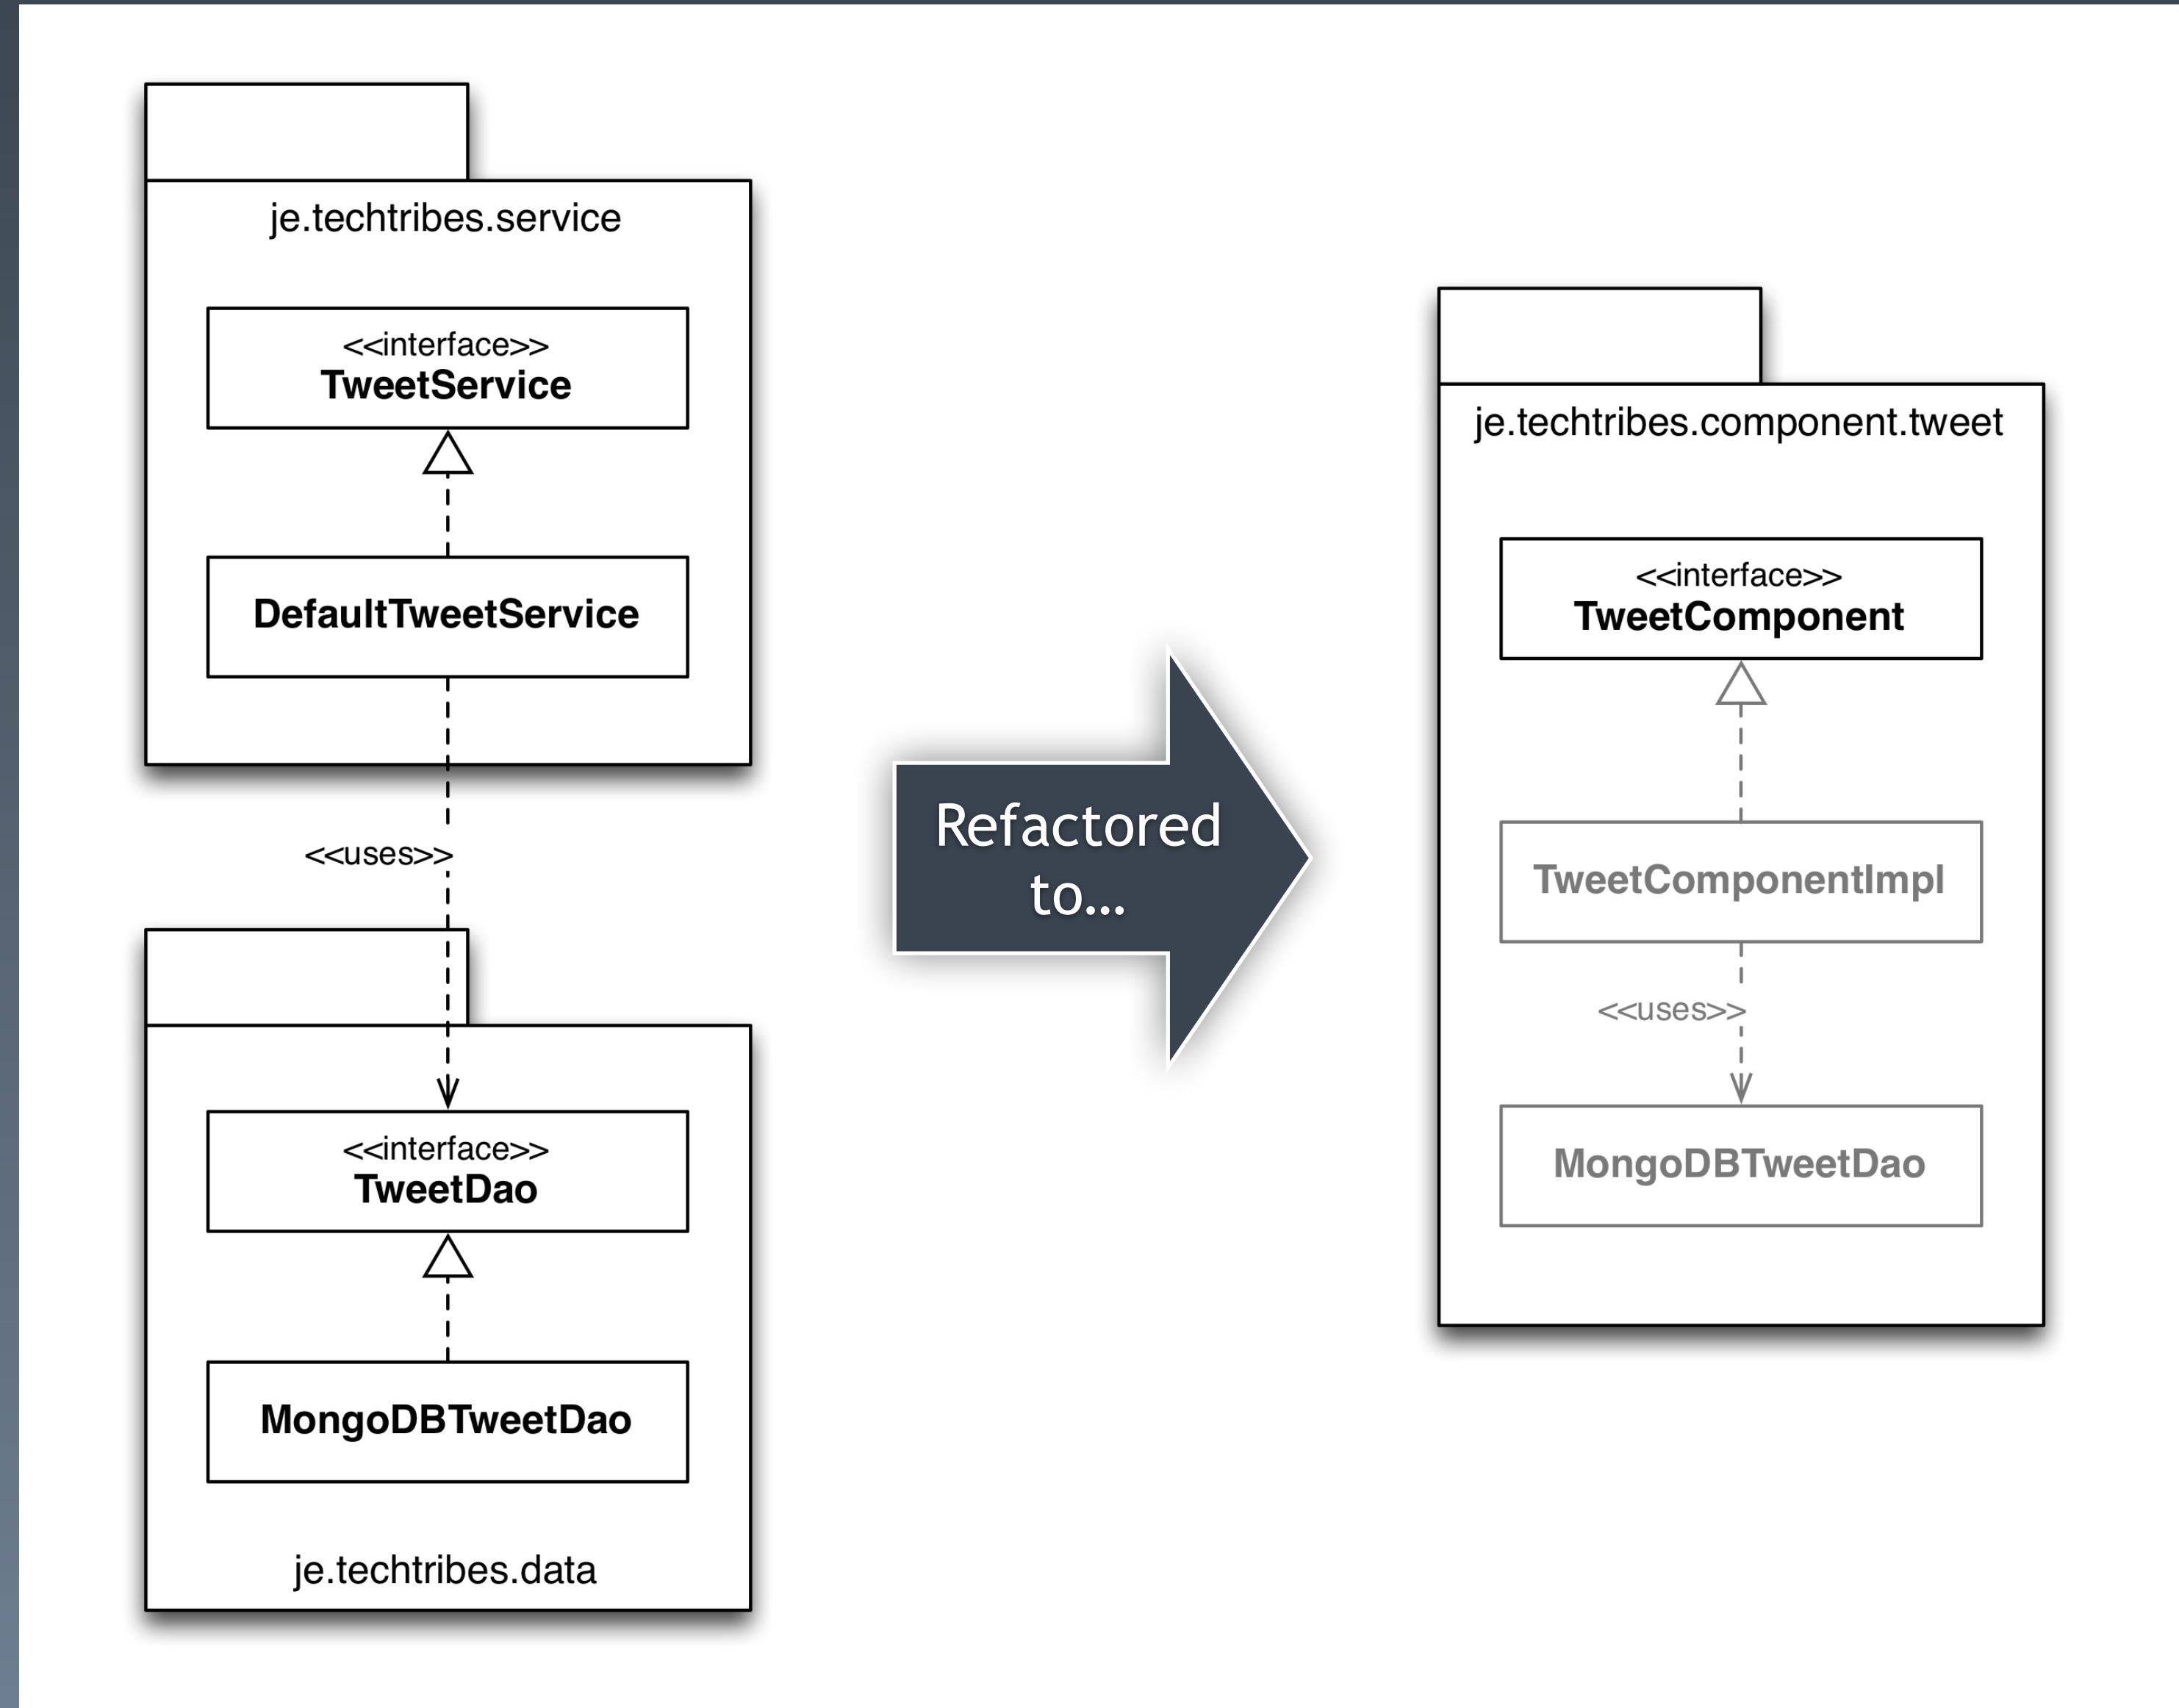
The screenshot shows a web browser window with the title "Test-induced design damage (DHH)" at the top. The URL in the address bar is "david.heinemeierhansson.com/2014/test-induced-design-damage.htm". The page content is an article titled "Test-induced design damage" by David Heinemeier Hansson, dated April 29, 2014. The text discusses the common TDD maxim that "Code that's hard to test in isolation is poorly designed" and argues against it, stating that such code is often a design smell and may be poorly designed, but it does not necessarily mean the code is always poorly designed.

## Test-induced design damage

By David Heinemeier Hansson on April 29, 2014

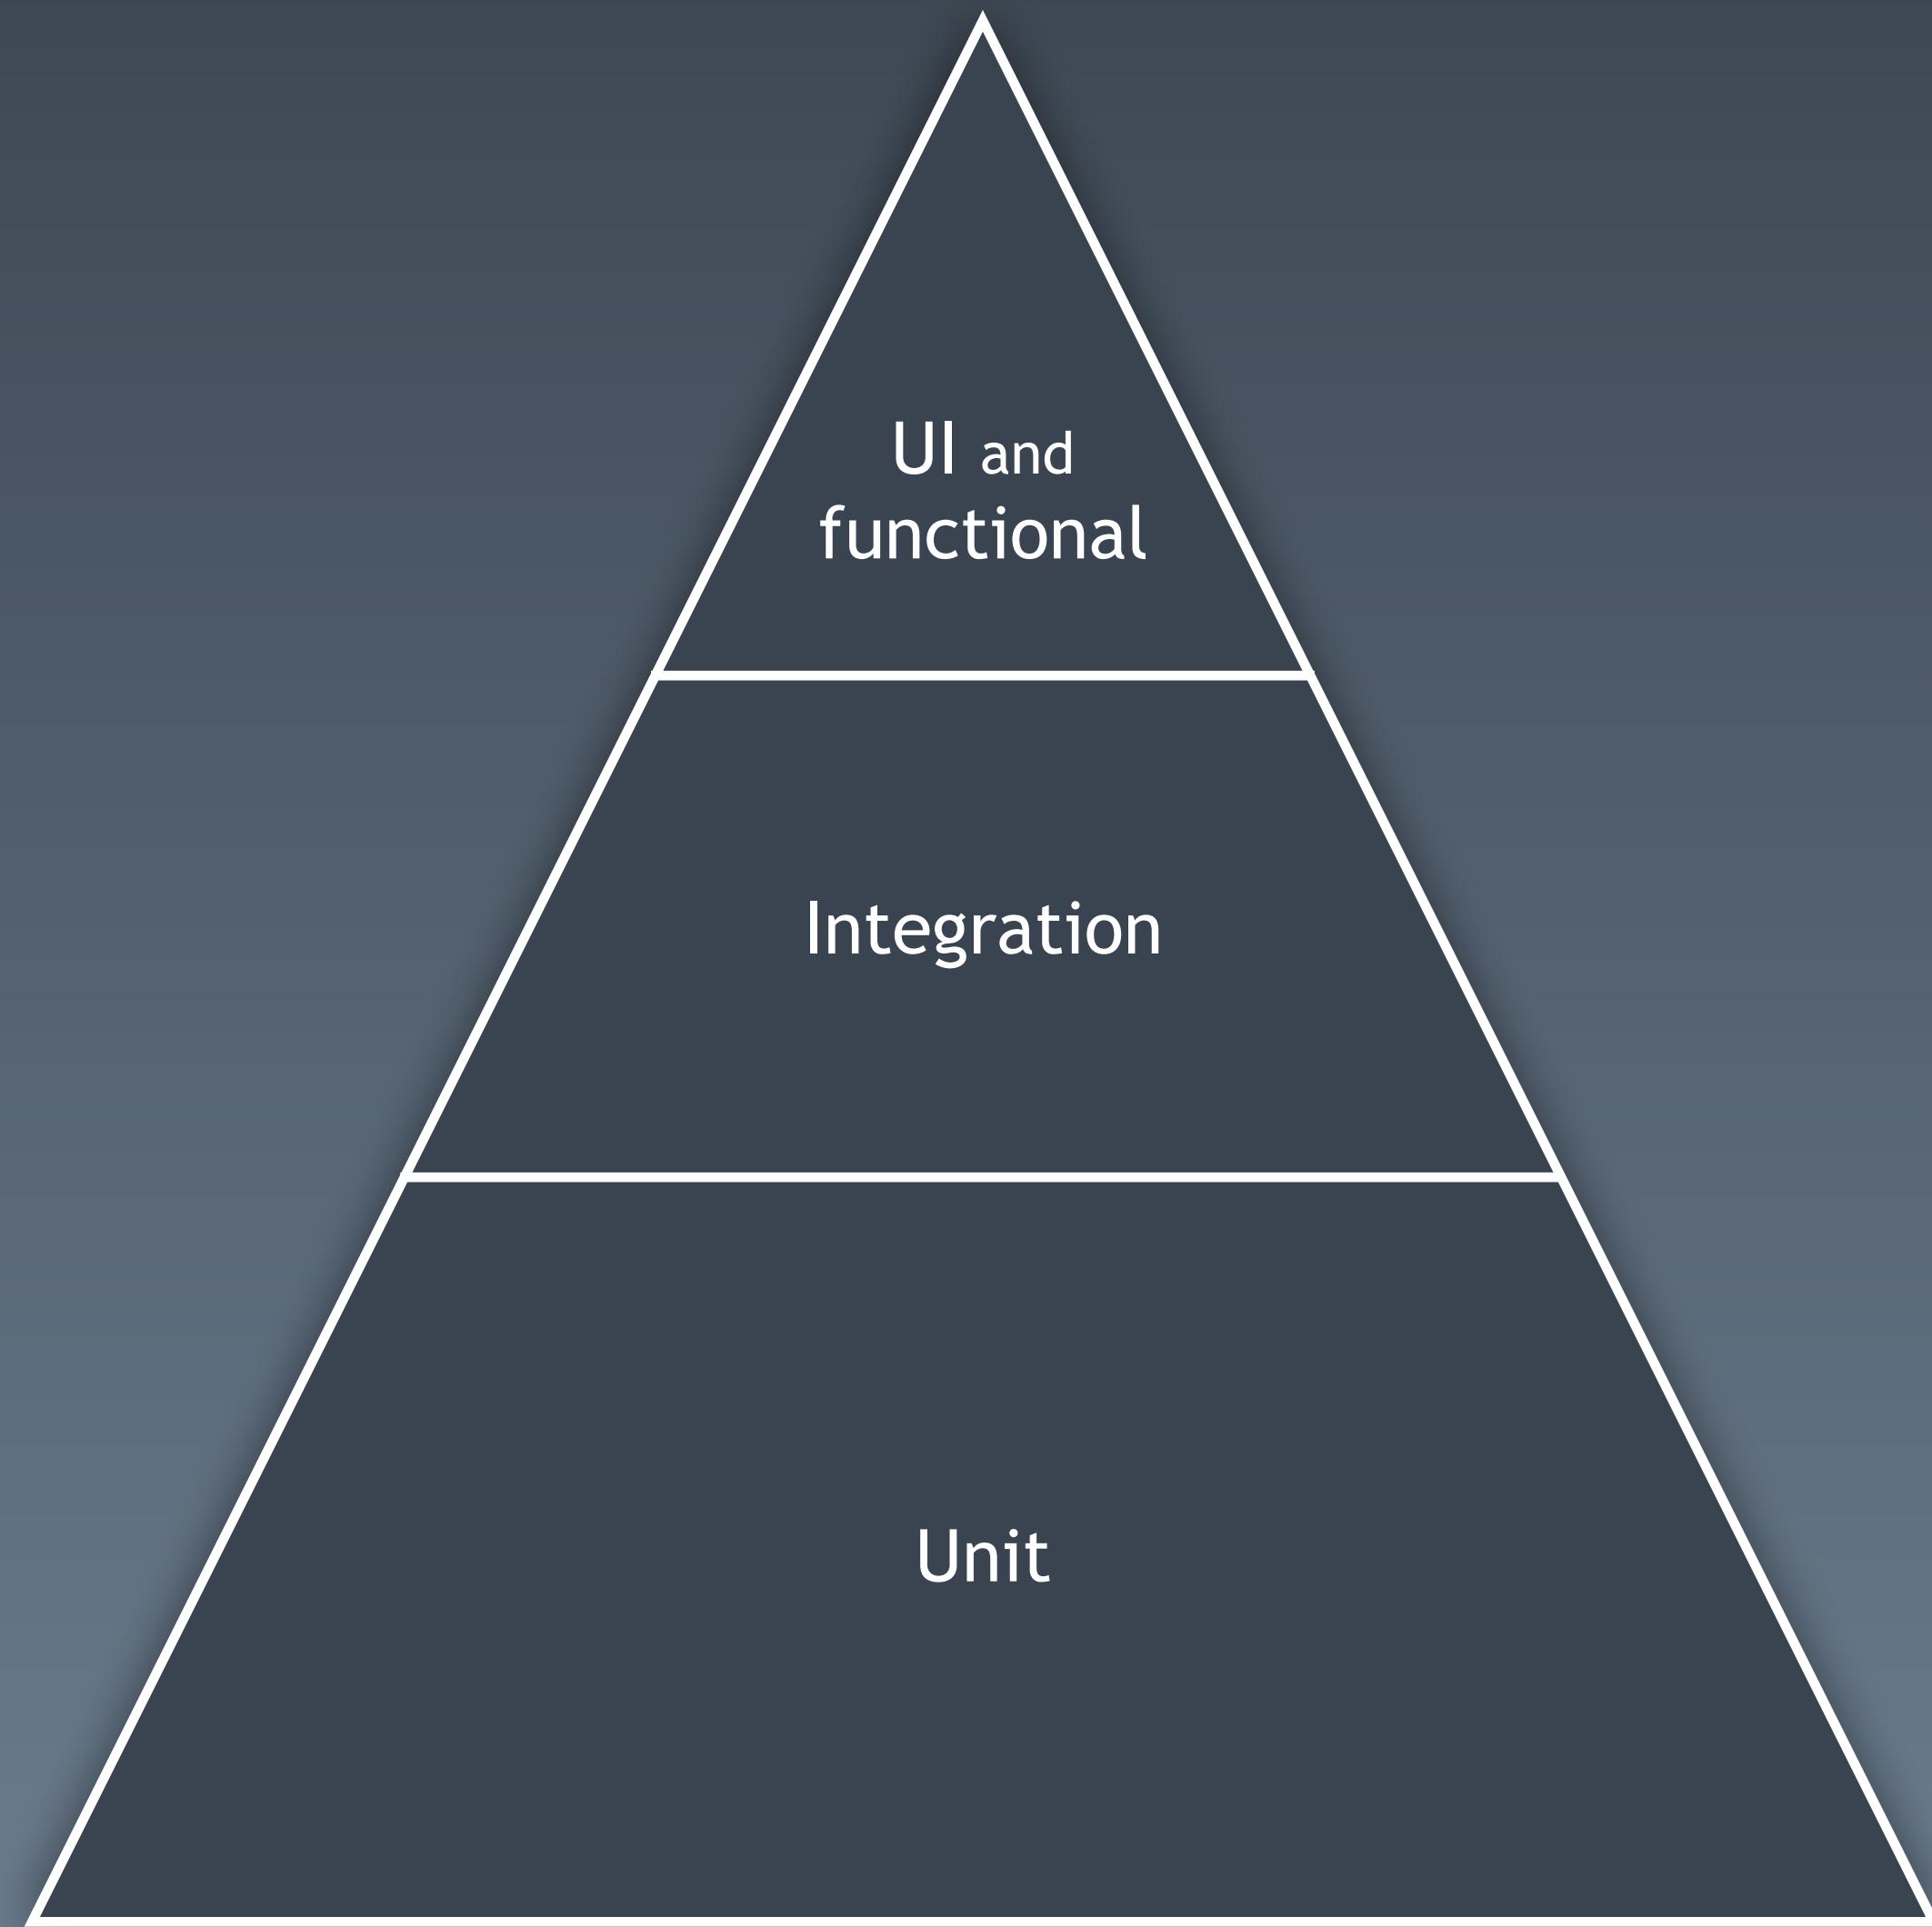
"Code that's hard to test in isolation is poorly designed", goes a common TDD maxim. Isolation meaning free of dependent context and separated from collaborators, especially "slow" ones like database or file IO. The prevalent definition of "unit" in unit testing (though not everyone agrees with this).

This is sometimes true. Some times finding it difficult to test points to a design smell. It may be tangled responsibilities or whatever. But that's a far ways off from declaring that hard-to-unit-test code is always poorly designed, and always in need of repair. That you cannot have well-designed code that is hard to unit test.



An example of an  
architecturally-evident coding style

Do we also  
cargo cult  
“unit testing”?



Do we need to rethink the testing pyramid?

Service

Uses

Repository

Component

Test the service in isolation?

(“unit tests”, with a mock repository)

Test the repository in isolation?

(“integration tests”, against the database)

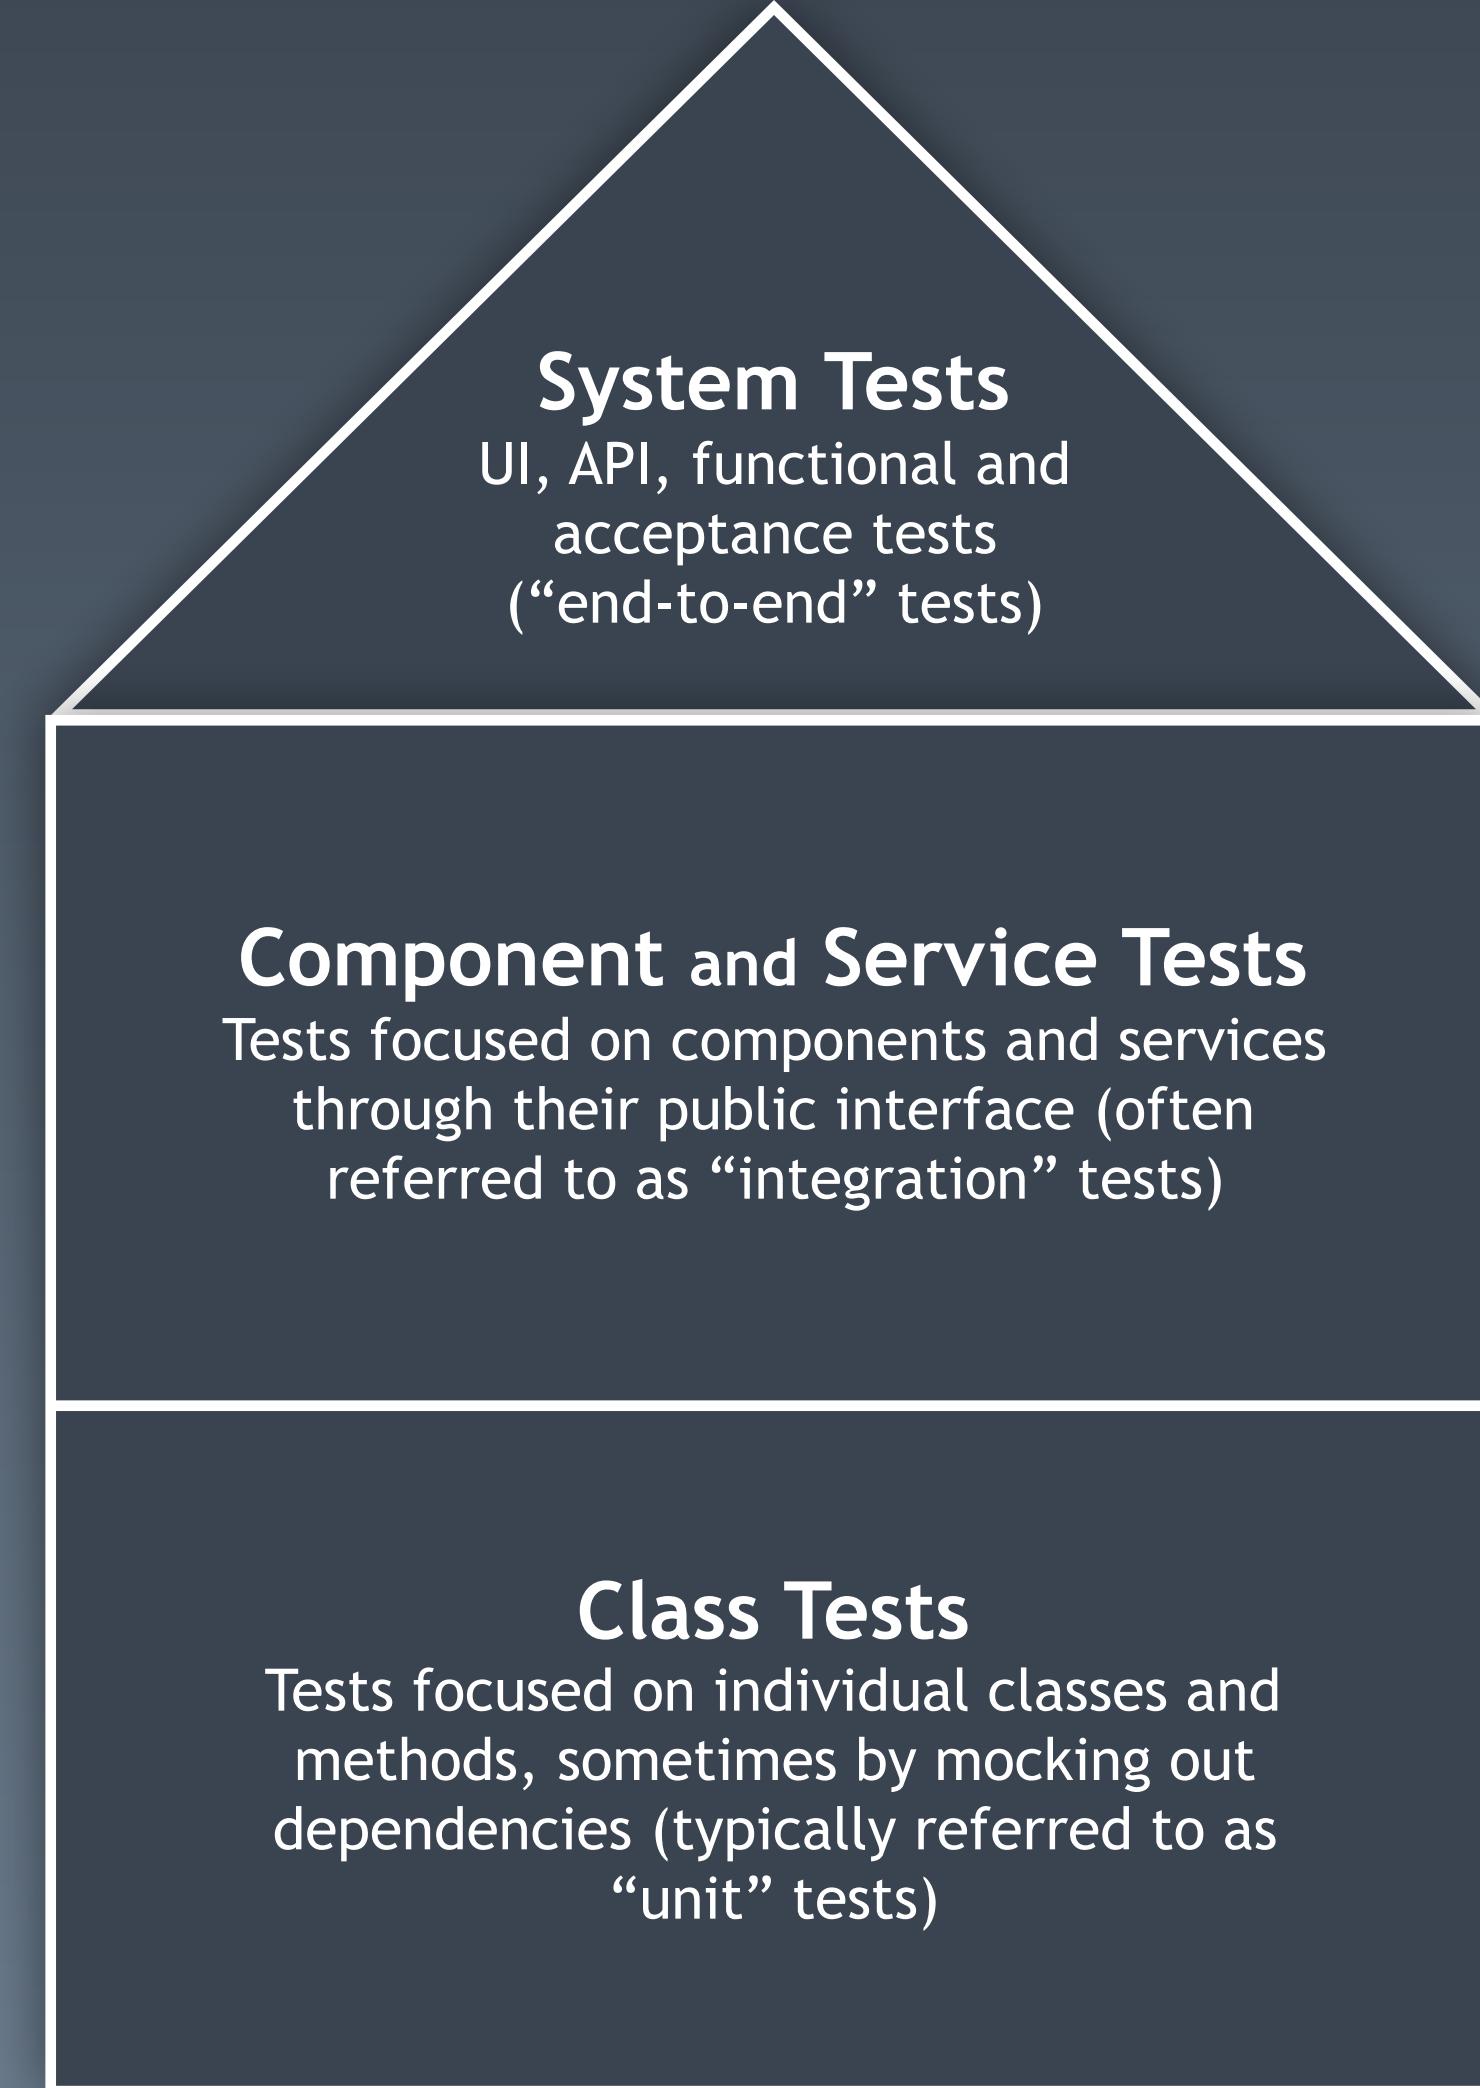
Or test the component through  
the public interface?

(“component tests”)

Instead of blindly unit testing everything, what about  
testing your significant structural  
elements as black boxes?

# Caveats apply!

Component testing doesn't necessarily suit asynchronous behaviour or integration with third-party systems outside of your control

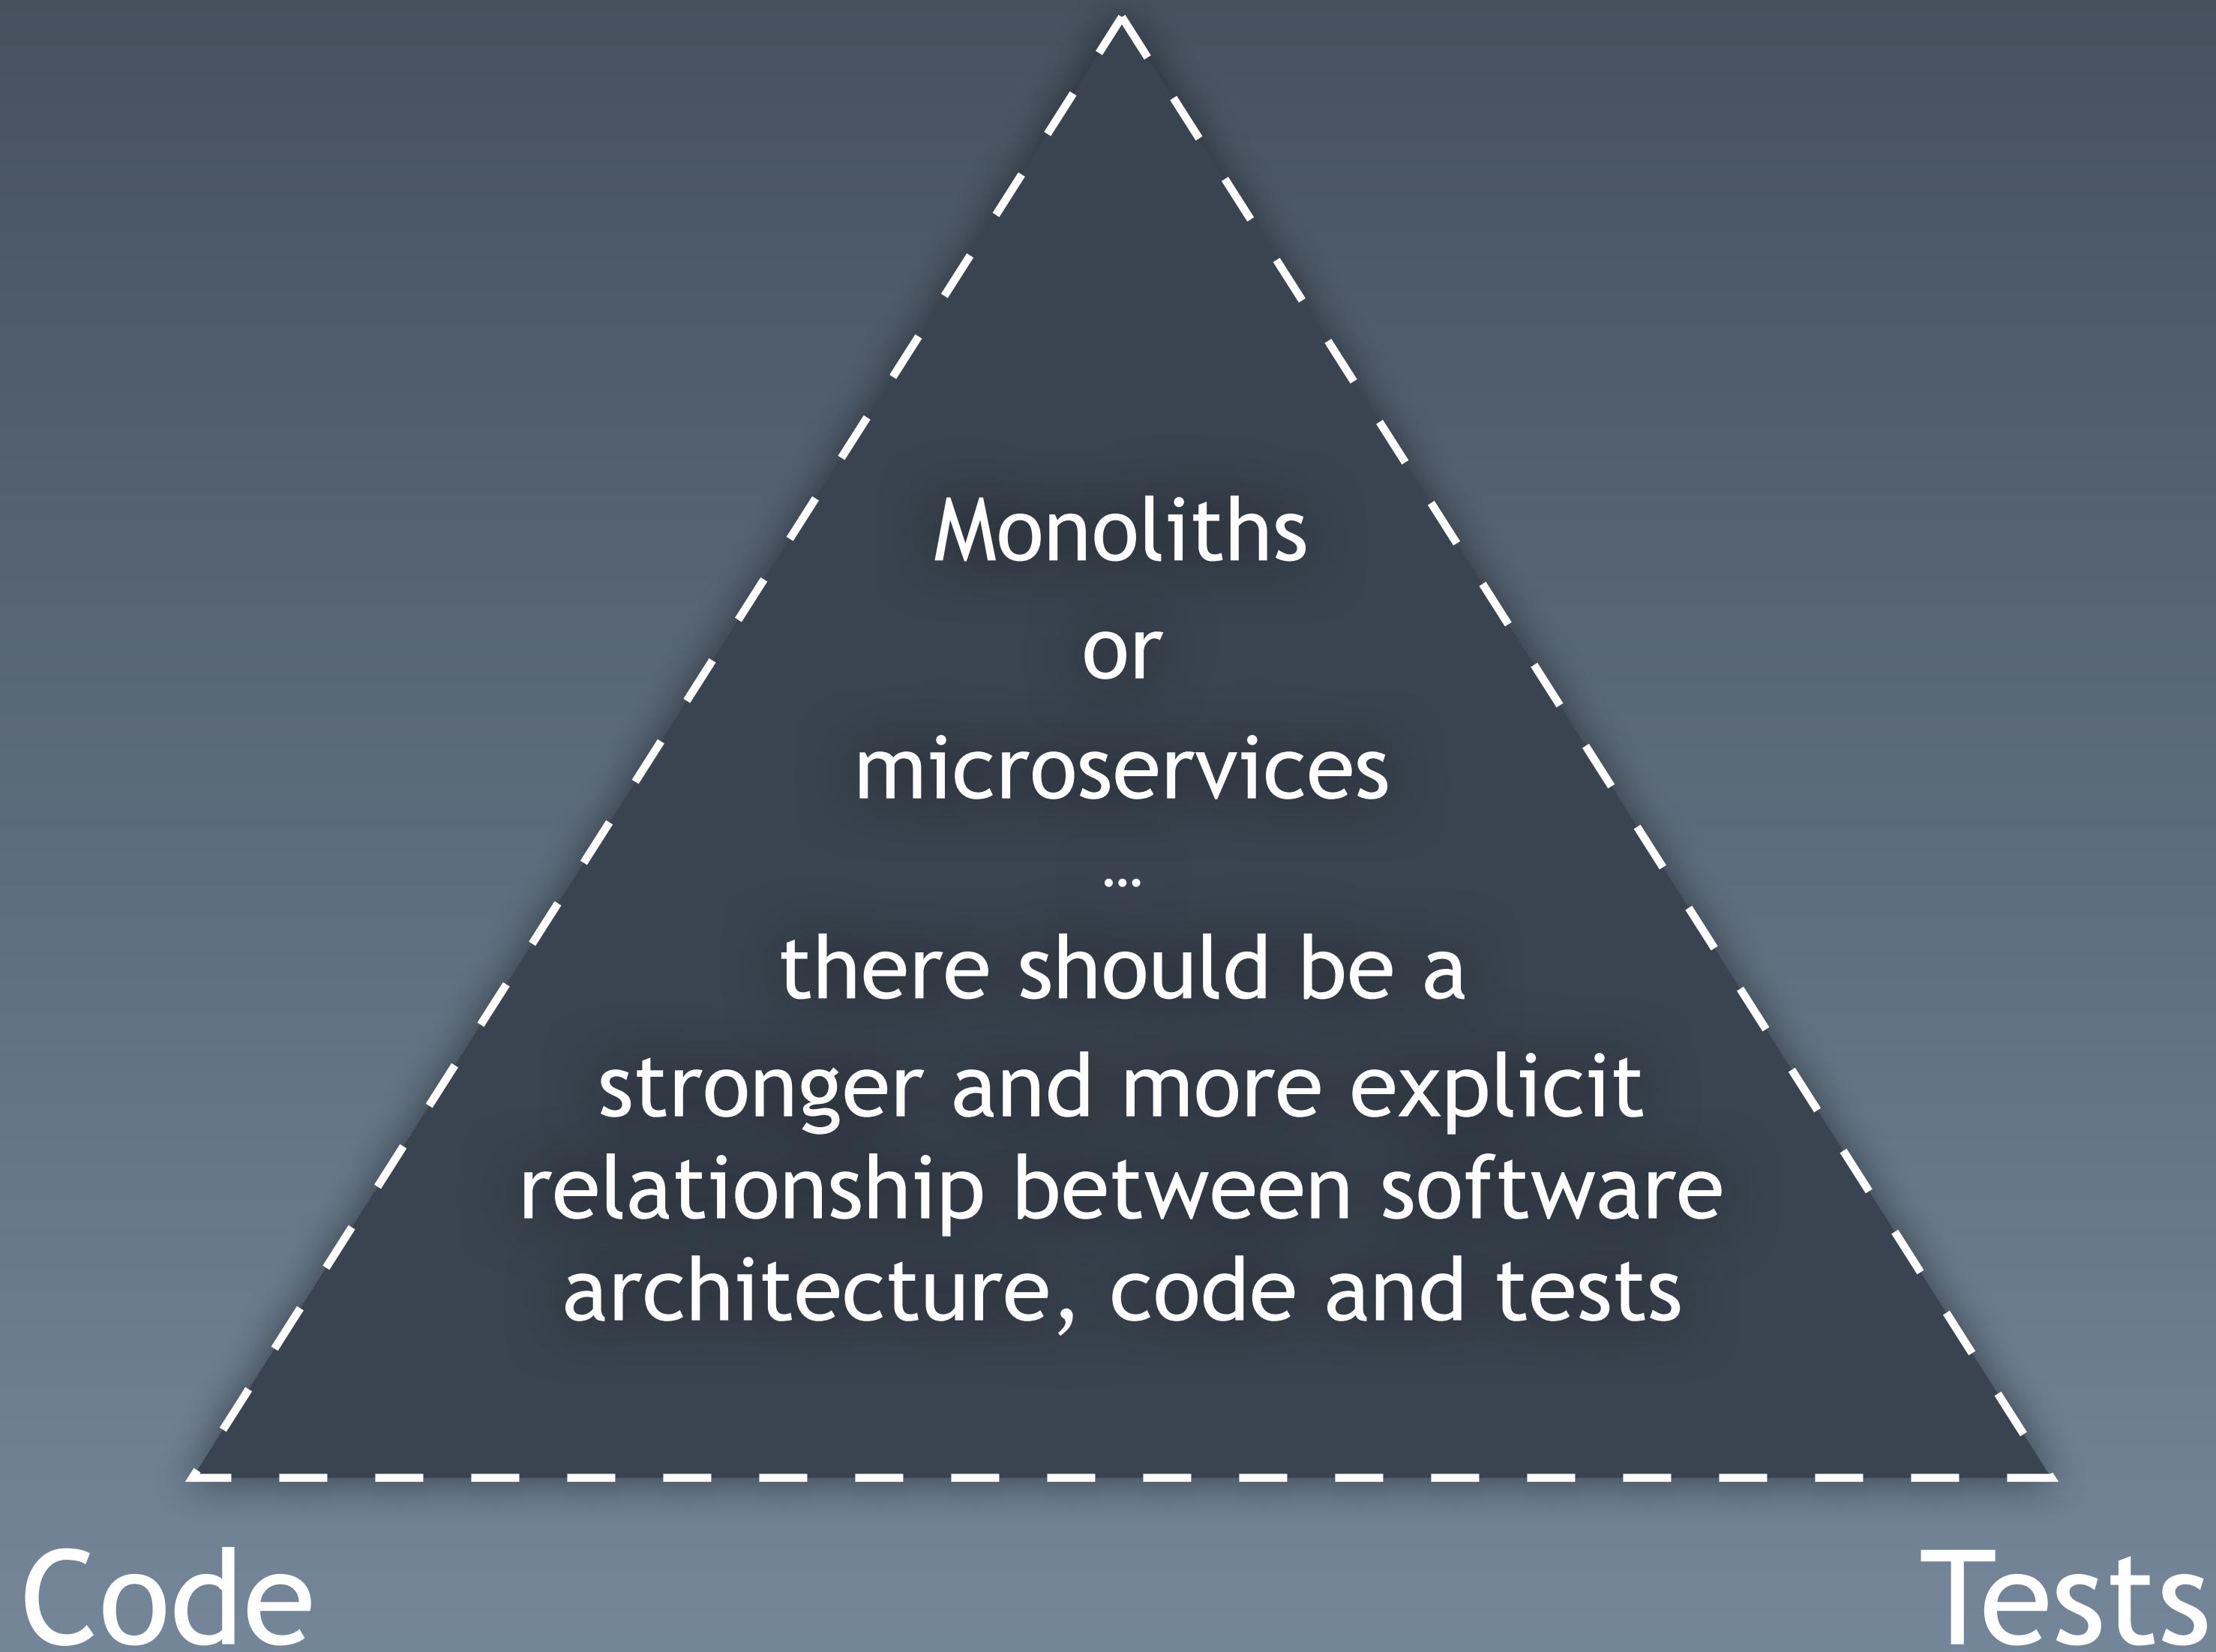


# Architecturally-aligned testing (and a reshaped testing pyramid)

I want to achieve maximum code coverage from a minimum amount of test code

(unit testing everything is one extreme,  
writing one gigantic acceptance test is the other  
... there is a sweet spot somewhere in-between)

# Software Architecture



# Why?

Maintainability is  
inversely proportional  
to the number of

public classes  
dependencies  
microservices

A good  
architecture  
enables

*agility*

*Agility*

is a

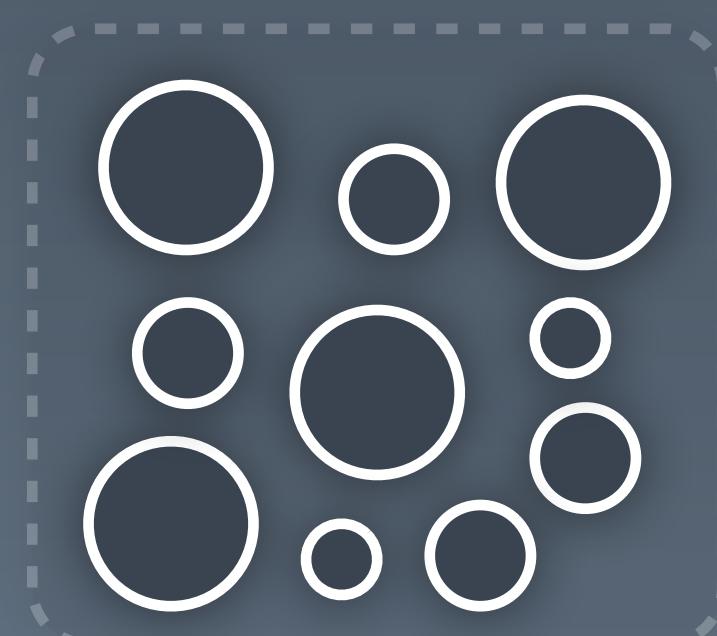
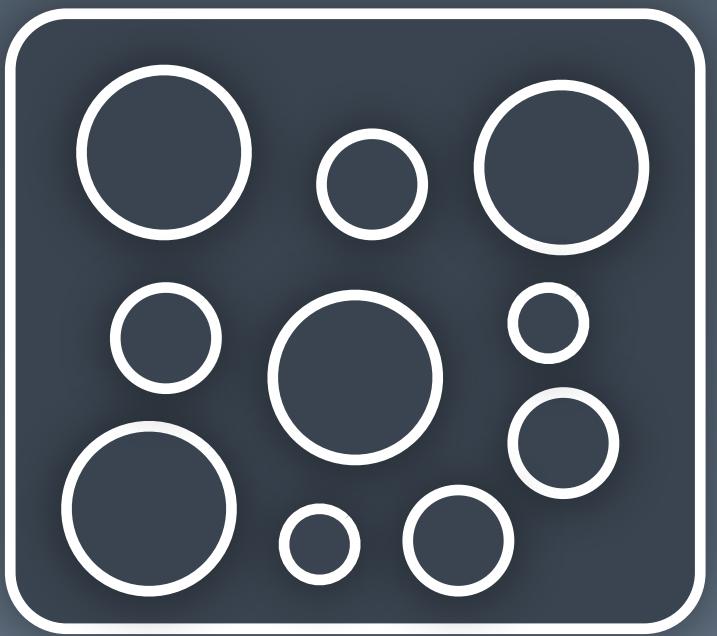
quality attribute

*Monolithic  
architecture*



*Service-based  
architecture*

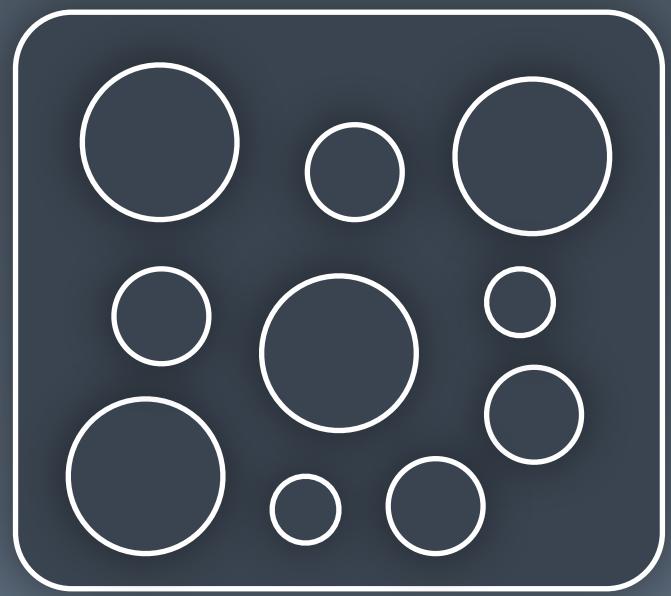
(SOA, microservices, etc)



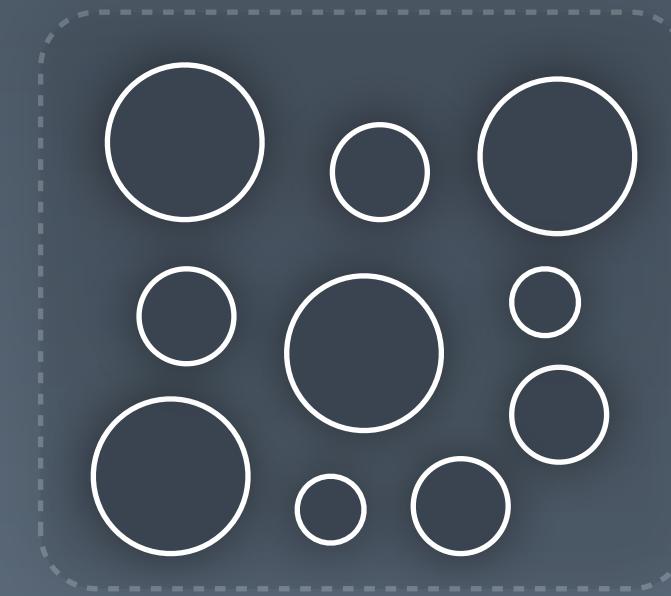
*Something in between  
(components)*



# Well-defined, in-process components is a stepping stone to out-of-process components (i.e. microservices)



From components  
to microservices



High cohesion  
Low coupling  
Focussed on a business capability  
Bounded context or aggregate  
Encapsulated data  
Substitutable  
Composable

<- All of that plus  
Individually deployable  
Individually upgradeable  
Individually replaceable  
Individually scalable  
Heterogeneous technology stacks

Choose  
microservices for  
the benefits,  
not because your monolithic  
codebase is a mess

Microservices ... also known as a  
distributed  
big ball of mud

If your software system  
is hard to work with,

*Change it!*

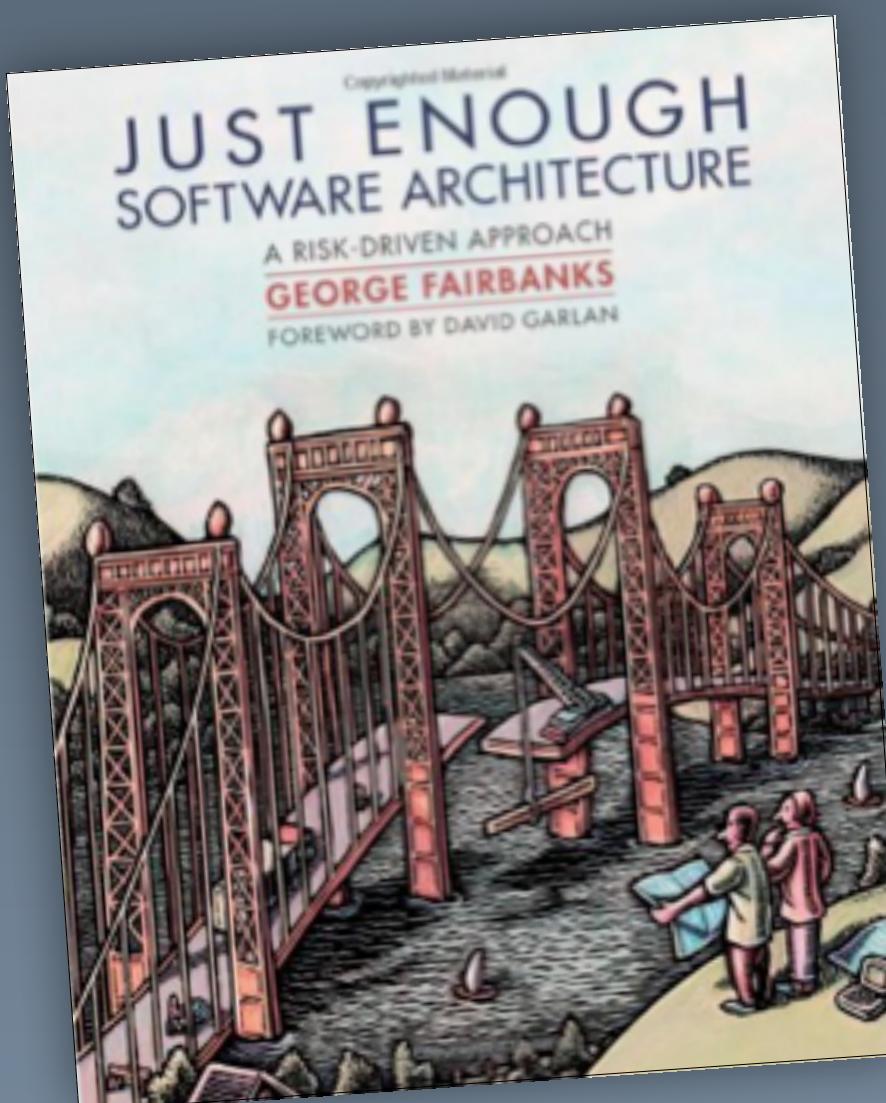
Think about how to align the

# software architecture

and the

code

Be conscious of the software  
architecture model ... adopt an  
architecturally-evident  
coding style



Stop making every class  
public

£1 charity donation  
every time you type

public class

without thinking :-)

If you can't build a  
modular  
monolith,  
what makes you think  
microservices is the answer?



[simon.brown@codingthearchitecture.com](mailto:simon.brown@codingthearchitecture.com)  
@simonbrown on Twitter